



ACV – Applied Computer Vision

Bachelor Medientechnik & Creative Computing

Matthias Zeppelzauer

matthias.zeppelzauer@fhstp.ac.at

Djordje Slijepcevic

djordje.slijepcevic@fhstp.ac.at

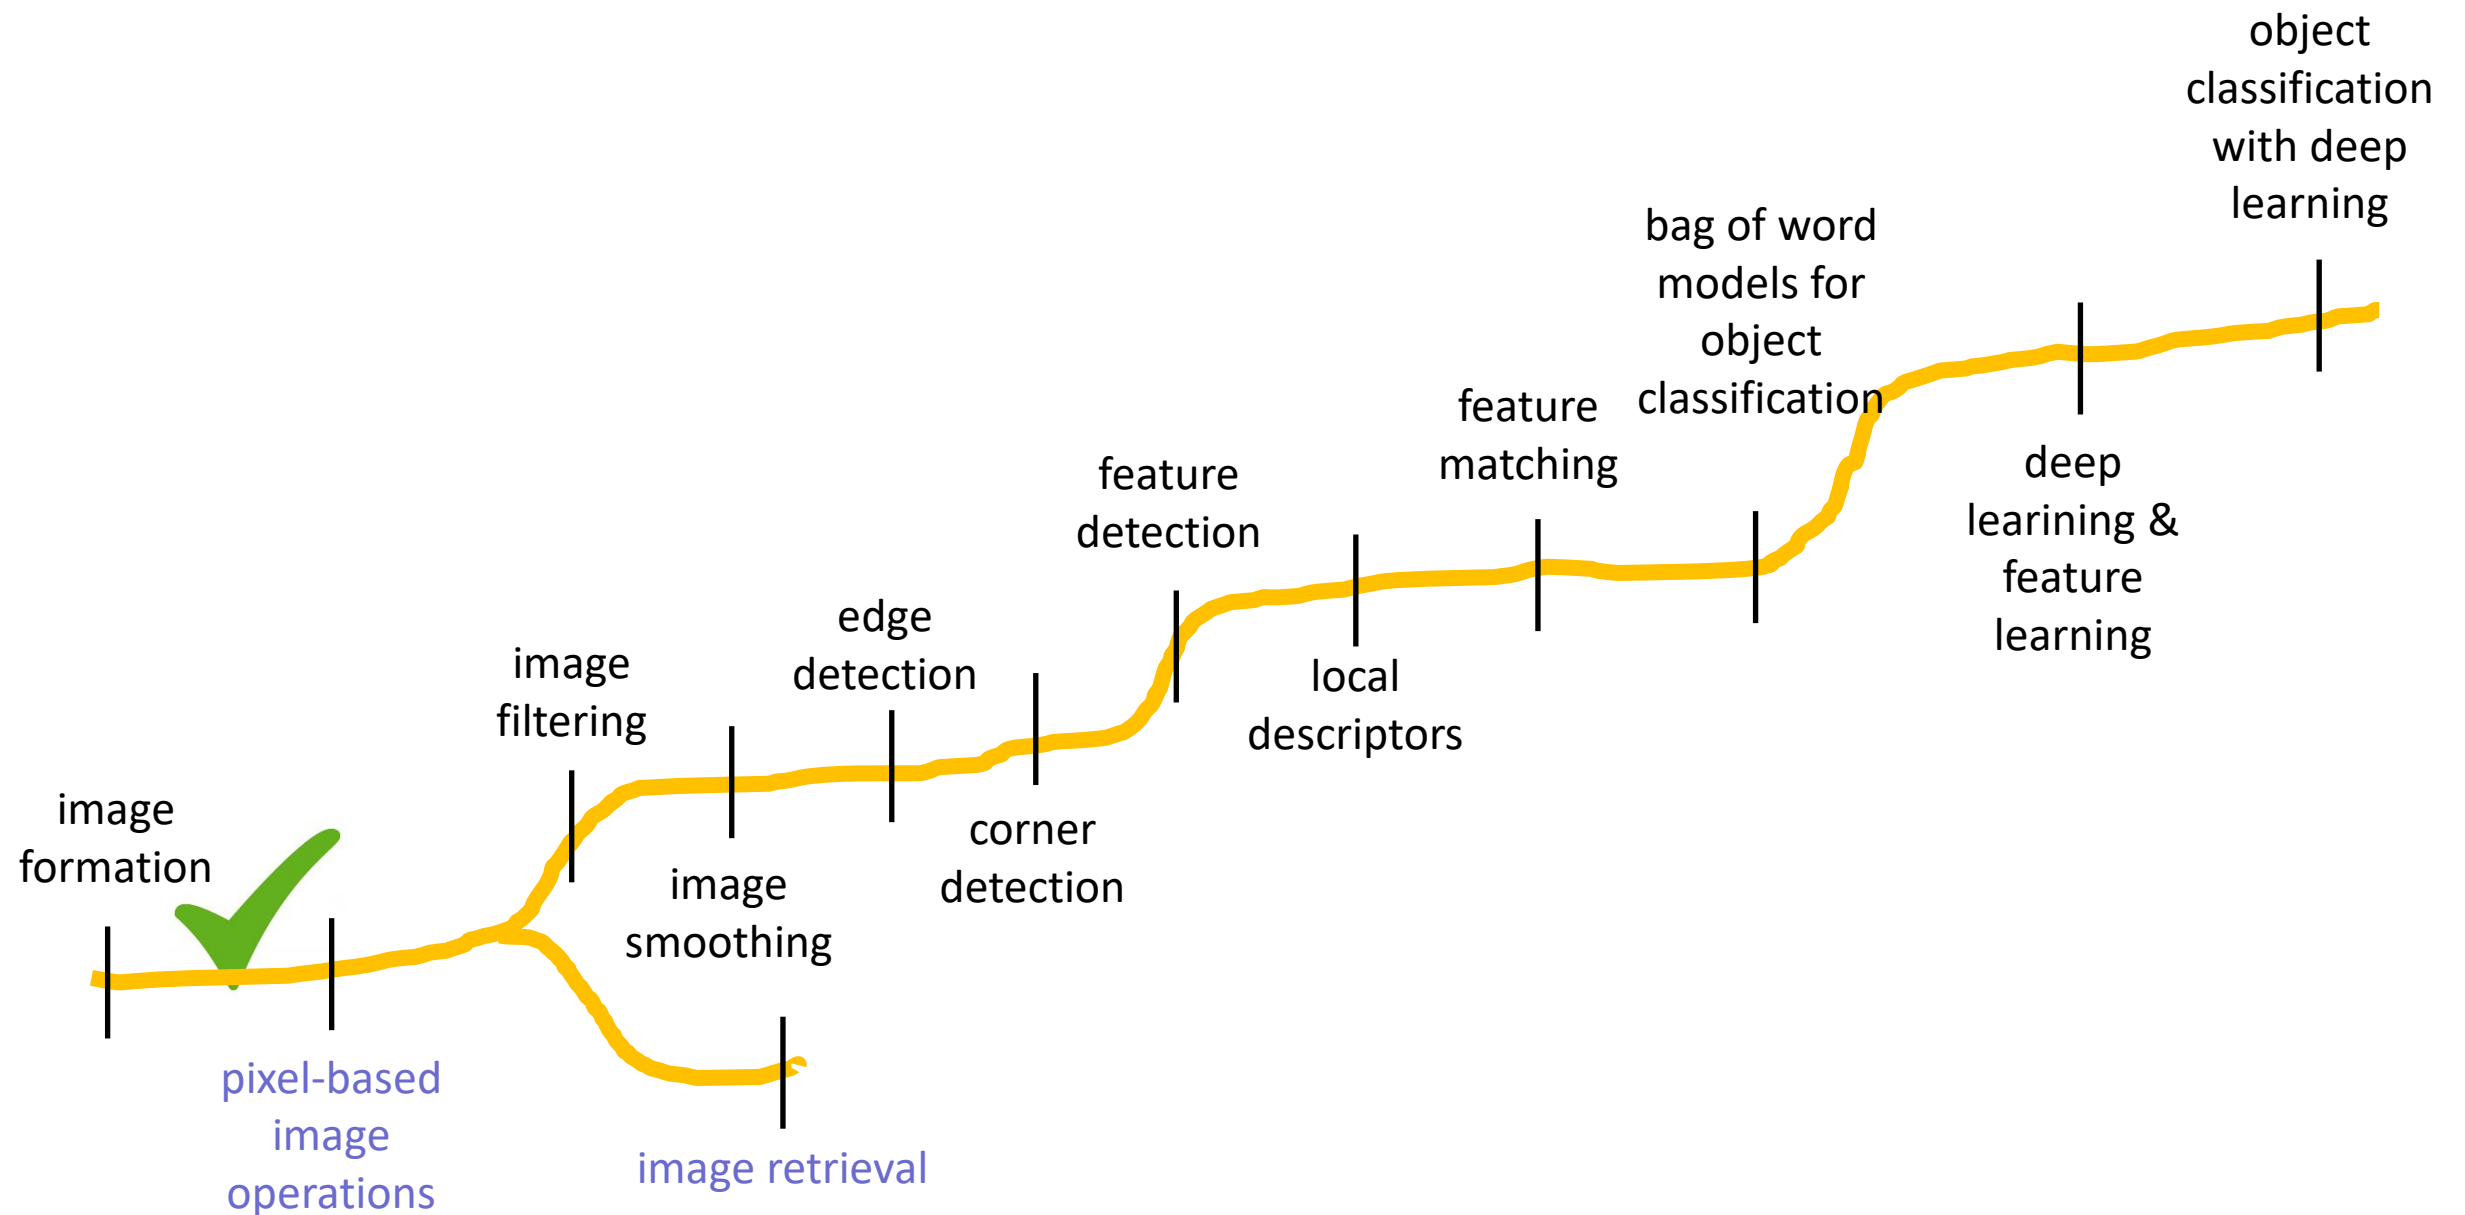
Credits

- Antonio Torralba
- Alexei Efros
- Leon Sigal
- Kristen Grauman
- Steve Seitz
- James Hays

Today

- Last time:
 - Challenges of computer vision
 - Image formation (pinhole camera etc.)
- Basic image operations
- Histograms
- Image Filtering

Roadmap






Basic Image Operations

Image Operations - Overview


- Arithmetic Operations
 - Add, subtract, multiply, divide, average
 - Not, and, or, xor
- Histograms
 - Applications: contrast improvement, threshold finding
- Thresholding
- Filtering
 - Applications: smoothing, sharpening, edge detection

Arithmetic Operations - Addition


- Addition:
 - Used to „blend“ 2 images
 - Pixels in the same coordinates are added to each other
- Python: `im3 = im2 + im1`
 - Values become larger, meaning the picture becomes brighter
 - Values could overflow (for example, if you have a range of 0 ... 255), so you may need to normalize after adding the images



+



=




Arithmetic Operations - Subtraction

- Substraction:
 - Pixels on the same coordinates are substracted from each other
 - Can be used to find differences in two images of the same scene
 - Areas that did not change turn black
 - Also used to detect motion in subsequent video frames
- Python: `im3 = np.abs(im2 - im1)`
 - It is important to take the „absolute“ value of the difference, because you cannot display negative values




Arithmetic Operations - Multiplication


- Multiplication:
 - It scales the grayscale values in the image
 - $C > 1$: Image becomes brighter, $C < 1$: Image becomes darker
 - $C > 1$: You might require clipping, to prevent values from going over the maximum value
- Python: `im2 = im1 * C`
- Using this method for correcting the brightness or contrast of an image is suboptimal, histogram-based approaches provide better results



Original



C=2



C=0.5

Arithmetic Operations - Division

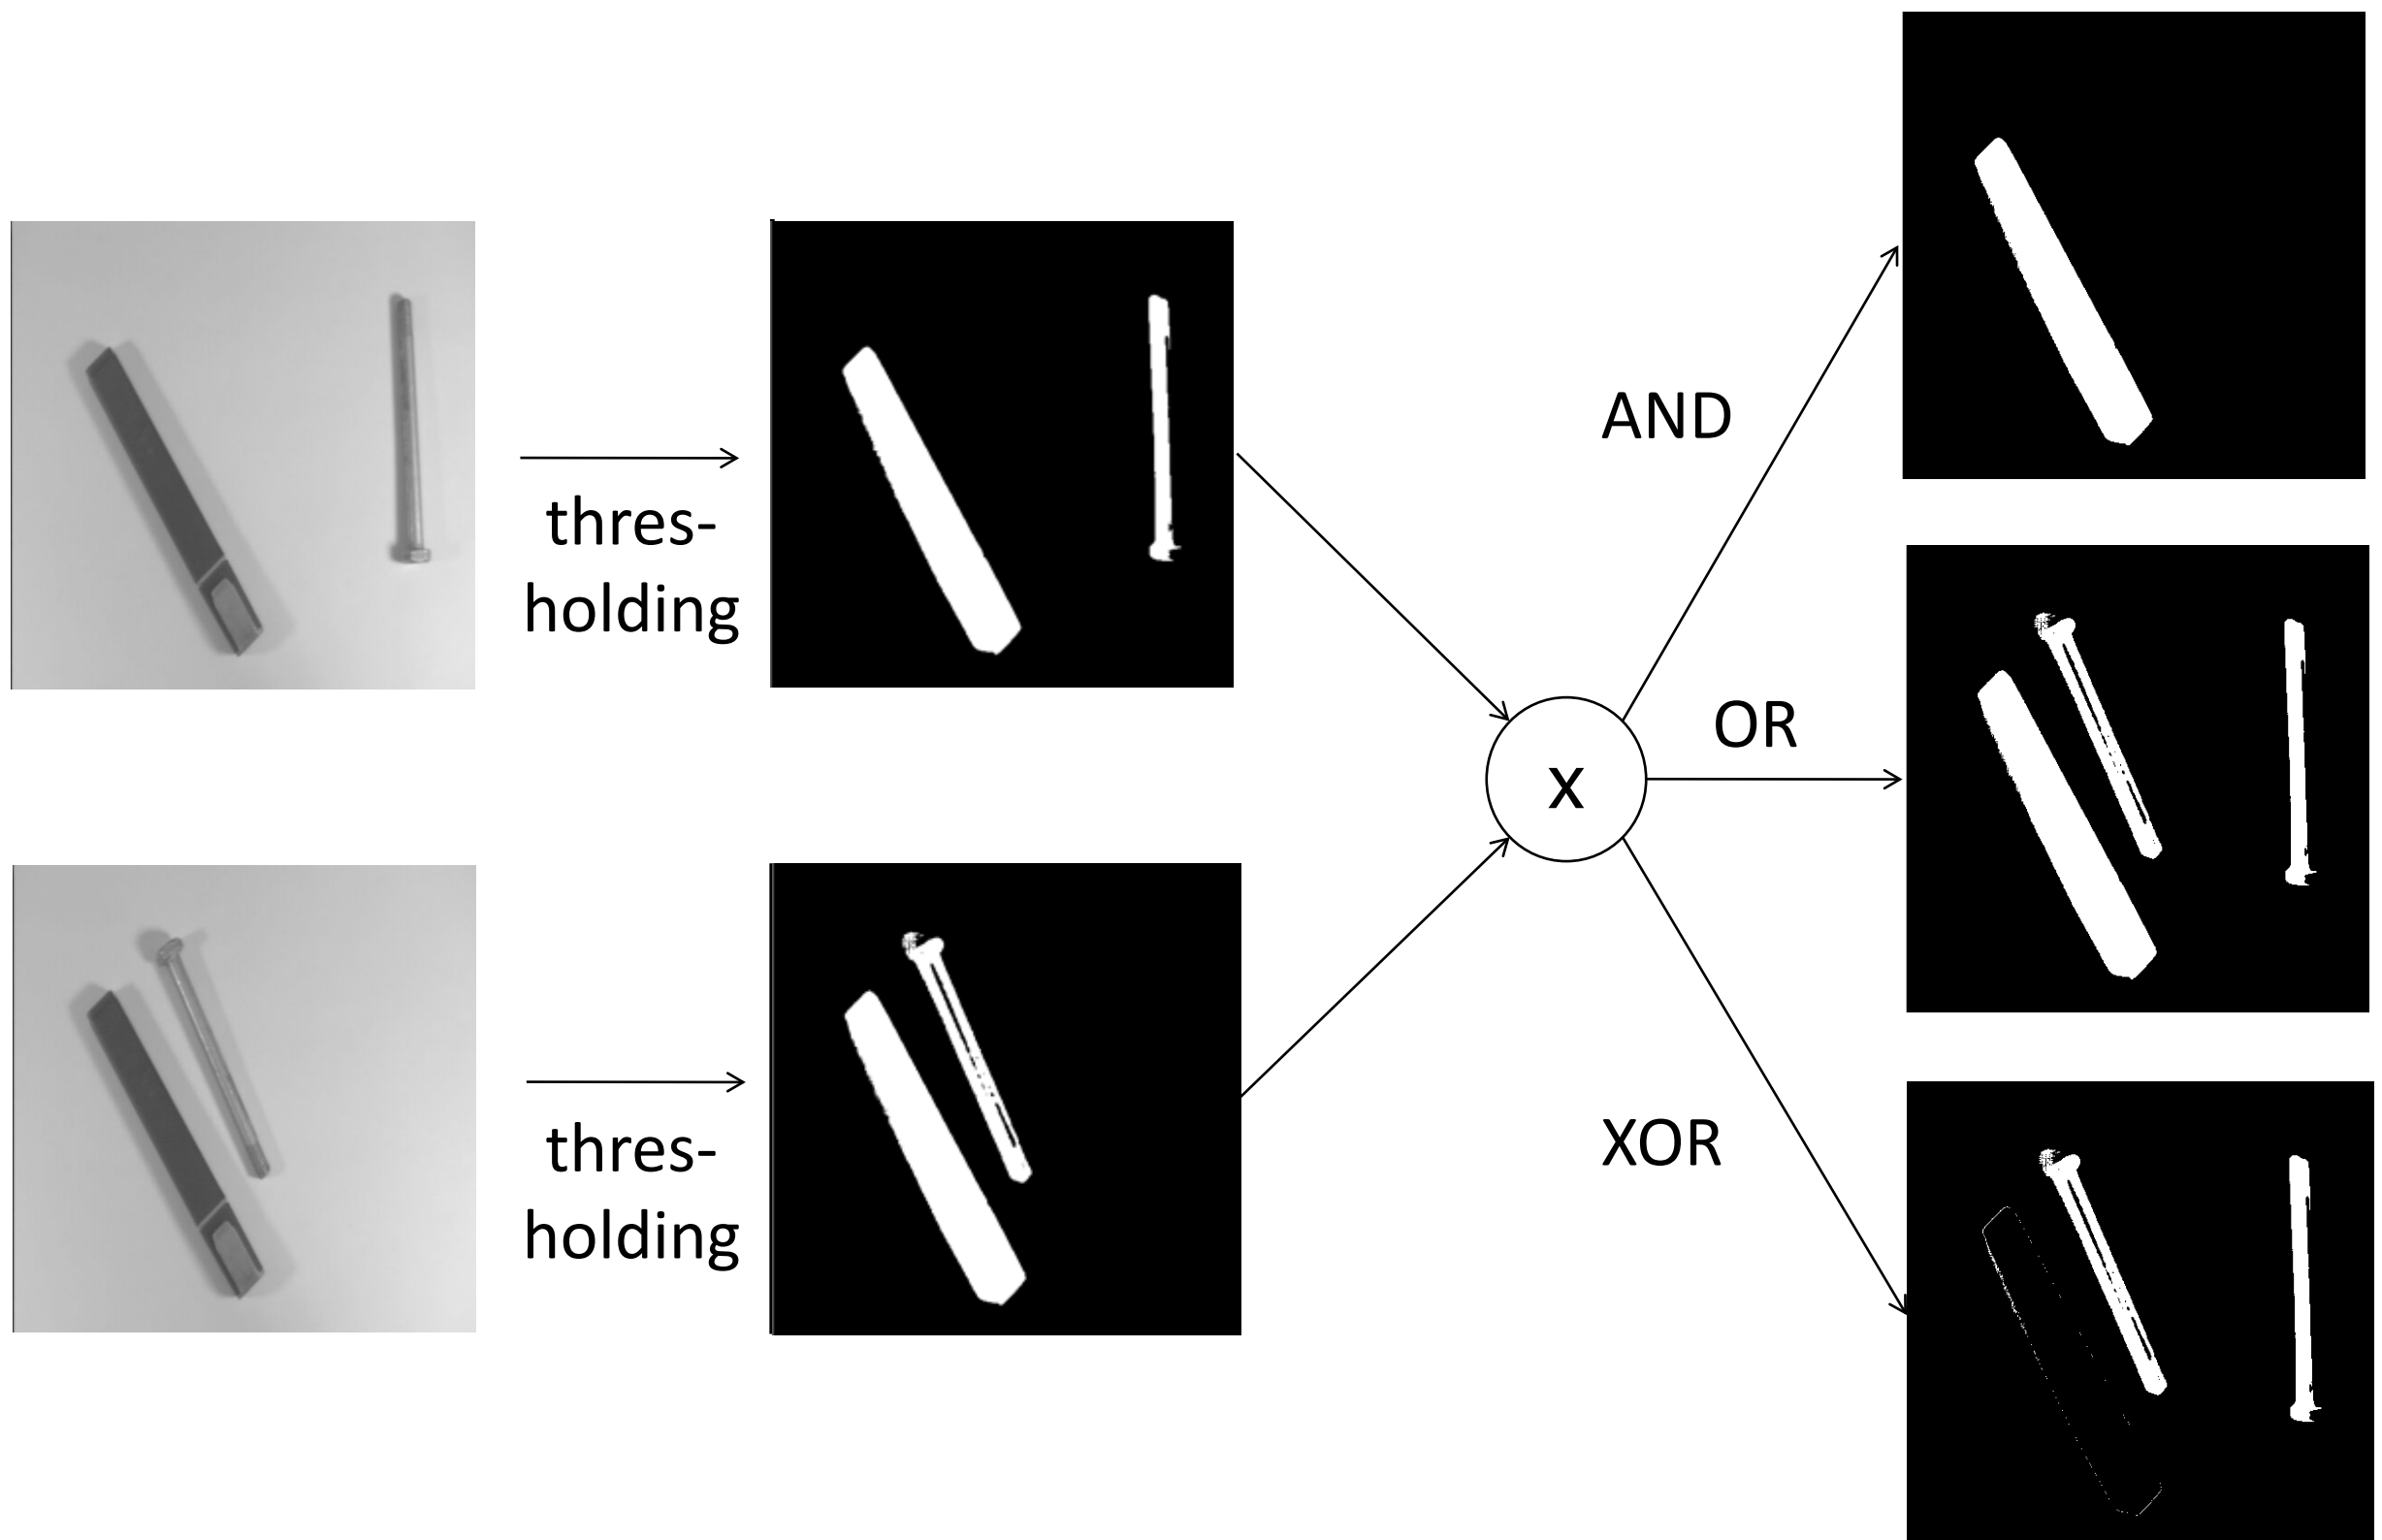
- Division:
 - Similar to subtraction, can be used to find changes or motion in the image
 - Areas that did not change get the value 1 → different effect on the image depending on the scale you use
 - Depending on the values of the pixels that are not 1, you can deduct which image has the higher value
 - $im1 > im2 \rightarrow [1 \ 255]$
 - $im1 < im2 \rightarrow [0 \ 1]$
- Python: $im3 = im1 / im2$



Logical Operations – AND, OR, XOR, NOT

- Logical Operations:
 - Pixelwise logic arithmetic on an image
 - Only makes sense for binary images, meaning each pixel is either 1 or 0
 - For example, you can mark objects in an image by setting all the pixels belonging to an object to 1 (using thresholding), and then check in the second image which pixels changes (to detect motion perhaps)
- NOT (`not`): Inversion
- AND (`&`, `and`) : Inner join of the images
- OR (`|`, `or`): Union of the images
- XOR (`^`): Detecting changes in the images

Examples – AND, OR, XOR



Questions

- How can arithmetic operations be used to detect motion in a video sequence?
- How can arithmetic operations be used to remove motion in a video sequence?


Image Averaging

- Motion in images can be removed by image averaging
- Averaging preserves only the constant image content
- Small movements do not prevail in the mean value image
- Example:
 - Webcam of e.g. Stephansplatz
 - Averaging the images over a longer time (1-2 minutes)
 - Result: Image of completely empty Stephansplatz!




Histograms

- Histogram:
 - 1D function that represents the occurrence frequencies of the gray values in an image
 - Normalized histogram represents occurrence probabilities of gray values (normalization factor = number of pixels); corresponds to density function!



`np.histogram(im)`



Histograms

- Gray level histogram represents the distribution of gray levels in an image
- Calculation: count occurrences of gray levels for a gray level range
- Optional: divide histogram by total number of pixels (normalization)
- Properties:
 - Resolution of histogram depends on size of histogram bins
 - Process is dimension-reducing and irreversible!
 - Location information of the image is completely lost
 - Histogram is rotation and translation invariant; normalized histogram also scale invariant
 - But: completely different images can have the same histograms

Histogram Applications

- **Contrast enhancement**
- Threshold search for image segmentation
- Search for similar images: Content-Based Image Retrieval (CBIR)


Histogram Applications

- Contrast enhancement of over- and underexposed images




overexposed

underexposed




Contrast enhanced



Contrast Enhancement - Histogram Stretch

- Low contrast narrow histogram not utilizing all gray levels.
- Histogram Stretch (Contrast Stretch):
 - Stretch gray level range linearly
 - Most common linear point operation
 - Utilize entire gray level range



Histogram Stretch

$A(x, y)$




$B(x, y)$



$$B(x, y) = \alpha A(x, y) - \beta$$

Scaling





Shift



$$\alpha = \frac{(K-1)}{\text{max} - \text{min}}$$

$$\beta = \frac{K-1}{\text{max} - \text{min}} \times \text{min}$$

Histogram Stretch - Example




Histogram Stretch - Summary

- Advantage:
 - Simple (linear)
 - Fast


- Disadvantage:
 - Hardly any improvement if gray value distribution is already wide.
 - Sensitive to outliers: Worst case: 1 white and one black pixel no scaling at all!

Histogram Equalization

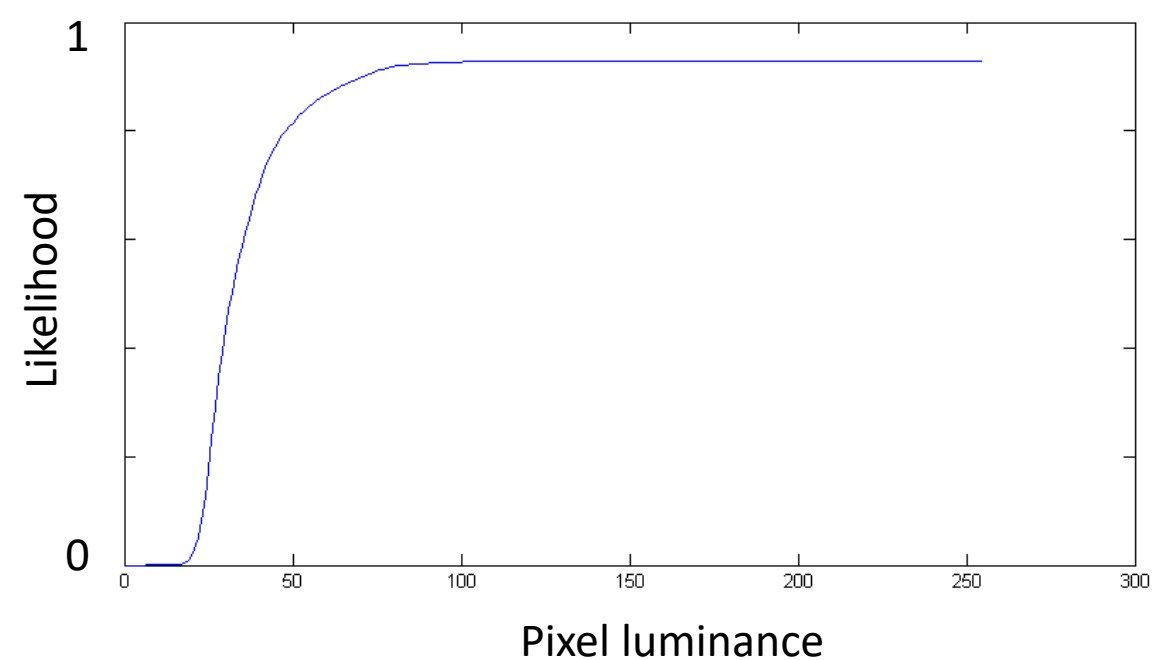
- Process for increasing contrast by spreading the histogram out to be approximately uniformly distributed
- Goal: use entire gray level range and a uniformly distributed (flat) histogram!
- Flat histogram maximizes the information content (cf. entropy)
- Nonlinear point operation → Changes shape of histogram!



Histogram Equalization



Histogram $h(x)$ /
Density function



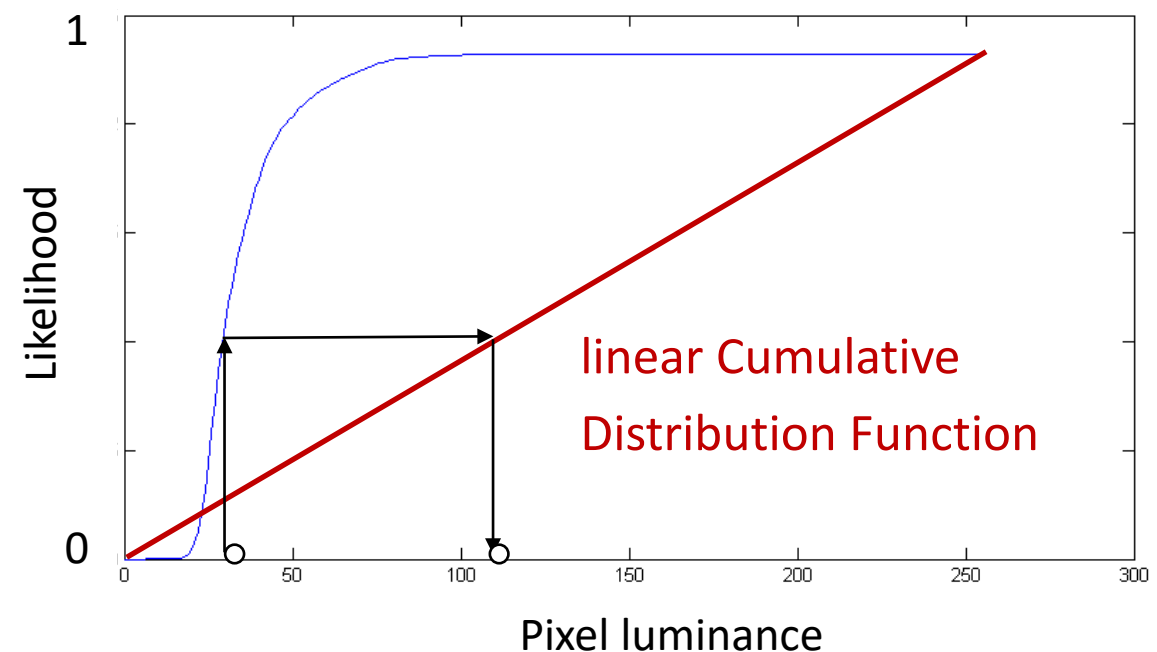
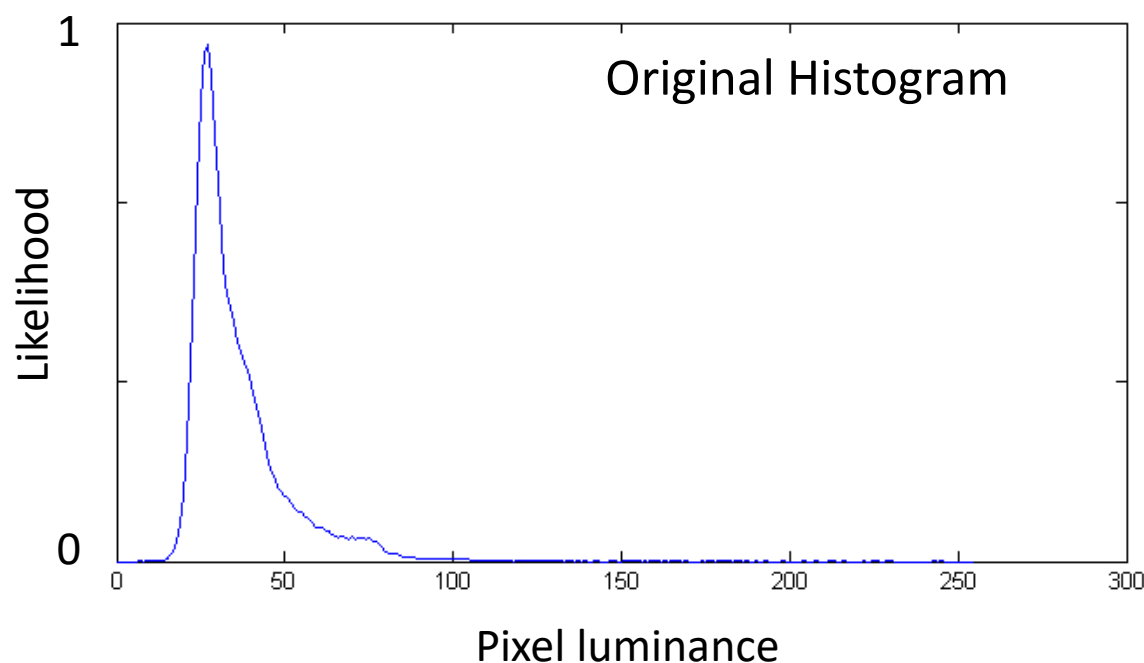
Cumulative histogram /
Cumulative distribution function:
 $c(i)$

$$c(i) = \int_0^i h(x) dx$$

Histogram Equalization

- Distribution function of a uniform distribution is linear
- Goal: make Distribution function $c(i)$ linear
- For linear distribution function holds:

$$c(i) = \frac{1}{N-1} f(i) \implies f(i) = c(i) * (N-1), \quad N = \#gray\ levels$$





- $f(i)$ is mapping from old gray level to new gray level --> replace gray level i in original image with $f(i)$ in new image

Histogram Equalization



- Distribution function of a uniform distribution is linear
- Goal: make Distribution function $c(i)$ linear
- For linear distribution function holds:

$$c(i) = \frac{1}{N-1} f(i) \implies f(i) = c(i) * (N-1), \quad N = \# \text{gray levels}$$





- $f(i)$ is mapping from old gray level to new gray level --> replace gray level i in original image with $f(i)$ in new image

Histogram Equalization - Result



Question

- Why does histogram equalization not work well in this example?



Histogram Equalization - Summary

- Resulting histograms are not really flat (because discrete distribution)
- Peaks in the histogram are preserved
- Histograms are non-linearly distorted. Compression where there are few gray levels, stretching where there are many gray levels
- N input gray levels can be assigned to 1 output gray level (reduction of image information, otherwise no problem)
- Generalization: Histogram Shaping (Histogram of arbitrary shapes)
- Python: `skimage.exposure.equalize_hist(im)`

Histogram Applications


- Contrast enhancement
- **Threshold search for image segmentation**
- Search for similar images: Content-Based Image Retrieval (CBIR)

Thresholding

- Conversion of gray scale image into binary image
- Quantization of gray levels to 2 values \rightarrow 0 and 1 using threshold T:







$$g(n) = \begin{cases} 1: f(n) \geq T \\ 0: f(n) < T \end{cases}$$

- **How to choose a reasonable threshold value?** Histogram can help if foreground and background have different gray level distributions:



Thresholding





- Application: Separate foreground and background



Histogram Applications





- Contrast enhancement
- Threshold search for image segmentation
- **Search for similar images: Content-Based Image Retrieval (CBIR)**

Color as a low-level cue for CBIR






- Color histograms:
Use distribution of colors to describe image
- No spatial info –
invariant to
translation, rotation,
scale

Color as a low-level cue for CBIR



Given two histogram vectors, sum the minimum counts per bin:

$$I(x, y) = \sum_{i=1}^n \min(x_i, y_i)$$



$$\sum_i \min(x_i, y_i) = 4$$

Color-based image retrieval


- Given collection (database) of images:
 - Extract and store one color histogram per image

- Given new query image:
 - Extract its color histogram
 - For each database image: Compute intersection between query histogram and database histogram
 - Sort intersection values (highest score = most similar)
 - Rank database items relative to query based on this sorted order





Example Retrieval Results

query







query



query





query






Example Retrieval Results



query



query



query



Final Remark on Histograms

- Q: What happens if we reshuffle all pixels within the images?



- A: Its histogram won't change!
→ Point-wise processing unaffected
- Need to measure properties relative to small *neighborhoods* of pixels

Roadmap

