



ACV – Applied Computer Vision

Bachelor Medientechnik & Creative Computing

Matthias Zeppelzauer

matthias.zeppelzauer@fhstp.ac.at

Djordje Slijepcevic

djordje.slijepcevic@fhstp.ac.at

Image Filtering

Roadmap

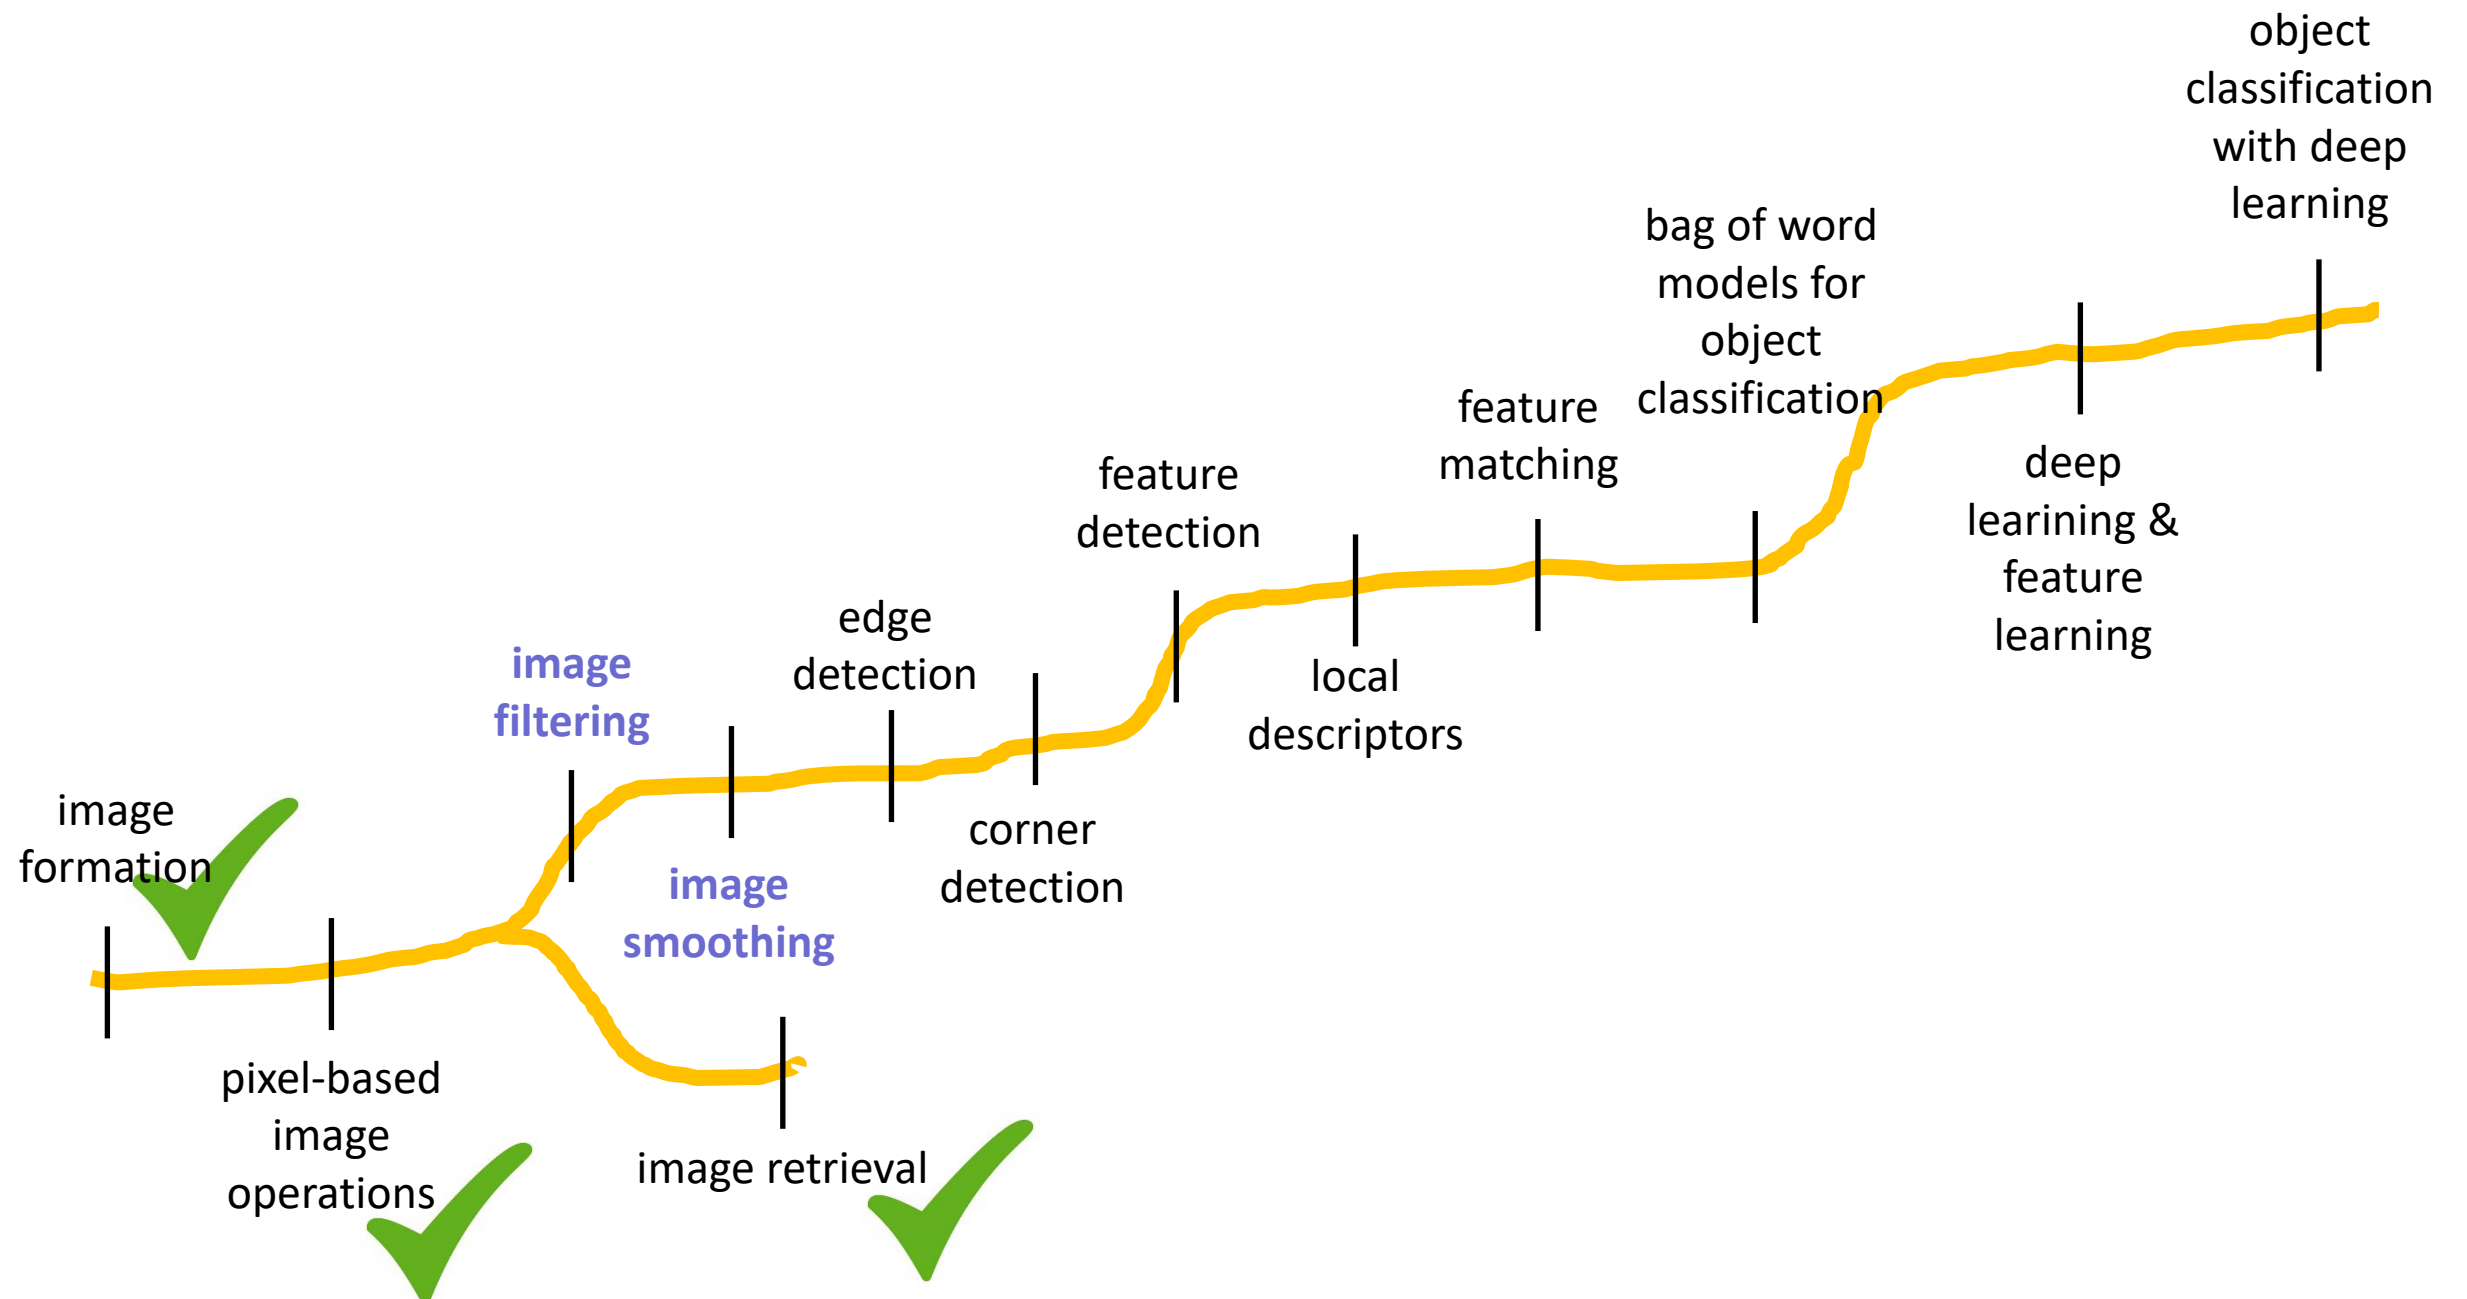


Image Filtering

- Compute a function of the **local neighborhood** at each pixel in the image
 - Function specified by a “filter” or mask saying how to combine values from neighbors.
- Uses of filtering:
 - Enhance an image (denoise, resize, etc)
 - Extract information (texture, edges, etc)
 - Detect patterns (template matching)

Image Filtering

- **Idea:** *take current pixel value plus values of its neighbors and compute something...*




Image Filtering

- **Idea:** *take current pixel value plus values of its neighbors and compute something...*
- Simplest case: linear filtering
 - Operation = Convolution
 - Sum of point-wise product between filter values and pixels



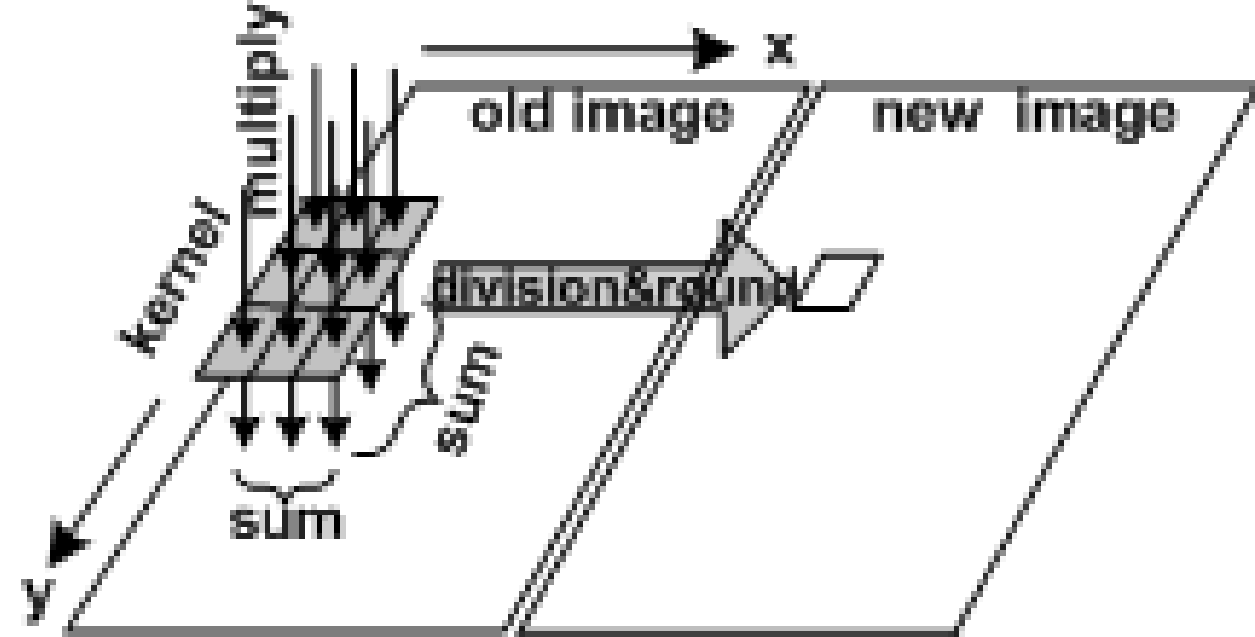
How is the computation scheme?

- Compute point wise product of filter weights and pixels
- Sum all products



Lets make an example

- Calculate by hand



$F[x, y]$

2	1	4	6	3
4	5	2	6	1
6	5	1	2	6
3	2	7	6	9
8	7	2	3	4

$H[u, v]$

-1	1	1
3	1	2
-1	1	3



=




$G[x, y]$

	?			

calculate the convolution on a piece of paper!

Boundary issues

- What is the size of the output?
- Output size options:
 - 'full': output size is sum of sizes of f and g
 - 'same': output size is same as f
 - 'valid': output size is difference of sizes of f and g




Filtering in Python

- `f` is the image, `h` is the filter.
 - `from scipy import signal`
 - `imConv = signal.convolve2d(f, h, mode='same')` : does the convolution
 - Attribute '`mode`':
 - `same`: zero padding, result image has same size as input image
 - `valid`: only locations are evaluated where the filter fits in completely → result image smaller than input image
 - `full`: filtering takes place at all possible locations. Missing values are replaced by zeros → result image is larger than input image

Mathematical Formulation

- At **each** location (i, j) in the image, compute:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k \underbrace{H[u, v]}_{\text{filter weights}} F[i + u, j + v]$$



- This is the result of convolution for **one** location
- Example for $k = 1$:
 - $u=-1, \dots, +1; v=-1, \dots, +1$, kernel size = $2*k+1 = 3$ (\rightarrow 3x3 filter)
 - H = filter kernel, $H(-1,-1)$: top left corner , $H(1,1)$: bottom right corner
 - In first iteration we get: $u=-1, v=-1 \rightarrow H(-1,-1)*F(i-1,j-1)$
 - In second iteration we get: $u=-1, v=0 \rightarrow H(-1,0)*F(i-1,j+0)$
 - How does the third iteration look like?

Mathematical Formulation

- At **each** location (i, j) in the image, compute:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k \underbrace{H[u, v]}_{\text{filter weights}} F[i + u, j + v]$$

- In Pseudocode (no handling boundary issues)

```
sum=0
```

```
k=1
```

```
For u = -1 to +1
```

```
    For v = -1 to +1
```

```
        product=H(u+k, v+k) *F(i+u, j+v)
```

```
        sum=sum+product
```




Mathematical Formulation

- At **each** location (i, j) in the image, compute:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k \underbrace{H[u, v]}_{\text{filter weights}} F[i + u, j + v]$$

- This is the result of convolution for **one** location
- Filtering an image: replace each pixel with a linear combination of its neighbors.
- The filter “kernel” or “mask” $H[u, v]$ is the prescription for the weights in the linear combination.



Applications of Filtering

- **Template Matching**
- Noise Reduction
- Edge Detection
- Corner Detection
- ...
- And: Representation Learning (Convolutional Neural Networks)


Template Matching

Find the chair in this image

This is a chair




Output of normalized correlation



Convolution yields high values where it finds similar patterns!

Motivation: noise reduction



- Even multiple images of the **same static scene** will not be identical.
- Reason: sensor noise

Common types of noise

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise






Impulse noise



Gaussian noise

Gaussian noise



$$f(x, y) = \underbrace{\hat{f}(x, y)}_{\text{Ideal Image}} + \underbrace{\eta(x, y)}_{\text{Noise process}}$$

Gaussian i.i.d. ("white") noise:
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

What is impact of the sigma?

Effect of
sigma on
Gaussian
noise

$\sigma=1$



Effect of
sigma on
Gaussian
noise

$\sigma=16$



Motivation: noise reduction





- Even multiple images of the same static scene will not be identical.
- **How can we remove the Gaussian noise from an image?**

First Attempt at a Solution

- Let's replace each pixel with an average of all the values in its neighborhood
- Assumptions:
 - Expect pixels to be like their neighbors
 - Expect noise processes to be independent from pixel to pixel, i.e. in average (over the neighborhood) the noise will cancel itself out

First Attempt at a Solution

- Let's replace each pixel with an average of all the values in its neighbors
- How does the filter has to look like?




“box filter”

“averaging filter”

Example: Averaging filter

- One of the simplest filters
- Why the multiplication with 1/9?



$$G = H \otimes F$$

Example: Averaging filter

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0								

- Moving average in 2D: compute average in each neighborhood and replace center value

Example: Averaging filter

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10							

Example: Averaging filter

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20						

Example: Averaging filter

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30					

Example: Averaging filter

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30	30				

Example: Averaging filter


$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	


Smoothing by averaging



depicts box filter:
white = high value, black = low value



original



filtered

What if the filter size was 5 x 5 instead of 3 x 3?

Gaussian filter

- What if we want nearest neighboring pixels to have the most influence on the output?

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0


$F[x, y]$

1	2	1
2	4	2
1	2	1



$H[u, v]$

This kernel is an approximation of a Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$

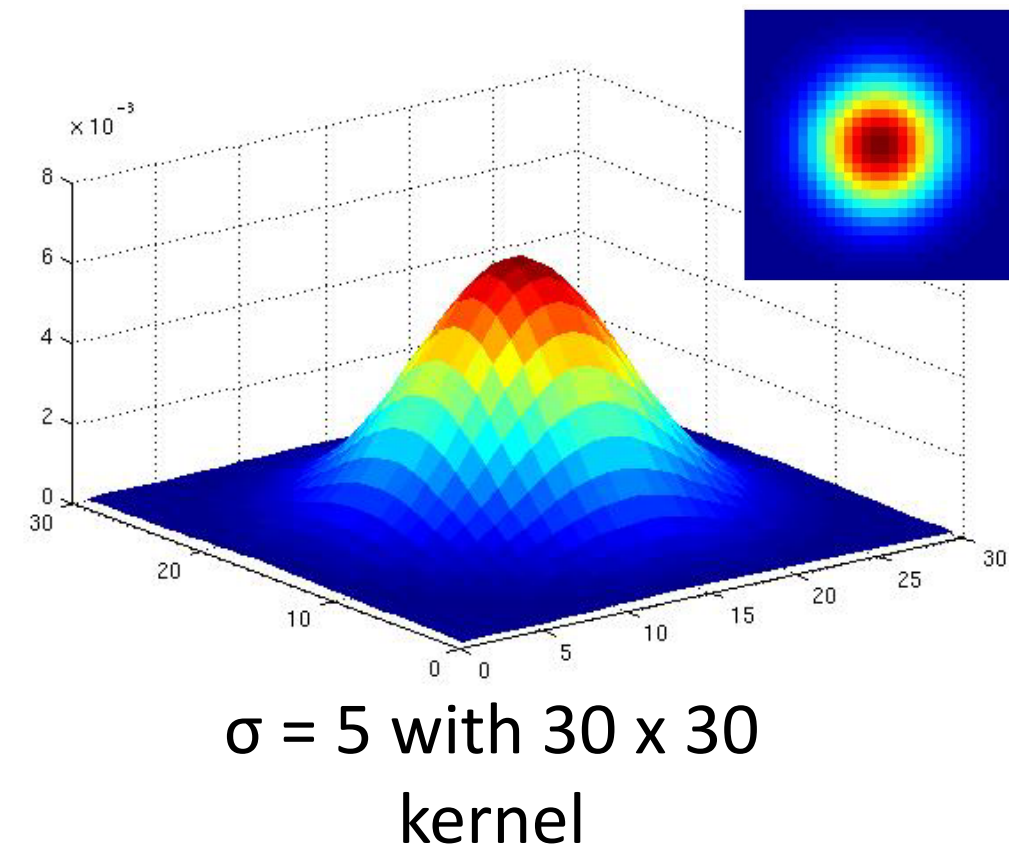



Smoothing with a Gaussian





Gaussian filters

- What parameters matter here?
- **Standard deviation (sigma)** of Gaussian: determines extent of smoothing




Smoothing with a Gaussian

- Parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



...



Properties of smoothing filters

- Values positive
- Sum to 1 \rightarrow constant regions same as input
- Amount of smoothing proportional to mask size
- Remove “high-frequency” components \rightarrow “low-pass” filter


$$\frac{1}{16}$$

1	2	1
2	4	2
1	2	1

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

Playing with filters



Original



0	0	0
0	1	0
0	0	0

=

?


Playing with filters



Original




0	0	0
0	1	0
0	0	0



Filtered
(no change)

Playing with filters



Original




0	0	0
0	0	1
0	0	0

=

?

Playing with filters



Original

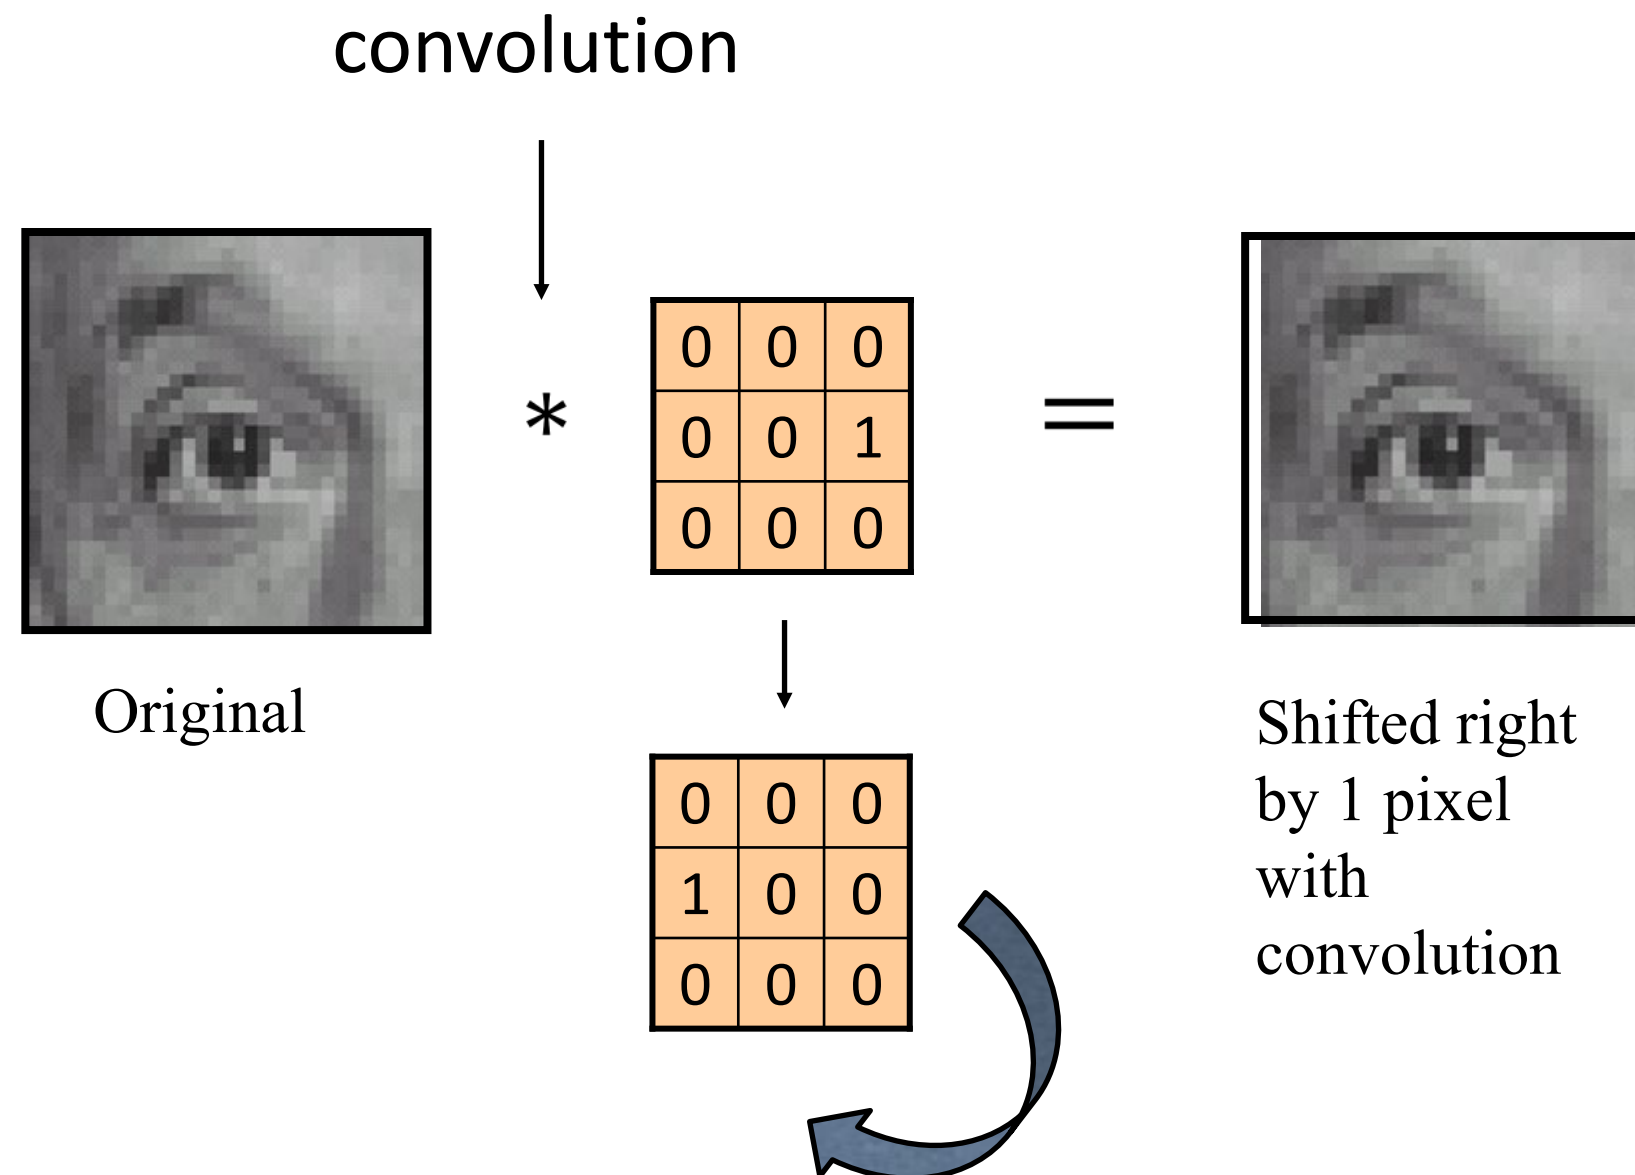


0	0	0
0	0	1
0	0	0




Shifted left
by 1 pixel
with
correlation

Playing with filters



Playing with filters



Original



$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

=

?

Playing with filters




Original



$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

=



Blur (with an averaging filter)

Playing with filters



Original



0	0	0
0	2	0
0	0	0

-


$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

=

?


Playing with filters

 \otimes


0	0	0
0	2	0
0	0	0

 $- \frac{1}{9}$

1	1	1
1	1	1
1	1	1


 $=$ 

Original




Sharpening filter:
accentuates differences with
local average

Sharpening Filter



before



after

Effect of Smoothing Filters

5x5



Additive Gaussian noise



Salt and pepper noise

- Works great for Gaussian noise
- But **fails** for salt and pepper noise – why?


Image Filtering

- **Idea:** *take current pixel value plus values of its neighbors and compute something...*
- Simplest case: linear filtering
- Alternative: non-linear filtering
 - Example: **median filter** (take median of all values in the neighborhood of a pixel)




Median Filter

- Note: this filter works different than the filters before!
- It is also a sliding window operation BUT: cannot be computed by correlation / convolution
- The median filter has no weights. It takes the values of the underlying images as inputs:





Median filter

Salt and
pepper
noise



Median
filtered



Plots of a row
of the image



Python:

```
scipy.ndimage.filters.median_filter(image, size=(3,3))
```


Median Filter vs. Average Filter



3x3 average
filter




3x3 median
filter



Median Filter

- Median filter preserves edges – it does not smooth over them
- No new pixel values introduced
- Removes spikes: good for salt & pepper noise
- Non-linear filter



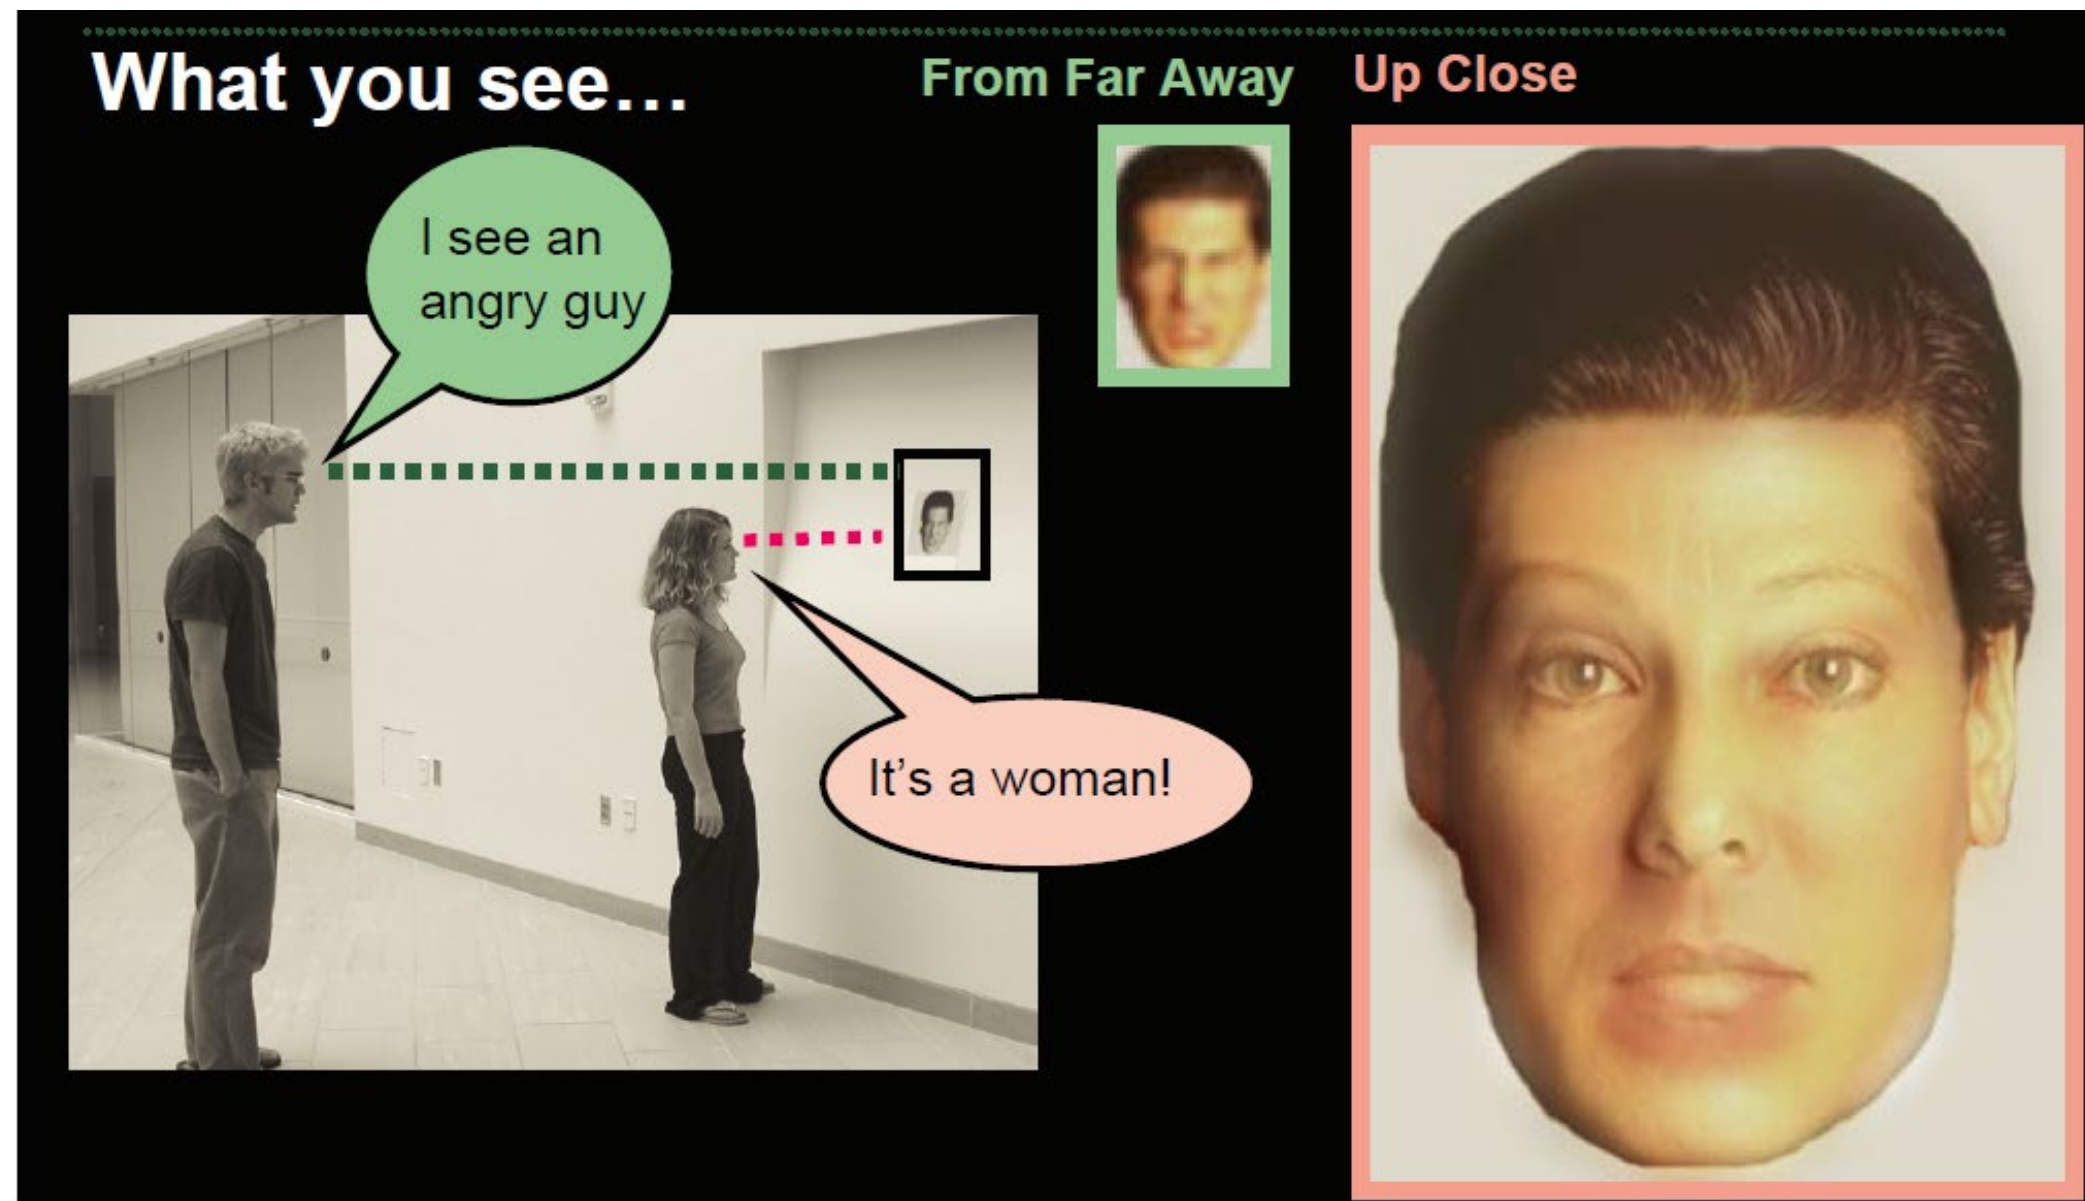
Applications of Filtering / Convolution

- Denoising: removing artifacts/noise/distortions in images
- Smoothing, Sharpening
- Template Matching
- Detection of basic image structures, such as edges and corners (next time)
- Image Effects, e.g. Hybrid Images...
- *Convolutional* Neural Networks 😊



Finally, let's have some fun with filtering

- How to create your own hybrid image?




A hybrid image is...

- <https://www.youtube.com/watch?v=OlumoQ05gS8>


Application: Hybrid Images

Gaussian Filter

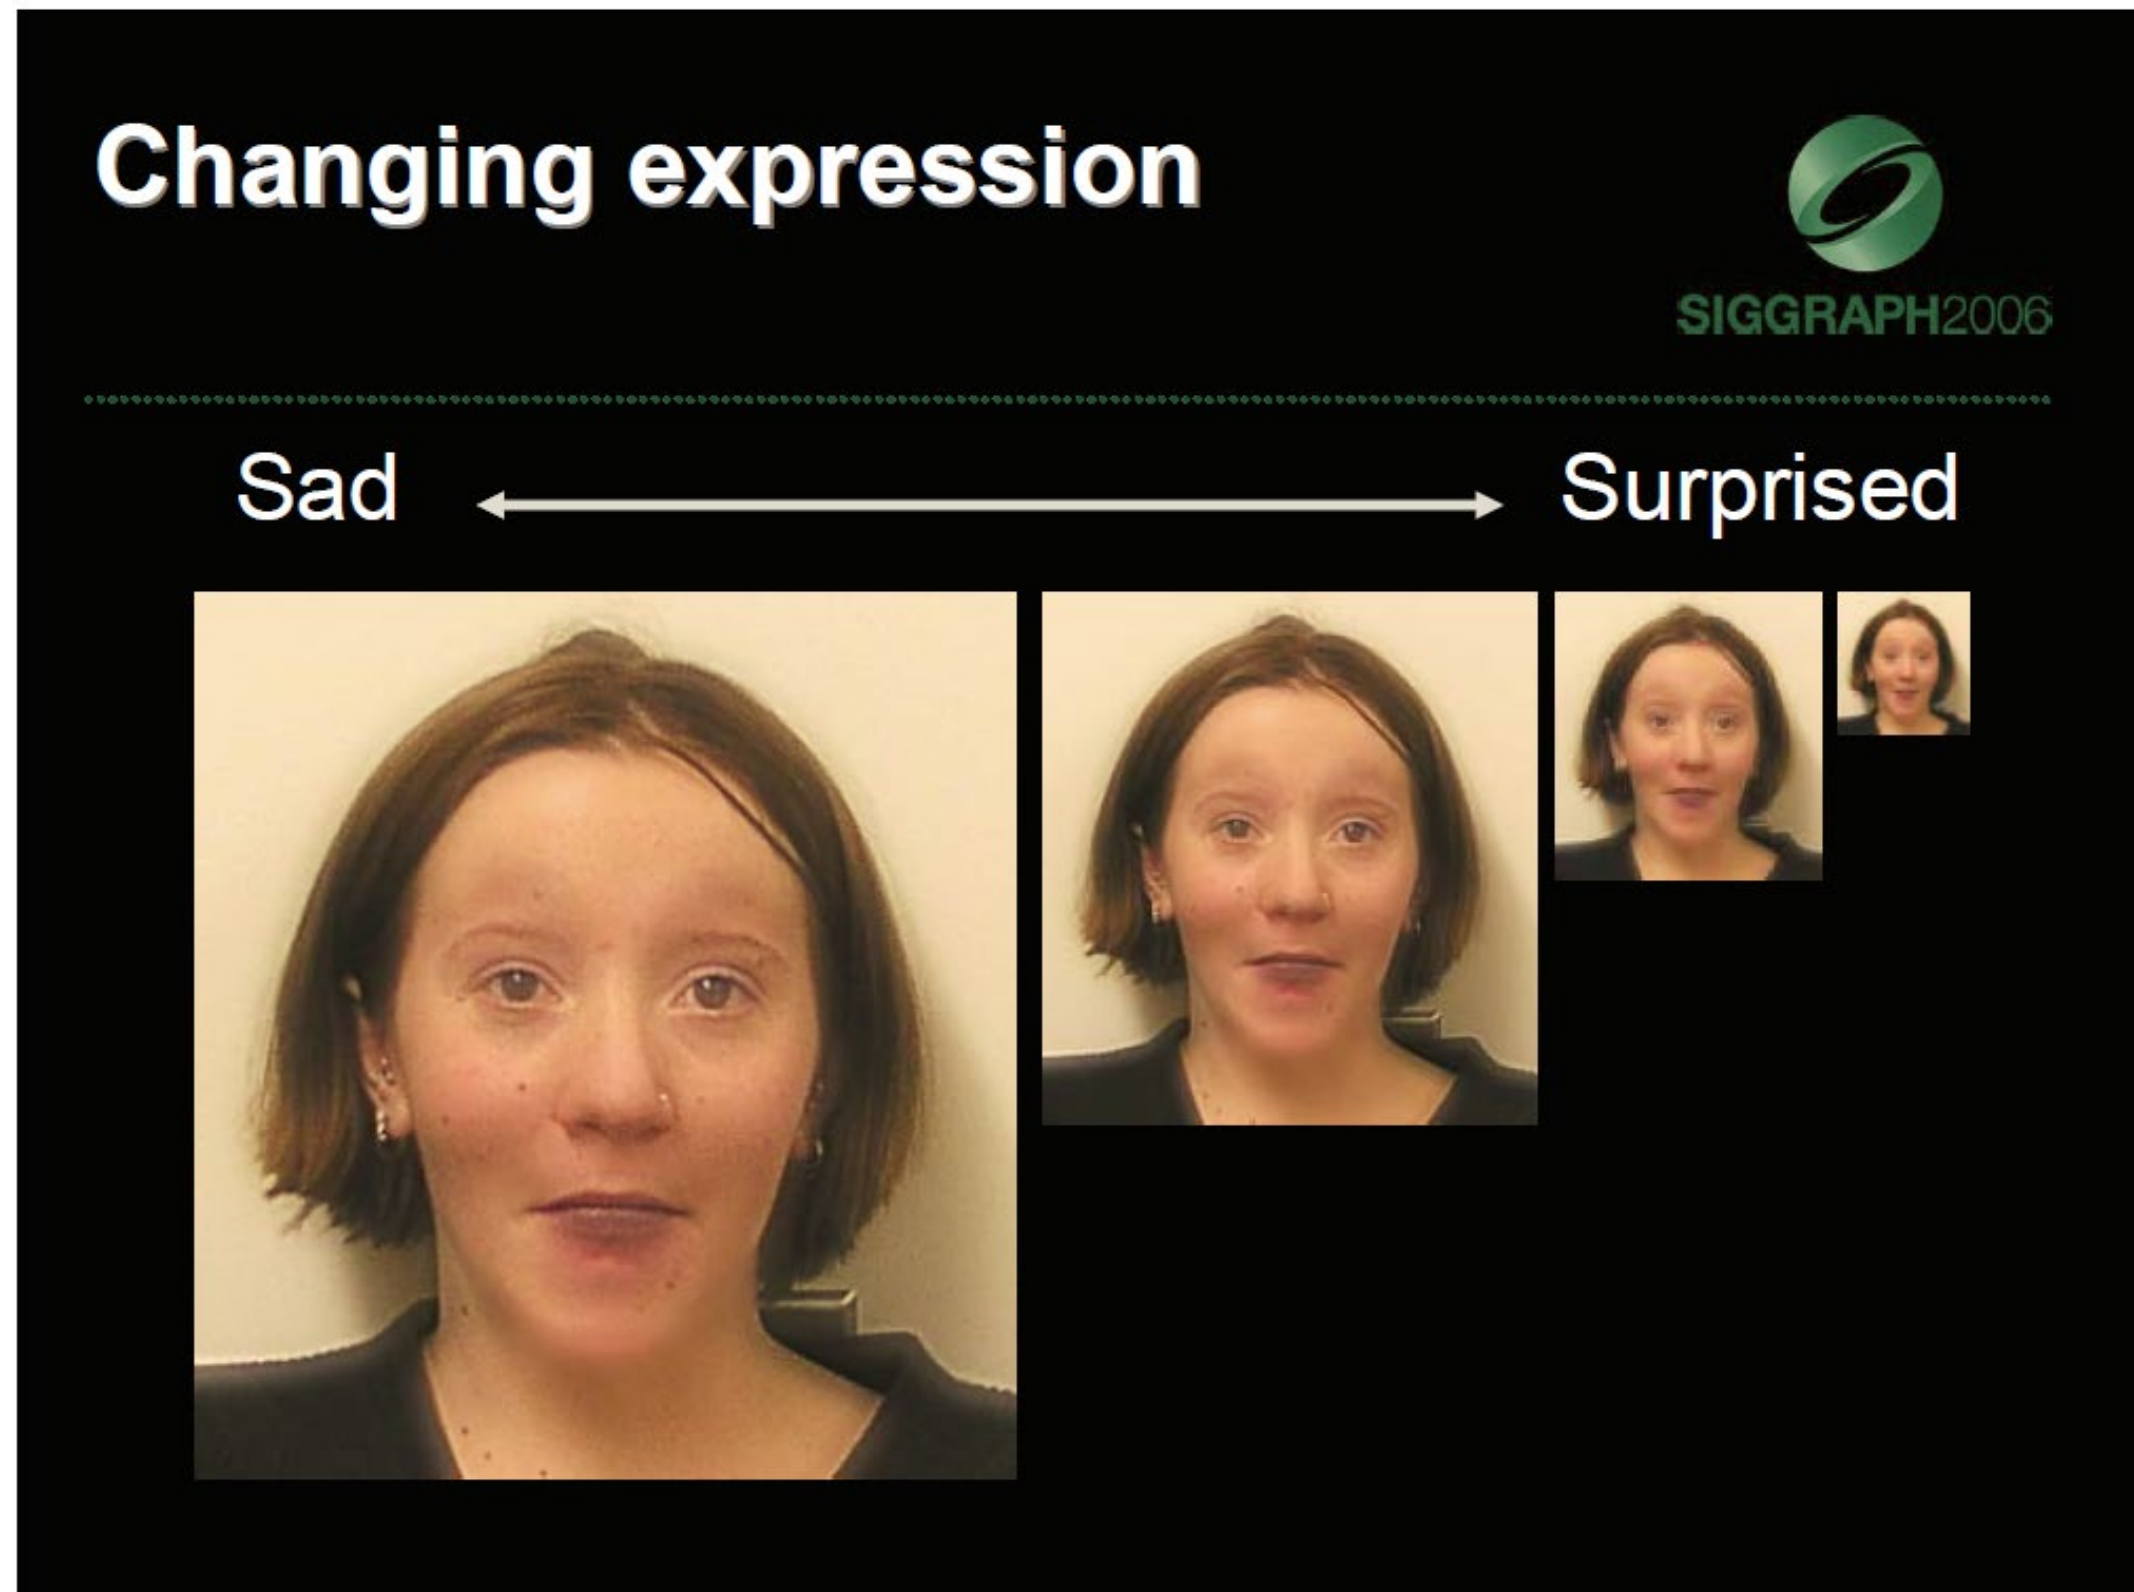


Laplacian Filter

Hybrid Image Result



Another Result




Summary

- Challenges of Computer Vision
- How are images generated
- How do we represent images digitally
- Point-wise image operations
 - Arithmetic operations (add, subtract, multiply...)
 - Logical operations (and, or,...)
 - Histograms
- Filtering
 - Convolution and Correlation filtering
 - Smoothing, Sharpening,
 - Median filters

Most important thing to remember: the process of filtering:

- Position filter on image (start top left corner)
- Move filter over image: row-by-row
- At each location: multiply values of filter with the values of the image
- Add all multiplied numbers together
- Optionally: divide sum by a constant (normalization)
- Add result at current location in new image



Roadmap

