



ACV – Applied Computer Vision

Bachelor Medientechnik & Creative Computing

Matthias Zeppelzauer
matthias.zeppelzauer@fhstp.ac.at

Djordje Slijepcevic
djordje.slijepcevic@fhstp.ac.at

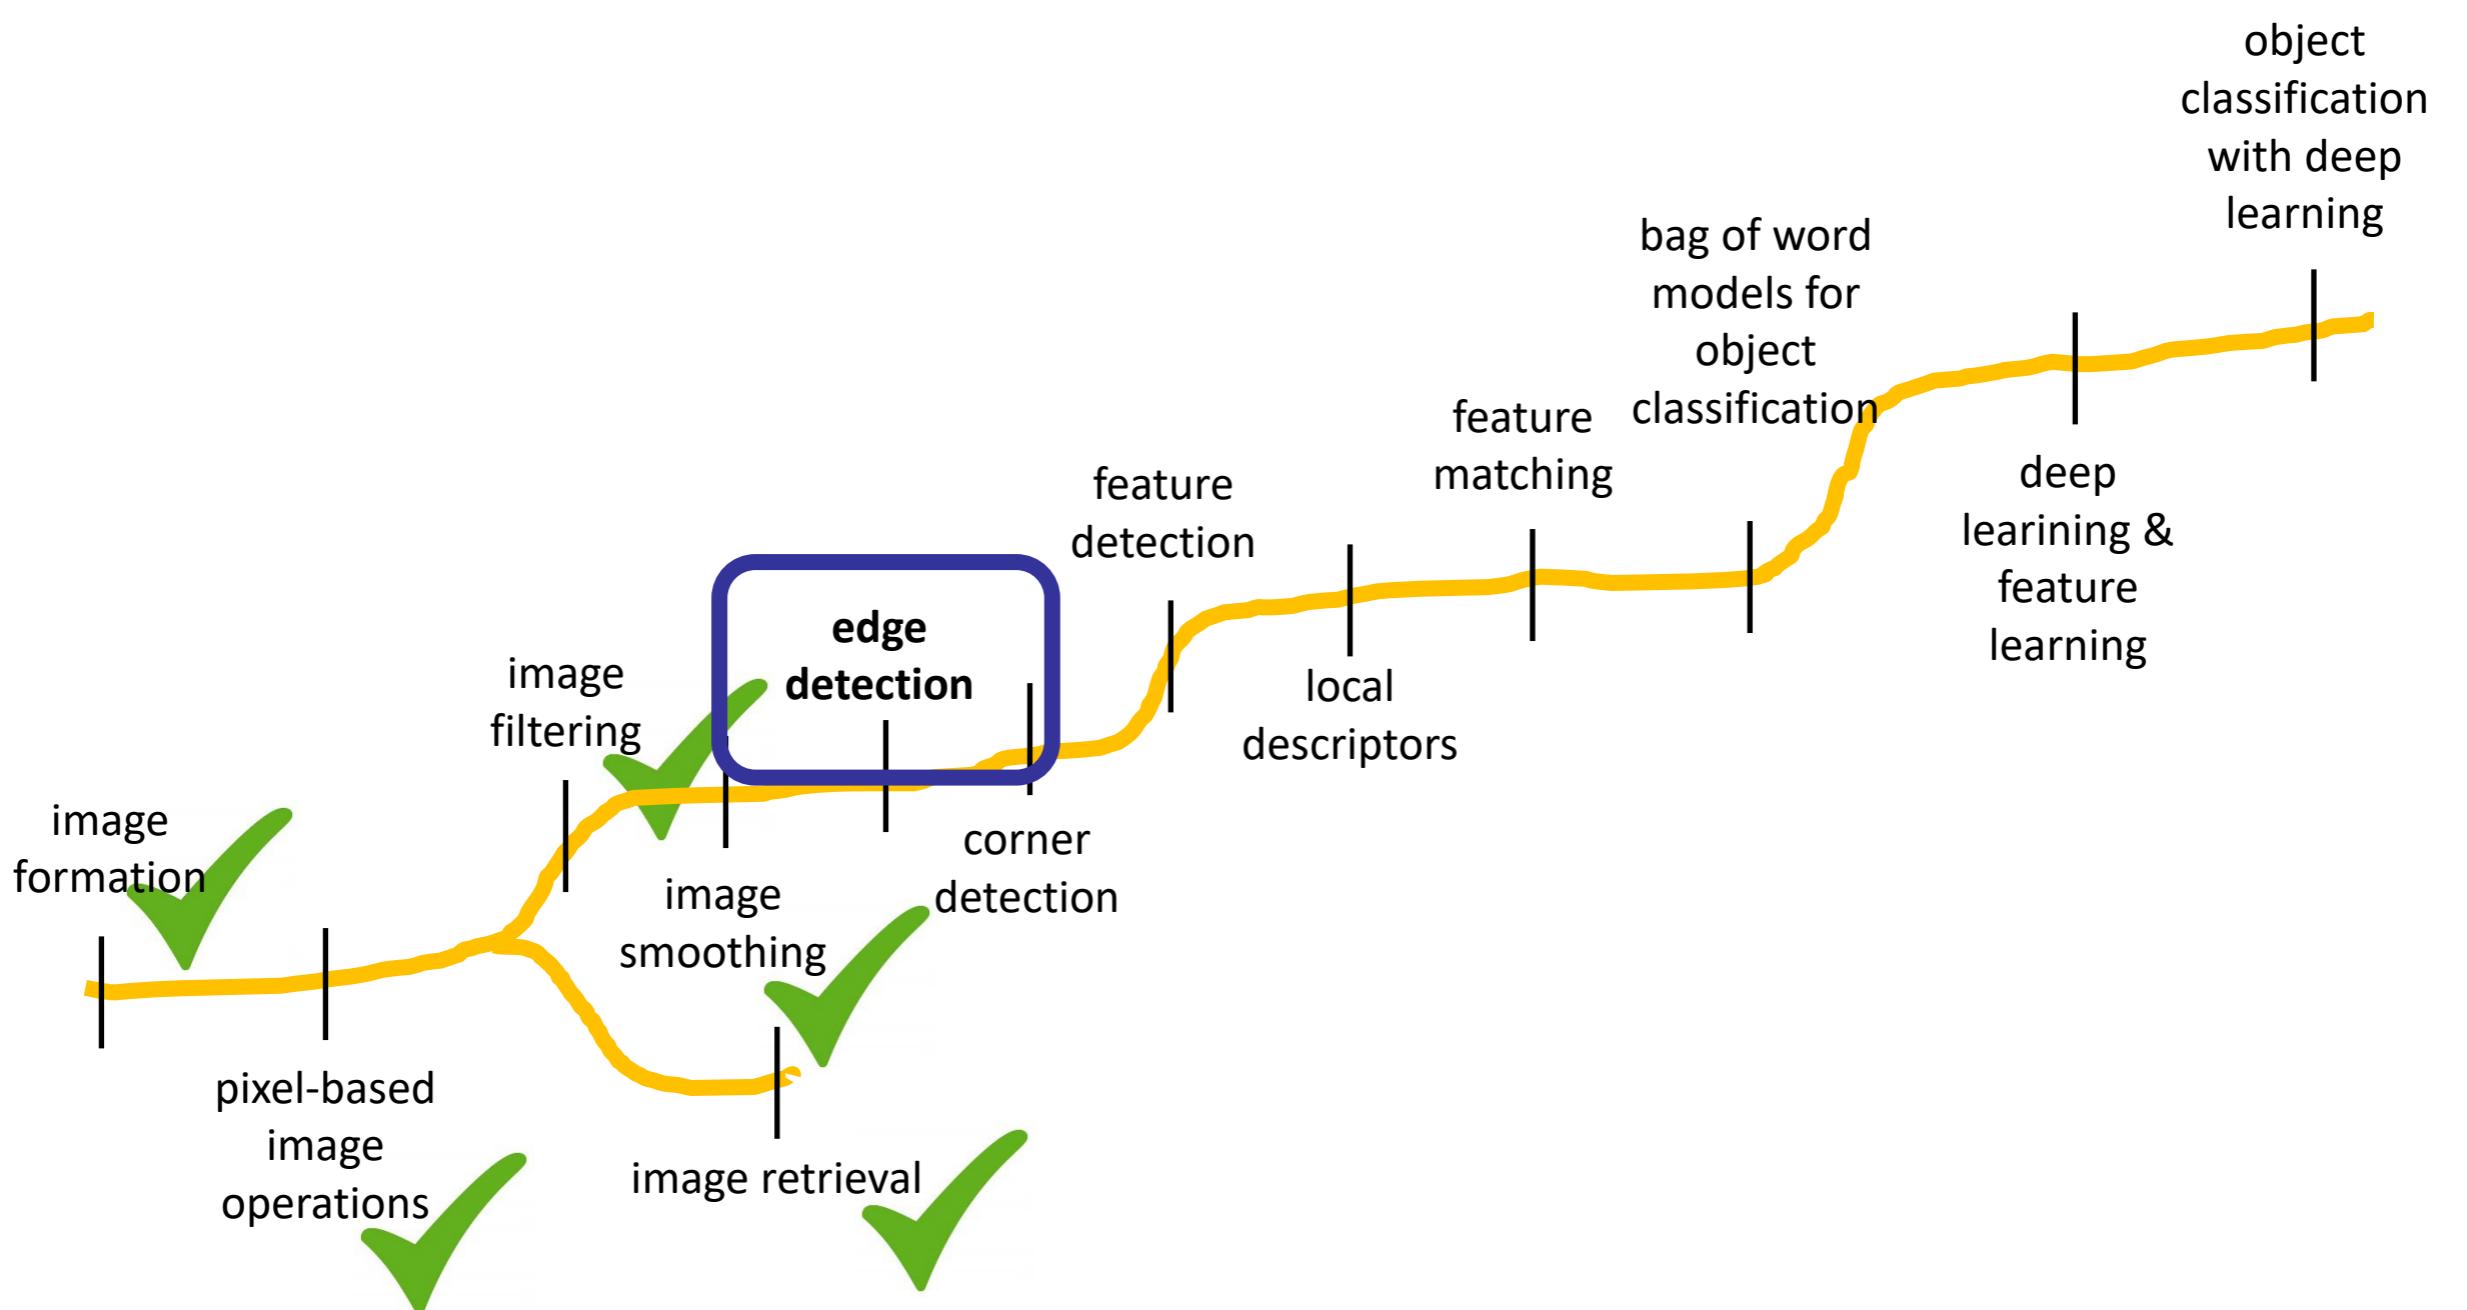
Slide Credits

- Kristen Grauman
- S. Seitz,
- Marschner,
- S. Lazebnik
- L. Fei Fei

Outlook

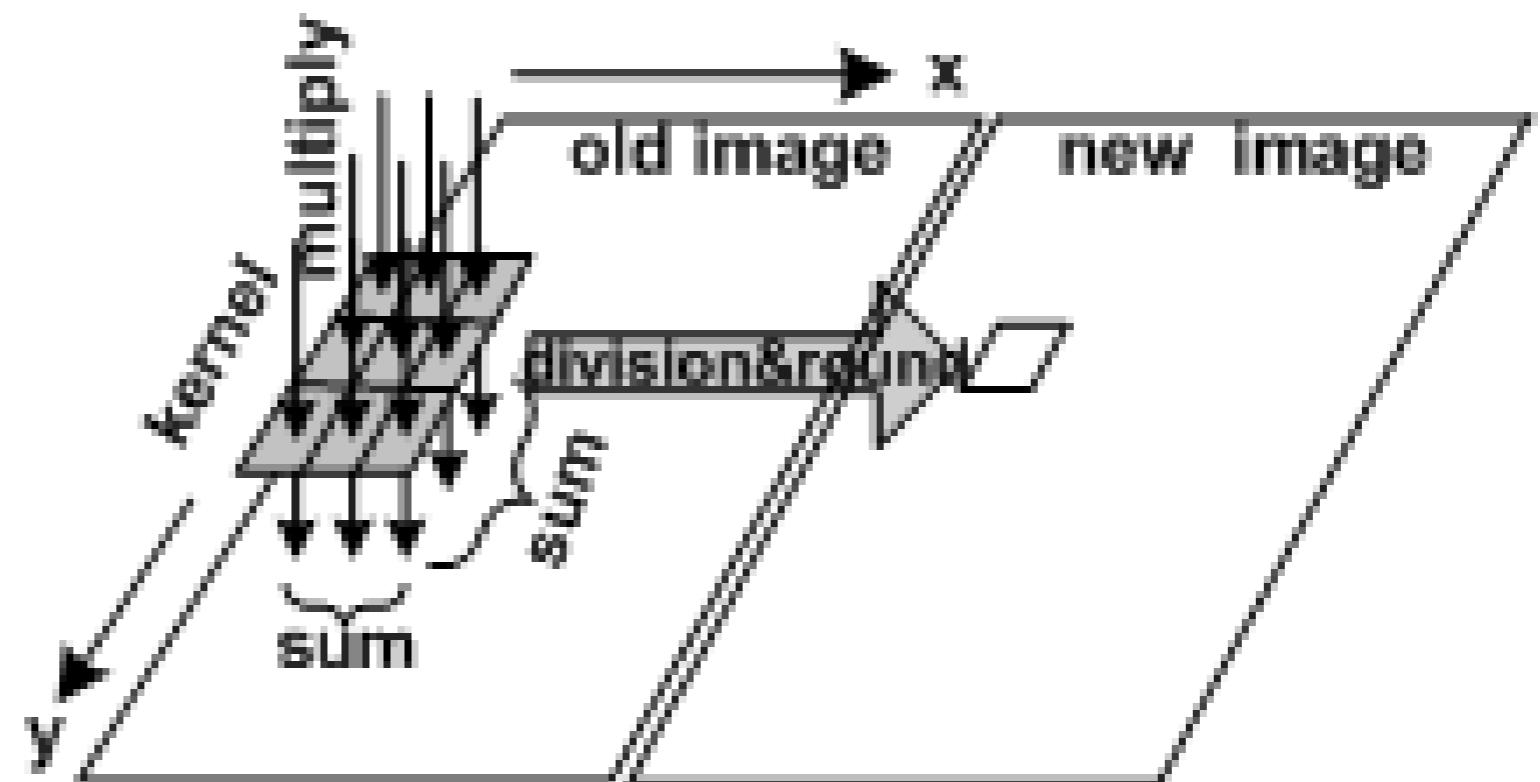
- So far:
 - Pixel-based operations on images
 - Filter (smoothing, sharpening)
- Today:
 - Finding structures in images!
 - (Template matching)
 - Edges

Roadmap



Remember Filtering

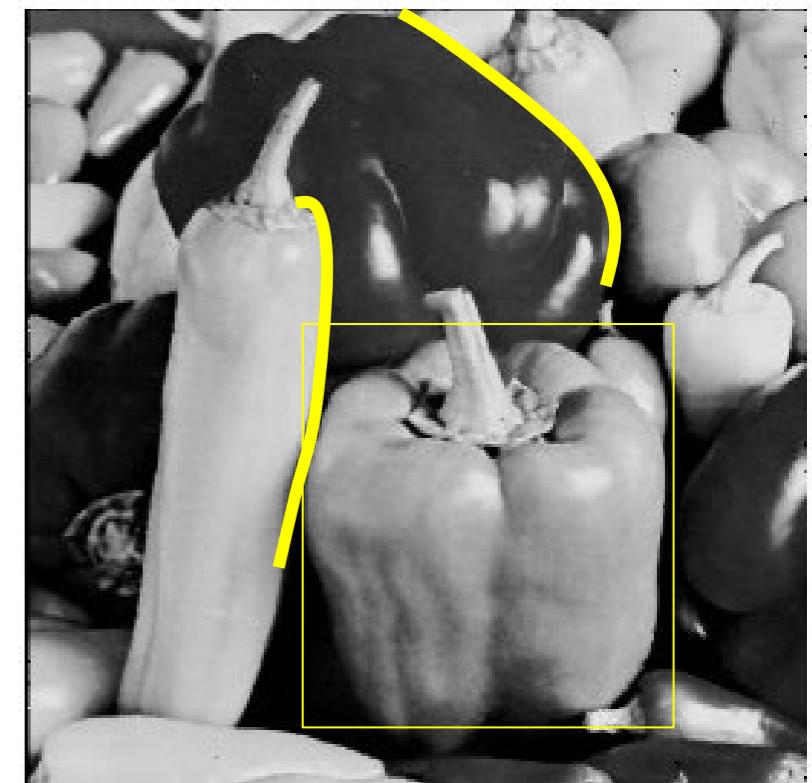
- Moving a filter across an image
- Computing the sum of products between filter and image



- Can we use this to find structures in the image?

Filters for features

- Previously, thinking of filtering as a way to remove or reduce **noise**
- Now, consider how filters will allow us to abstract higher-level “**features**”.
 - Map raw pixels to an **intermediate representation** that will be used for subsequent processing
 - Goal: **reduce amount of data, discard redundancy, preserve what's useful**



Edge detection

- **Goal:** map image from 2d array of pixels to a set of curves or line segments or contours.
- **Why?**

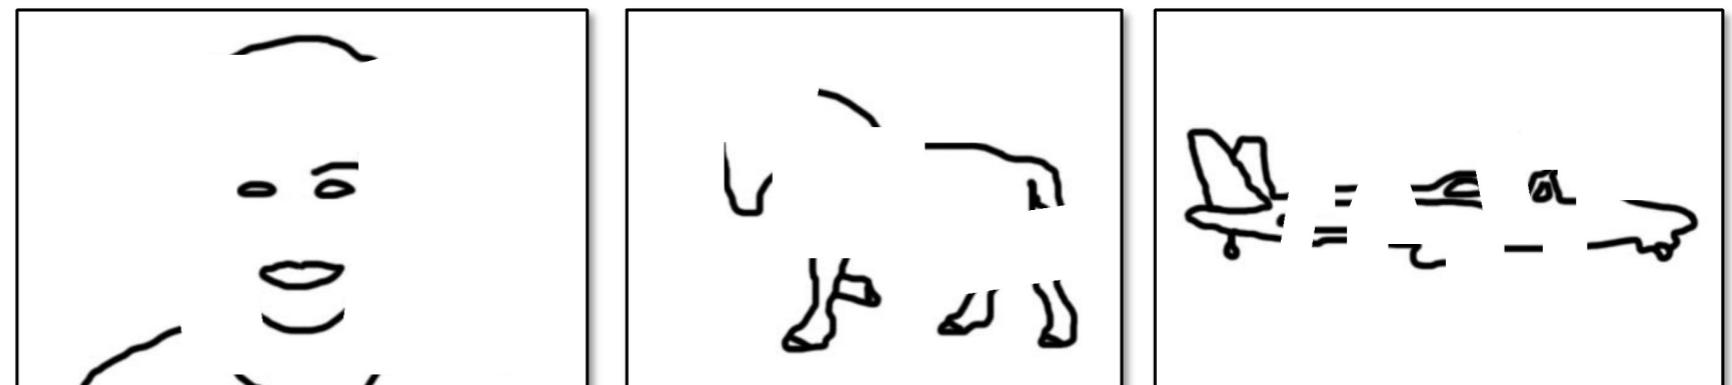


Figure from J. Shotton et al., PAMI 2007

- **Main idea:** look for strong gradients, post-process

Edge Detection

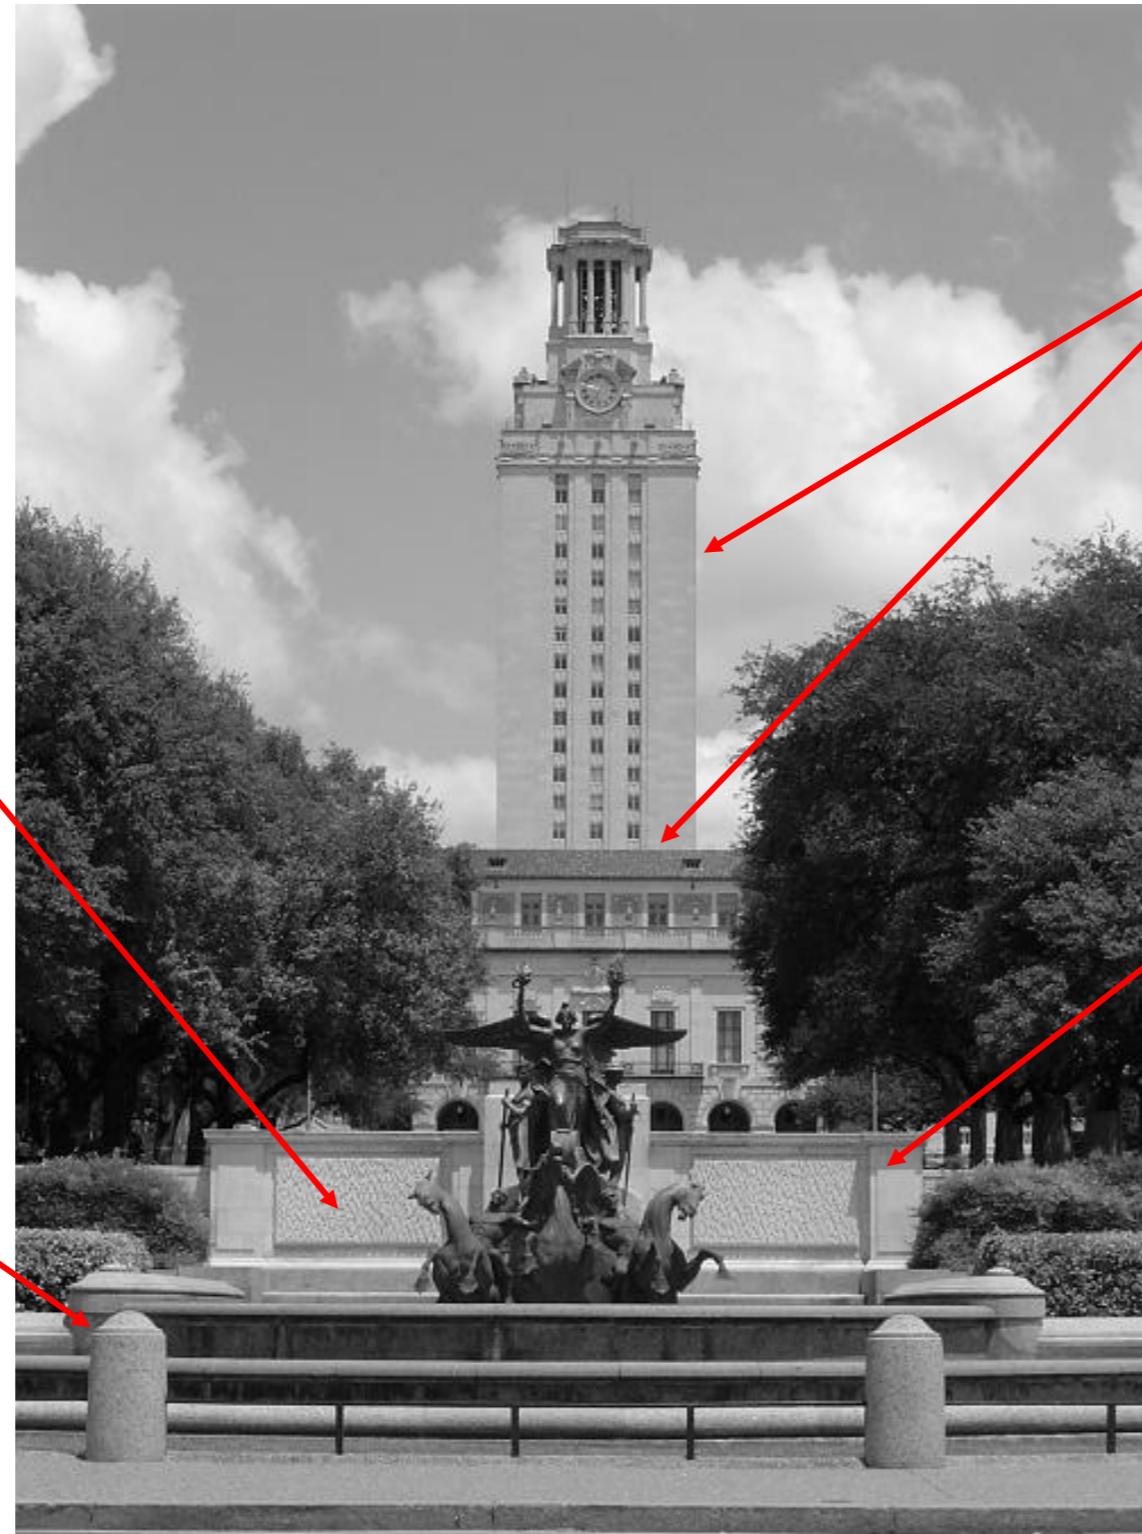
- **Goal:** Identify sudden changes (discontinuities) in an image
 - Intuitively, most semantic and shape information from the image can be encoded in the edges
 - More compact than pixels
- **Ideal:** artist's line drawing (but artist is also using object-level knowledge and world knowledge)



What can cause an edge?

Reflectance change:
appearance
information, texture

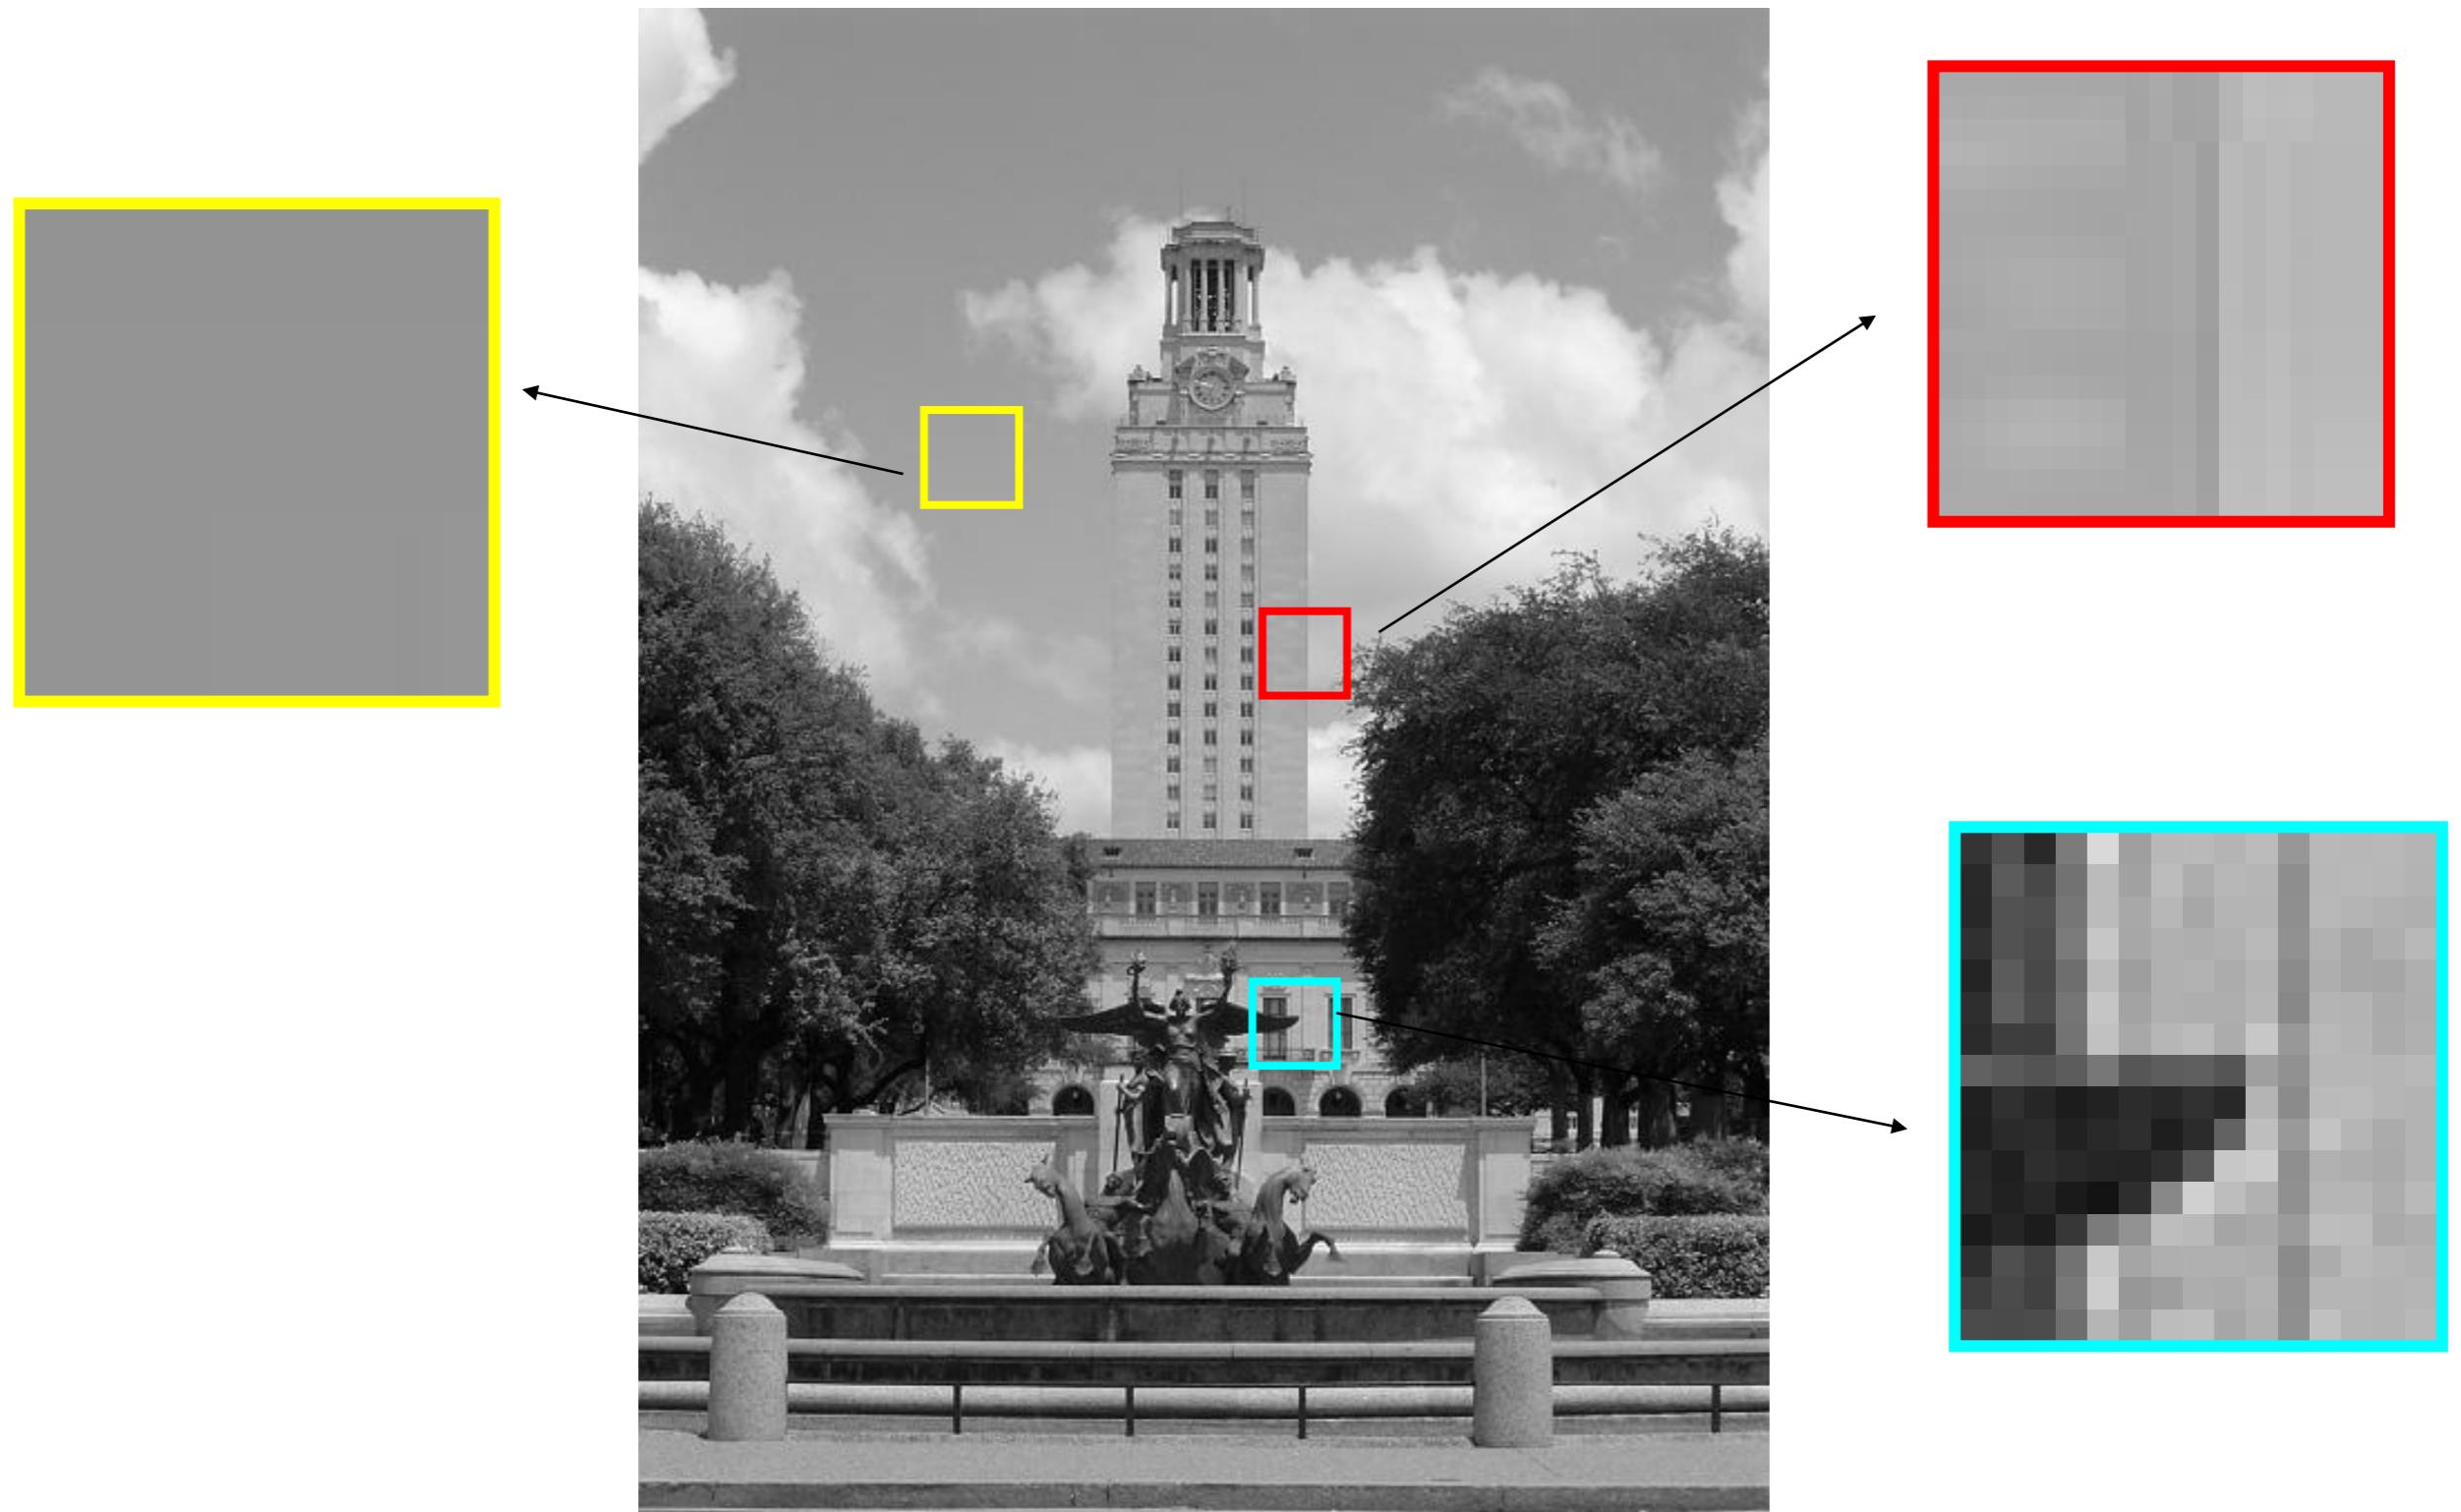
Change in surface
orientation: shape



Depth discontinuity:
object boundary

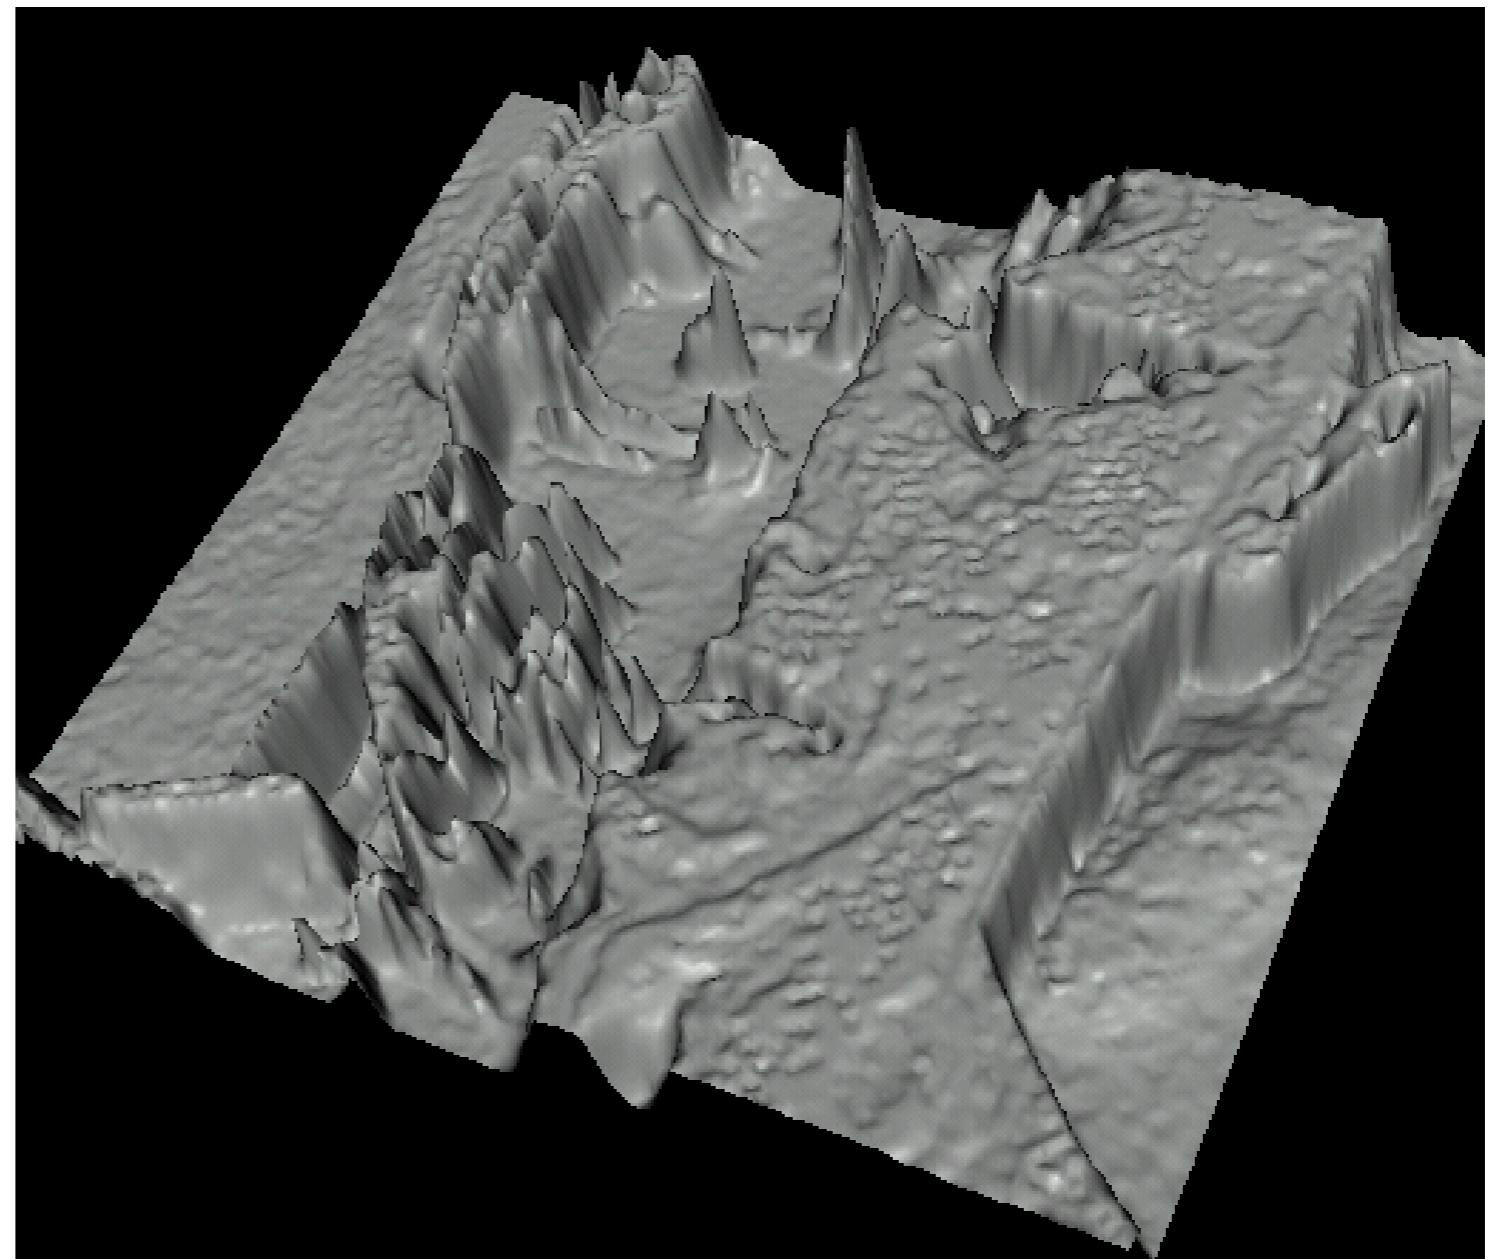
Cast shadows

Contrast and invariance



Recall : Images are functions

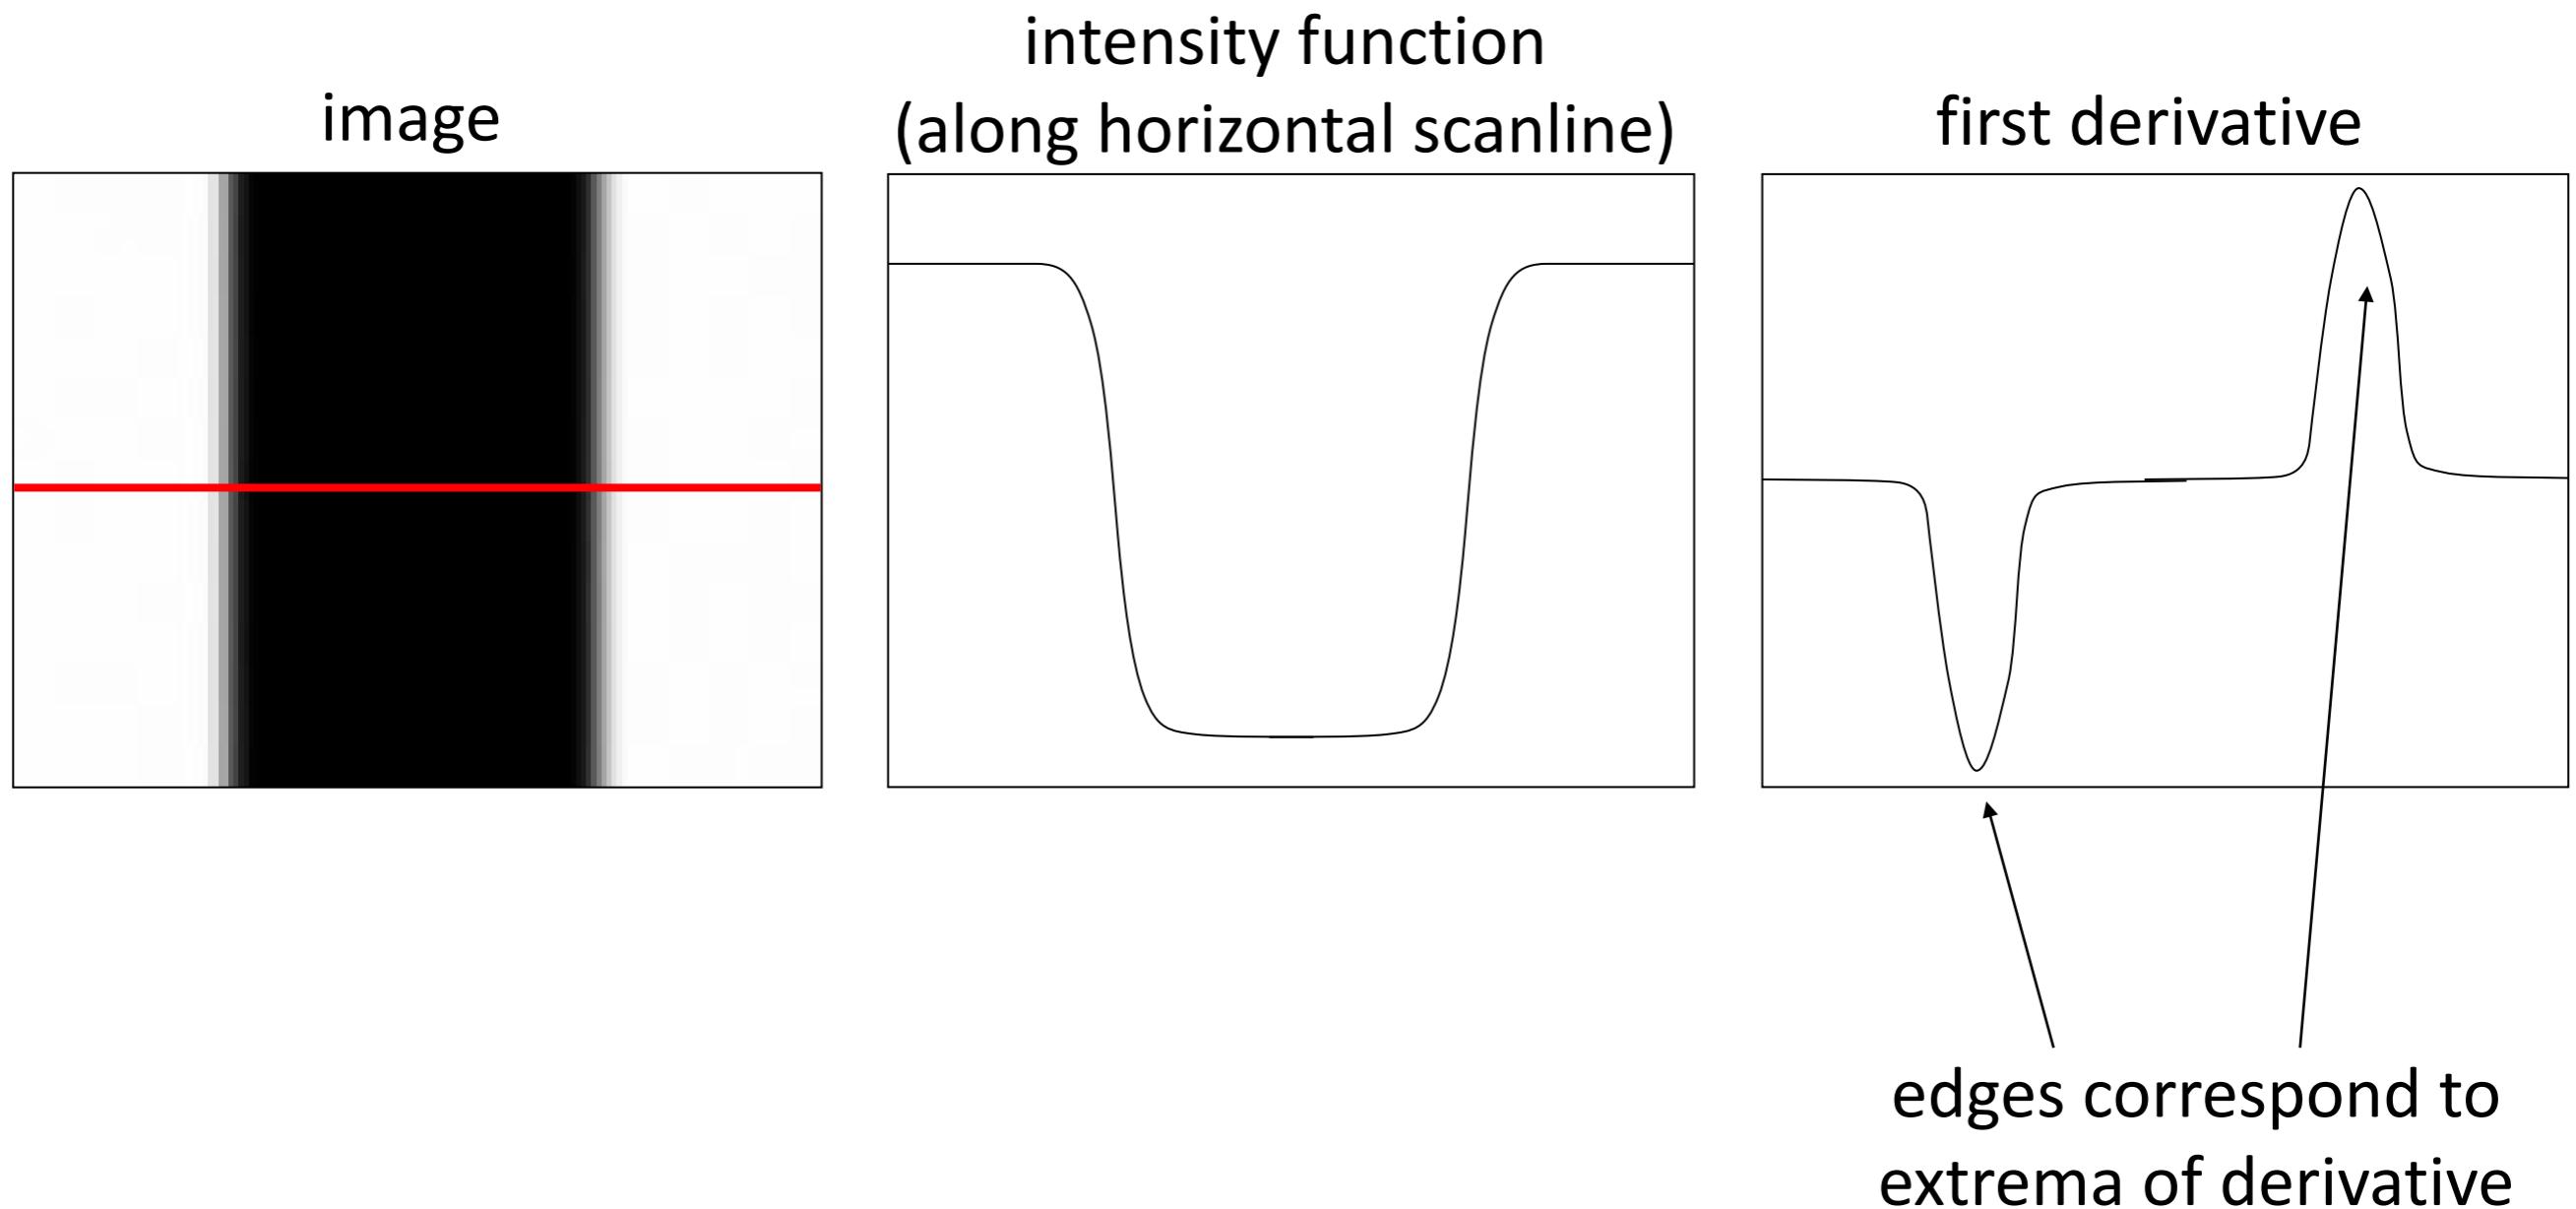
- Edges look like steep cliffs



Source: S. Seitz

Derivatives and edges

- An edge is a place of rapid change in the image intensity function.



Differentiation

- For discrete 2D function, $f(x,y)$, the partial derivative is:

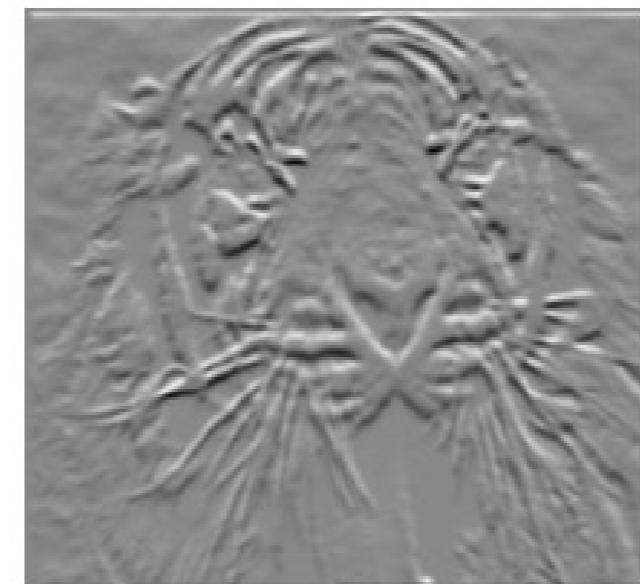
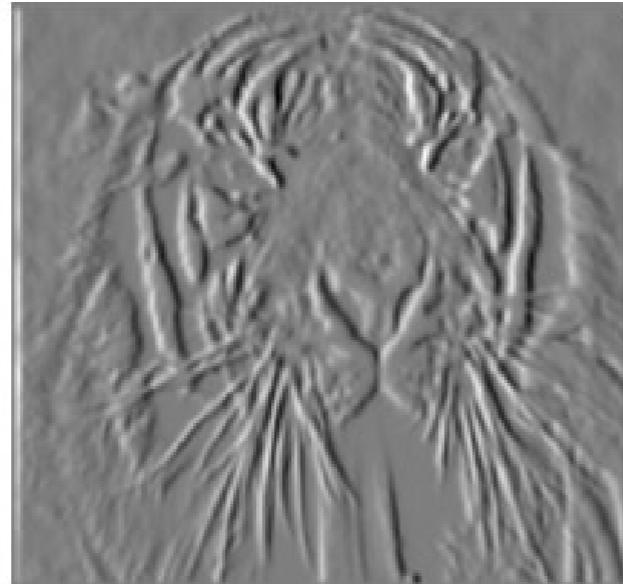
$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x+1, y) - f(x, y)}{1}$$

- To implement above as convolution, what would be the associated filter?

Partial derivatives of an image

$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

-1
1

Which shows changes with respect to x?

Typical Edge filters

Prewitt: $M_x =$

-1	0	1
-1	0	1
-1	0	1

; $M_y =$

1	1	1
0	0	0
-1	-1	-1

Sobel: $M_x =$

-1	0	1
-2	0	2
-1	0	1

; $M_y =$

1	2	1
0	0	0
-1	-2	-1

Roberts: $M_x =$

0	1
-1	0

; $M_y =$

1	0
0	-1

Typical Edge filters

Prewitt: $M_x =$

-1	0	1
-1	0	1
-1	0	1

; $M_y =$

1	1	1
0	0	0
-1	-1	-1

Sobel: $M_x =$

-1	0	1
-2	0	2
-1	0	1

; $M_y =$

1	2	1
0	0	0
-1	-2	-1

Roberts: $M_x =$

0	1
-1	0

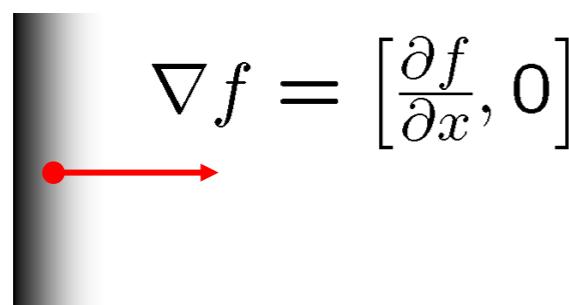
; $M_y =$

1	0
0	-1

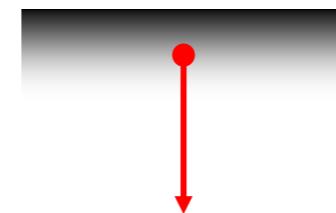
In general: the filters yield strong response on image locations which look similar – what does that mean?

Image gradient

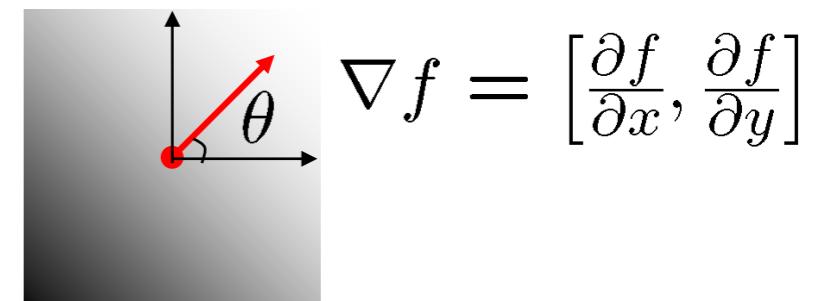
- The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$
- The gradient points in the direction of most rapid change in intensity



$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$



$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

- The gradient direction (orientation of edge normal) is given by:

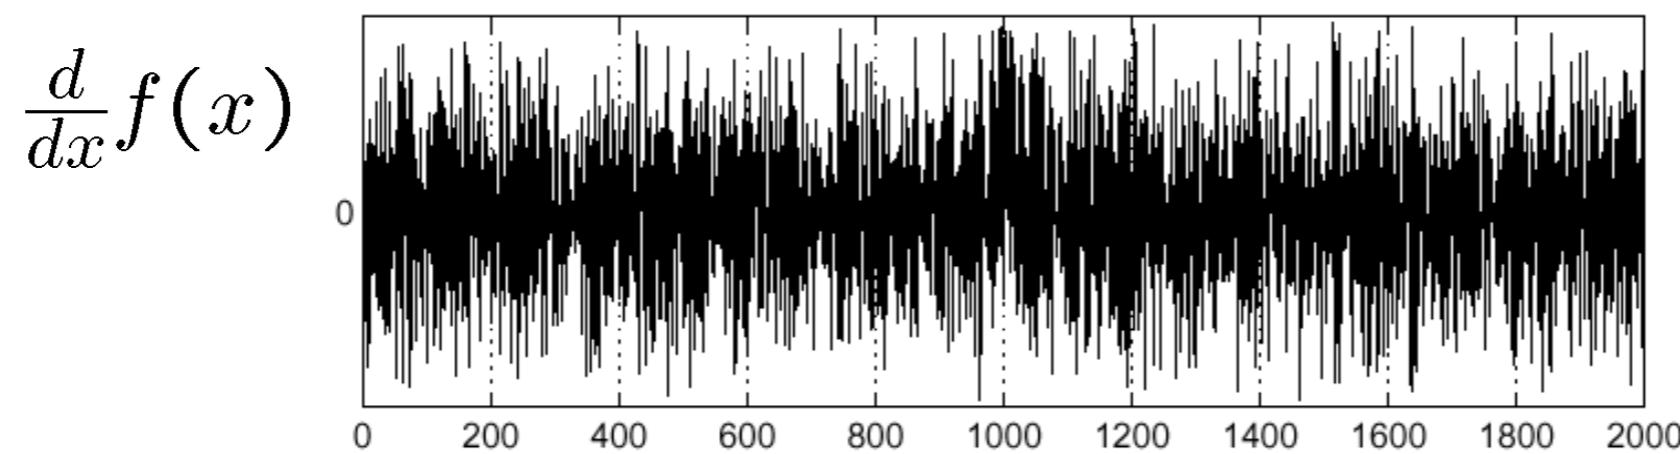
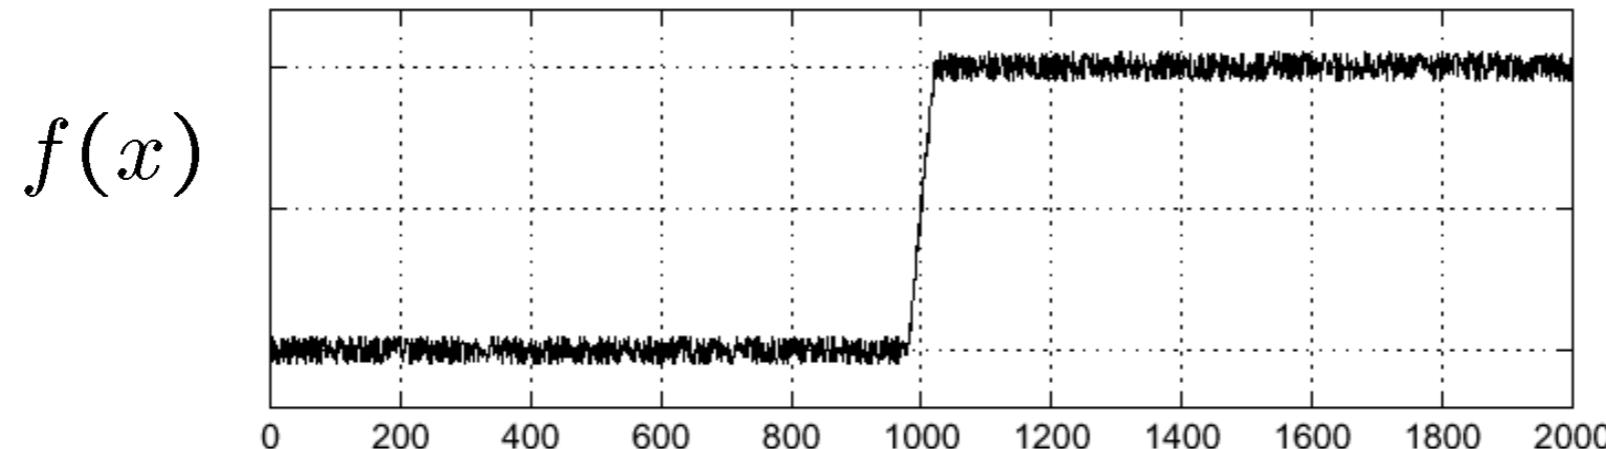
$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

- The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Effects of noise on first derivation (gradient)

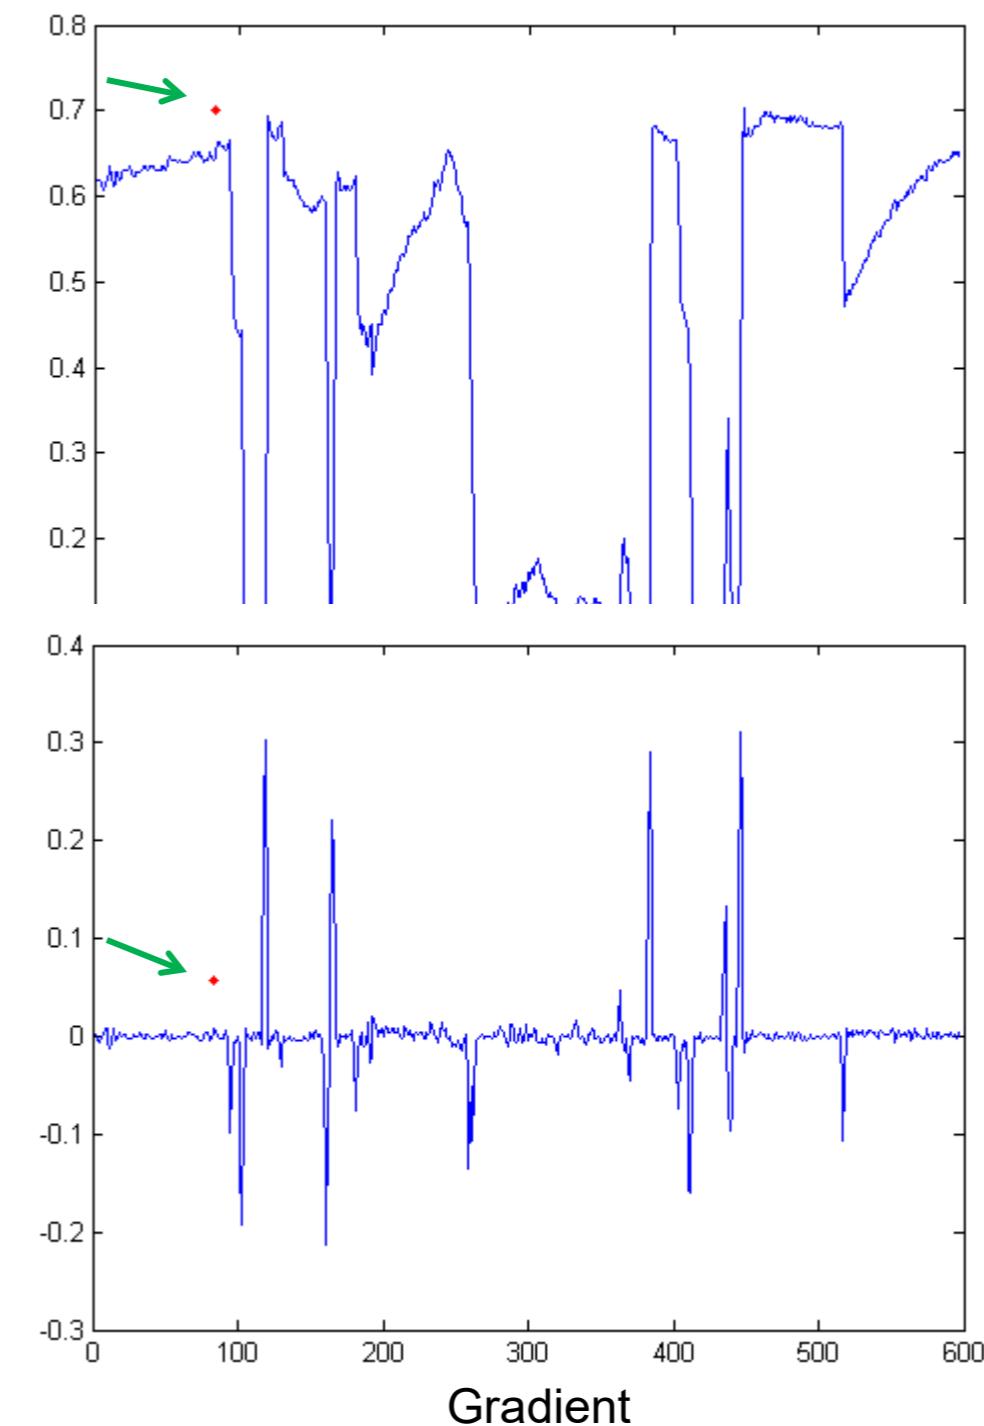
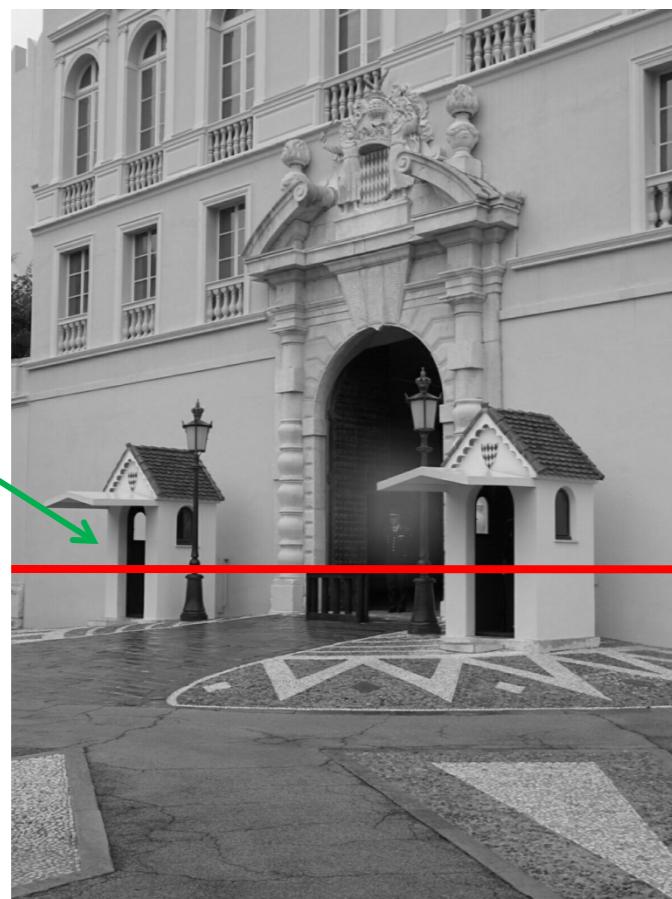
- Consider a single row or column of the image
- Plotting intensity as a function of position gives a signal



- Where is the edge?

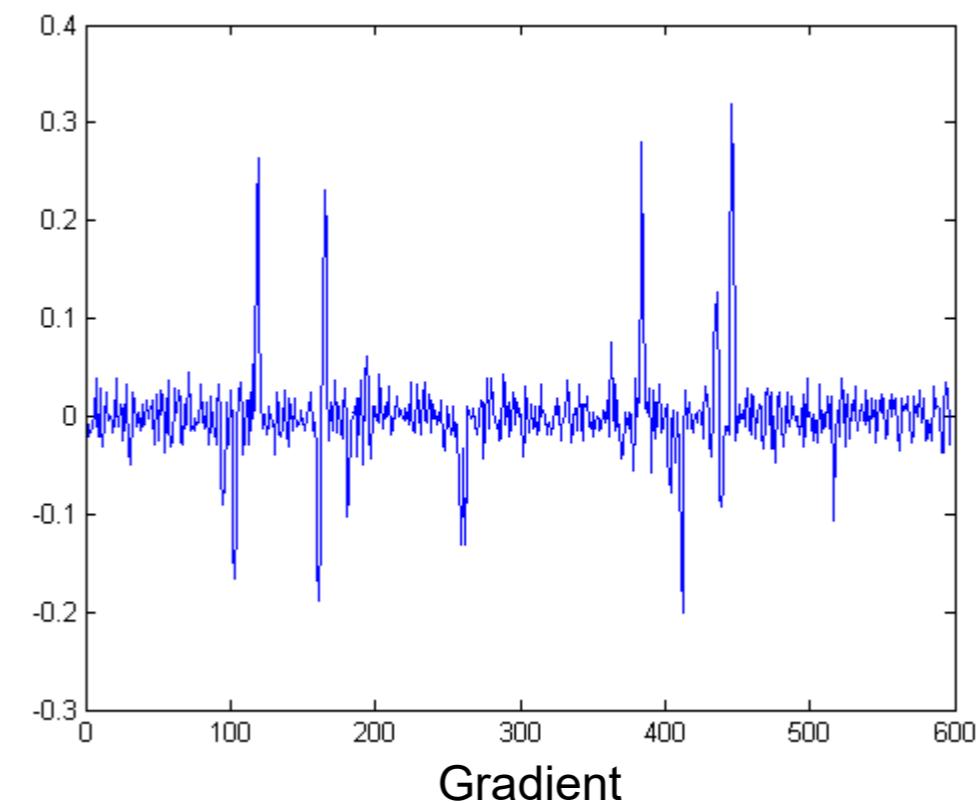
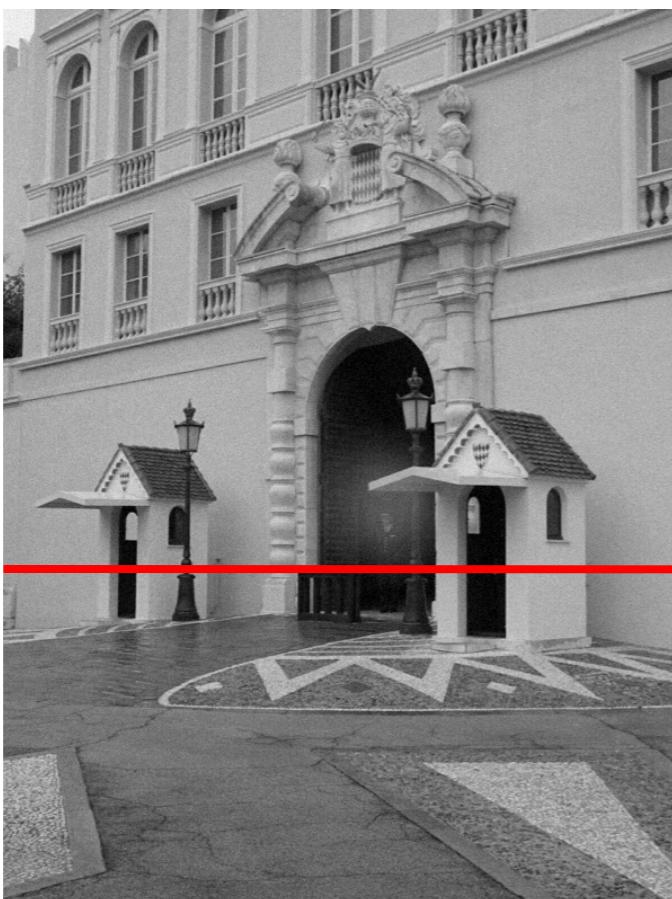
Example

- Intensity profile:

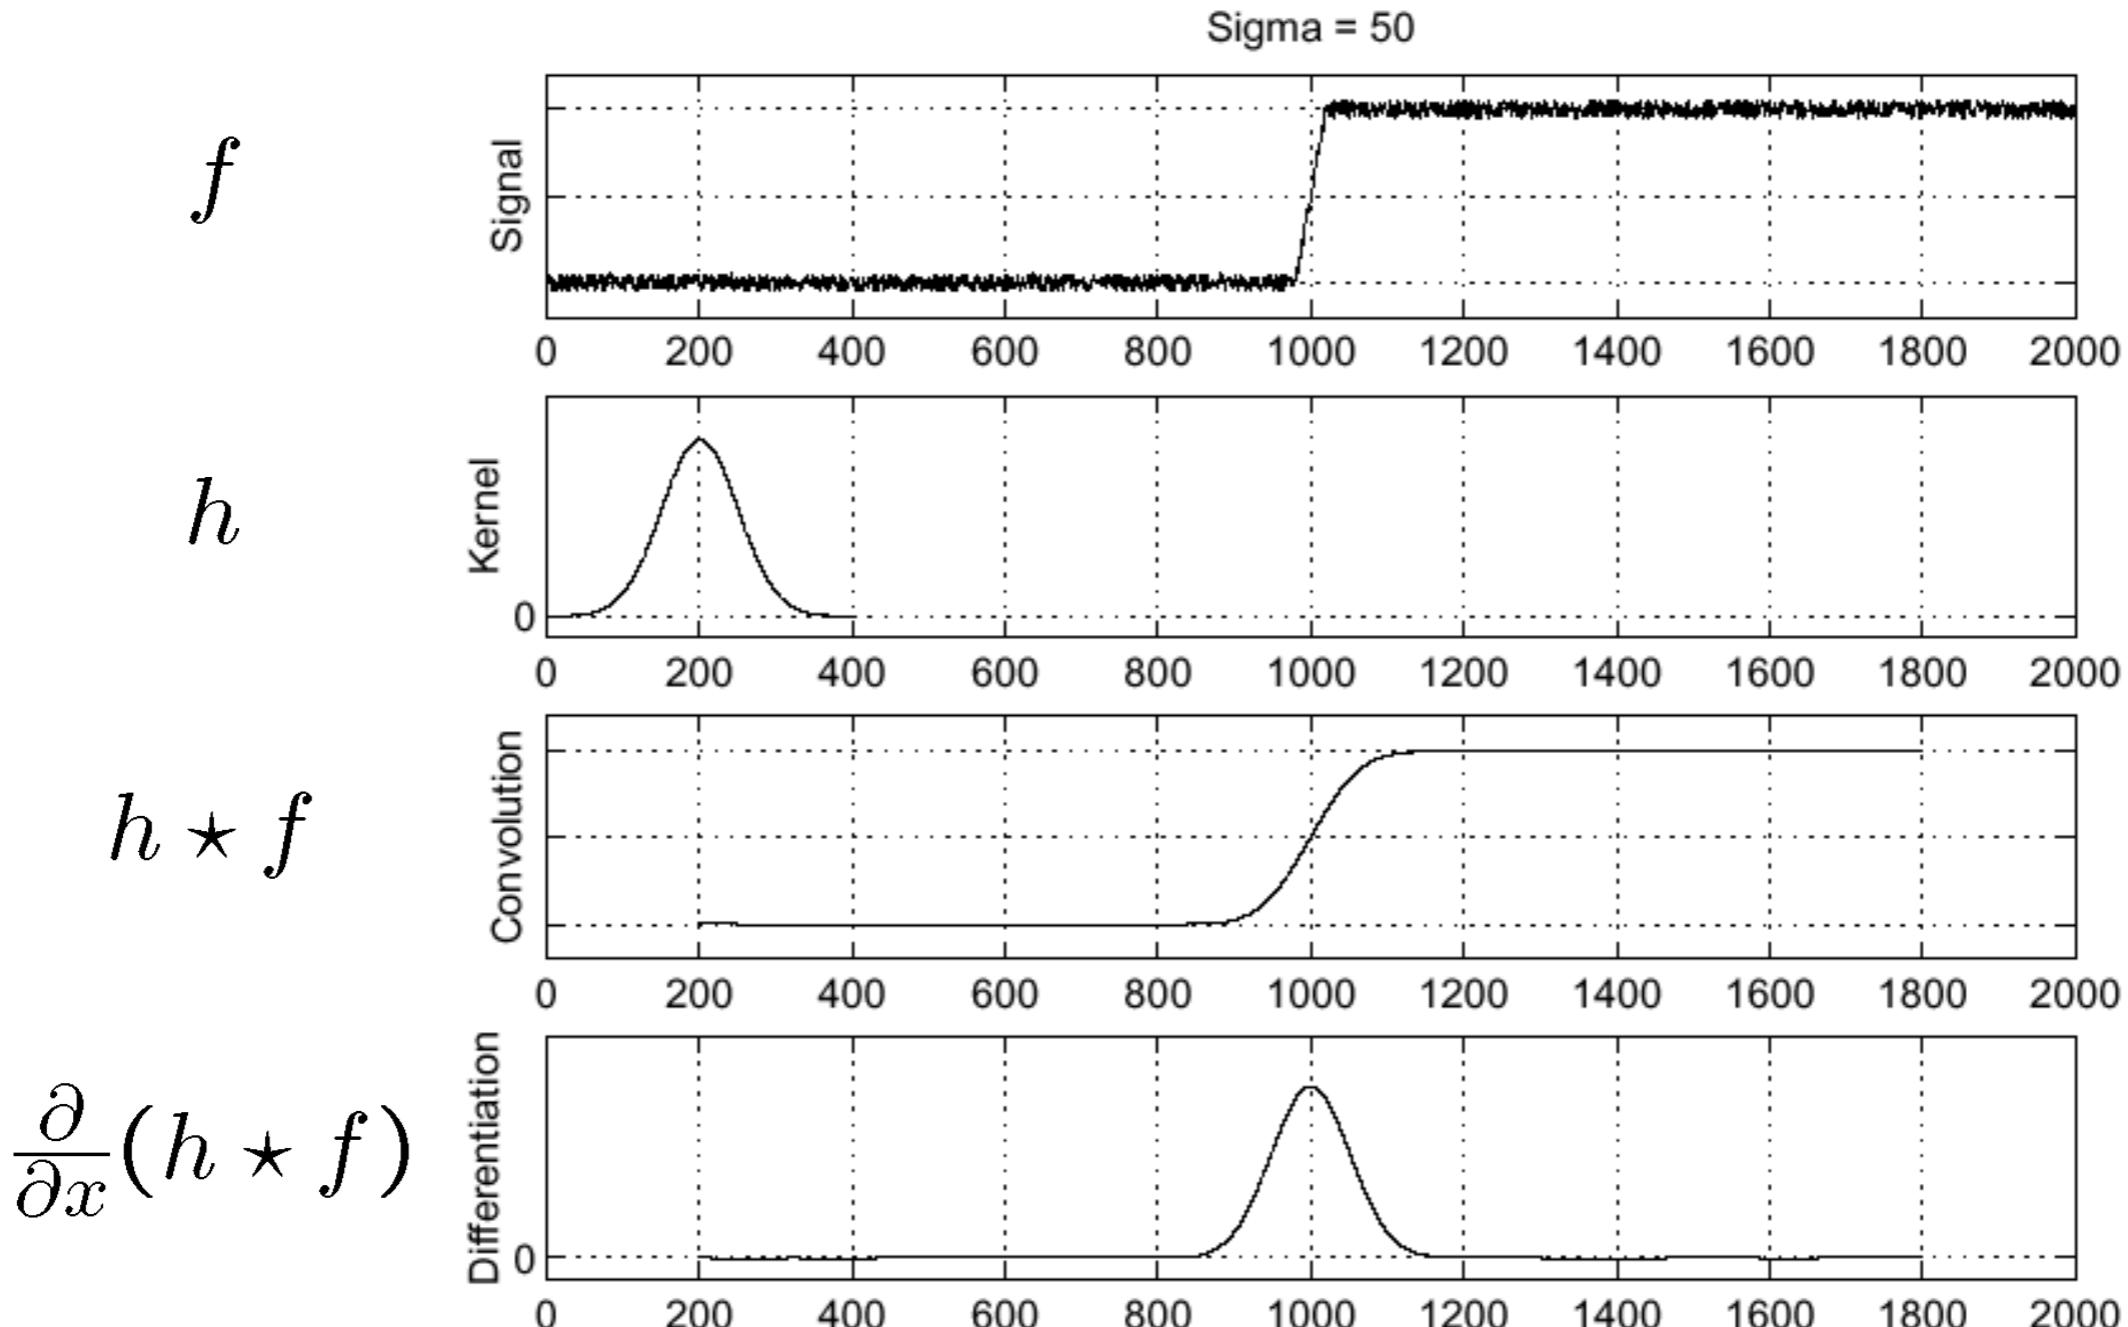


Example with added Noise

- Intensity profile:



Solution: smooth first to remove noise



Where is the edge?

Look for peaks in

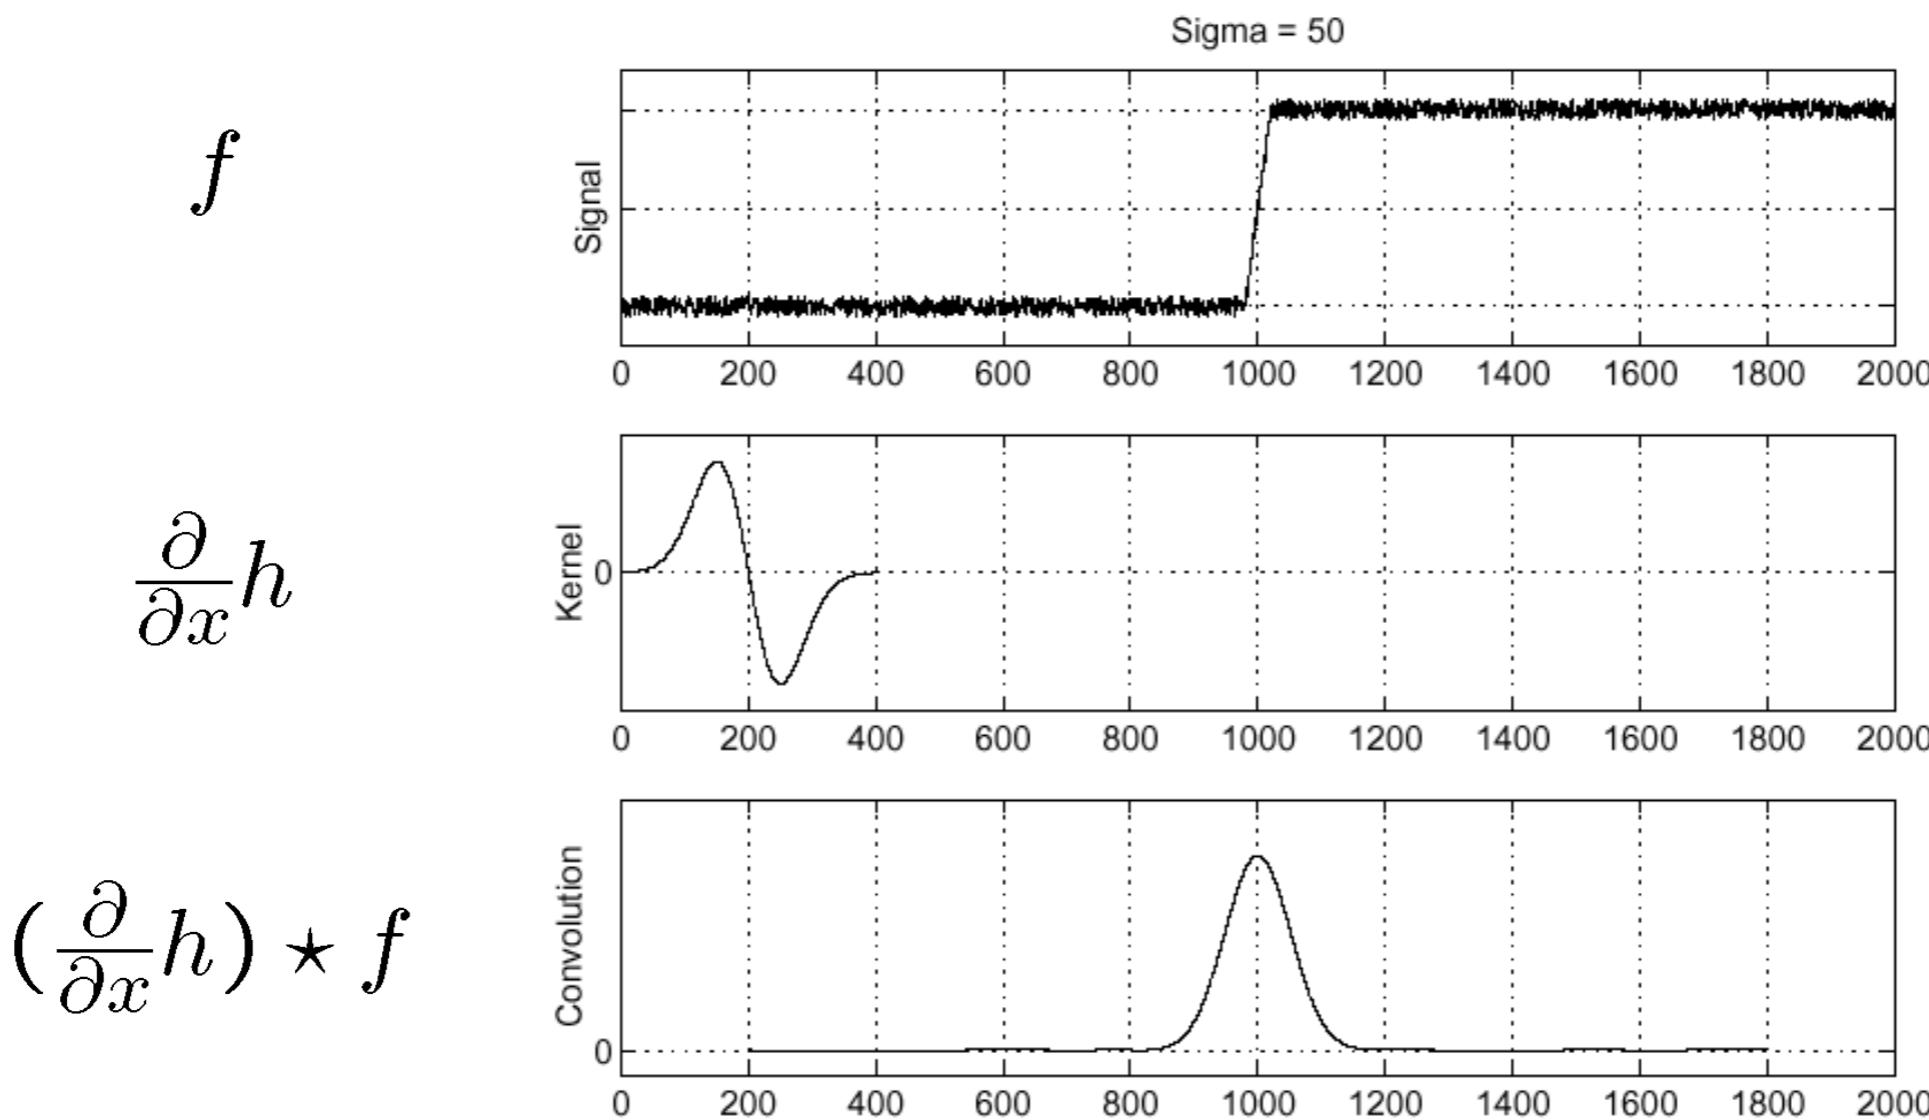
$\frac{\partial}{\partial x}(h \star f)$

Derivative theorem of convolution

- Combine the filters first, i.e. take Gaussian and compute 1st derivation

$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

- And filter then → same result!

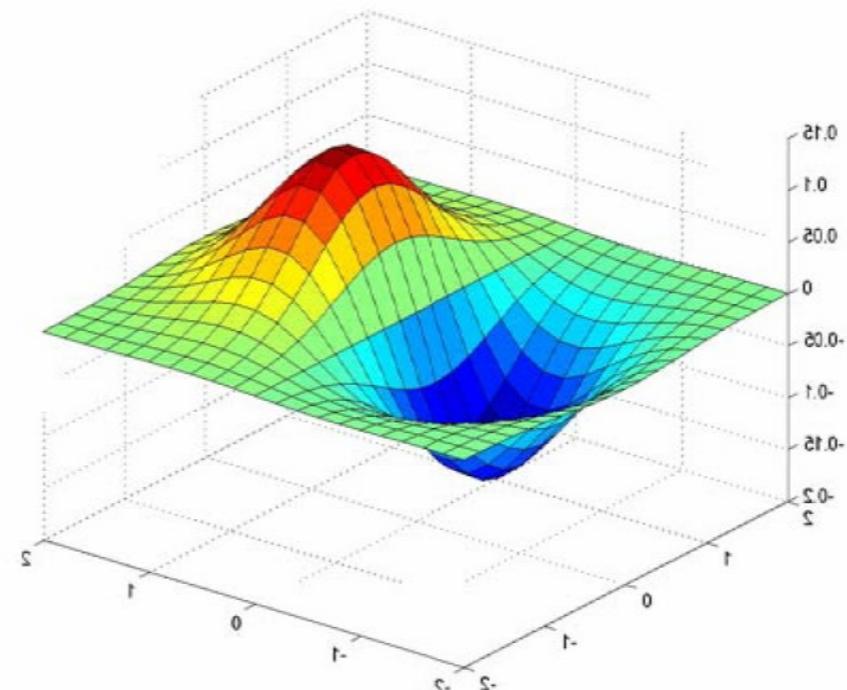


Derivative of Gaussian filter

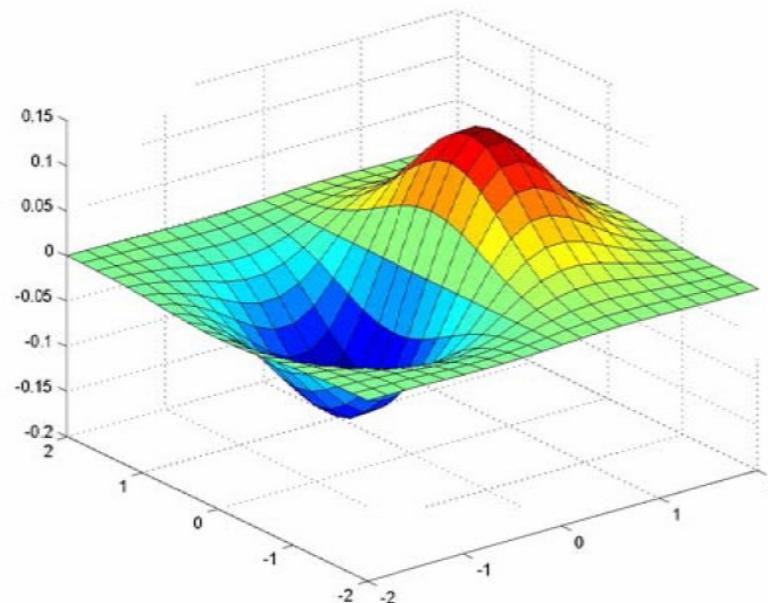
$$(I \otimes g) \otimes h = I \otimes (g \otimes h)$$

$$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix}$$

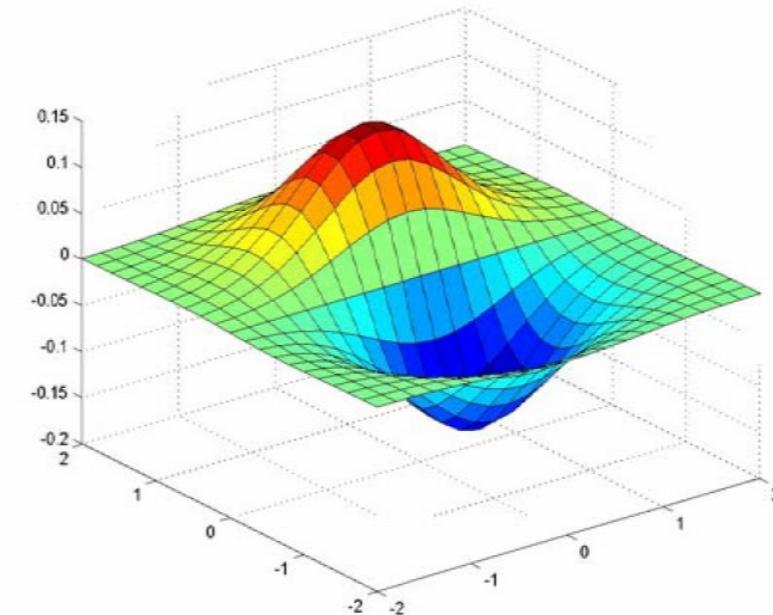
$$\otimes \begin{bmatrix} 1 & -1 \end{bmatrix}$$



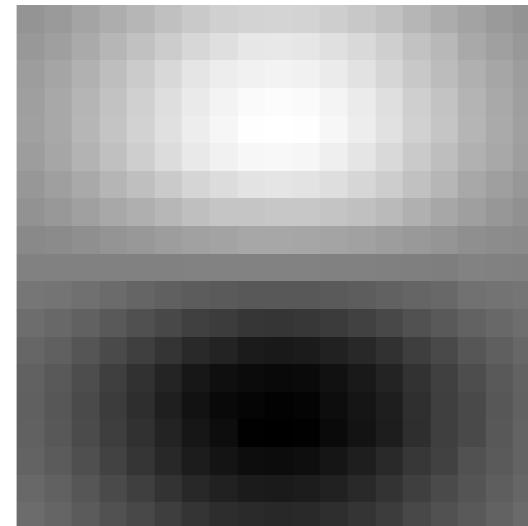
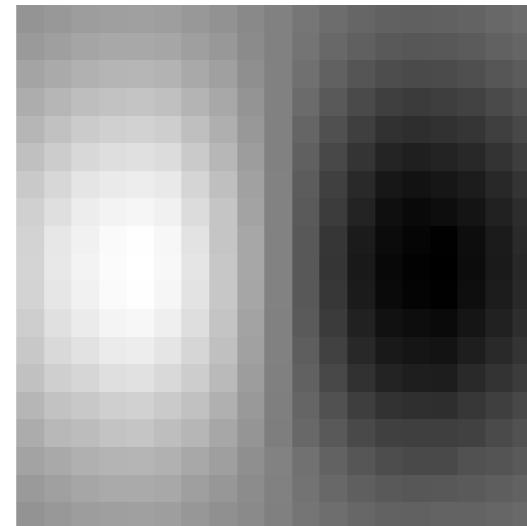
Derivative of Gaussian filters



x-direction



y-direction



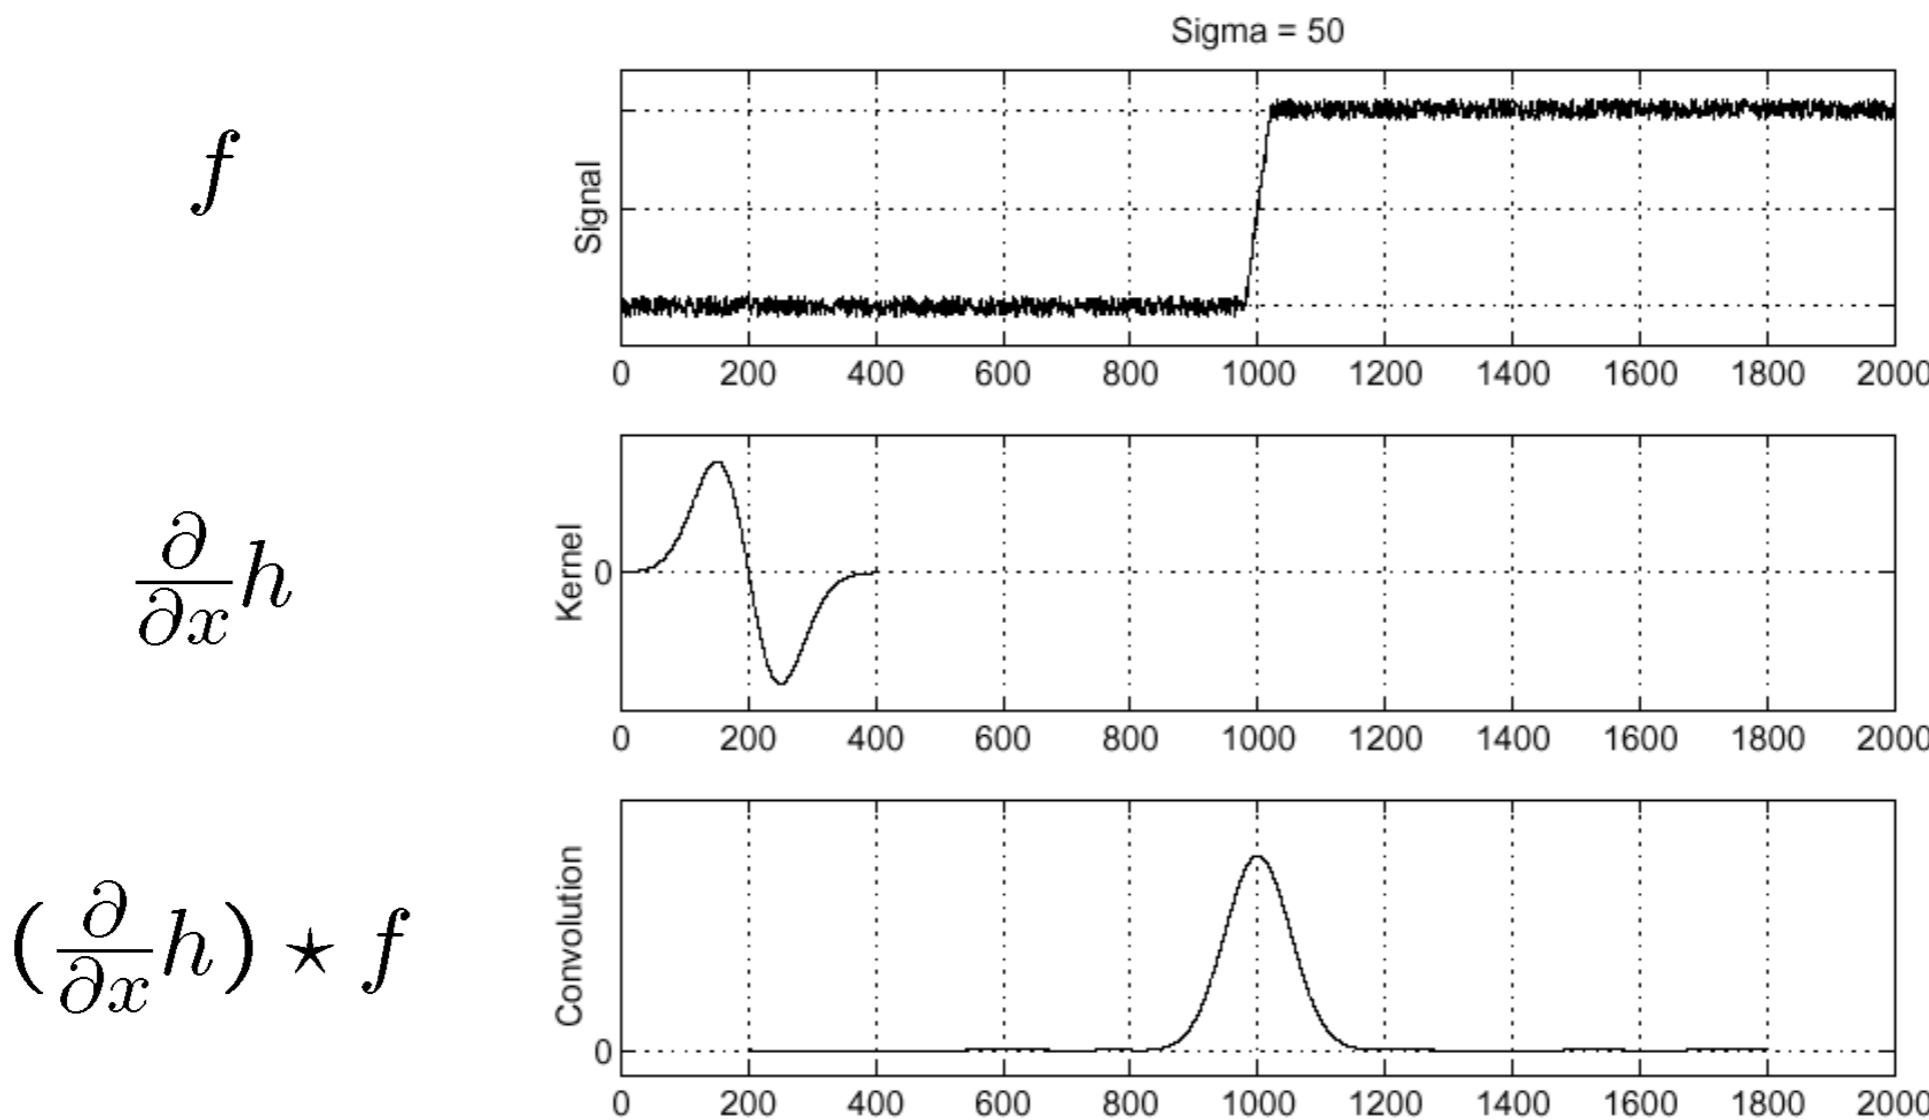
- Nice (noise robust) filters for horizontal and vertical edges!
- But what about other edge directions?

Derivative theorem of convolution

- Combine the filters first, i.e. take Gaussian and compute 1st derivation

$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

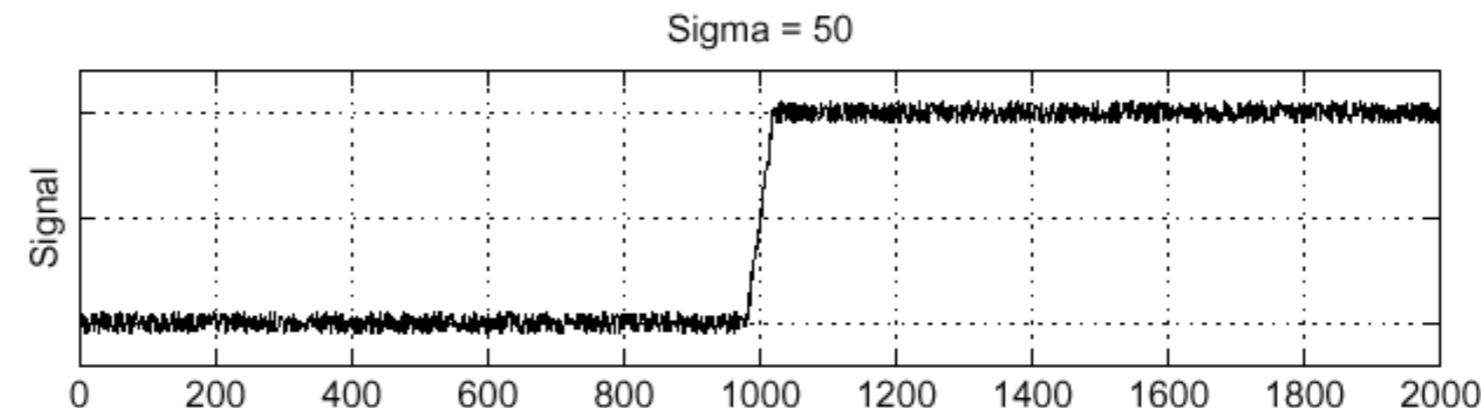
- And filter then → same result!



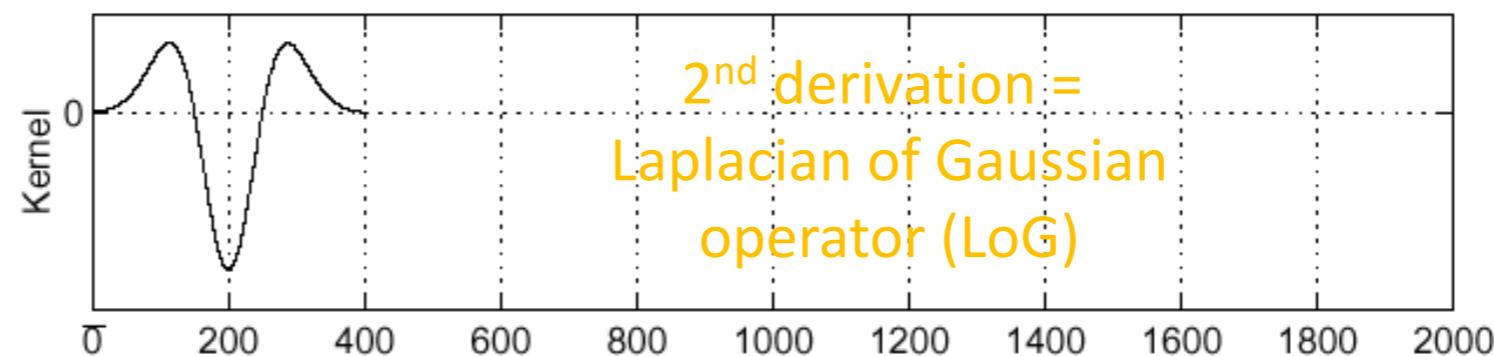
In 1D: Laplacian of Gaussian

- Consider $\frac{\partial^2}{\partial x^2}(h \star f)$

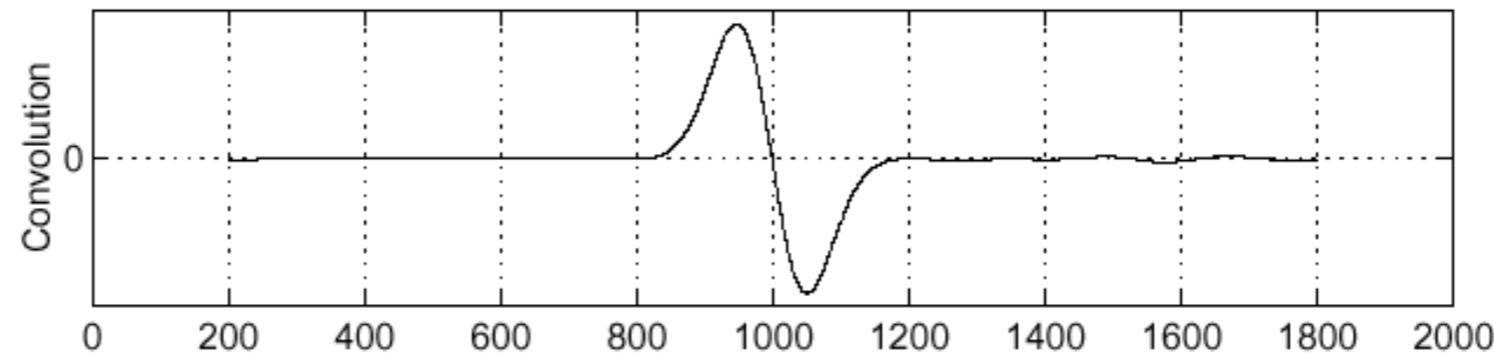
f



$\frac{\partial^2}{\partial x^2} h$



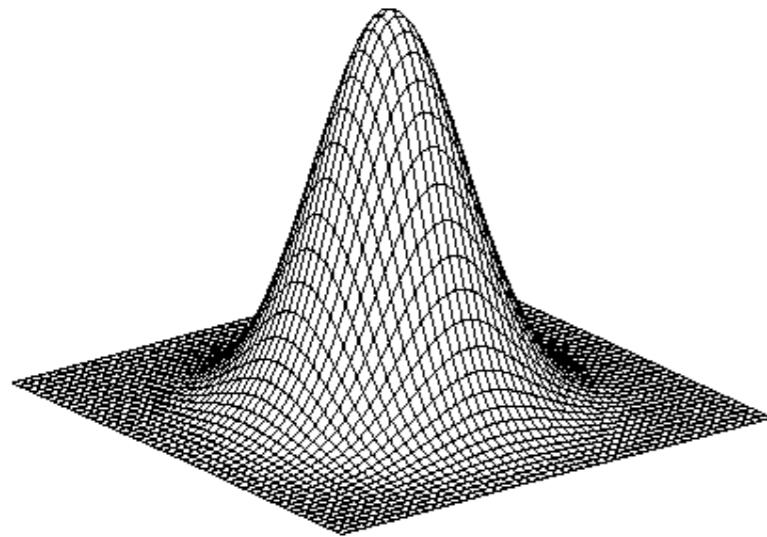
$(\frac{\partial^2}{\partial x^2} h) \star f$



Where is the edge?

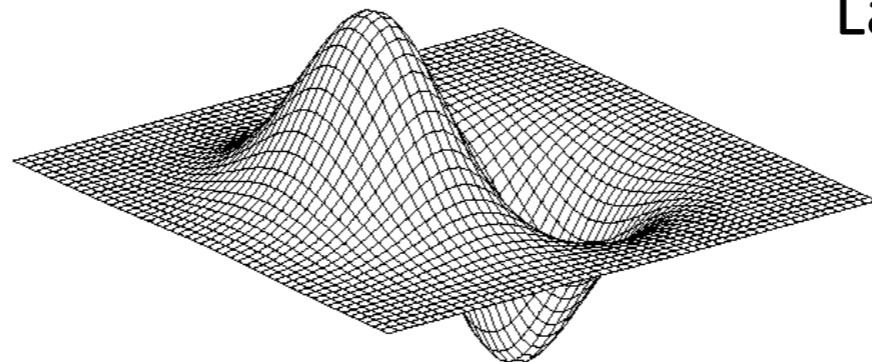
Zero-crossings of bottom graph

2D edge detection filters



Gaussian

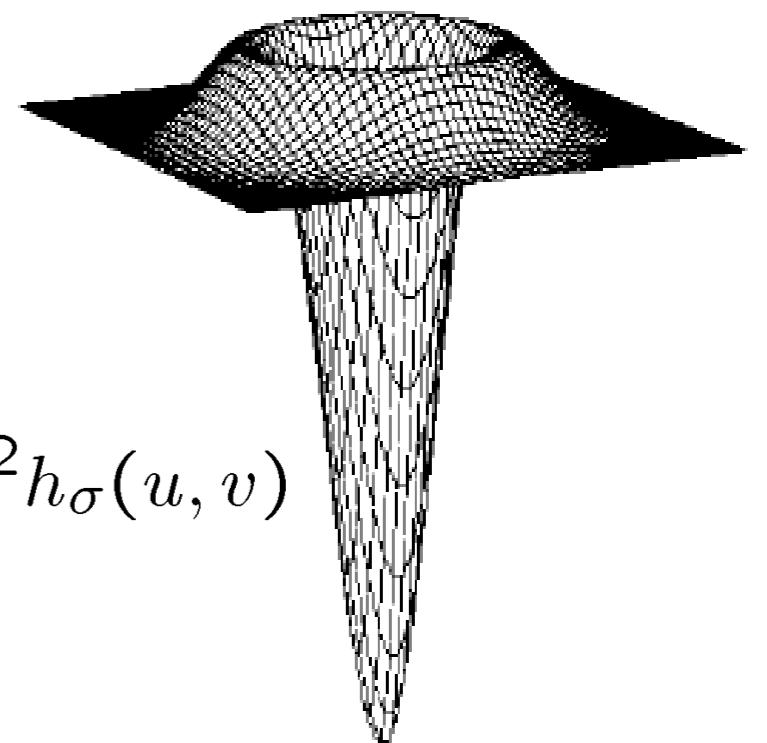
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

Laplacian of Gaussian (LoG)



$$\nabla^2 h_\sigma(u, v)$$

∇^2 is the Laplacian operator:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

3x3 LoG:

0	1	0
1	-4	1
0	1	0

Filter properties

- Smoothing
 - All values positive
 - Sum to 1 → constant regions same as input
 - Amount of smoothing proportional to filter size
 - Remove “high-frequency” components → “low-pass” filter

- Edges
 - Opposite signs used to get high response in regions of high contrast
 - Sum to 0 → no response in constant regions
 - High absolute value at points of high contrast
 - Remove “low-frequency” components → “high-pass” filter

- Filters act as templates
 - Highest response for regions that “look the most like the filter”
 - Dot product as correlation

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$$\begin{matrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{matrix}$$

Computation of Gradients

- Apply the LoG filter to the image
- Result is an intensity map
 - High values: strong gradients → potentially important edges (high contrast regions)
 - Low values: weak or no gradients → rather homogenous regions



Input image

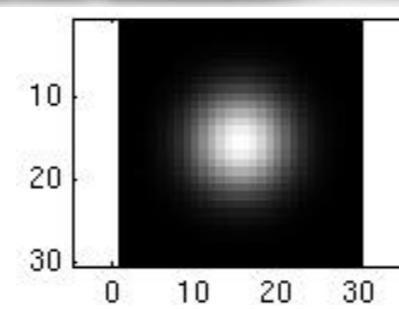
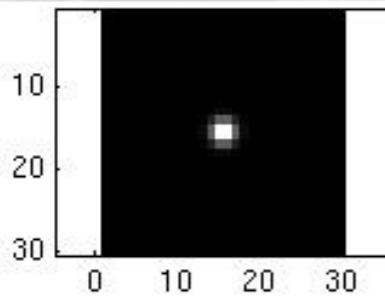
LoG
→



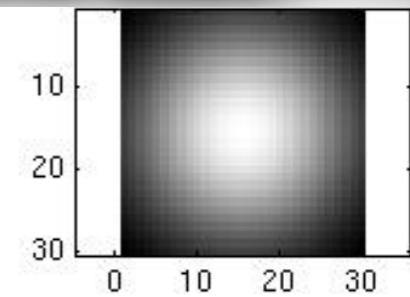
Gradients

Influence on Amount of Smoothing on Edge Detection (LoG Filter)

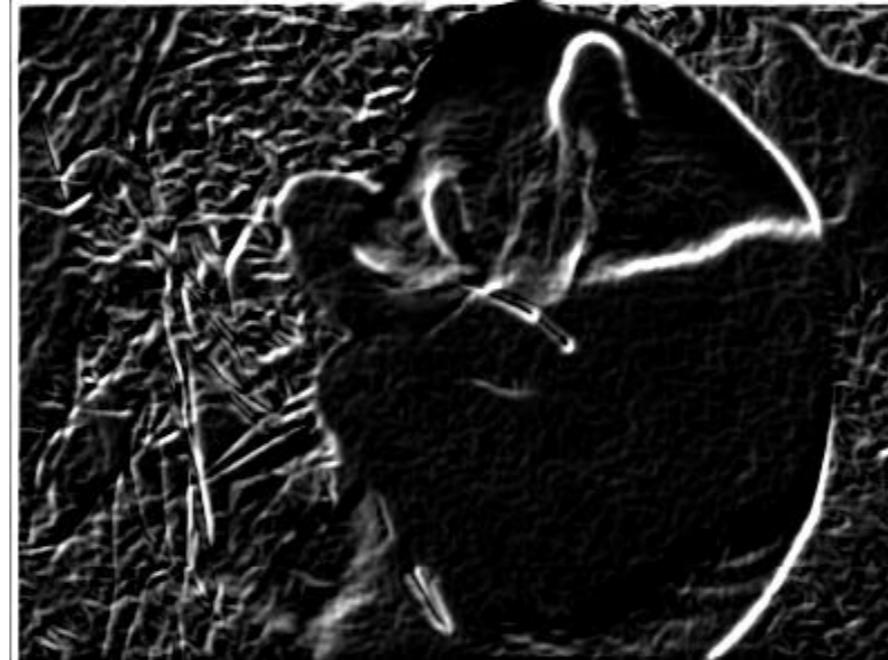
- Recall: parameter σ is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



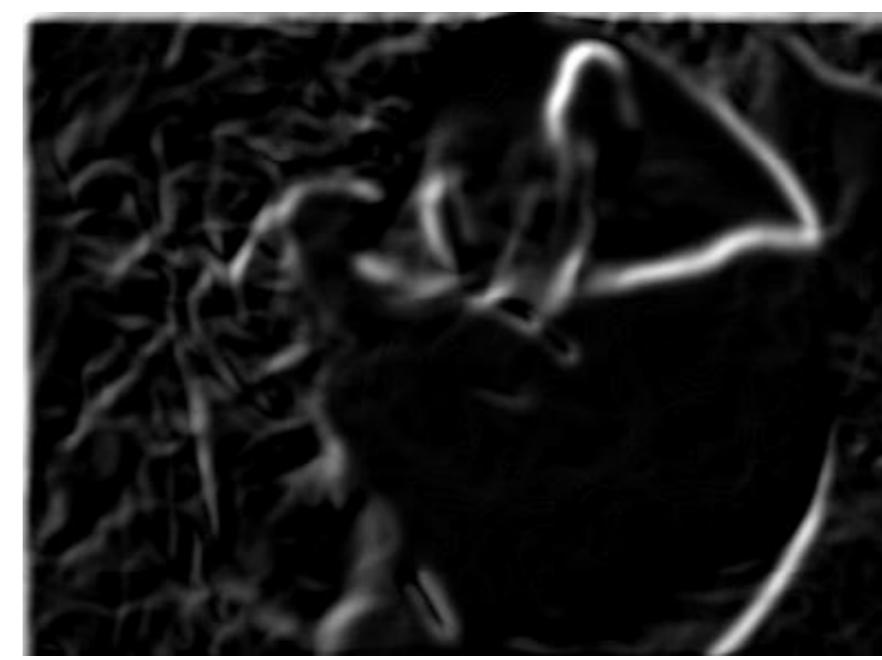
...



Influence on Amount of Smoothing on Edge Detection (LoG Filter)



$\sigma = 1$ pixel



$\sigma = 3$ pixels

- The apparent structures differ depending on Gaussian's scale parameter.
- Larger values: larger scale edges detected
- Smaller values: finer features detected

So, what scale to choose?

- It depends what we are looking for...



- Too fine of a scale... can't see the forest for the trees.
- Too coarse of a scale...can't tell the maple grain from the cherry.

A Simple Edge Detection:

- Choose a threshold value t
- Set any pixels less than t to zero (off)
- Set any pixels greater than or equal to t to one (on)

Original image



Gradient magnitude image



Thresholding gradient with a lower threshold



Thresholding gradient with a higher threshold



Are we happy with that?

Canny edge detector

- Filter image with derivative of Gaussian (e.g. LoG)
- Find magnitude and orientation of gradient
- **Non-maximum suppression:**
 - Thin multi-pixel wide “ridges” down to single pixel width
- Linking and thresholding (**hysteresis**):
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them
- Python:
 - `edges2 = feature.canny(im, sigma=3)`
 - see: http://scikit-image.org/docs/dev/auto_examples/edges/plot_canny.html

The Canny edge detector



The Canny edge detector



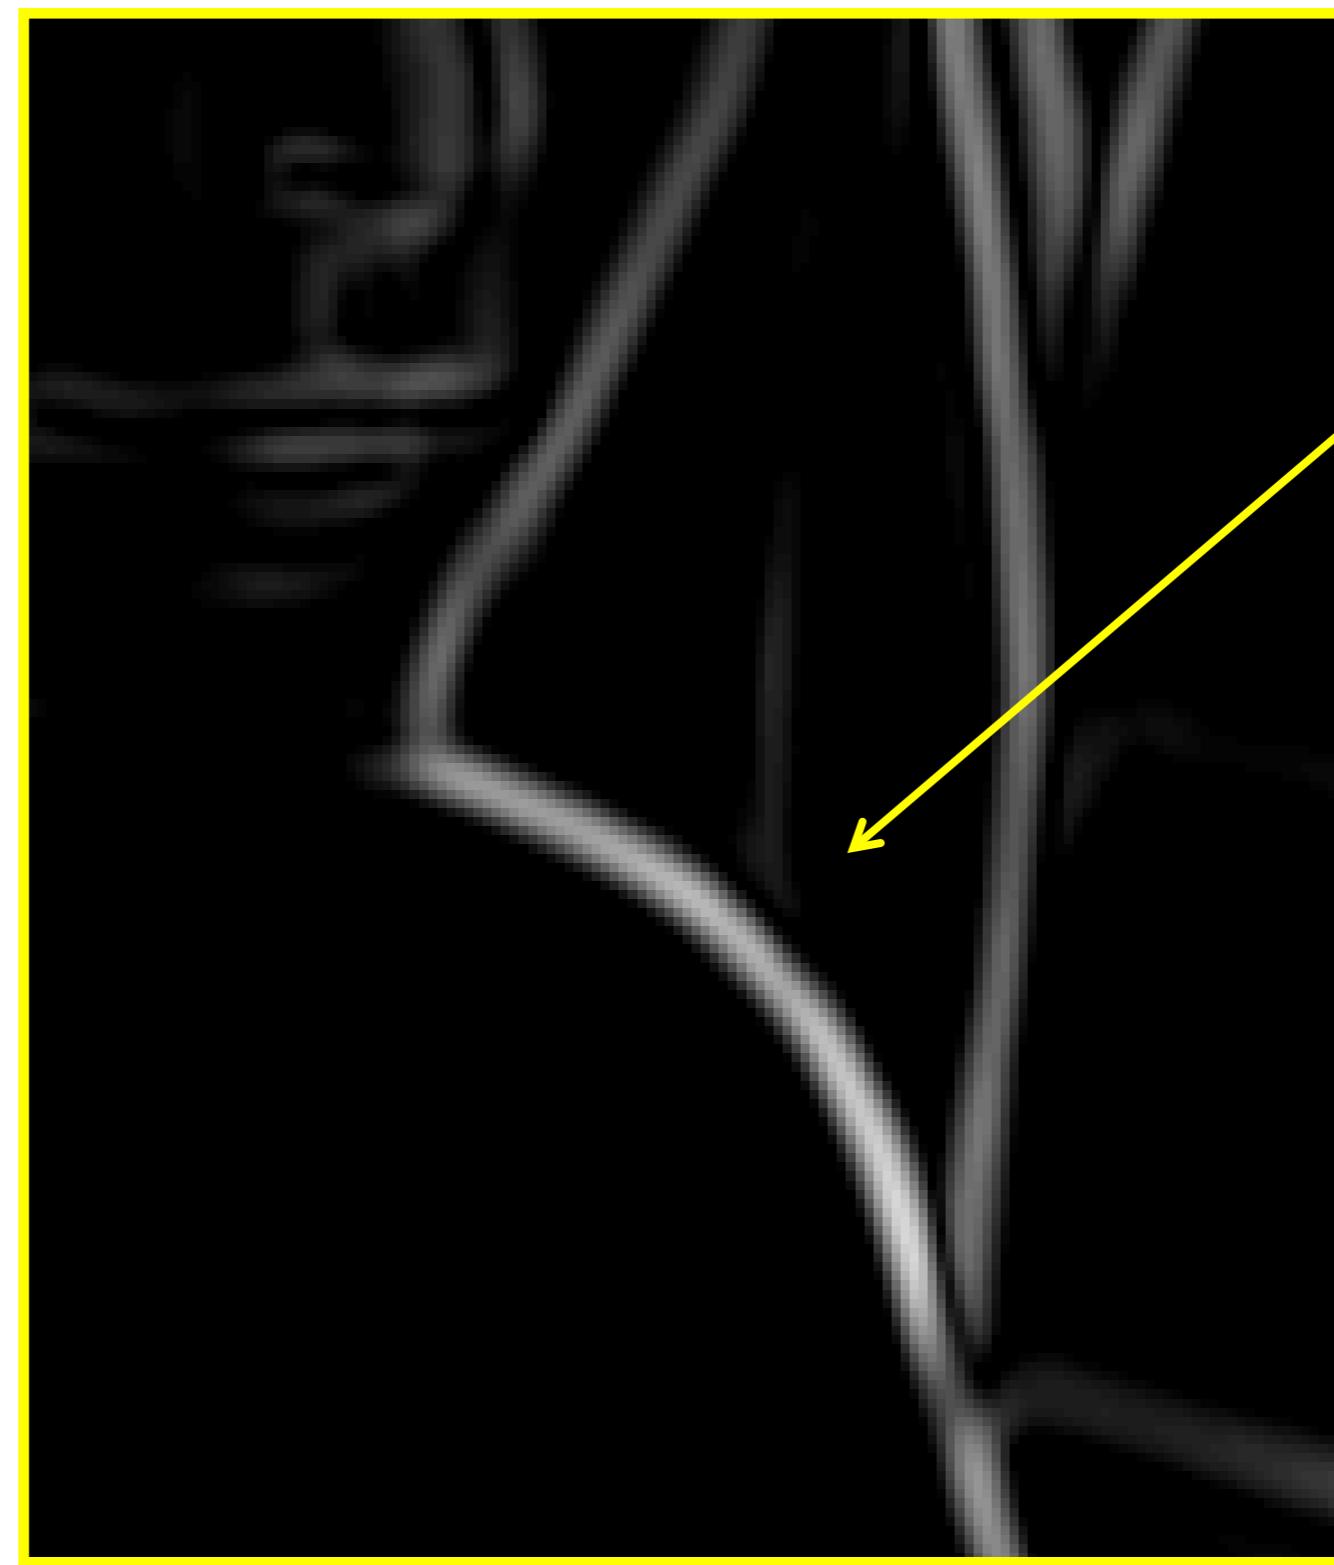
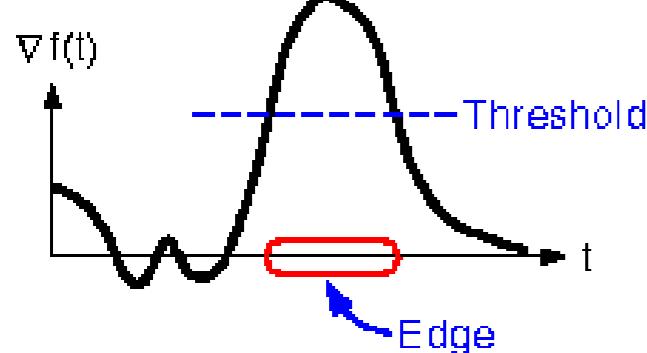
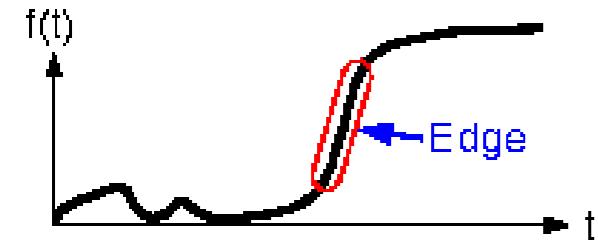
gradient map

The Canny edge detector



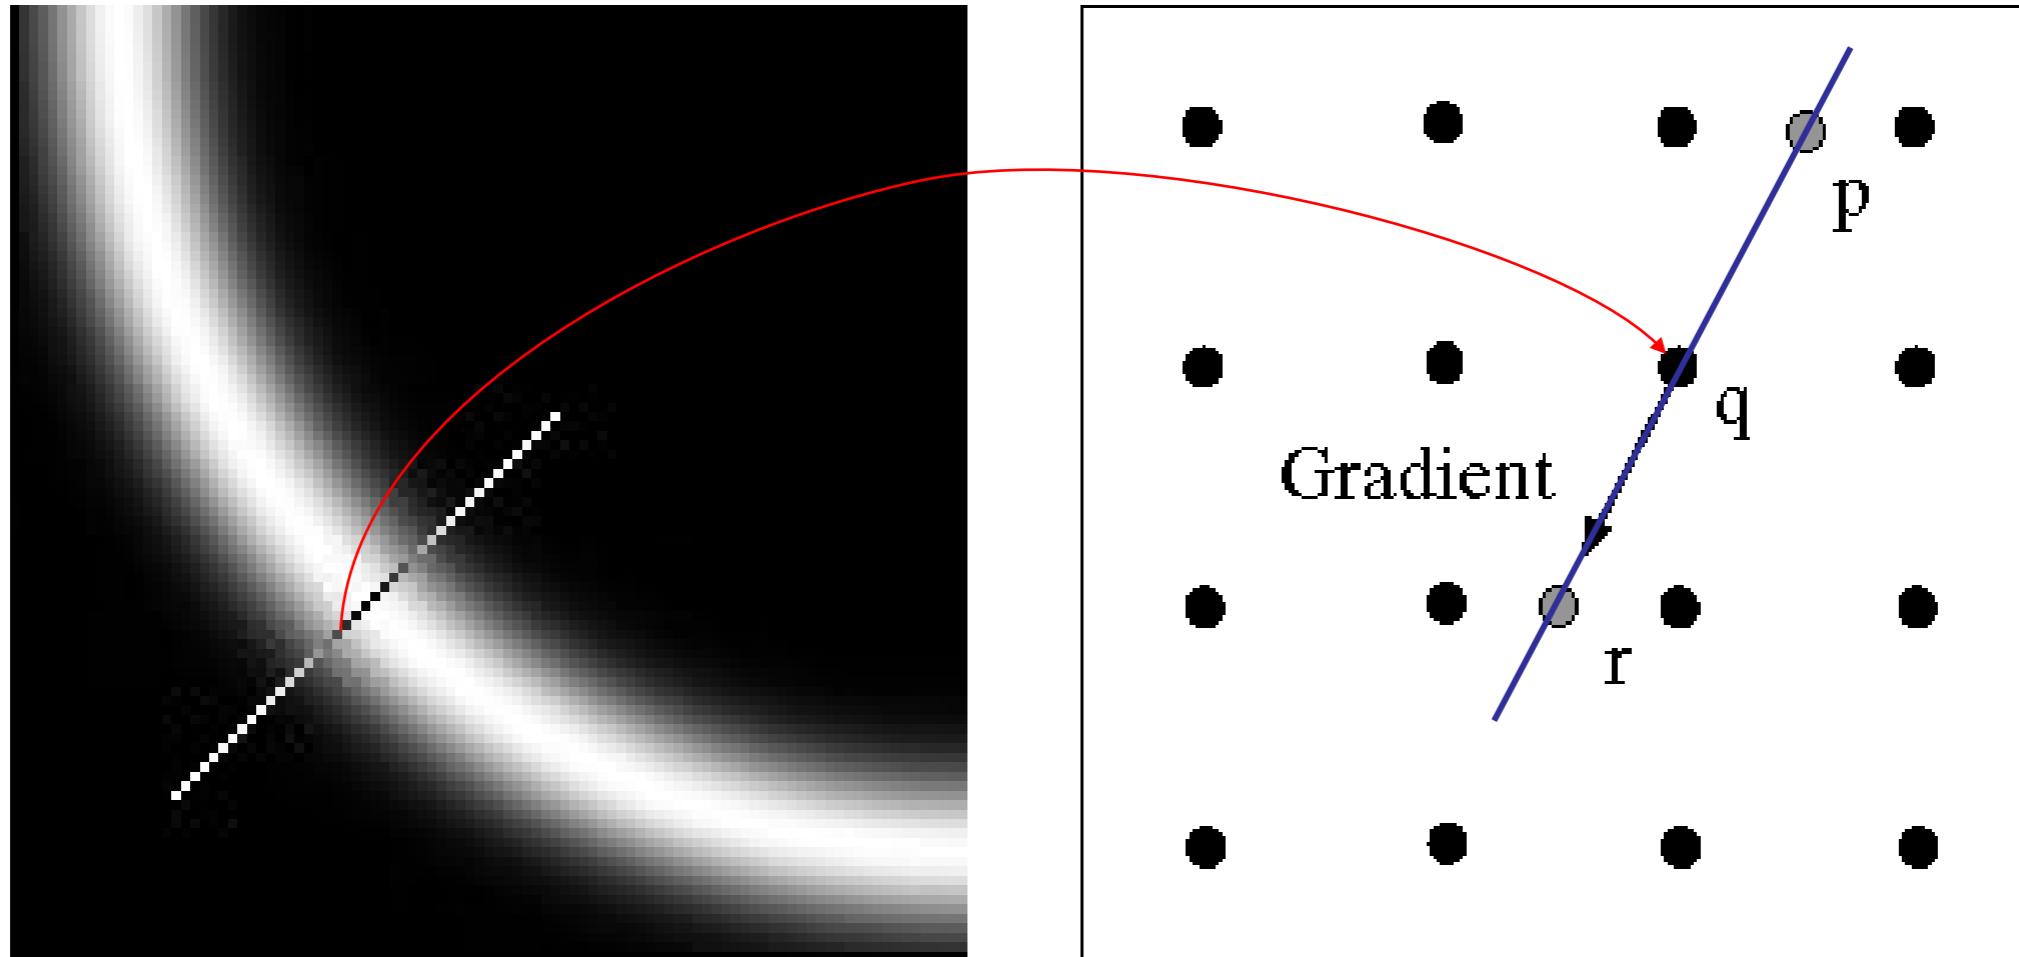
thresholding

The Canny edge detector



How to turn
these thick
regions of the
gradient into
curves?

Non-maximum suppression



- Check if a given pixel q is local maximum along gradient direction, select single max across width of the edge
 - The “gradient line” does not intersect with the pixels! Points p and r are not valid pixels.
 - What are the values of p and r ? Estimate via bilinear interpolation
 - If $q > r$ and $q > p$ then declare q is a local maximum

The Canny edge detector

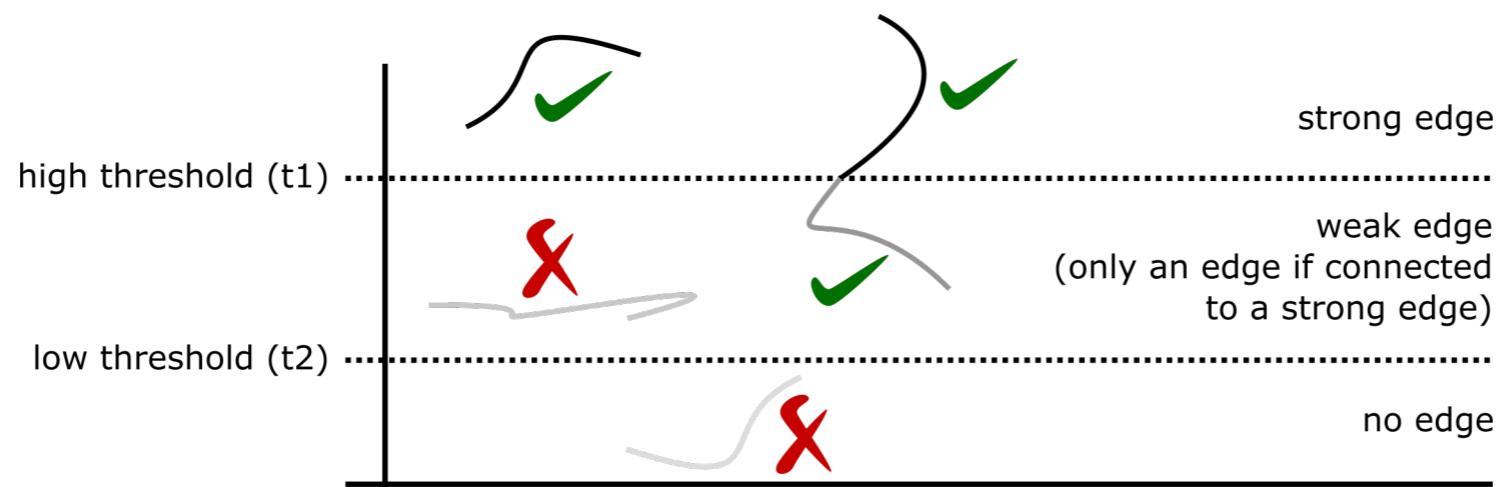


thinning
(non-maximum suppression)

Problem: pixels along this edge didn't survive the thresholding

Hysteresis thresholding

- Check that maximum value of gradient value is sufficiently large
 - Use “Hysteresis”:
 - Define two thresholds t_1 and t_2 , where t_1 is larger than t_2
 - Look for strong edge pixels which are larger than t_1
 - Go along edge until edge strengths falls under $t_2 \rightarrow$ stop edge
 - Idea:
 - Weak edges that are not connected to strong edge pixels and considered noise and are removed
 - Weak edges that are connected to strong edge pixels are kept



Hysteresis thresholding



original image



high threshold
(strong edges)



low threshold
(+weak edges)



hysteresis threshold

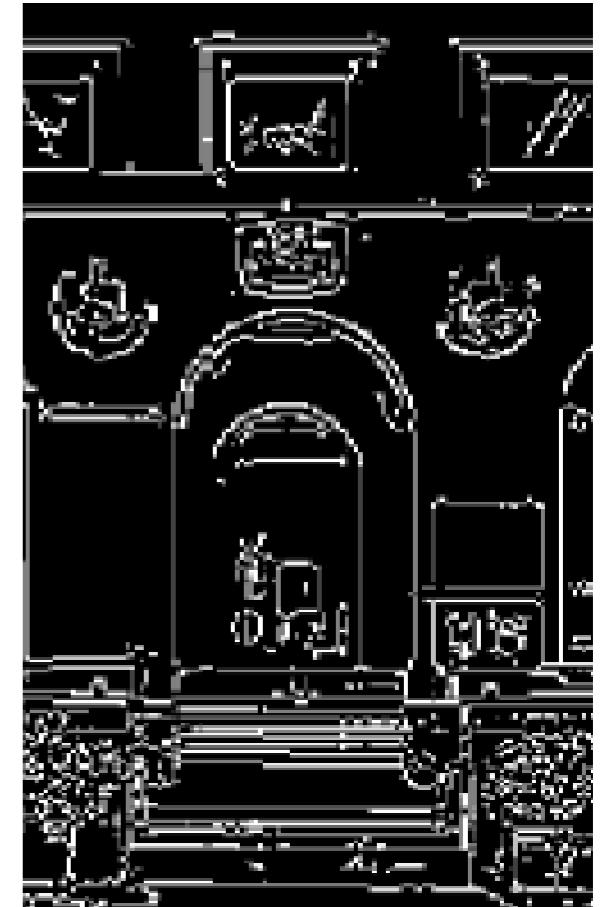
Hysteresis thresholding



**high threshold
(strong edges)**

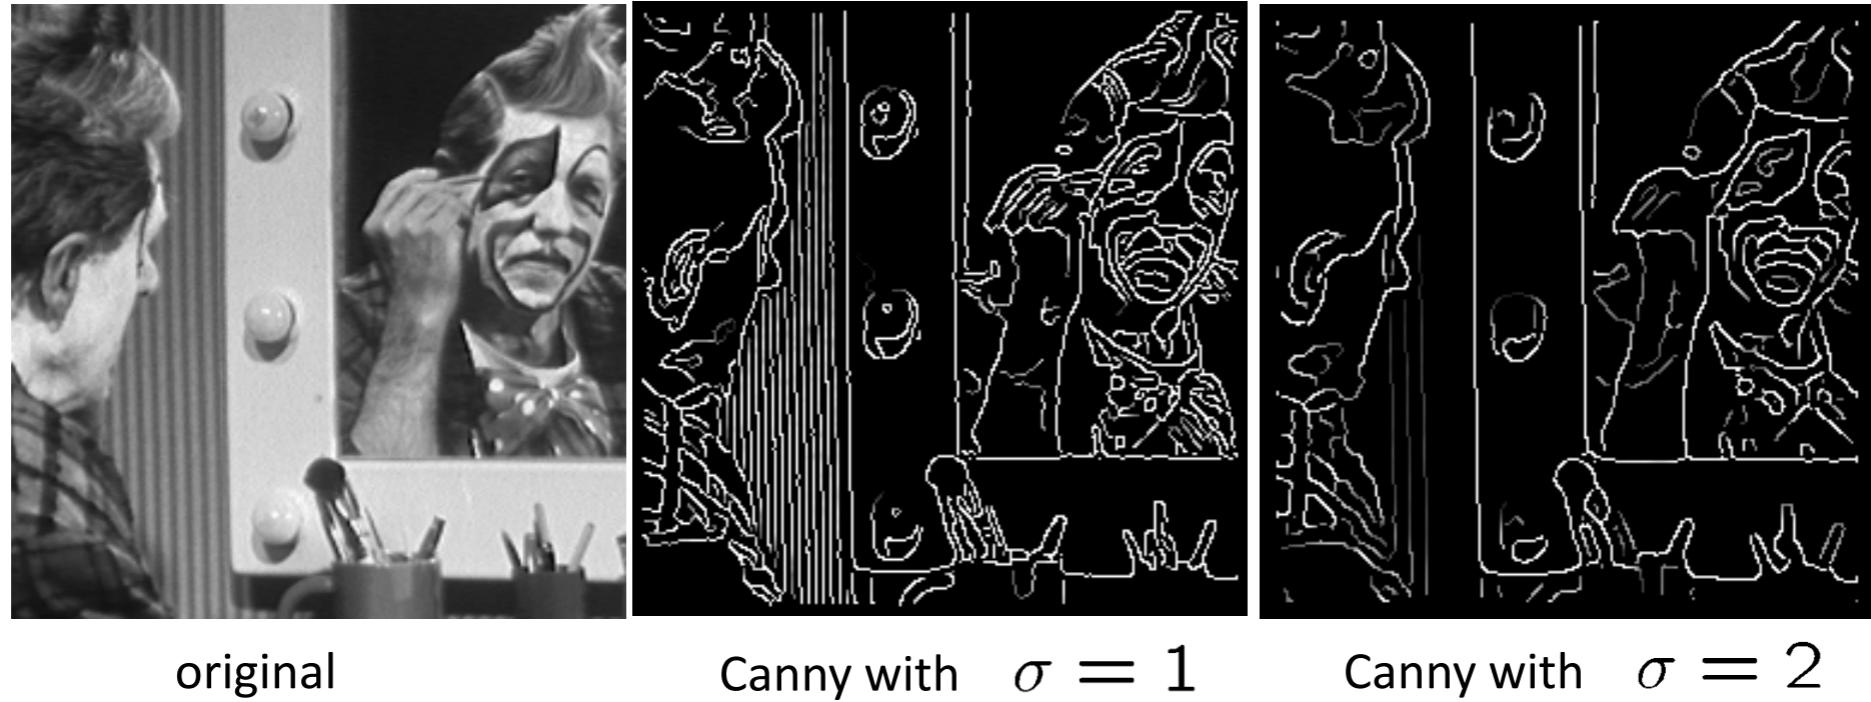


**low threshold
(weak edges)**



hysteresis threshold

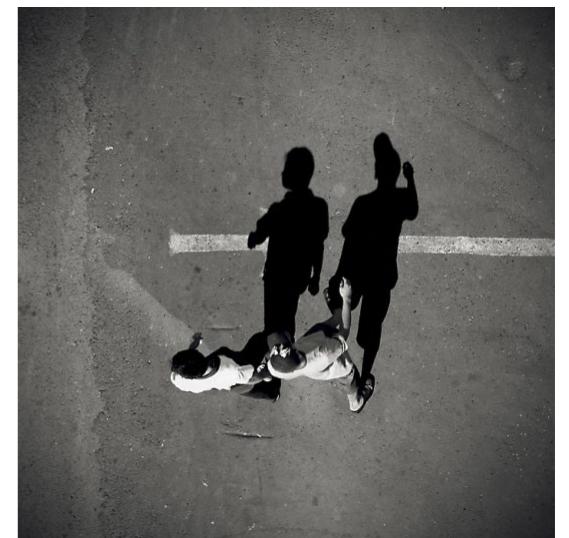
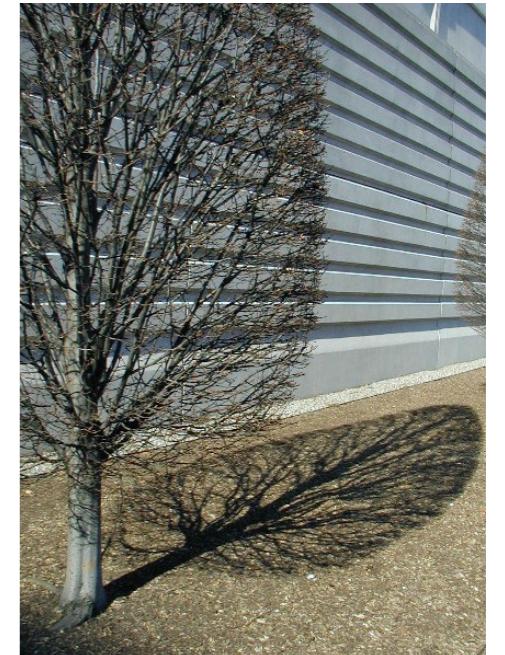
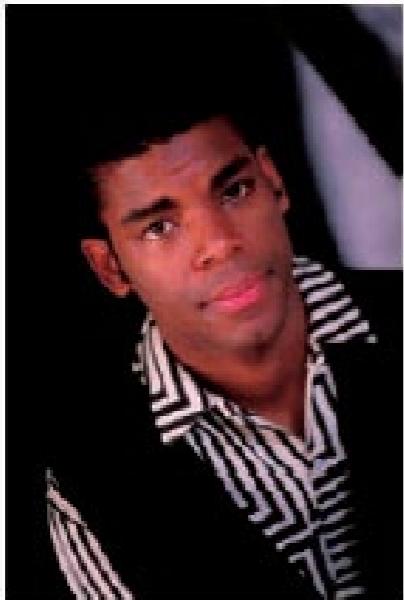
Effect of σ (Gaussian kernel spread/size)



The choice of σ depends on desired behavior

- large σ detects large scale edges
- small σ detects fine features

Object boundaries vs. edges

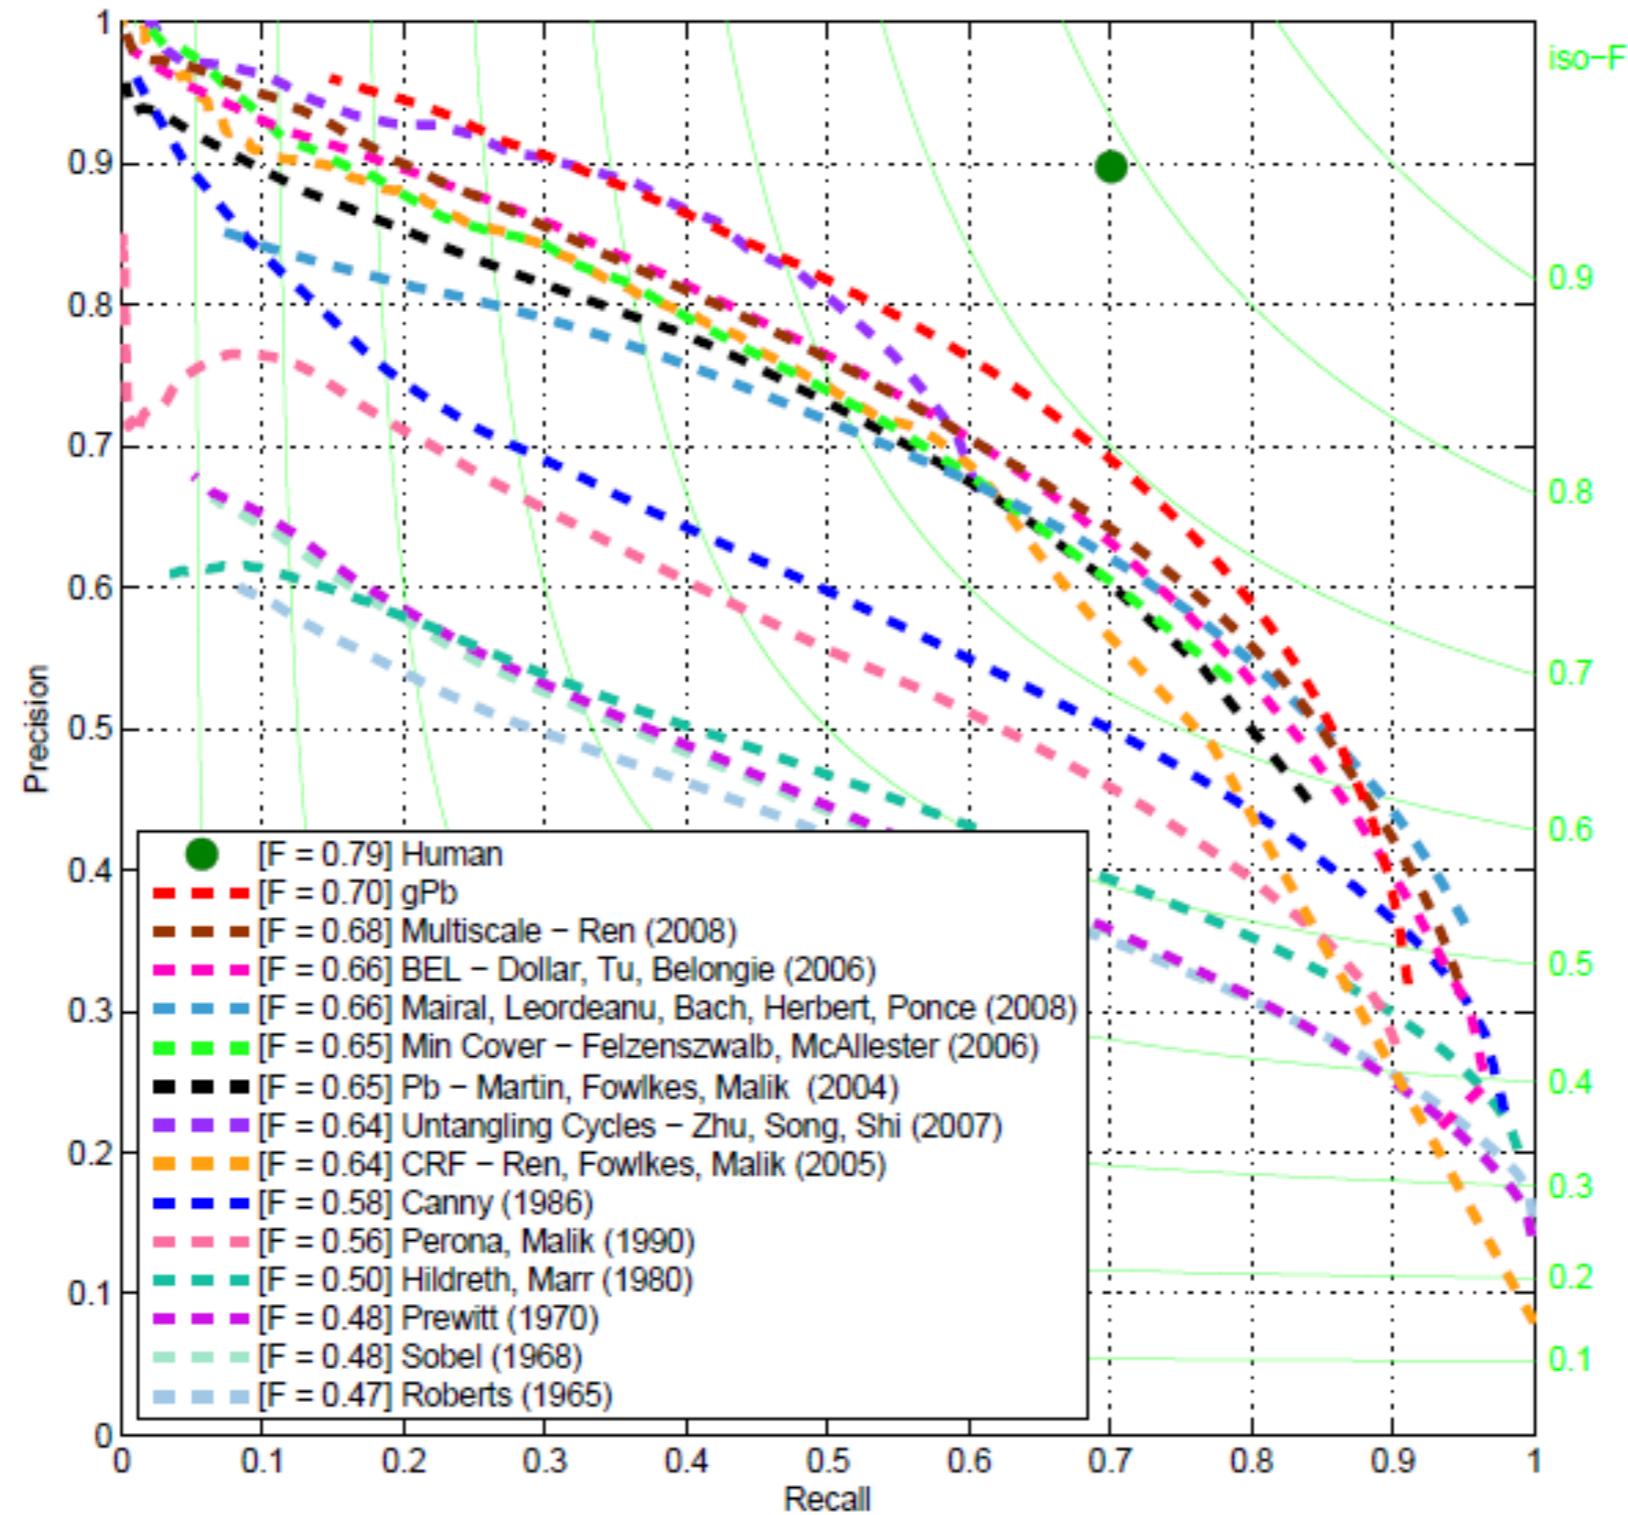


Background

Texture

Shadows

45 years of edge/boundary detection



There are limits...



What information is
missing for the
algorithm?



For more:
[http://www.eecs.berkeley.edu/Research/Projects/CS/
vision/bsds/bench/html/108082-color.html](http://www.eecs.berkeley.edu/Research/Projects/CS/vision/bsds/bench/html/108082-color.html)

Semantic Segmentation (by Facebook Research)

- Segmentation by Deep Networks (SharpMask Algorithm)



Summary on Edge Detection

- Filters allow local image neighborhood to influence our description and features
 - Smoothing to reduce noise
 - Derivatives to locate high contrast areas and gradients
- Filters have highest response on neighborhoods that “look like” it; can be thought of as template matching.
- Edge detection uses the image gradient to find curves, or chains of edges.