



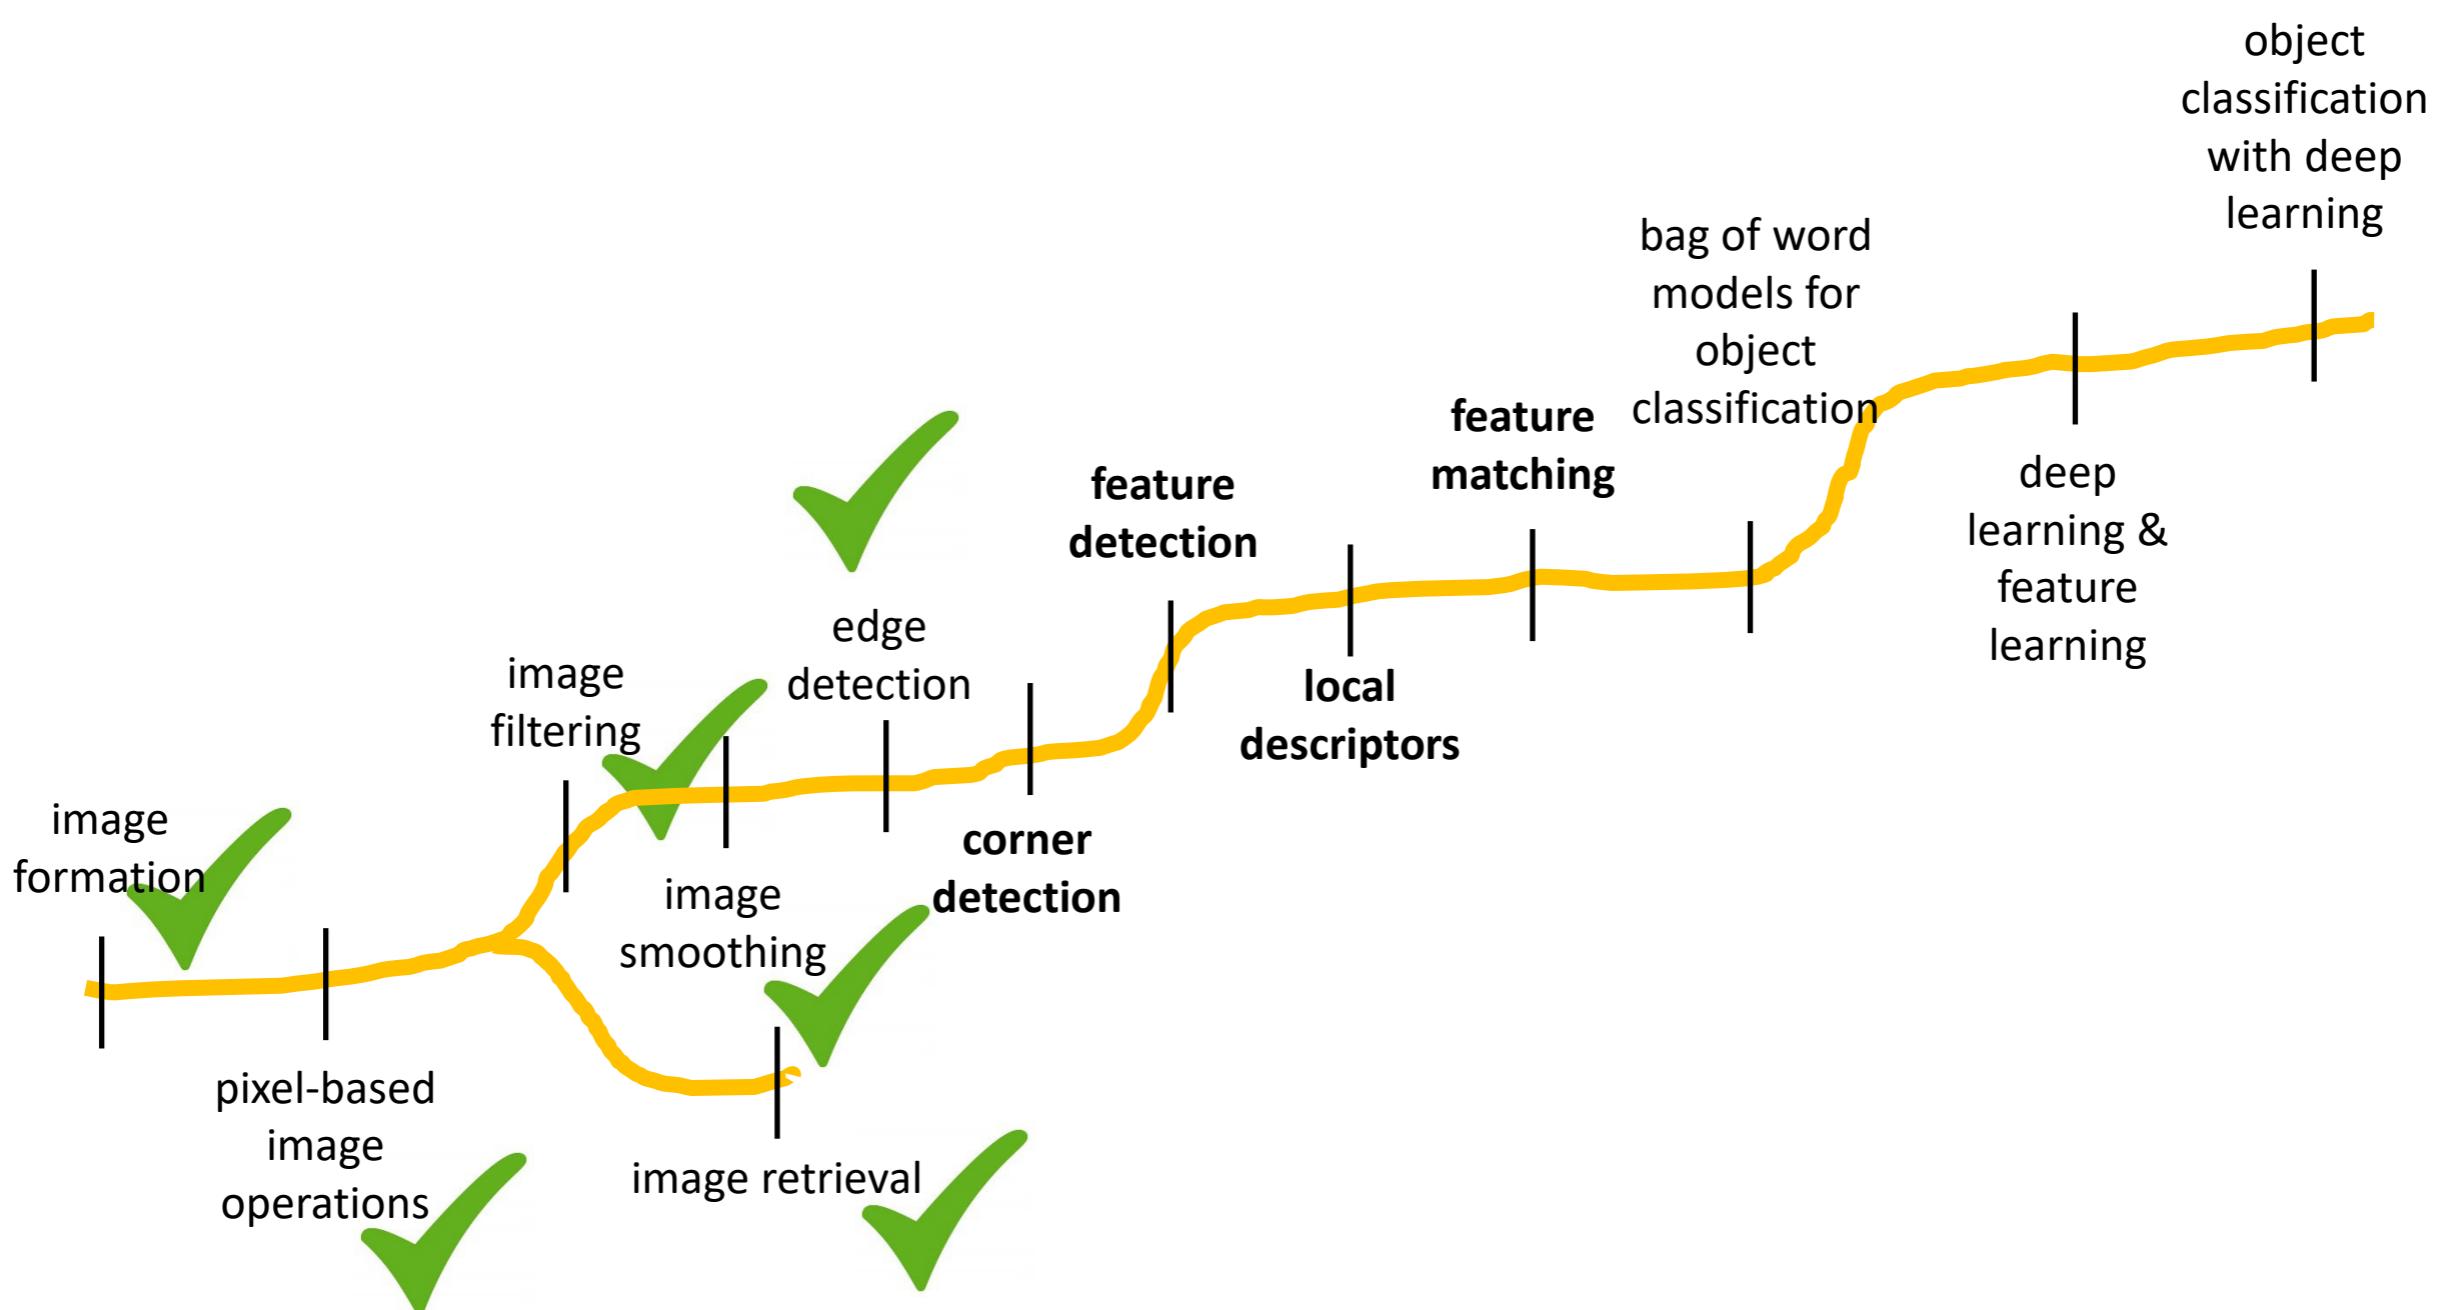
# ACV – Applied Computer Vision

Bachelor Medientechnik & Creative Computing

Matthias Zeppelzauer  
[matthias.zeppelzauer@fhstp.ac.at](mailto:matthias.zeppelzauer@fhstp.ac.at)

Djordje Slijepcevic  
[djordje.slijepcevic@fhstp.ac.at](mailto:djordje.slijepcevic@fhstp.ac.at)

# Roadmap

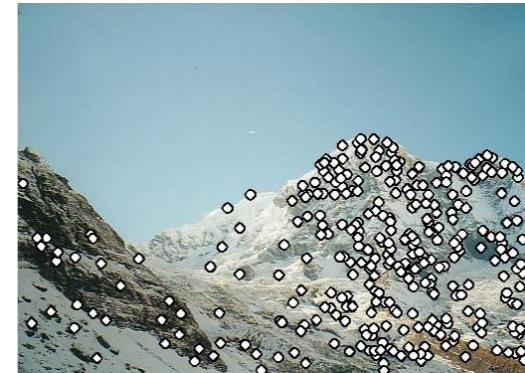


# Outlook

- So far:
  - Pixel-based operations on images
  - Filter (smoothing, sharpening)
  - Edge detection
- Today:
  - Corner detection
  - Local features and descriptors
  - Feature matching

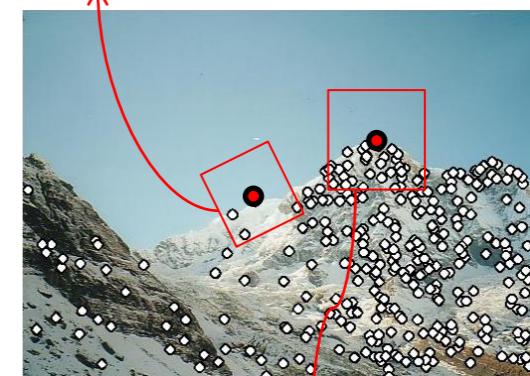
# Overview – The Big Picture

- **1. Detection:** identify the interest points



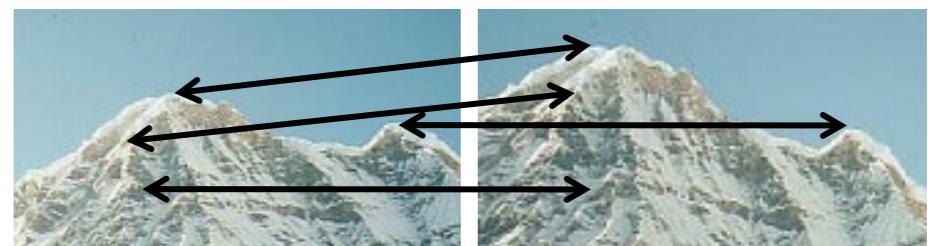
$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$

- **2. Description:** extract a vector (feature descriptor) that represents the surrounding of each feature point



$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$

- **3. Matching:** determine correspondence between descriptors in two views

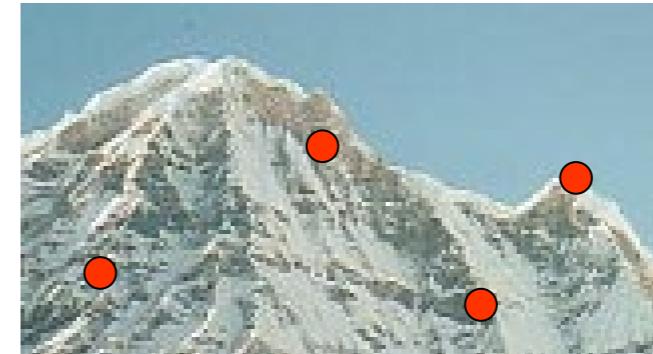


# Let's start with feature detection...

- Note: Local features are often also referred to as:
  - Keypoints
  - Feature points
  - Interest Points
  - Salient Points
  - or simply “features”

# Finding reliable local features

- Problem:
  - Detect the *same* points *independently* in both images



no chance to match!

We need a repeatable *detector*

# What Points would you choose?

- Where would you tell your friend to meet you?



# What Points would you choose?

- Where would you tell your friend to meet you?



Source: James Hays

# Properties of the ideal feature

- Local: features are local, so robust to occlusion and clutter (no prior segmentation necessary)
- Invariant to **scale, rotation, translation**, affine and perspective transforms
- Invariant to **brightness variation** in different images (different illuminations)
- Robust to **noise**, blur, discretization, compression etc.
- Quantity: many features can be generated for even small objects
- Accurate: precise **localization**
- Efficient: close to **real-time** performance

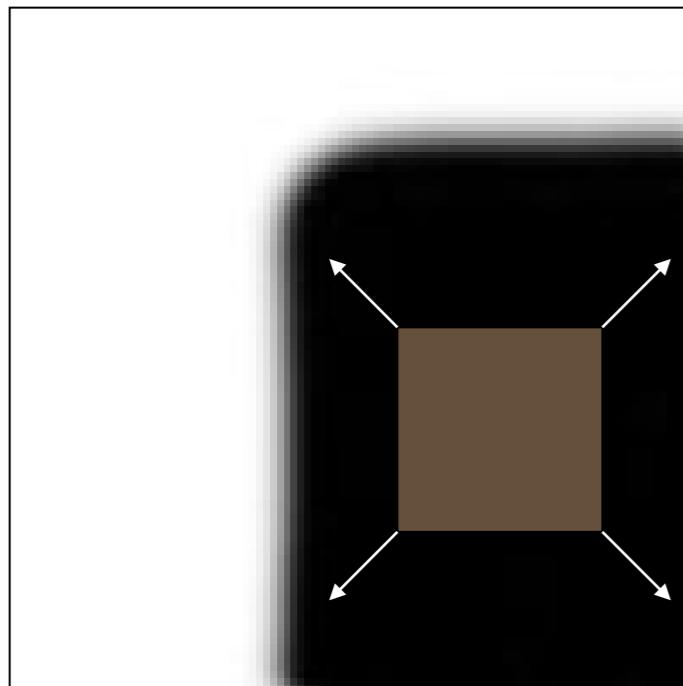
# Good Candidates for Local Features are *Corners*:

Are edges  
good  
feature  
locations?

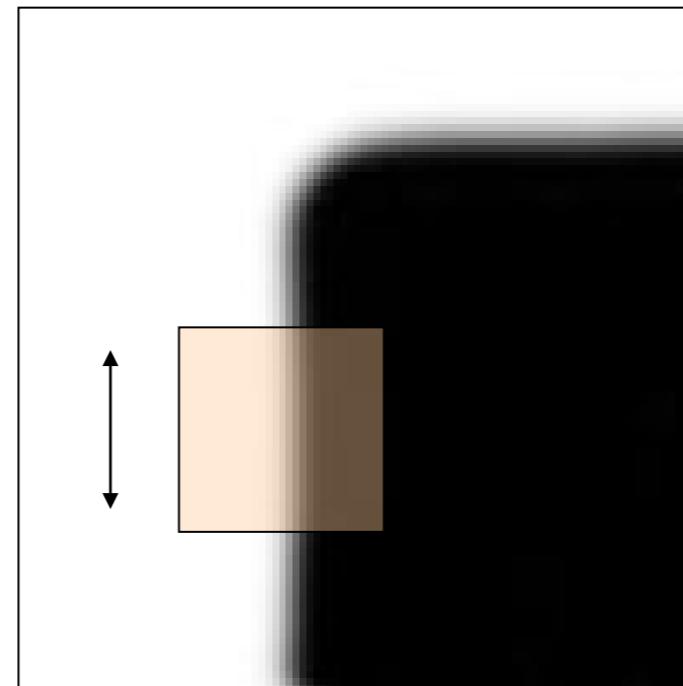


# Why Corners?

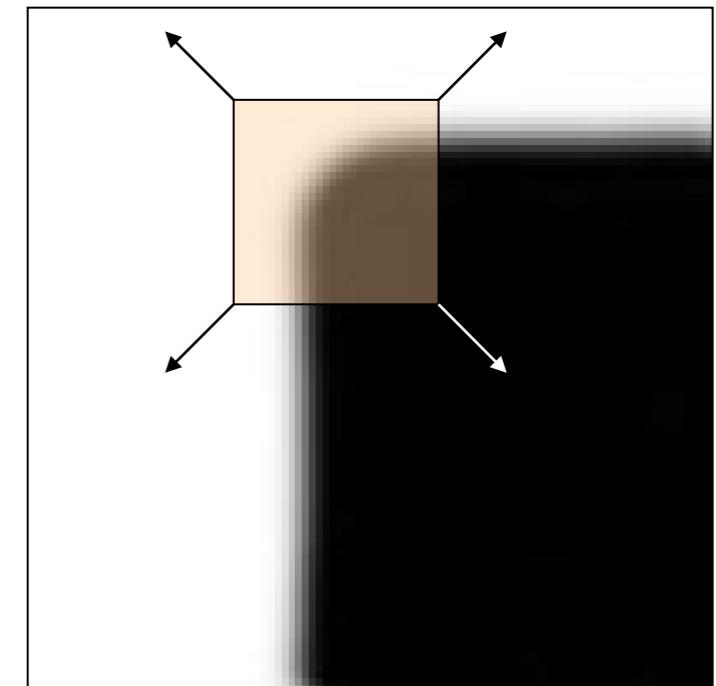
- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give a *large change* in intensity



“flat” region:  
no change in all  
directions



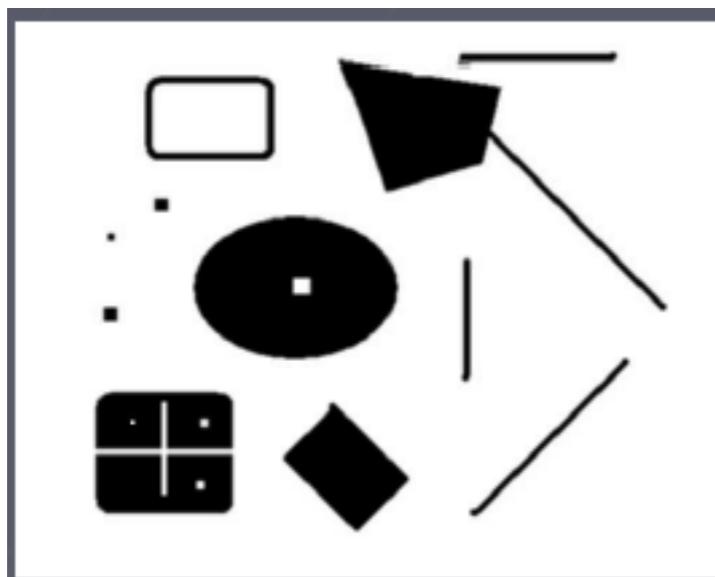
“edge”:  
no change along  
the edge direction



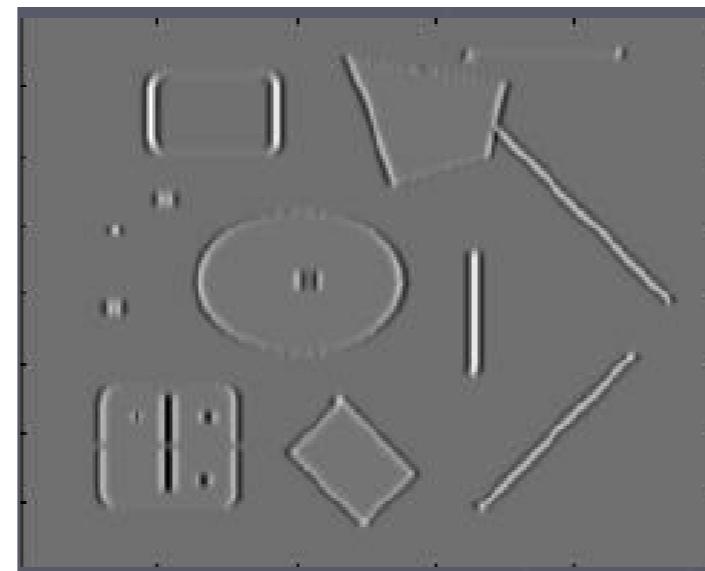
“corner”:  
significant change  
in all directions

# The Harris Corner Detector

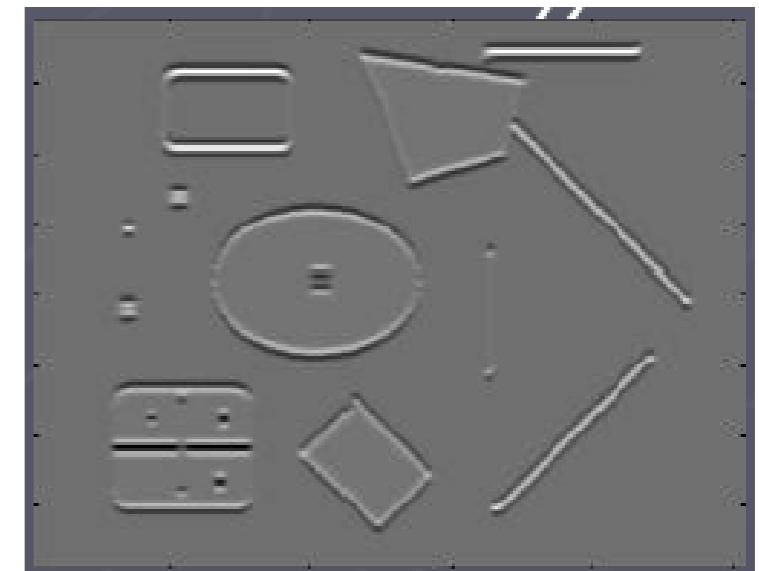
- A robust corner detector that is based on the **image gradient**
- Idea: find locations in the image where the gradient is strong in at least two directions
- So back to gradients and partial derivations:



Notation:



$$I_x \Leftrightarrow \frac{\partial I}{\partial x}$$



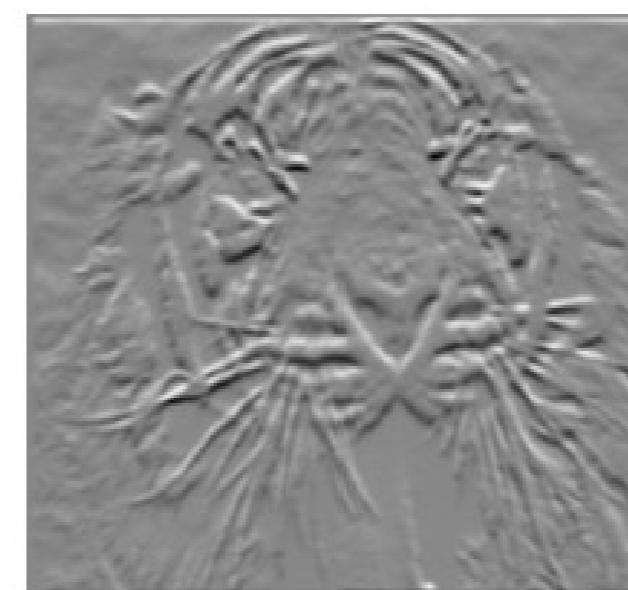
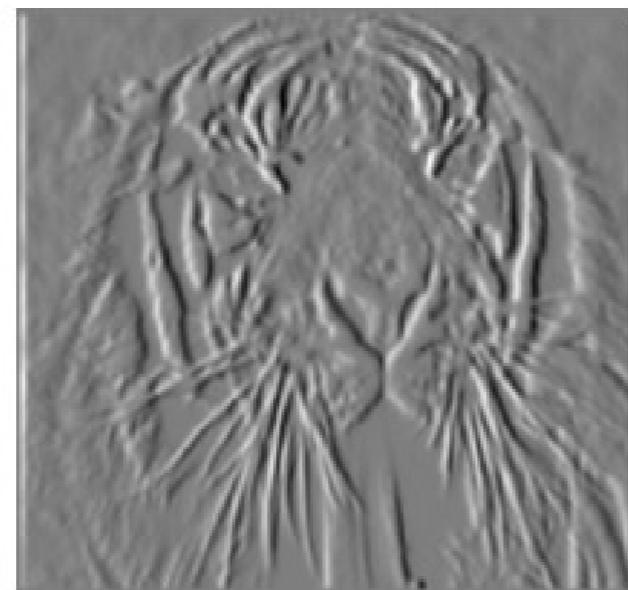
$$I_y \Leftrightarrow \frac{\partial I}{\partial y}$$

C.Harris and M.Stephens. "[A Combined Corner and Edge Detector.](#)" *Proceedings of the 4th Alvey Vision Conference*: pages 147--151.

# Recap: how to compute the derivation in X- or Y-direction?

$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

-1	1
----	---

# Recap: how to compute the derivation in X- or Y-direction?

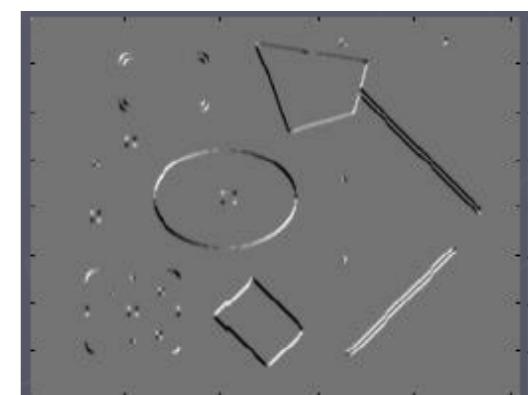
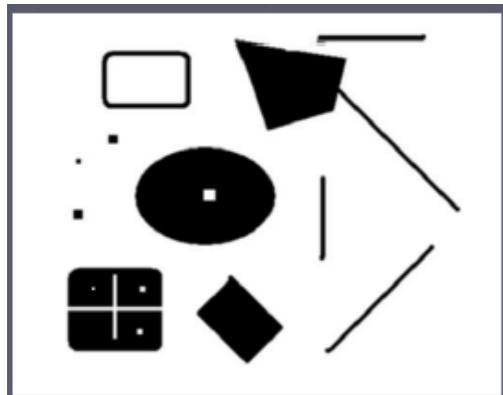
- The gradient map is then simply the product of horizontal derivative and vertical derivative of the image:

$$\left| \frac{\partial f(x, y)}{\partial x} * \frac{\partial f(x, y)}{\partial y} \right|$$



- That's all we need to detect Harris corners...

# Harris Corner Detection



Notation:

$$I_x \Leftrightarrow \frac{\partial I}{\partial x}$$

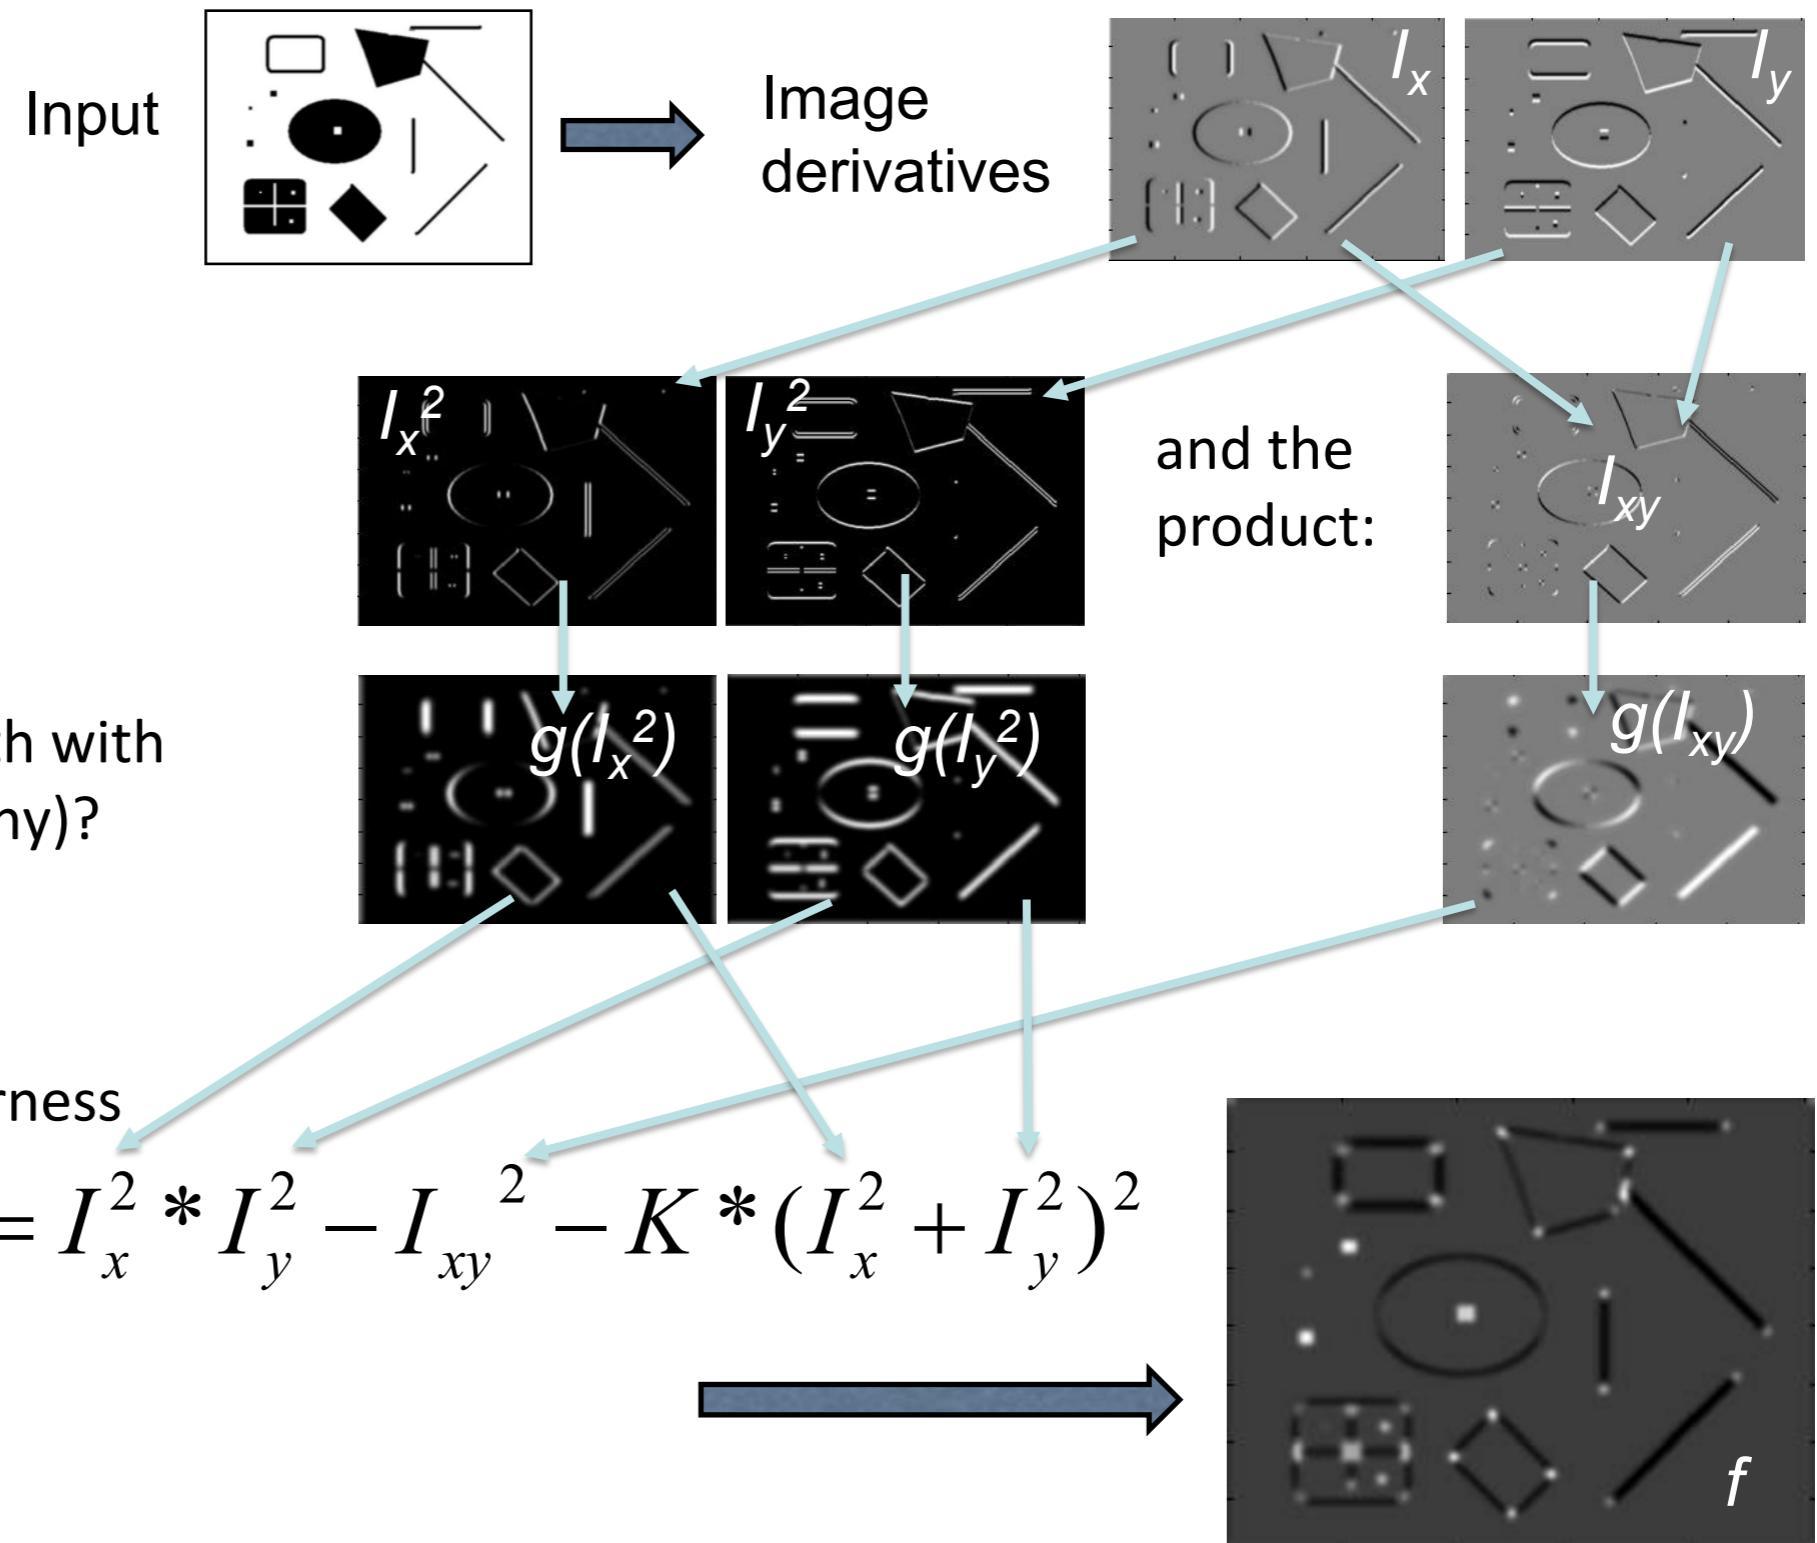
$$I_y \Leftrightarrow \frac{\partial I}{\partial y}$$

$$I_x I_y \Leftrightarrow \frac{\partial I}{\partial x} \frac{\partial I}{\partial y}$$

# Harris Corner Detection

- Computation:

- 1. Compute image derivatives
- 2. Compute square of derivatives
- 3. Optionally smooth with Gaussian filter  $g$  (why)?
- 4. Compute „cornerness“ function



## Harris corner detector

- 1) Compute *cornerness* score  $R$  (=an image)



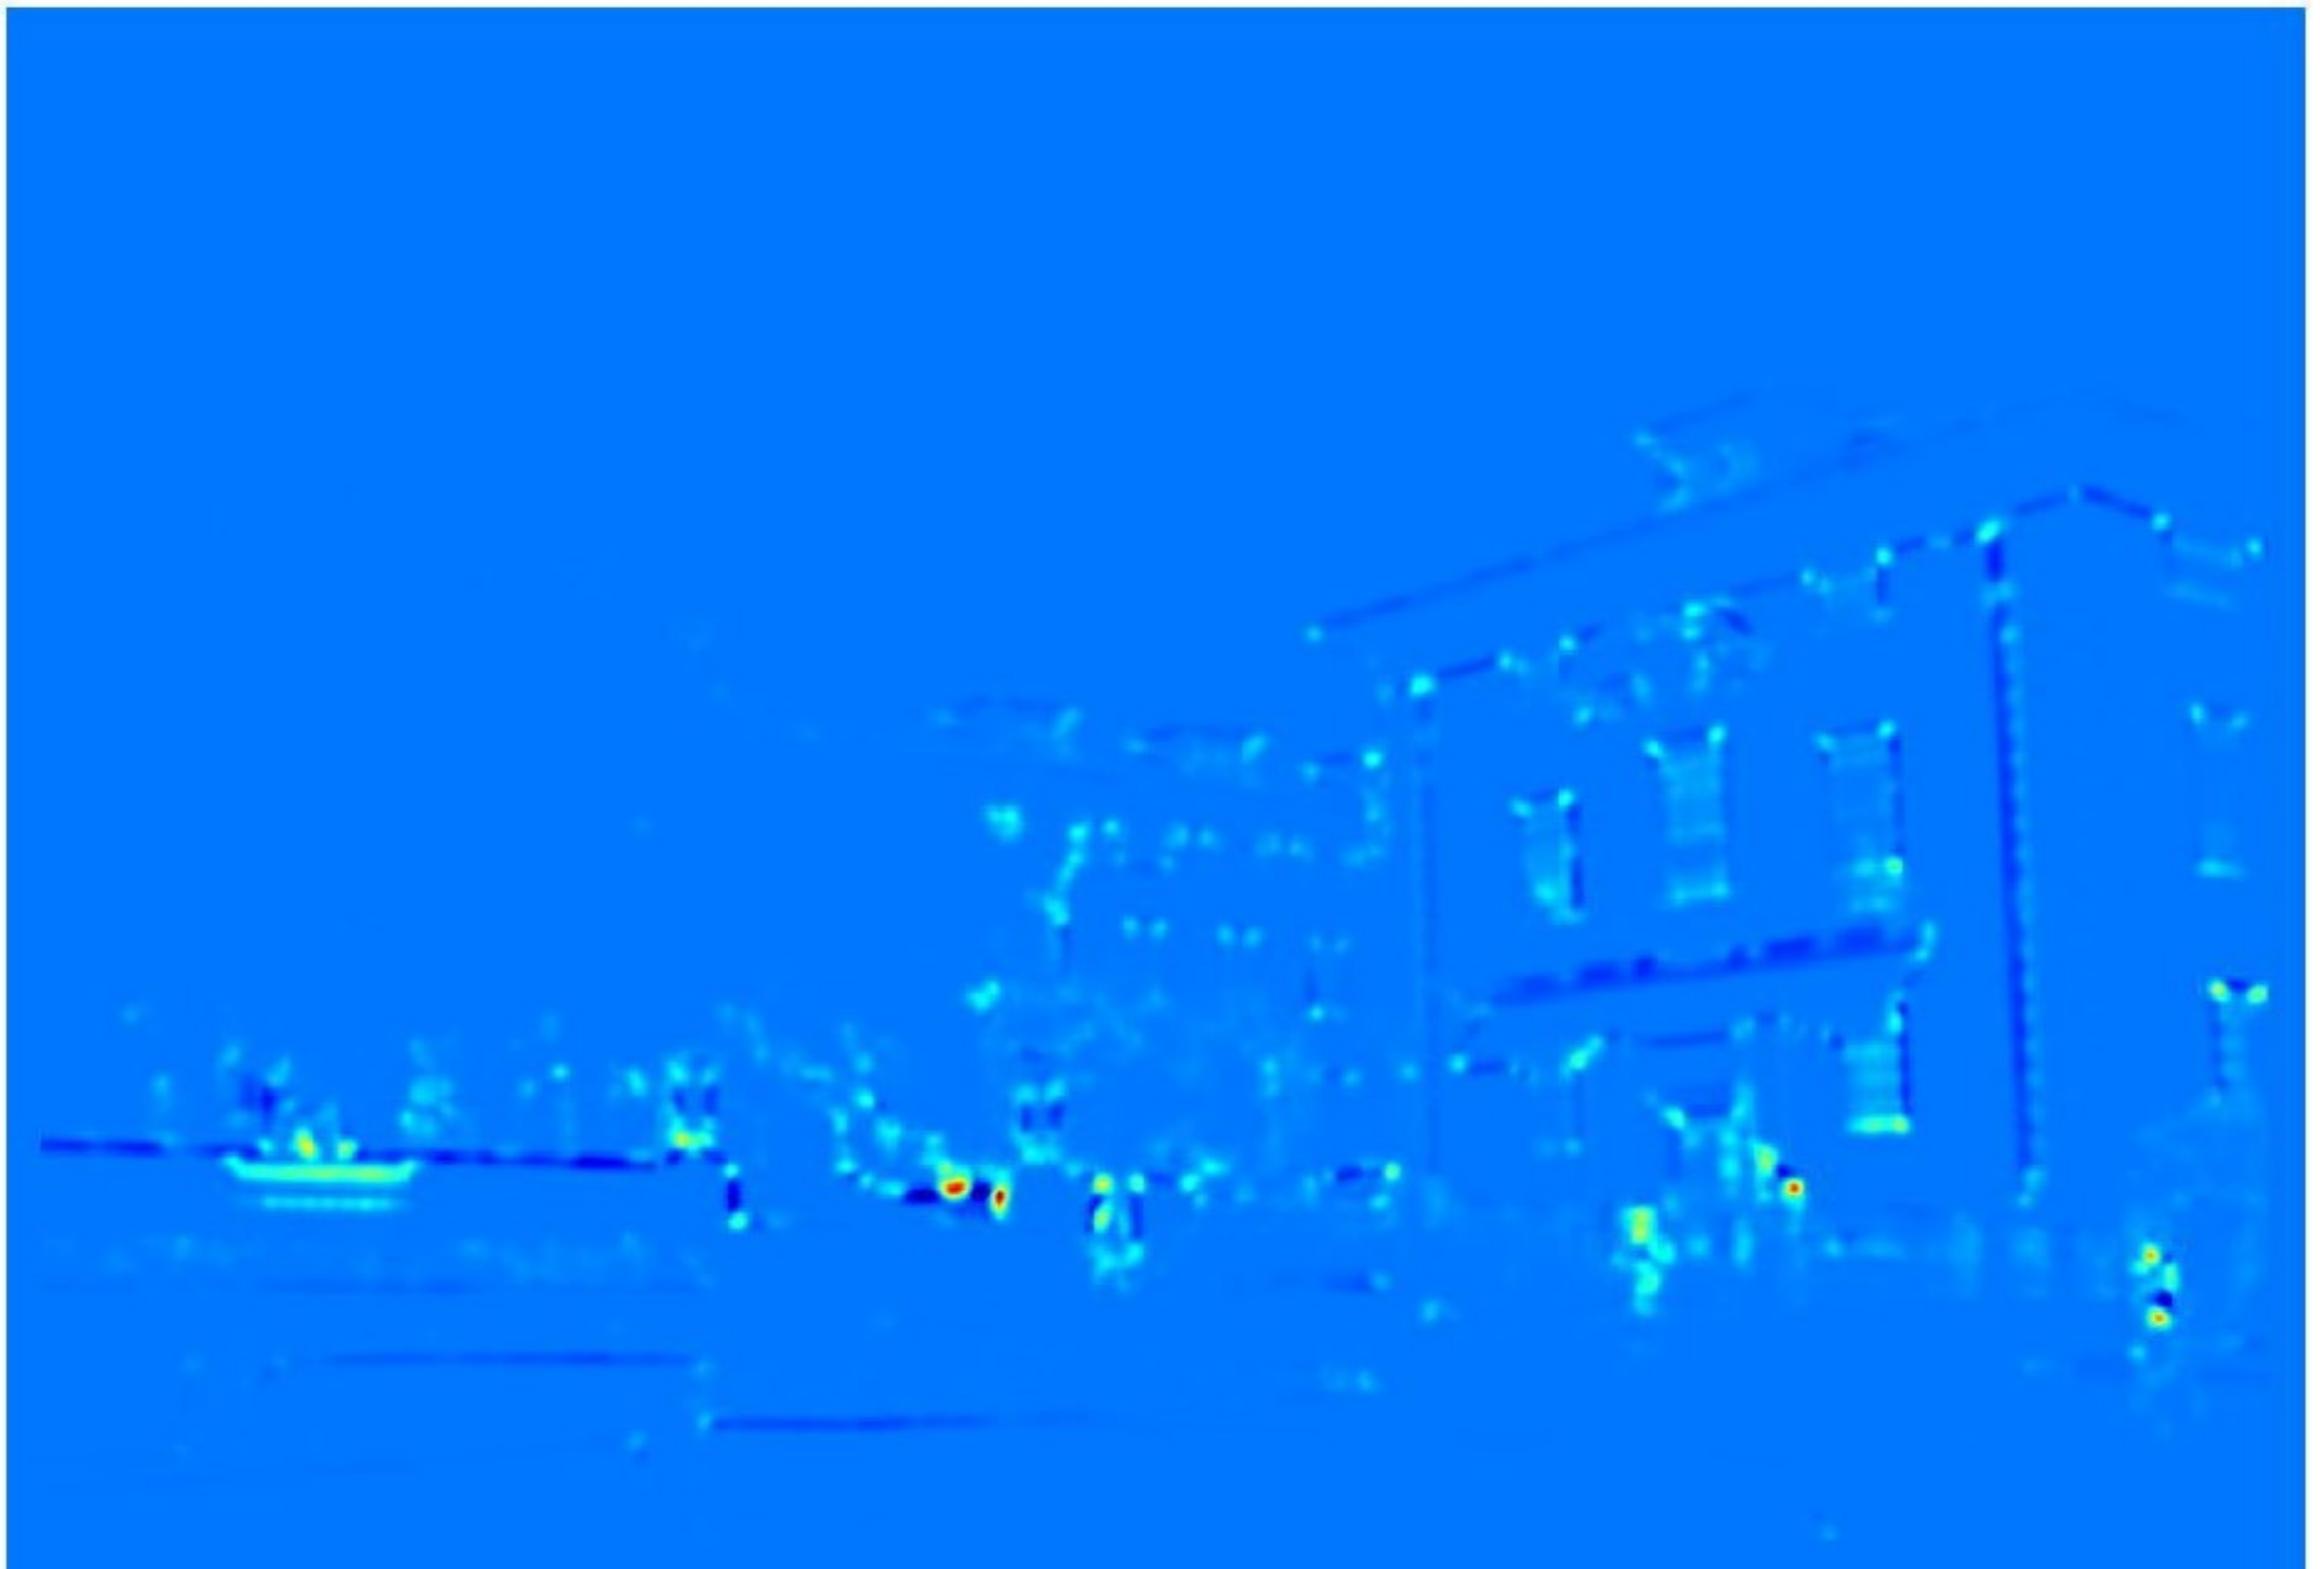
- 2) Find peaks in  $R$ , i.e. points which are surrounded by smaller values and take the points of local maxima as corners
- 3) Or find the  $N$  strongest points by selecting all points  $p$  by sorting all intensities and selecting the first  $N$  entries.

# Example of Harris application



# Example of Harris application

- Compute corner response at every pixel.



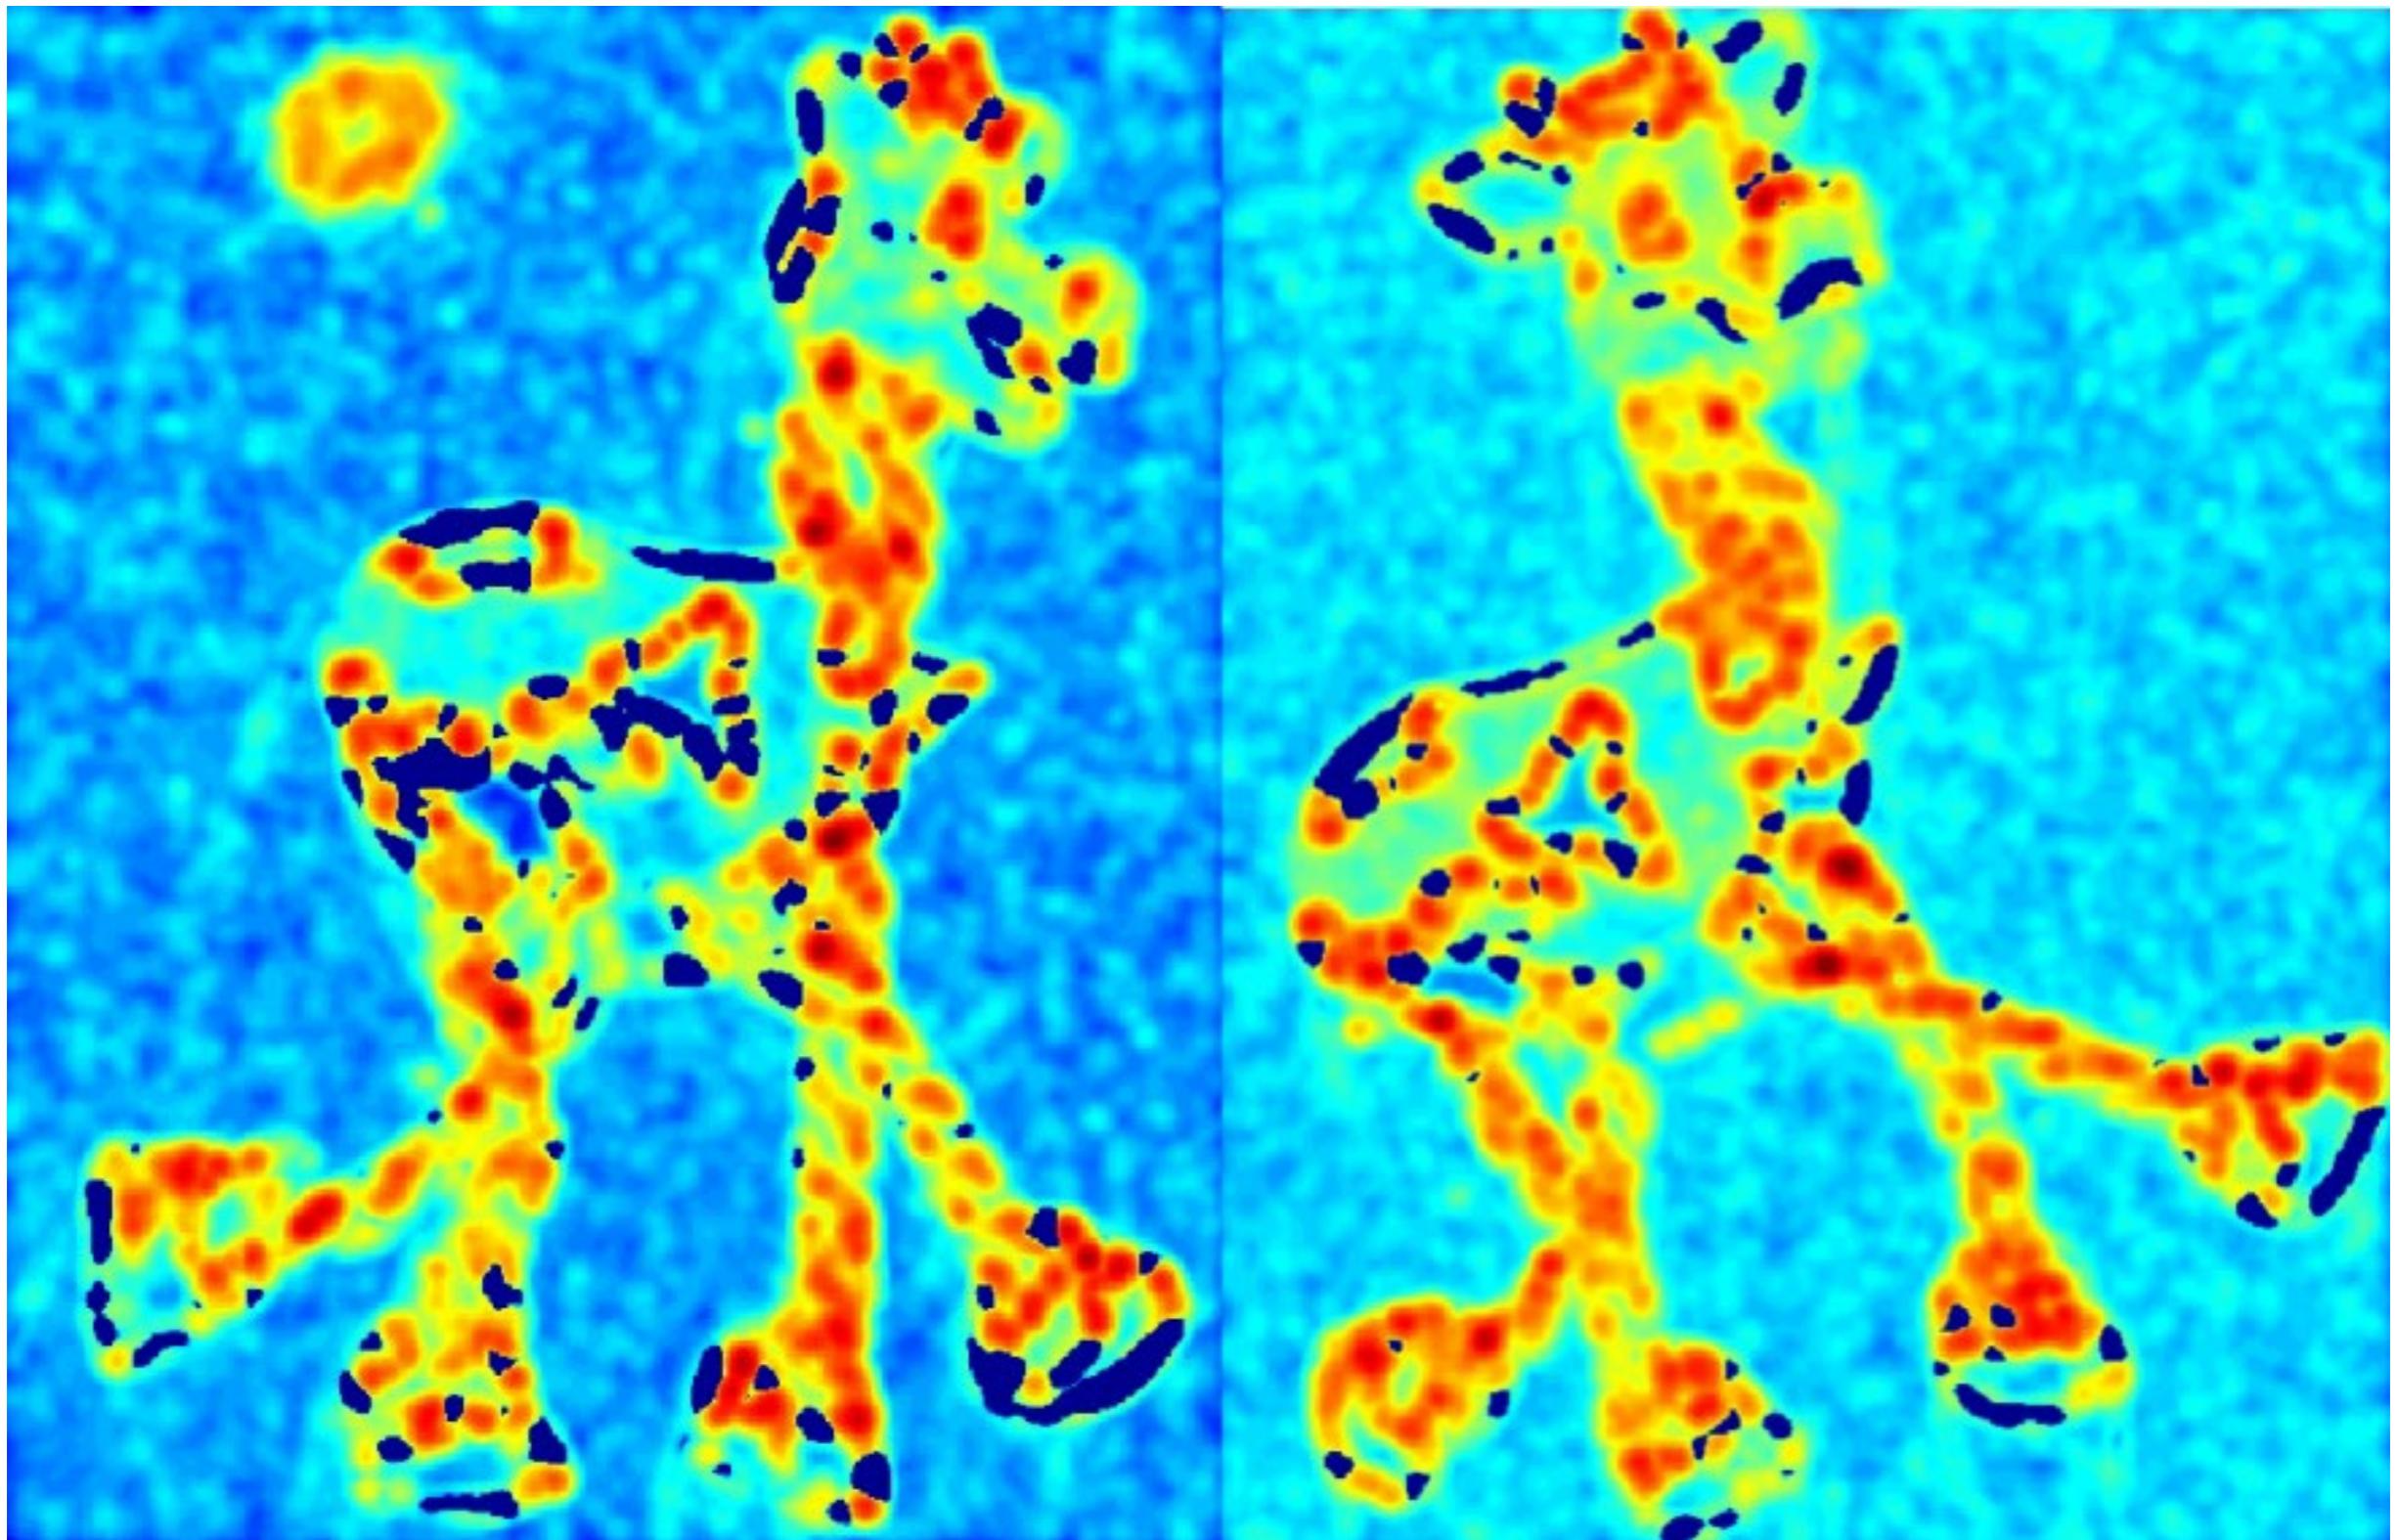
# Example of Harris application



# Harris detector example – robustness?



## Cornerness function R (red high, blue low)



## Threshold ( $f > \text{value}$ )



# Search for maxima in Cornerness Function



# Find local maxima of f

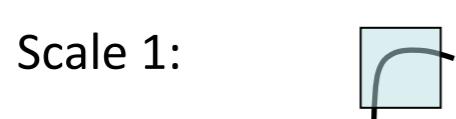


## Harris features (in red)



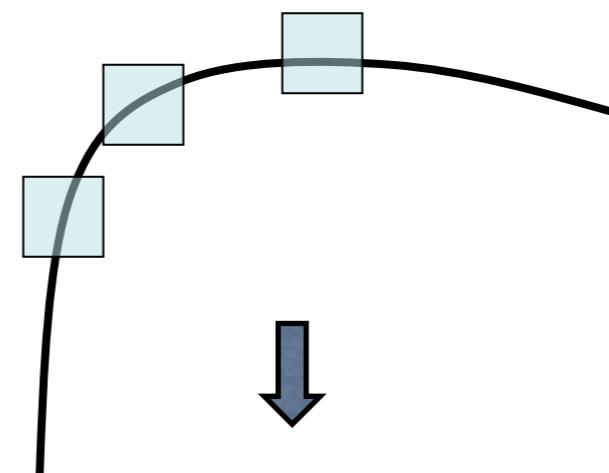
# Harris Corner Detector - Summary

- Harris corner detector is a frequently and popular detector
- Easy to compute (from image derivatives only)
- Harris corners are rotation invariant and translation invariant
- Harris corner detector is, however, **not scale invariant**:



Corner !

Scale 2:



The detector sees only  
edges but no corners!

# Corners vs. Local Features

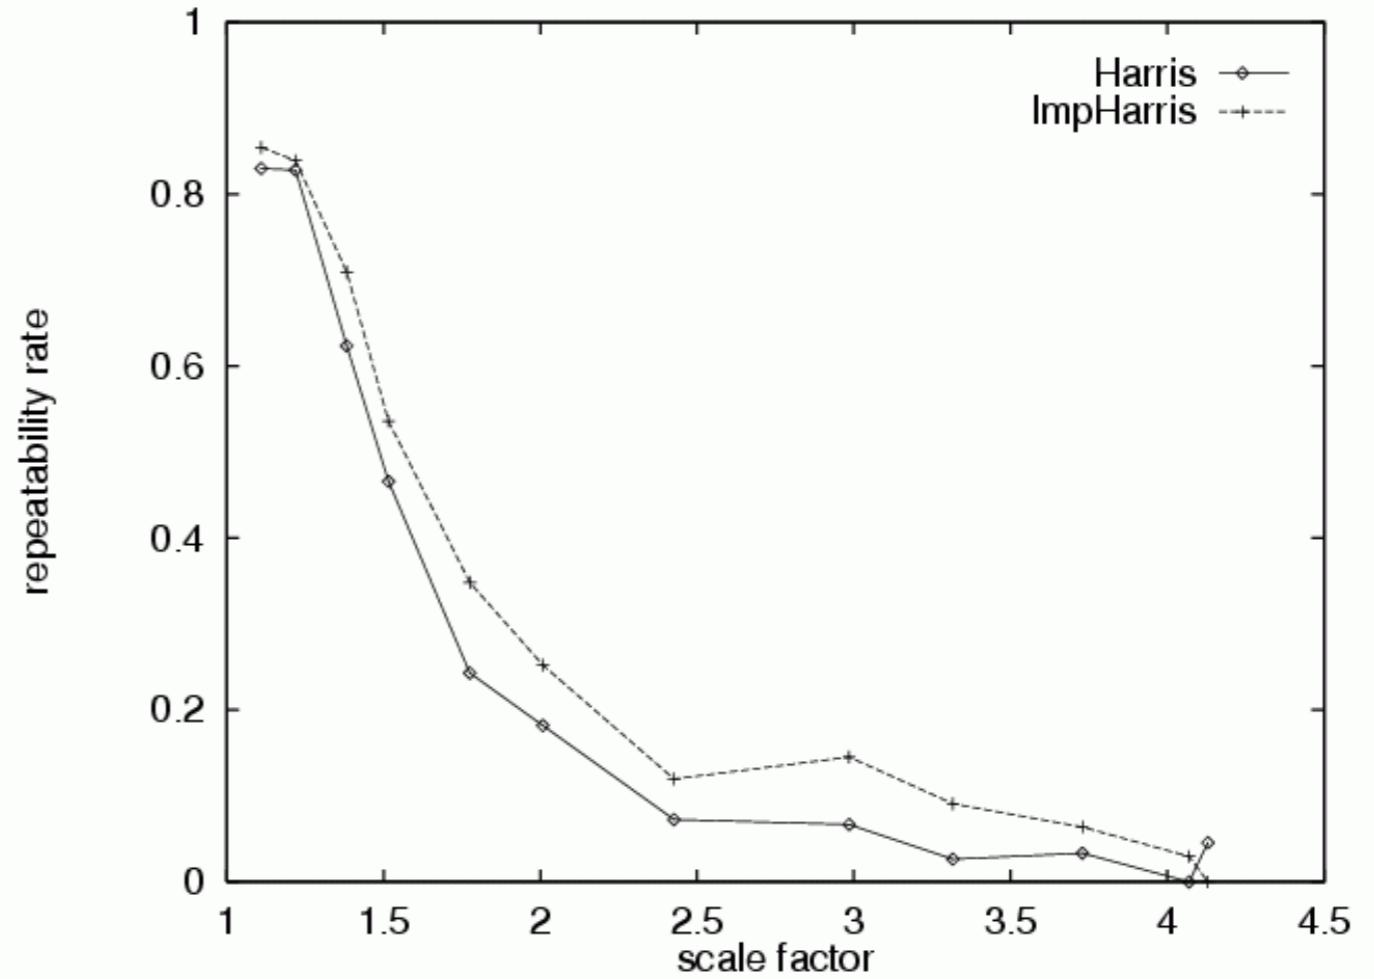
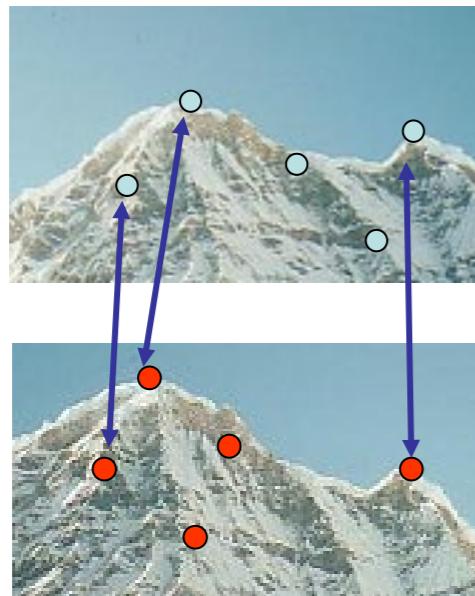
- A “local feature” is more than a “corner”
- Corner: has an x and y coordinate, i.e. is a point in the image
- Local feature
  - has also an x and y coordinate, but also has:
  - **scale** and
  - **orientation**
- Scale: the size of the structure that generates the corner (e.g. a corner in a window has a smaller scale than a corner produced by the roof of a building)
- Orientation: imagine a corner related to the rim (Felge) of a car wheel. The same corner in different images of the wheel might be differently rotated. Orientation helps to compensate this rotation.

# Harris Detector: Some Properties

- Quality of Harris detector for different scale changes

Repeatability rate:

$$\frac{\# \text{ correspondences}}{\# \text{ possible correspondences}}$$



# How to get a scale invariant detector?

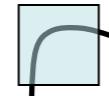
- To match (or stitch) these images we need a feature detector whose detections are *repeatable* across different scales



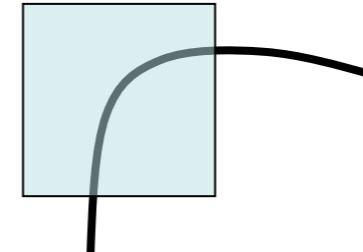
# How to get a scale invariant detector?

- Apply corner detection with different window sizes

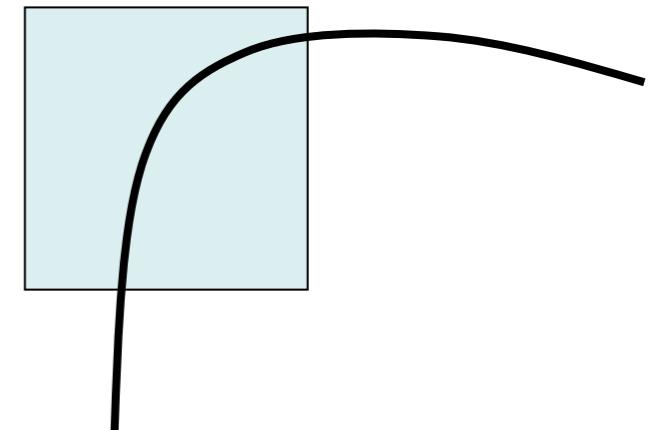
Scale 1:



Scale 2:



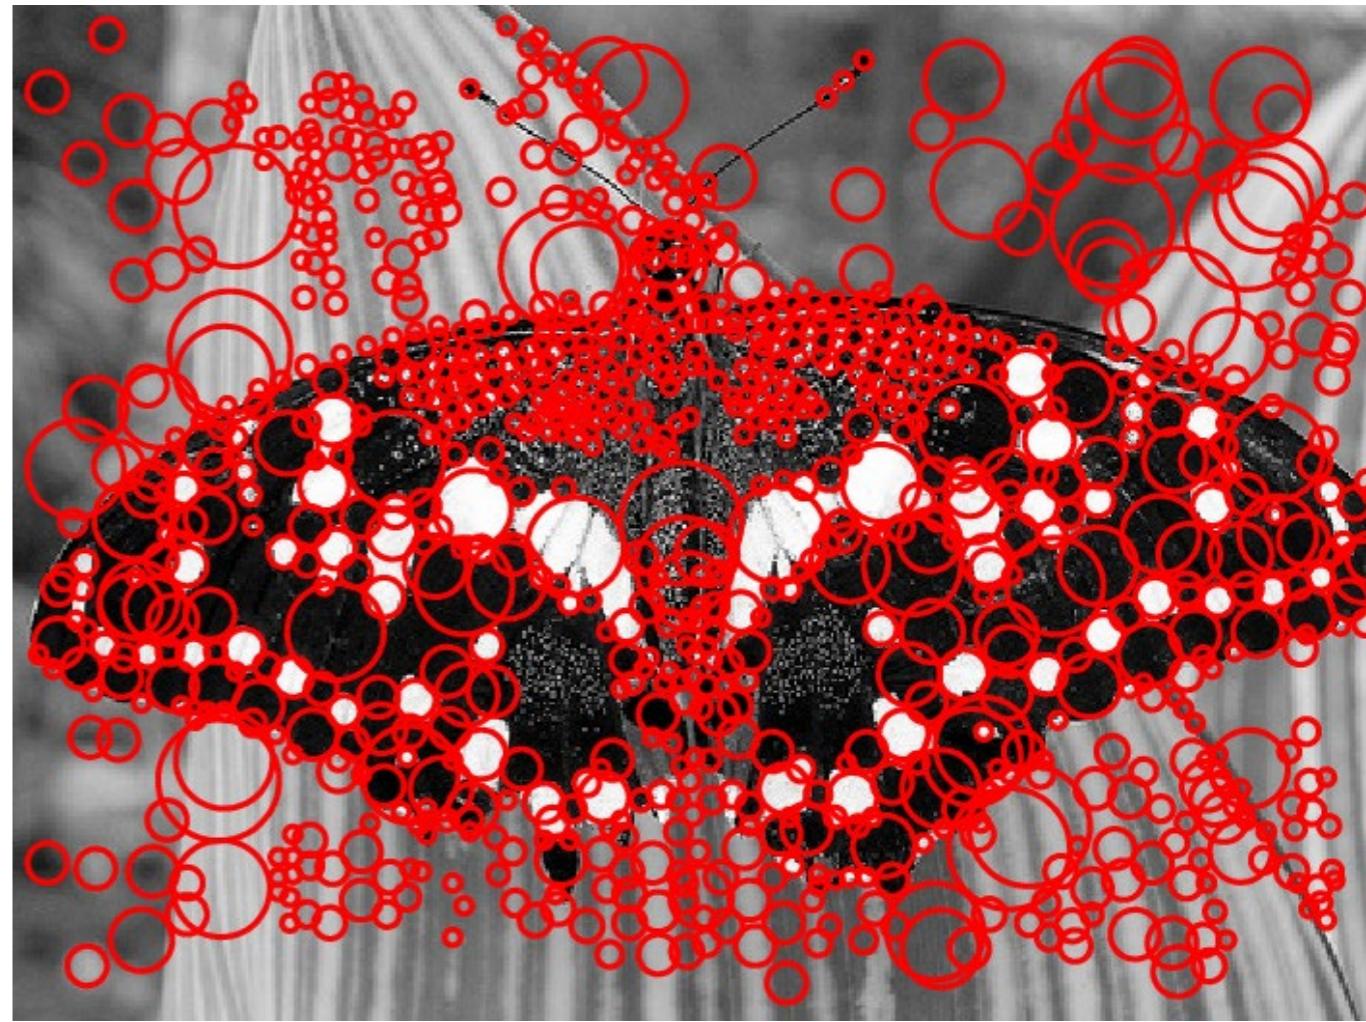
Scale 3:



- Apply same corner detection to different sizes of images! → „image pyramid“



# Multi-scale corner detection

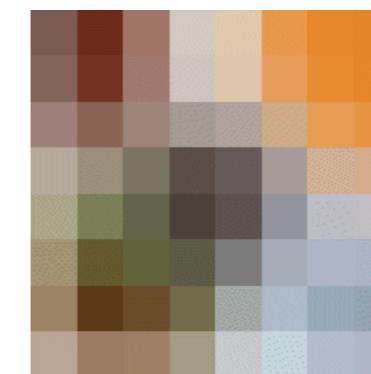
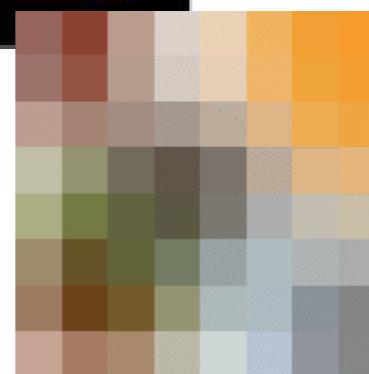
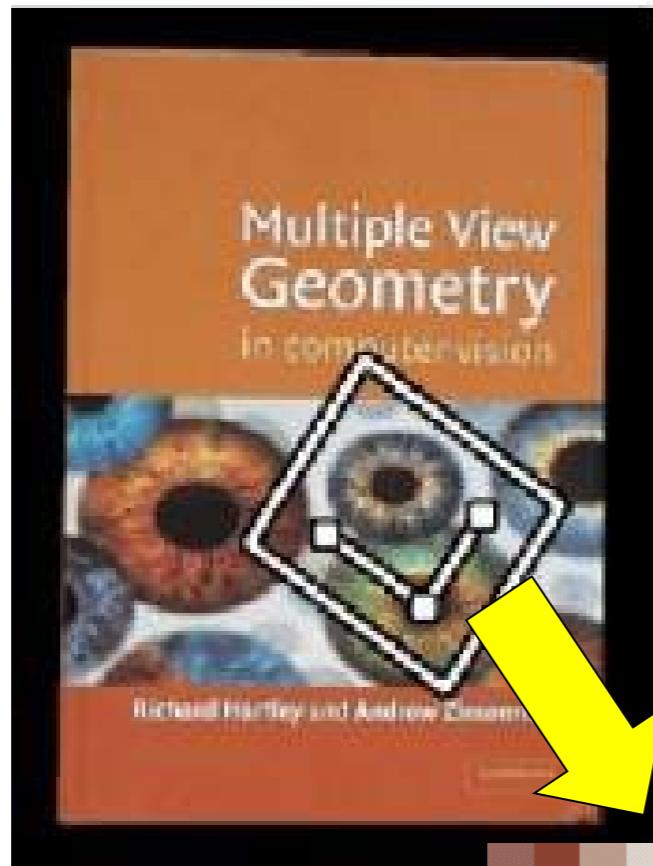


# Orientation Normalization

- Now every feature point has a **scale** (a “size”)
- But what about the **orientation** of the point?

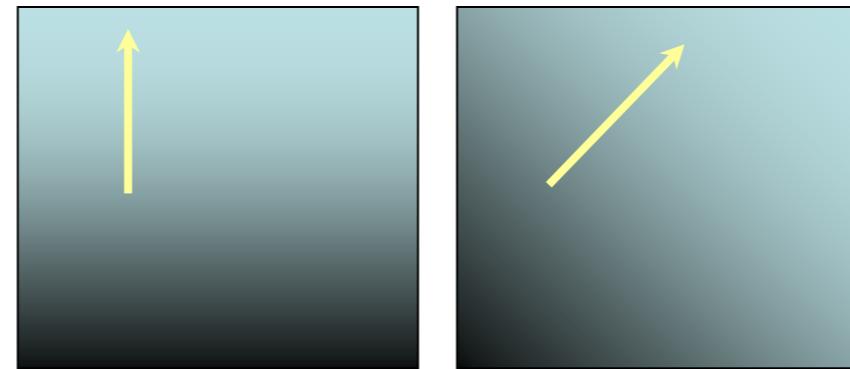
# Orientation Normalization

- Now every feature point has a **scale** (a “size”)
- But what about the **orientation** of the point?



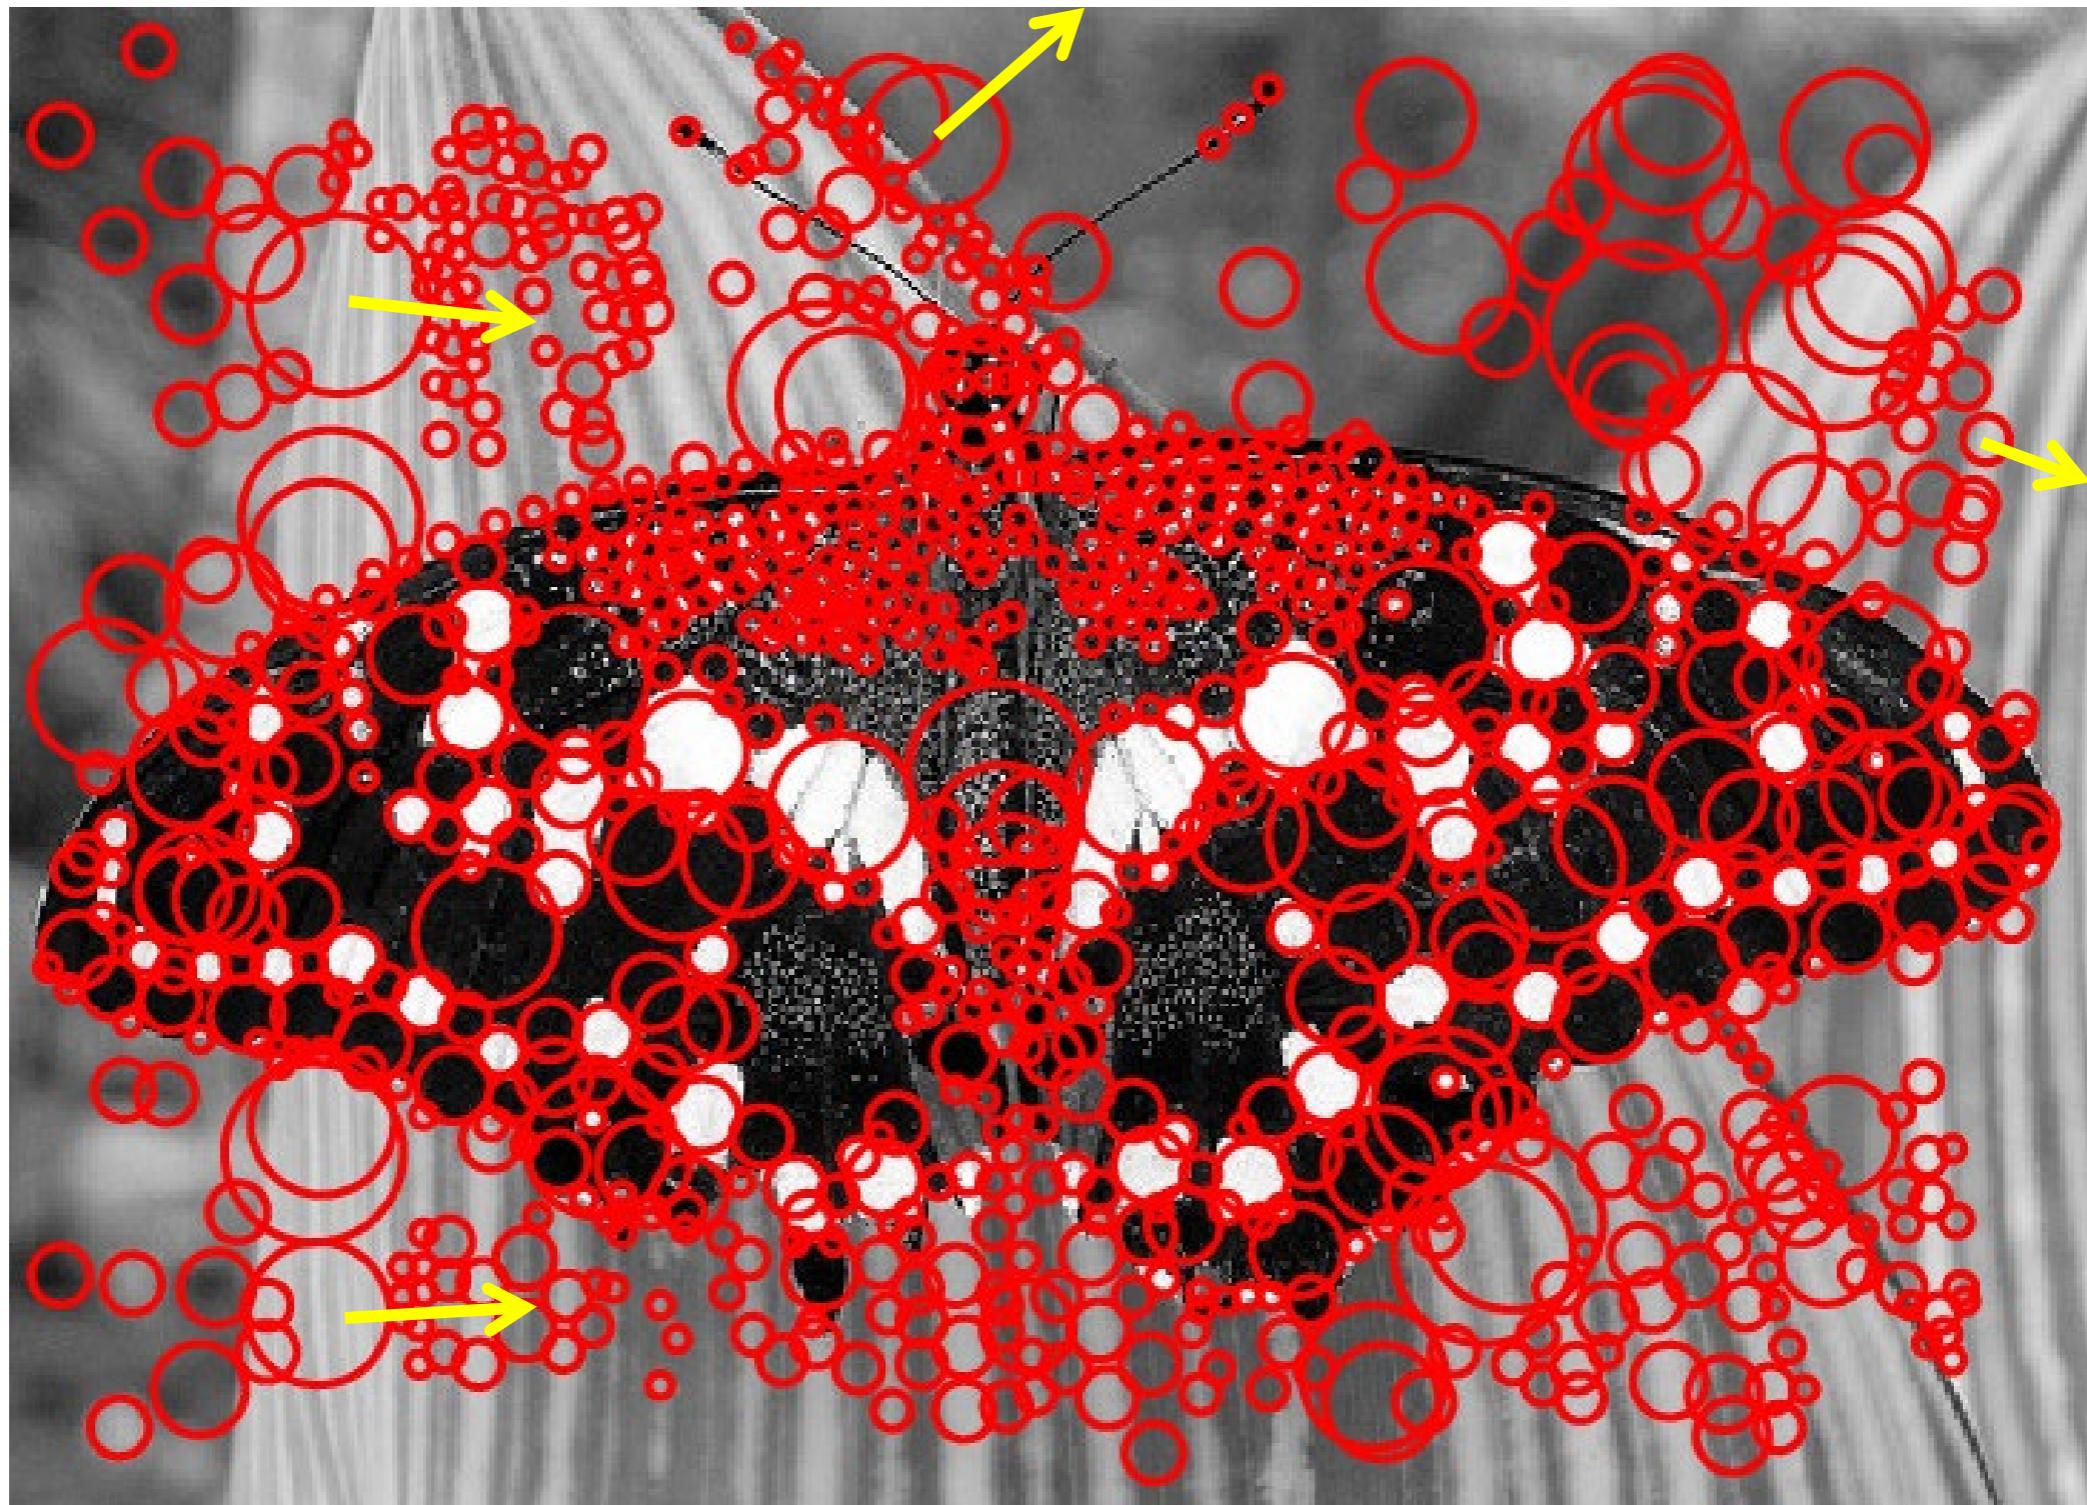
# Find local orientation of feature point:

- Use the dominant direction of the gradient for the entire image patch



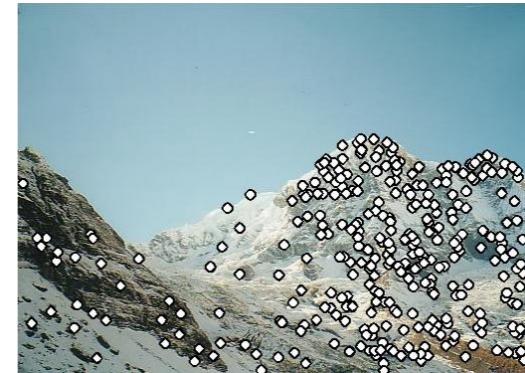
- Rotate feature point according to this angle
- Now all feature points can be compared in a rotation normalized way. In other words: differently rotated local features can be compared directly

# Local features: Each corner gets a scale and an orientation



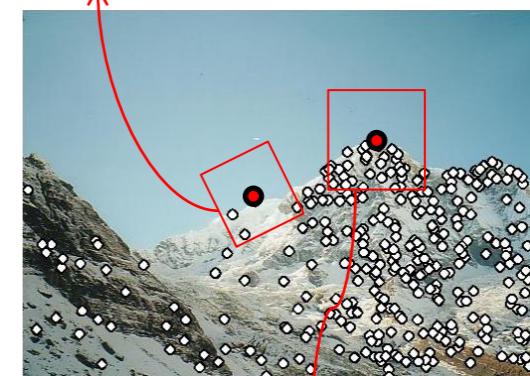
# Overview – The Big Picture

- **1. Detection:** identify the interest points



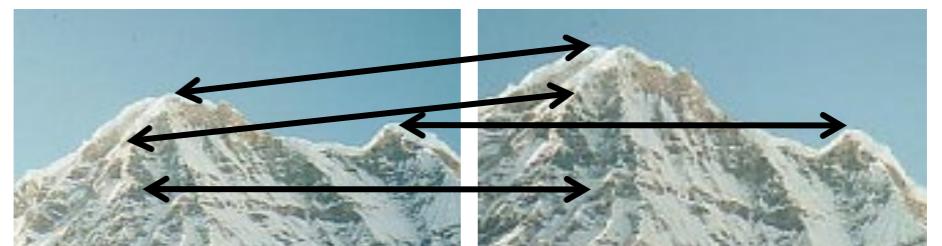
$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$

- **2. Description:** extract a vector (feature descriptor) that represents the surrounding of each feature point



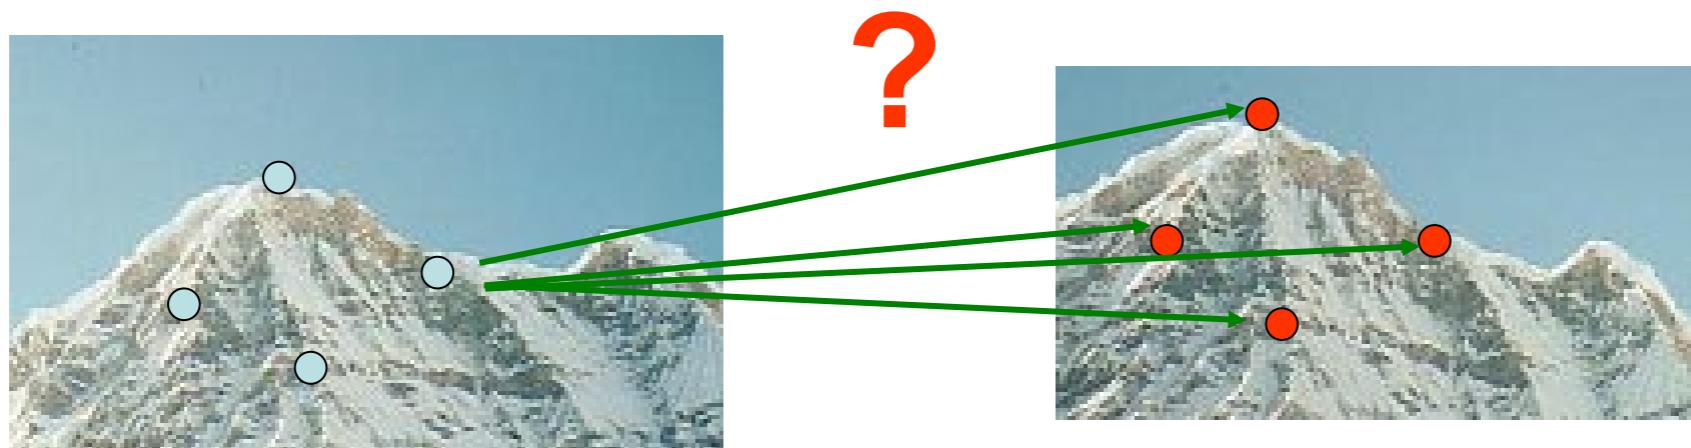
$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$

- **3. Matching:** determine correspondence between descriptors in two views



# Major Challenge:

- For each point correctly recognize the corresponding one



We need a reliable and distinctive *descriptor*

# Example

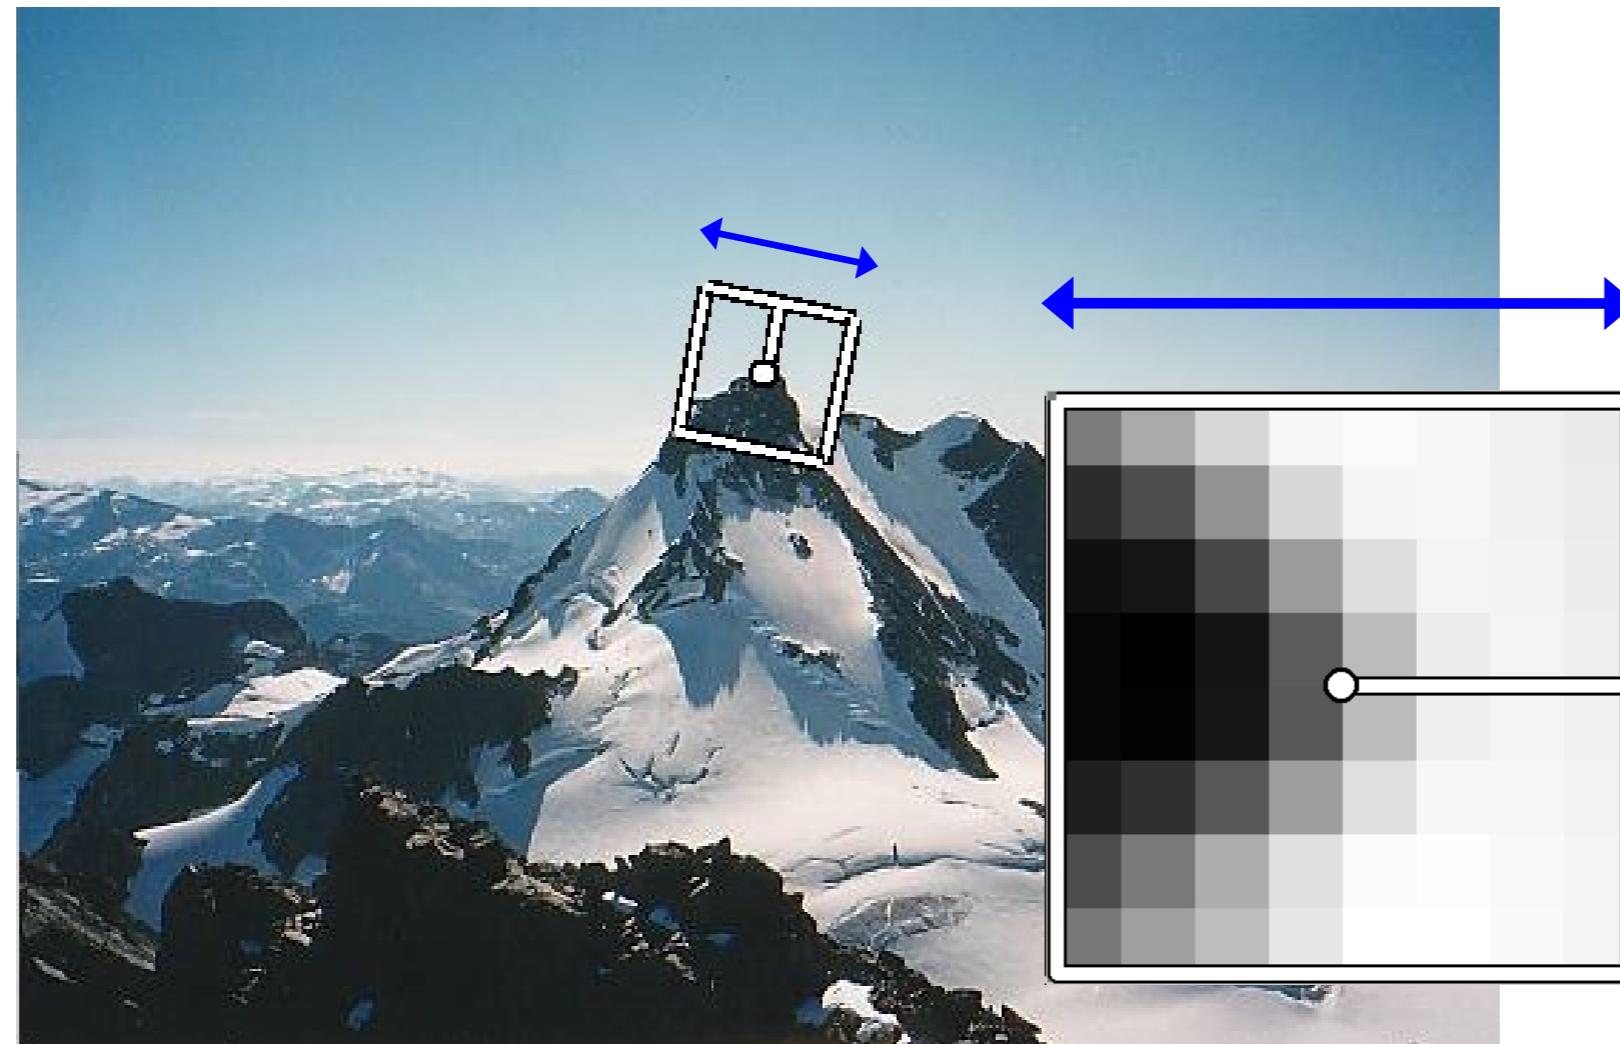


Image from Matthew Brown

# Detections at multiple scales and with orientation

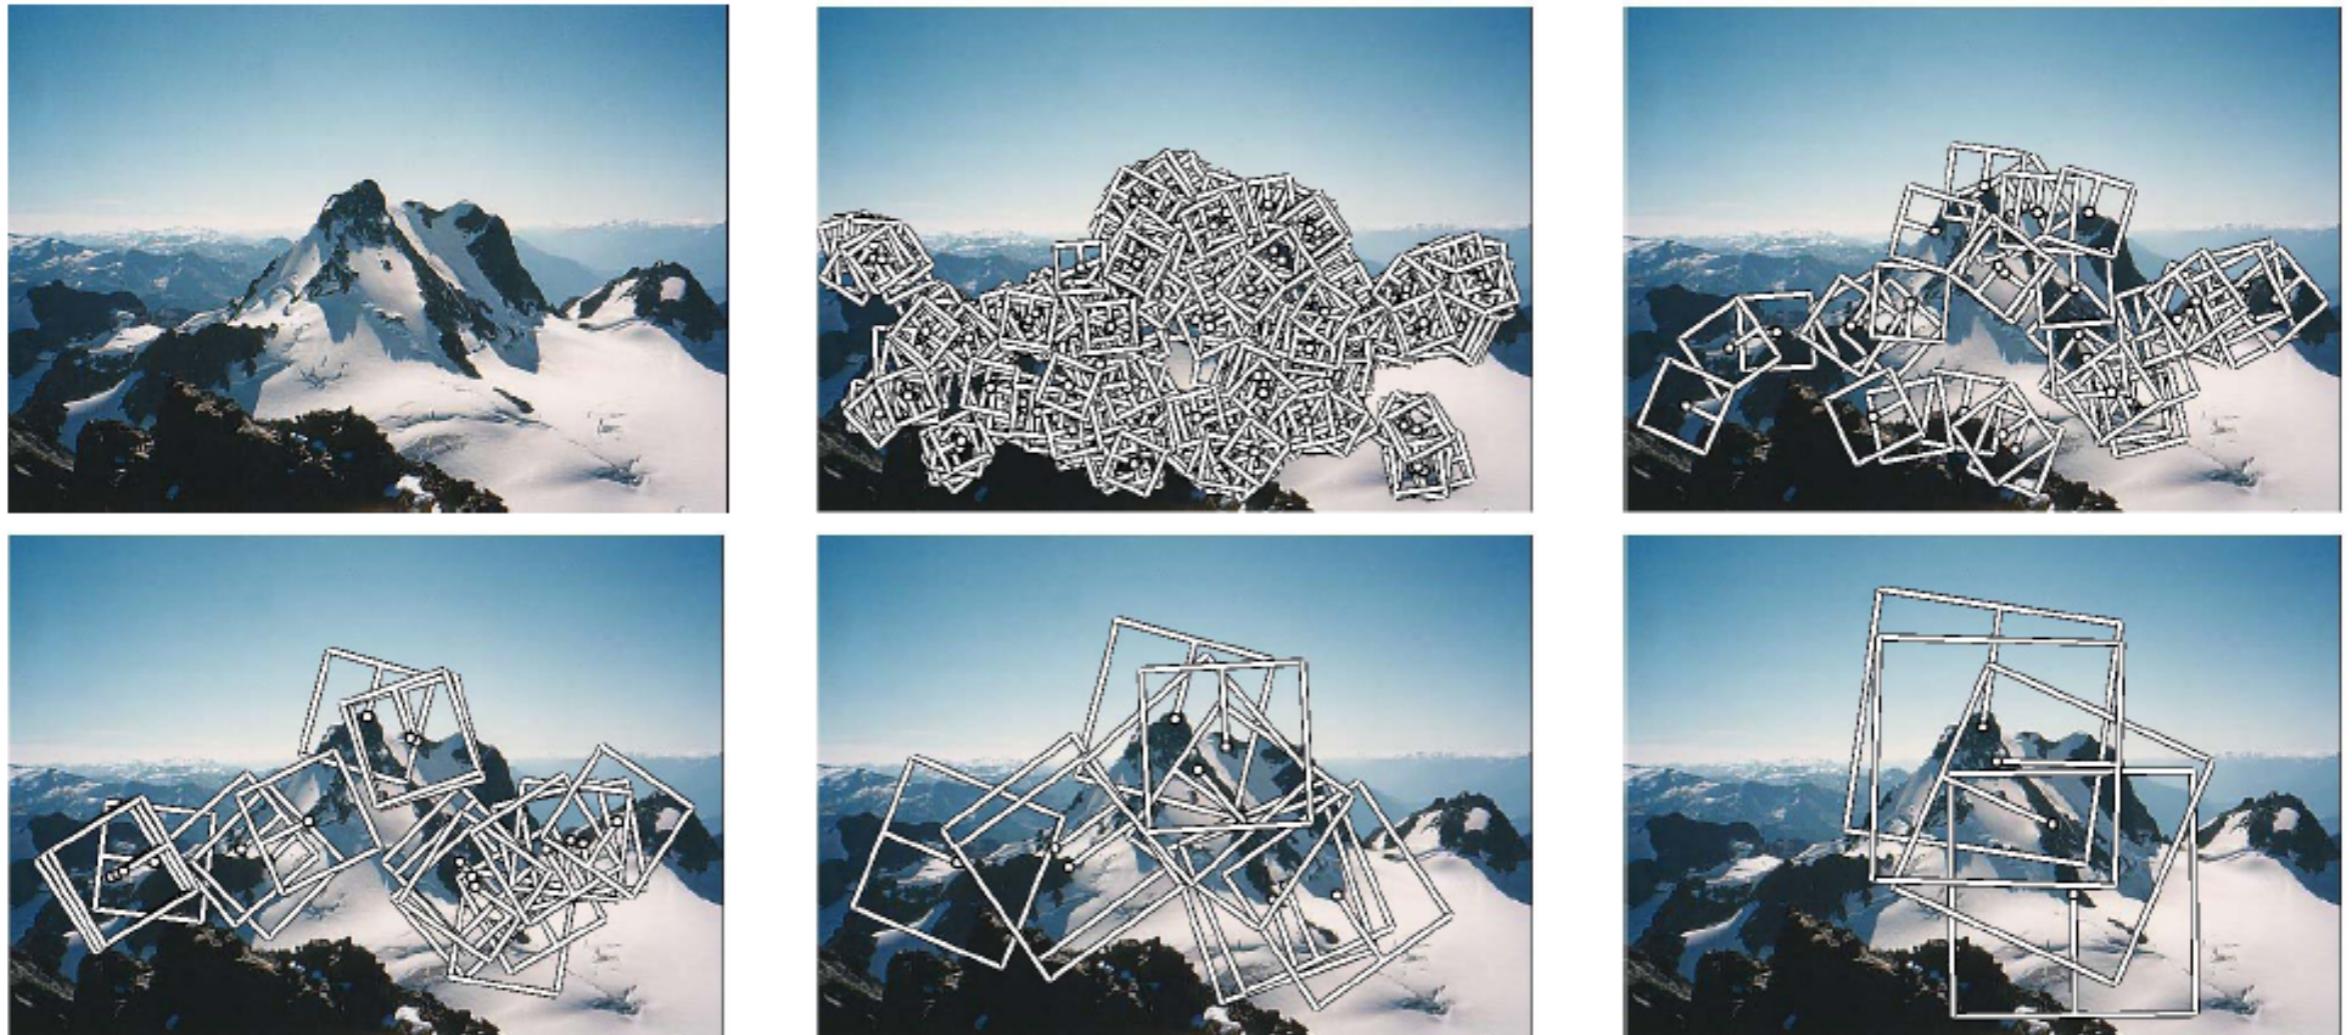
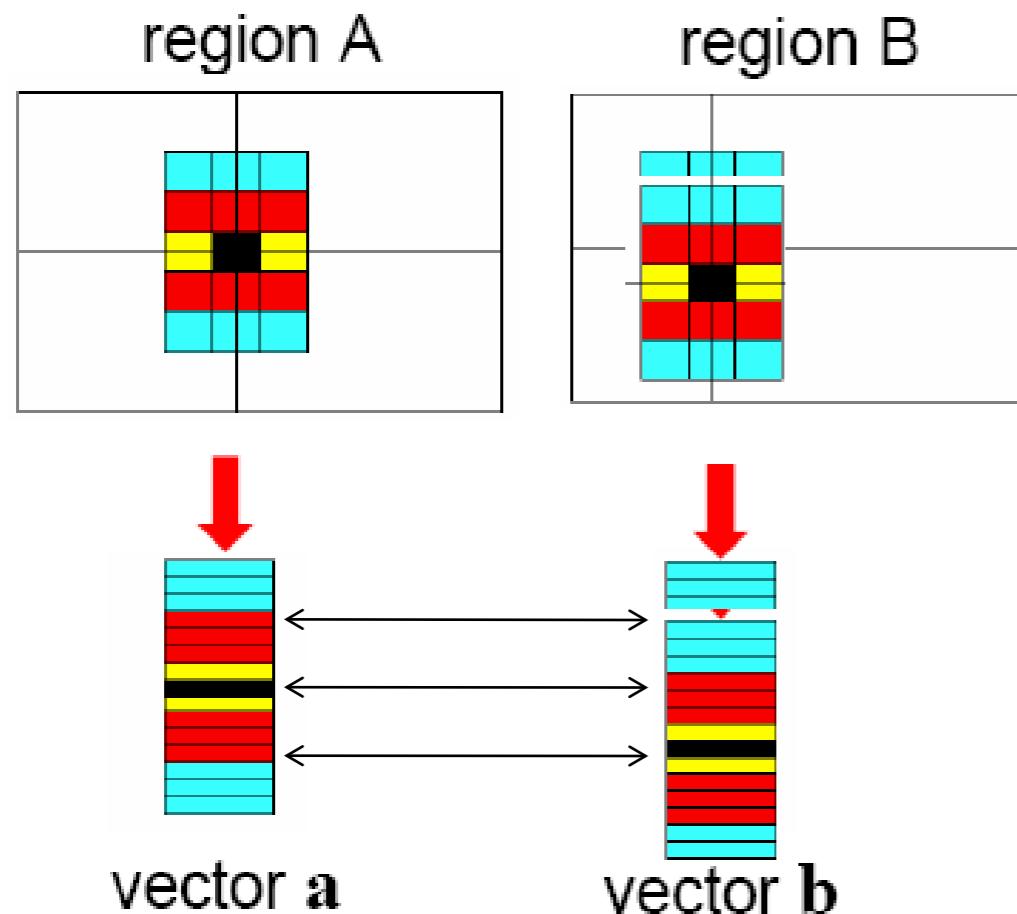


Figure 1. Multi-scale Oriented Patches (MOPS) extracted at five pyramid levels from one of the Matier images. The boxes show the feature orientation and the region from which the descriptor vector is sampled.

# How to describe the surrounding of a feature point?

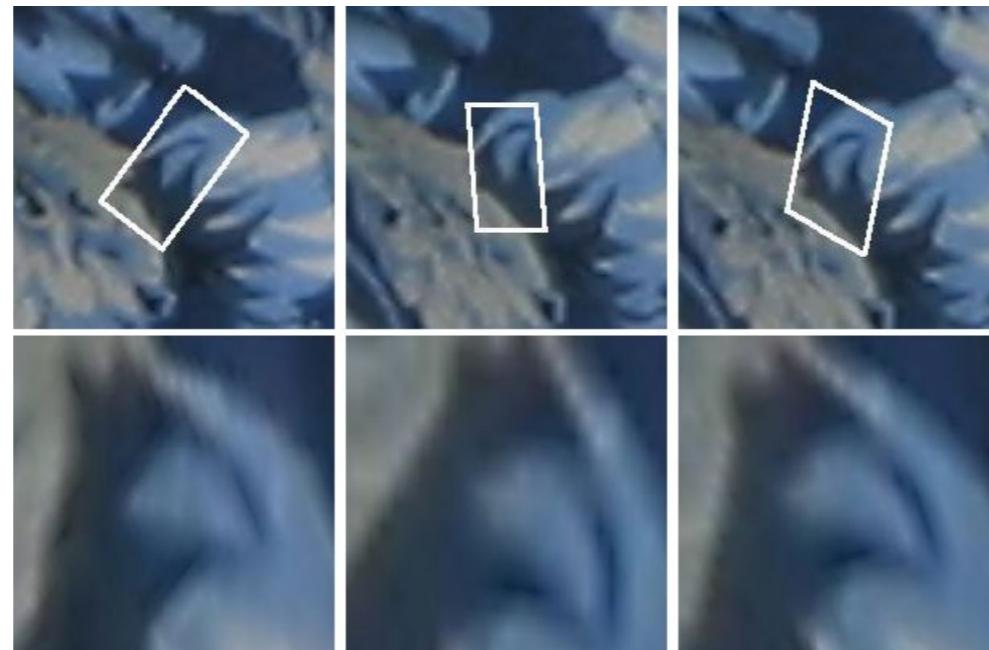


The simplest way to describe the neighborhood around an interest point is to write down the list of intensities to form a feature vector.

But this is very sensitive to even small shifts, rotations.

# Feature descriptors

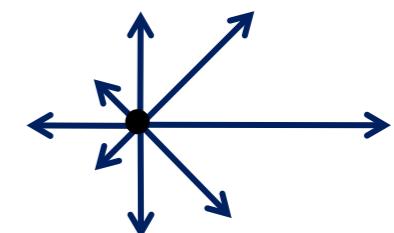
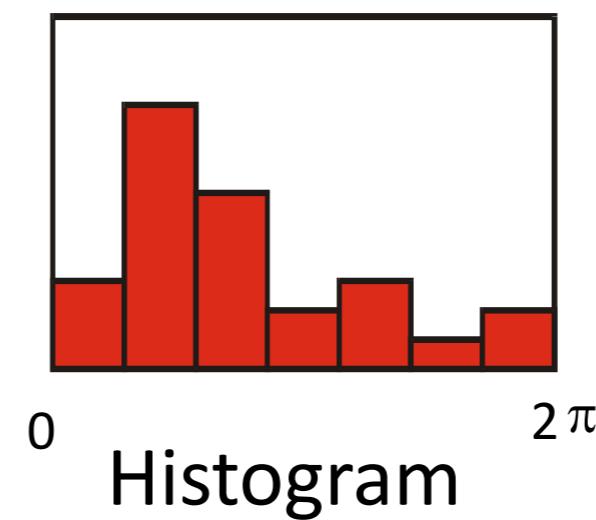
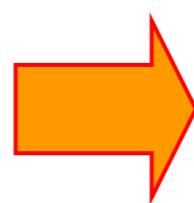
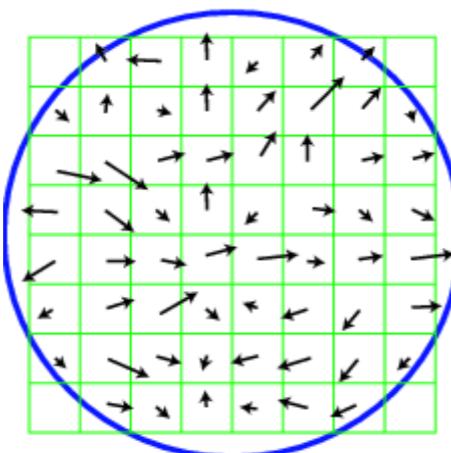
- Disadvantage of patches as descriptors: Small shifts can affect matching score a lot



- List of pixels is a bad idea..

# Feature descriptors

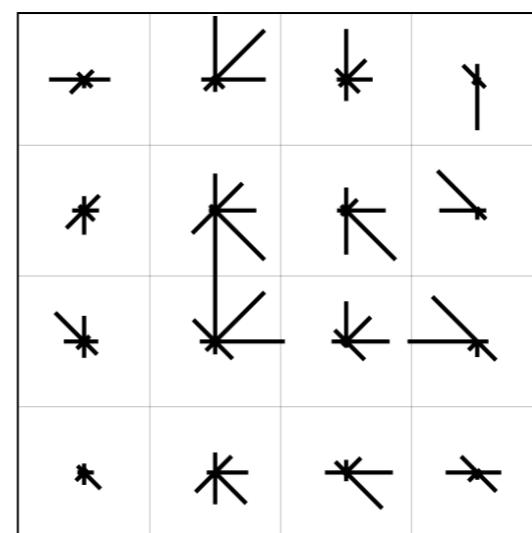
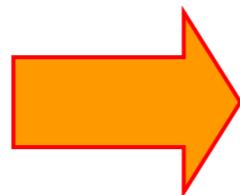
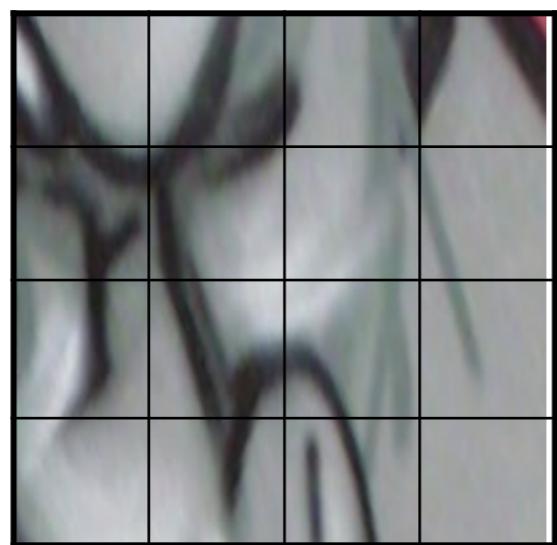
- Better: not pure pixels but image gradients (edge directions)
  - Advantage: more robust to illumination variations
- Aggregate over several pixels using histogram
  - To compensate small shifts
- Histograms of gradients



Gradient Lengths

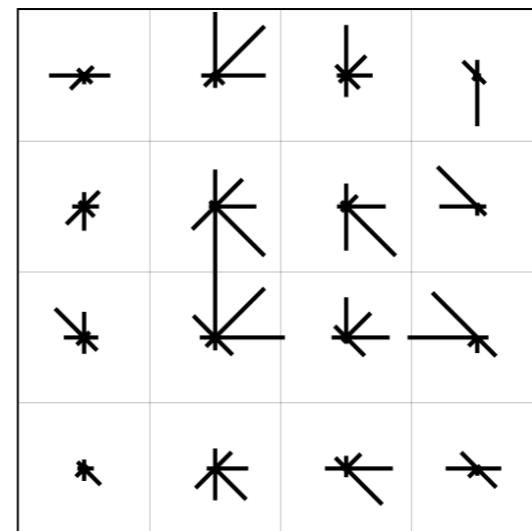
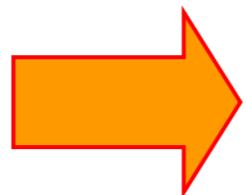
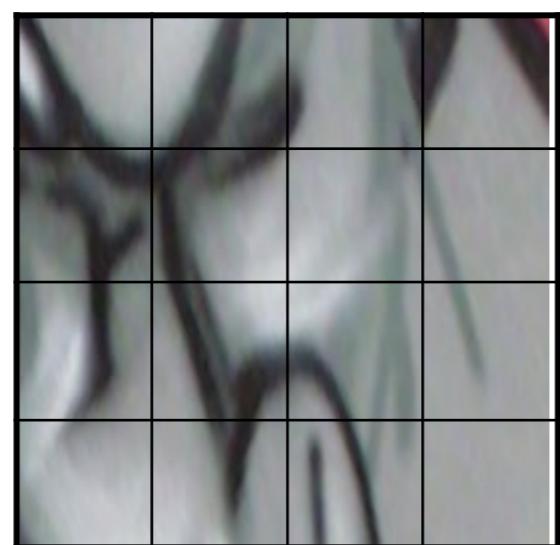
# Feature descriptors: SIFT

- Scale Invariant Feature Transform
- Descriptor computation:
  - Divide patch into  $4 \times 4$  sub-patches: 16 cells
  - Compute histogram of gradient orientations (8 reference angles) for all pixels inside each sub-patch. Includes also magnitude!
  - Resulting descriptor:  $4 \times 4 \times 8 = 128$  dimensions



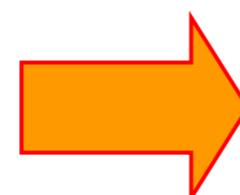
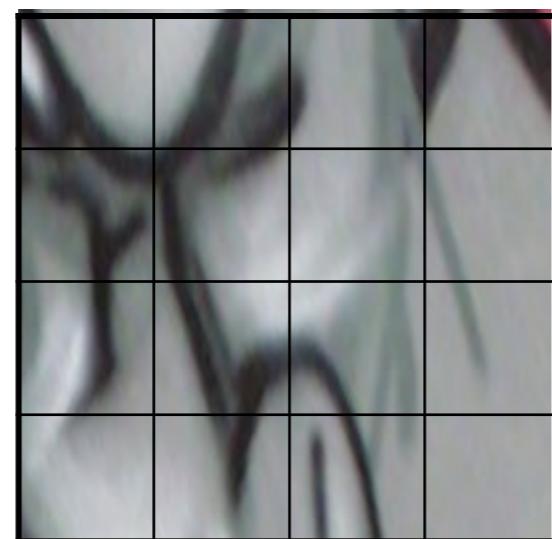
# Feature descriptors: SIFT

- Scale Invariant Feature Transform
- Descriptor computation (cont.):
  - Normalize vector to length of 1 ( $\rightarrow$  robustness to lightning variations)
  - Remove entries  $< 0.2$  (very small gradients – probably noise)
  - Renormalize vector again to length of 1



# Feature descriptors: SIFT

- Why subpatches?
- What invariances do we get?
  - Illumination?
  - Translation?
  - Scale?
  - Rotation?



*	↖	*	↑
*	↗	↑	↘
*	↙	↖	→
*	↖	*	↗

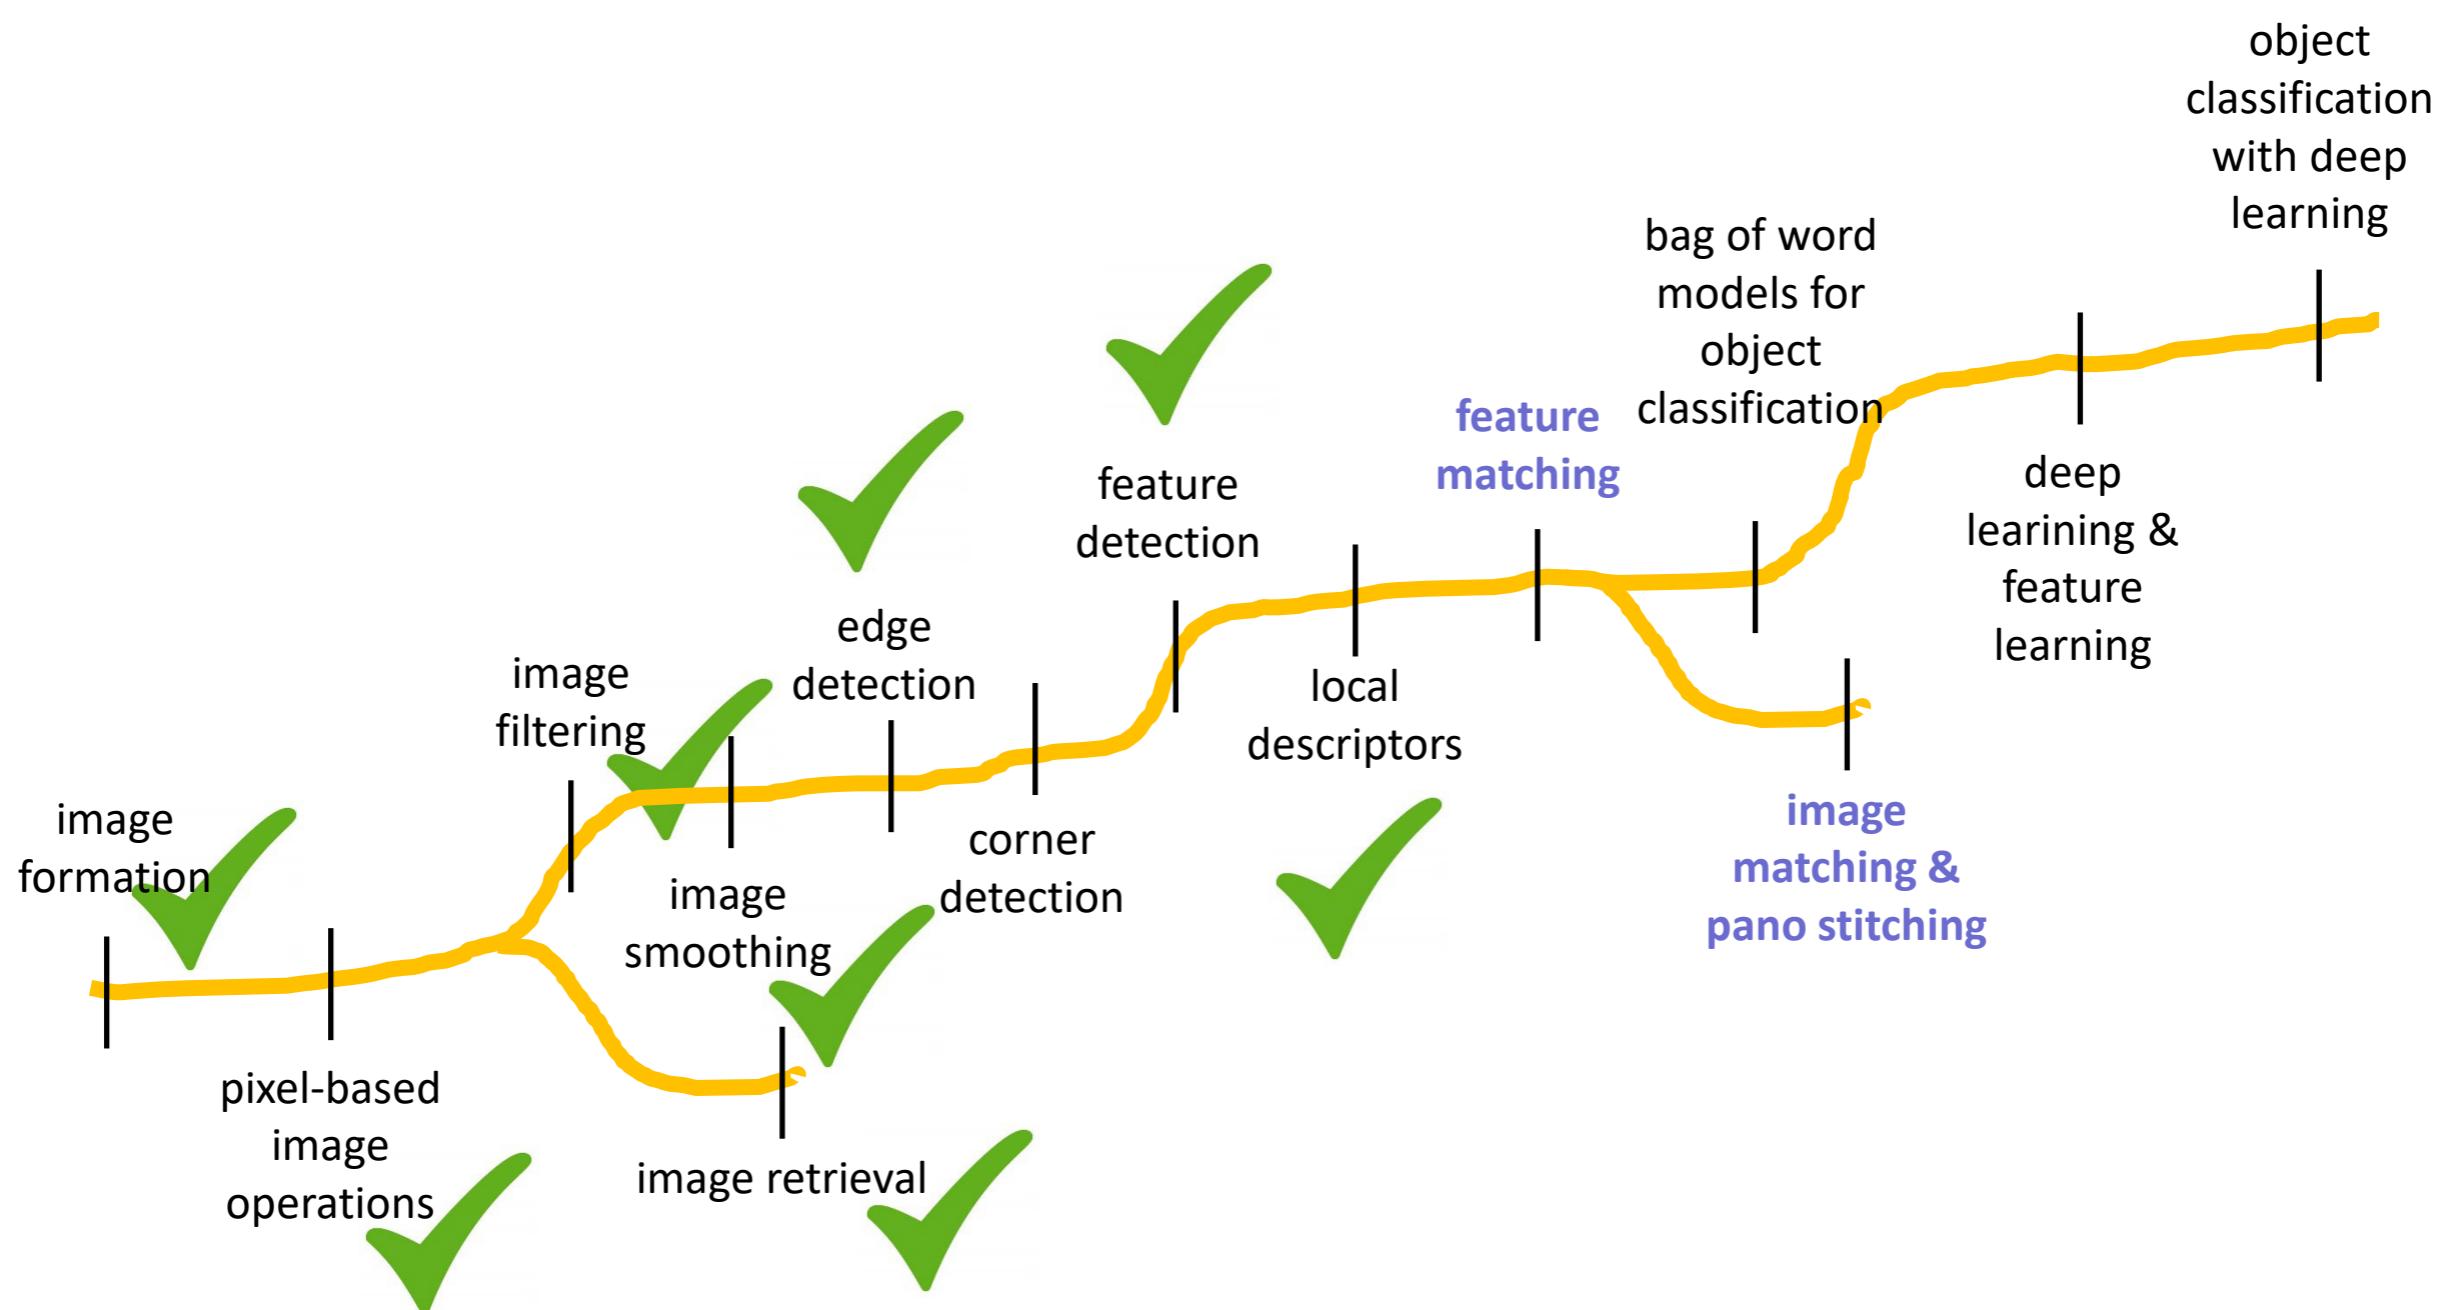
# Feature descriptors: SIFT

- Extraordinarily robust descriptor
  - Can handle changes in viewpoint: Up to about 60 degree out of plane rotation
  - Can handle significant changes in illumination: sometimes even day vs. night (below)
  - Fast and efficient
  - Lots of code available (however patented!)
  - Today many different types of descriptors exist: **ORB**, GLOH, HOG, ...



Steve Seitz

# Roadmap



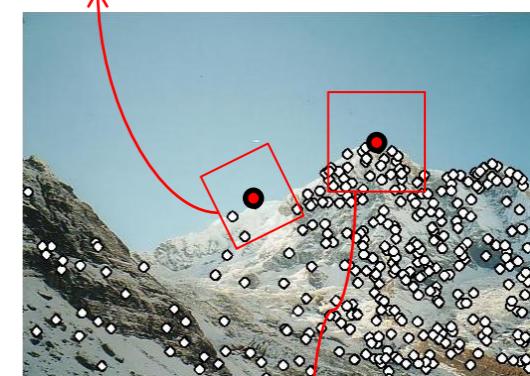
# Overview – The Big Picture

- 1. **Detection:** identify the interest points

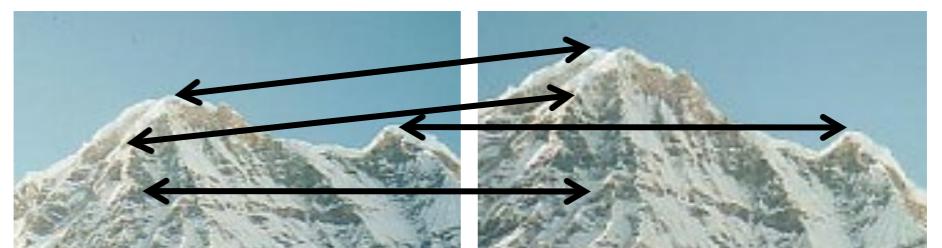


- 2. **Description:** extract a vector (feature descriptor) that represents the surrounding of each feature point

$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$

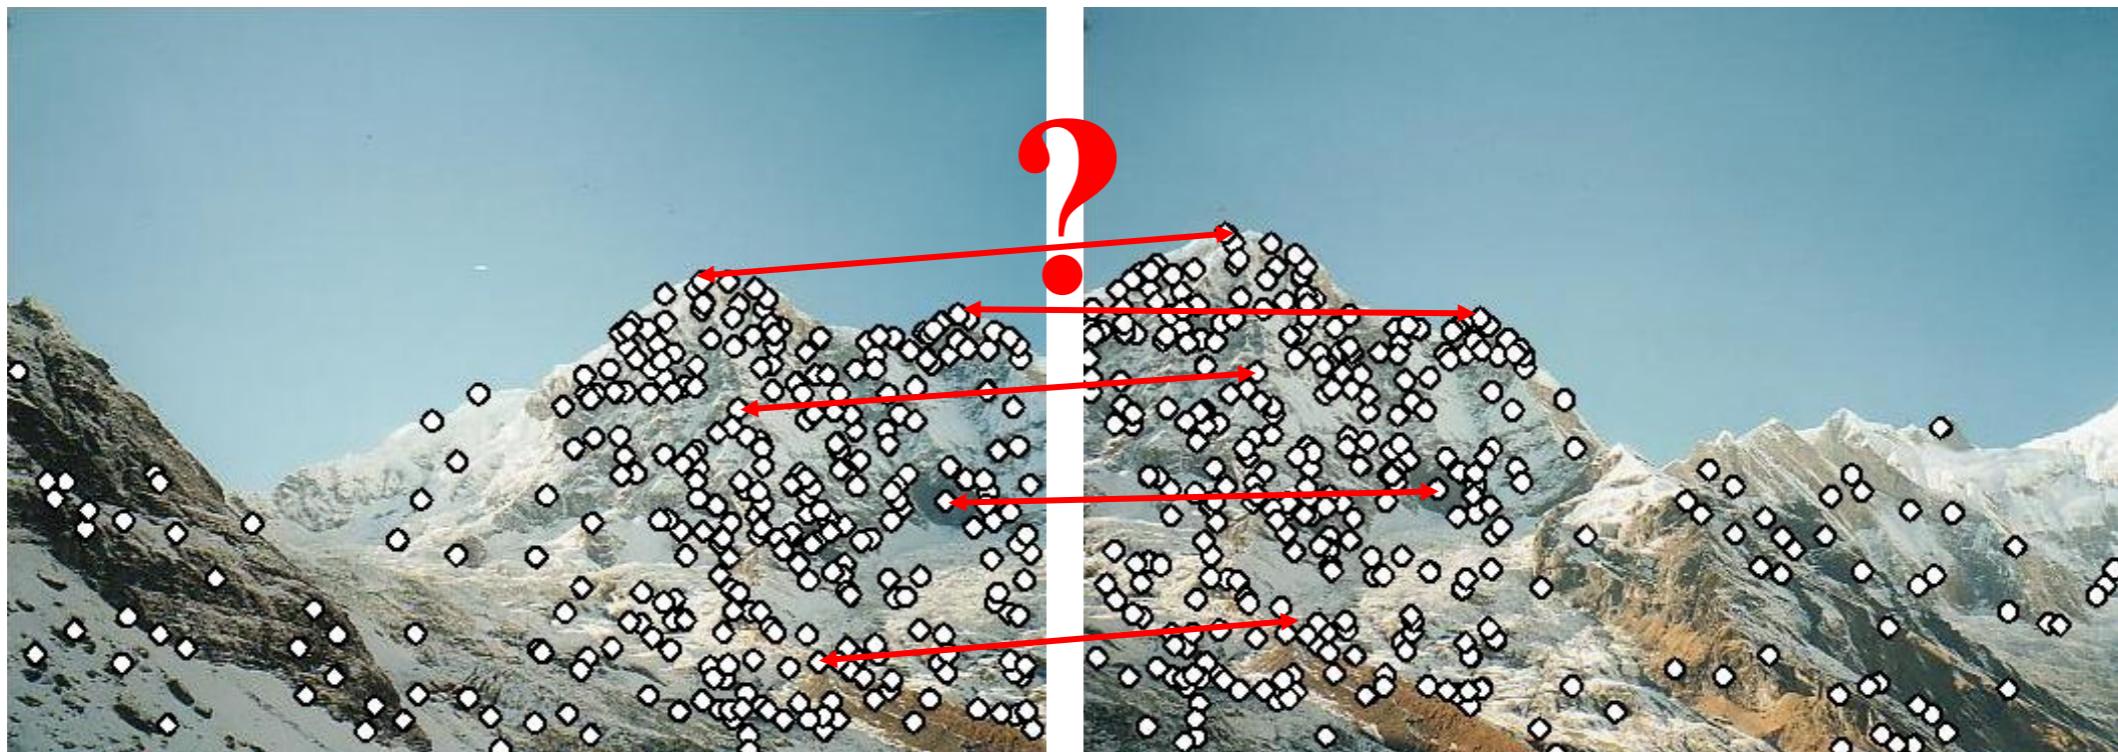


- 3. **Matching:** determine correspondence between descriptors in two views



# Finding corresponding points

- We know now how to detect and describe local feature points. How can we use them now to match images?



# Feature matching

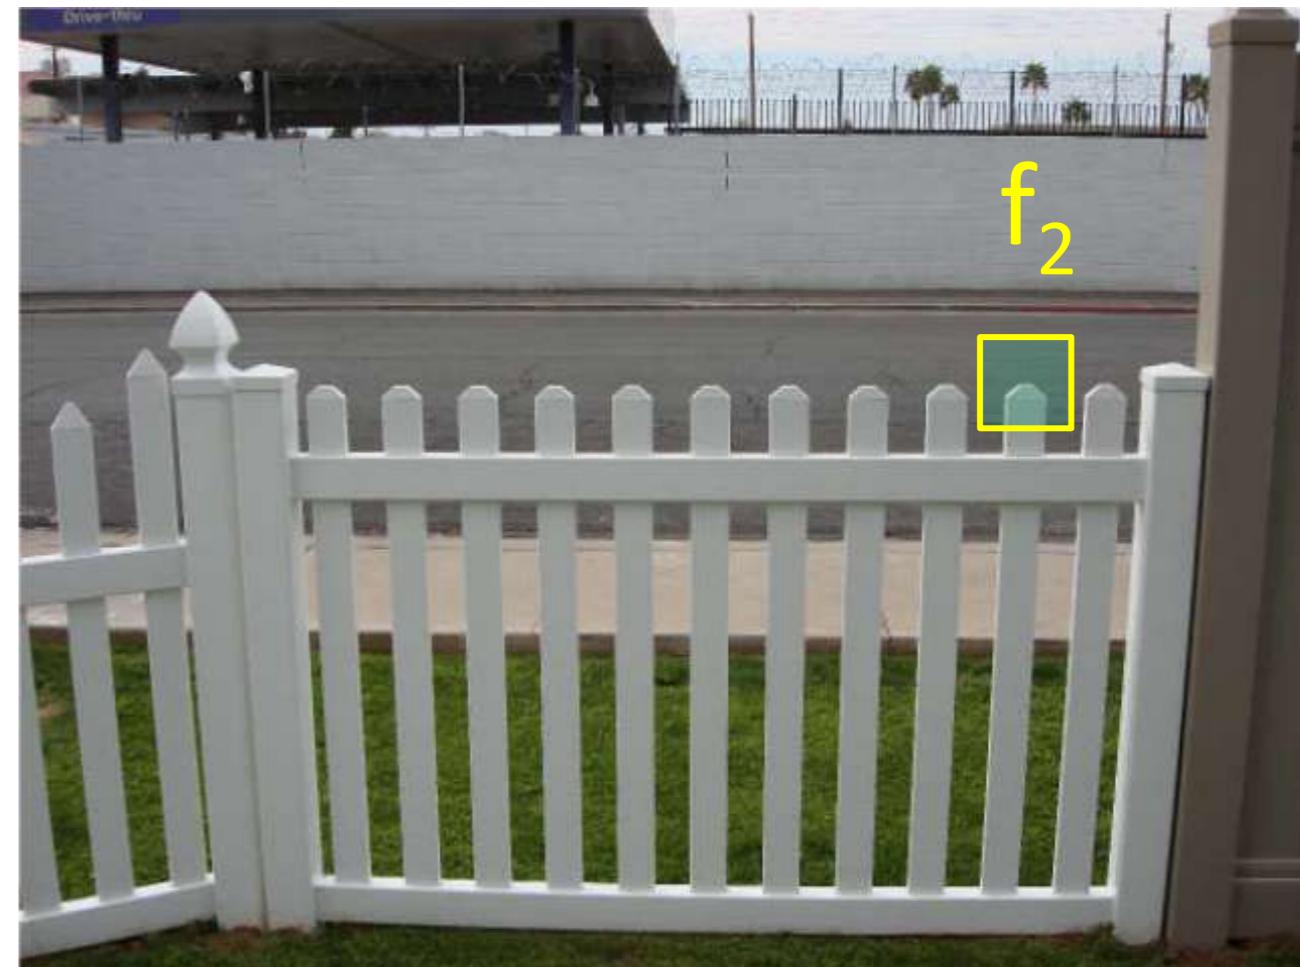
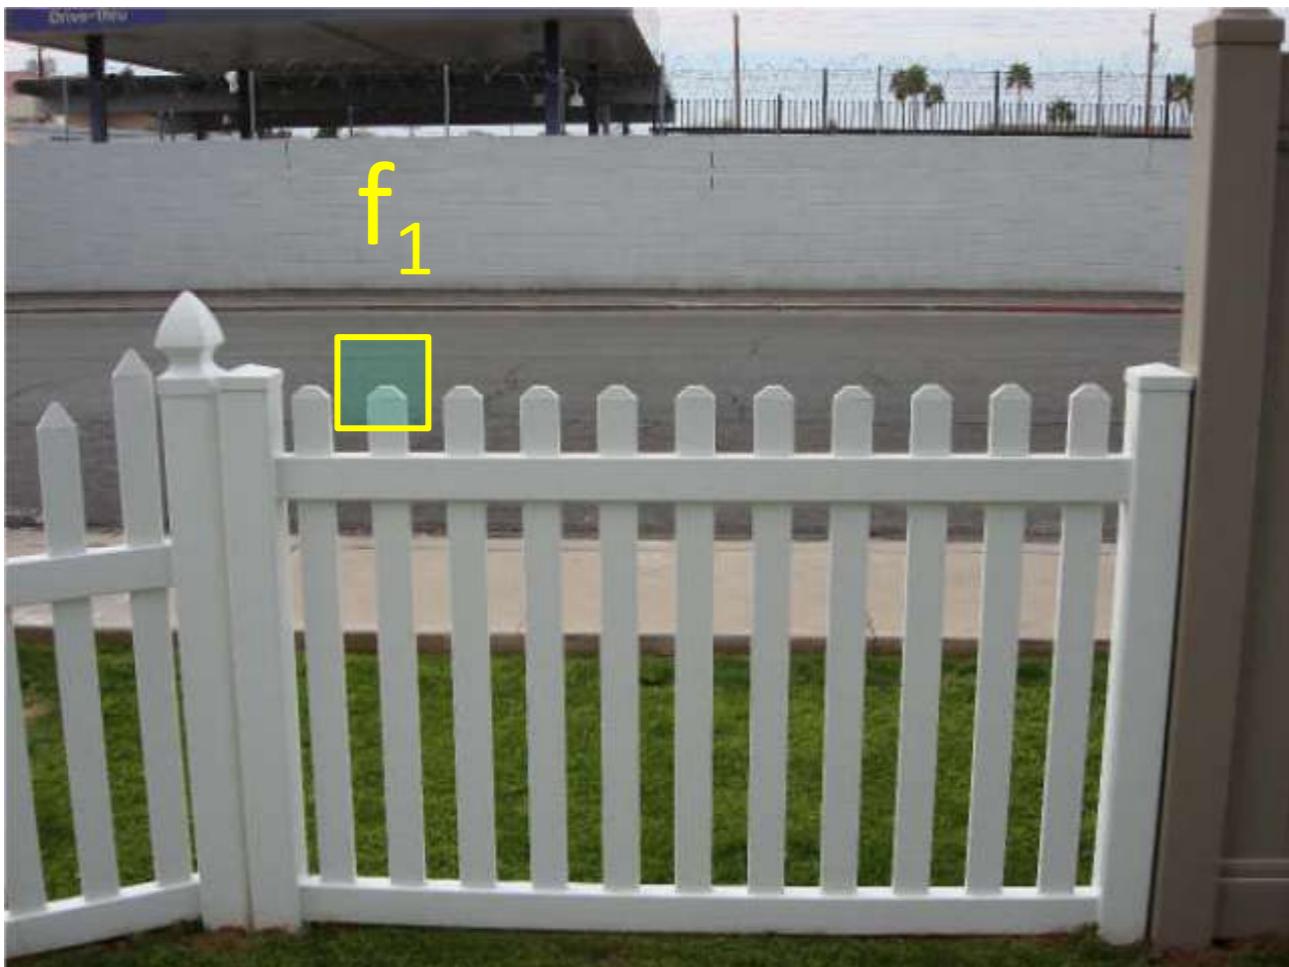
Given a feature in image  $I_1$ , how to find the best match in image  $I_2$ ?

1. Define distance function that compares two **descriptors**
2. Test all the features in  $I_2$ , find the one with **minimum distance**

# Feature distance

How to define the difference between two features  $f_1, f_2$ ?

- Simple approach is  $\text{SSD}(f_1, f_2)$ 
  - **Sum of square differences** between entries of the two descriptors
  - Can give good scores to very ambiguous (bad) matches



$I_1$

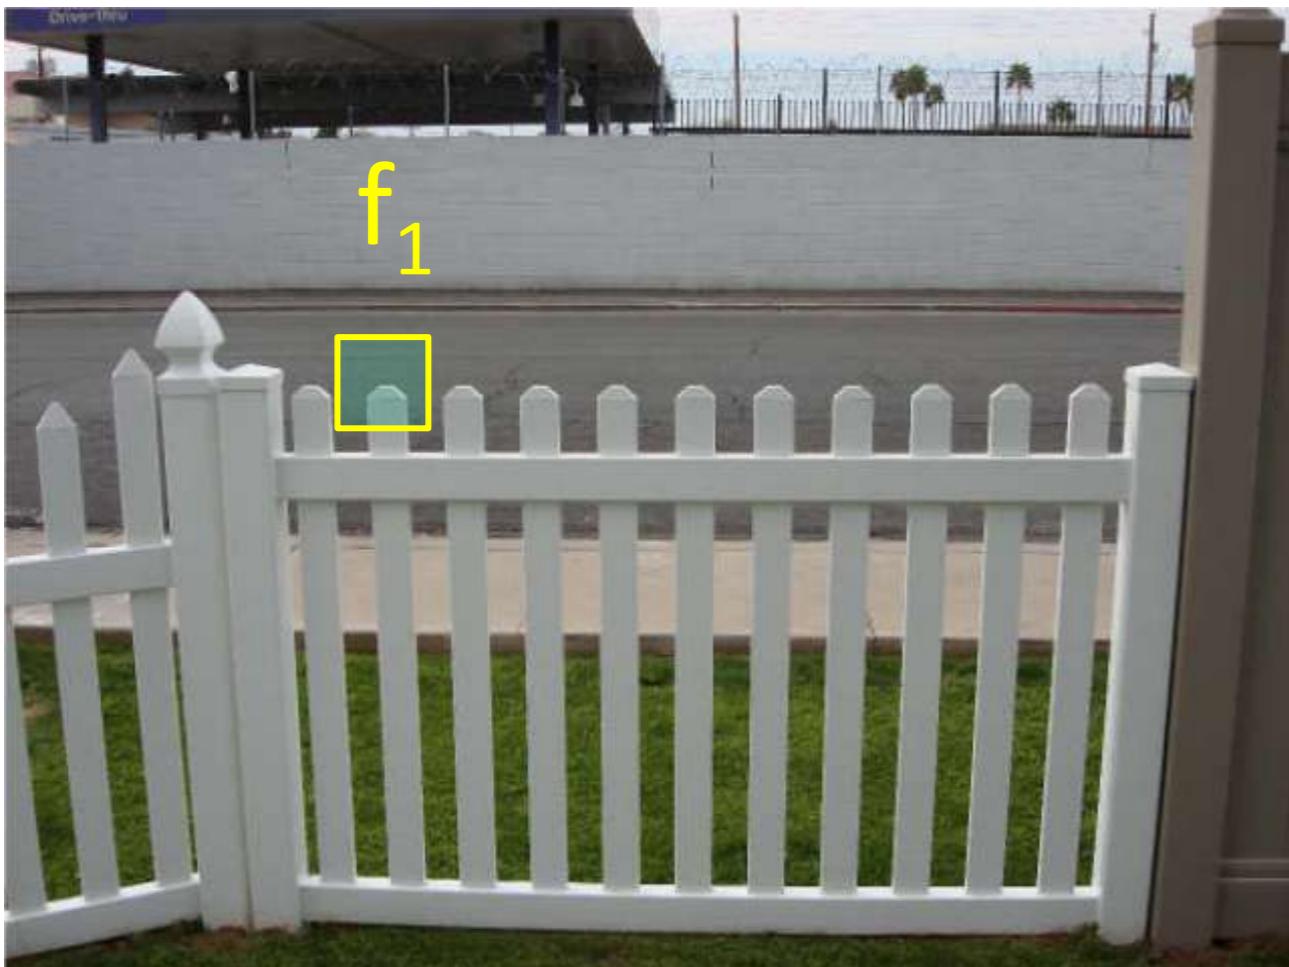
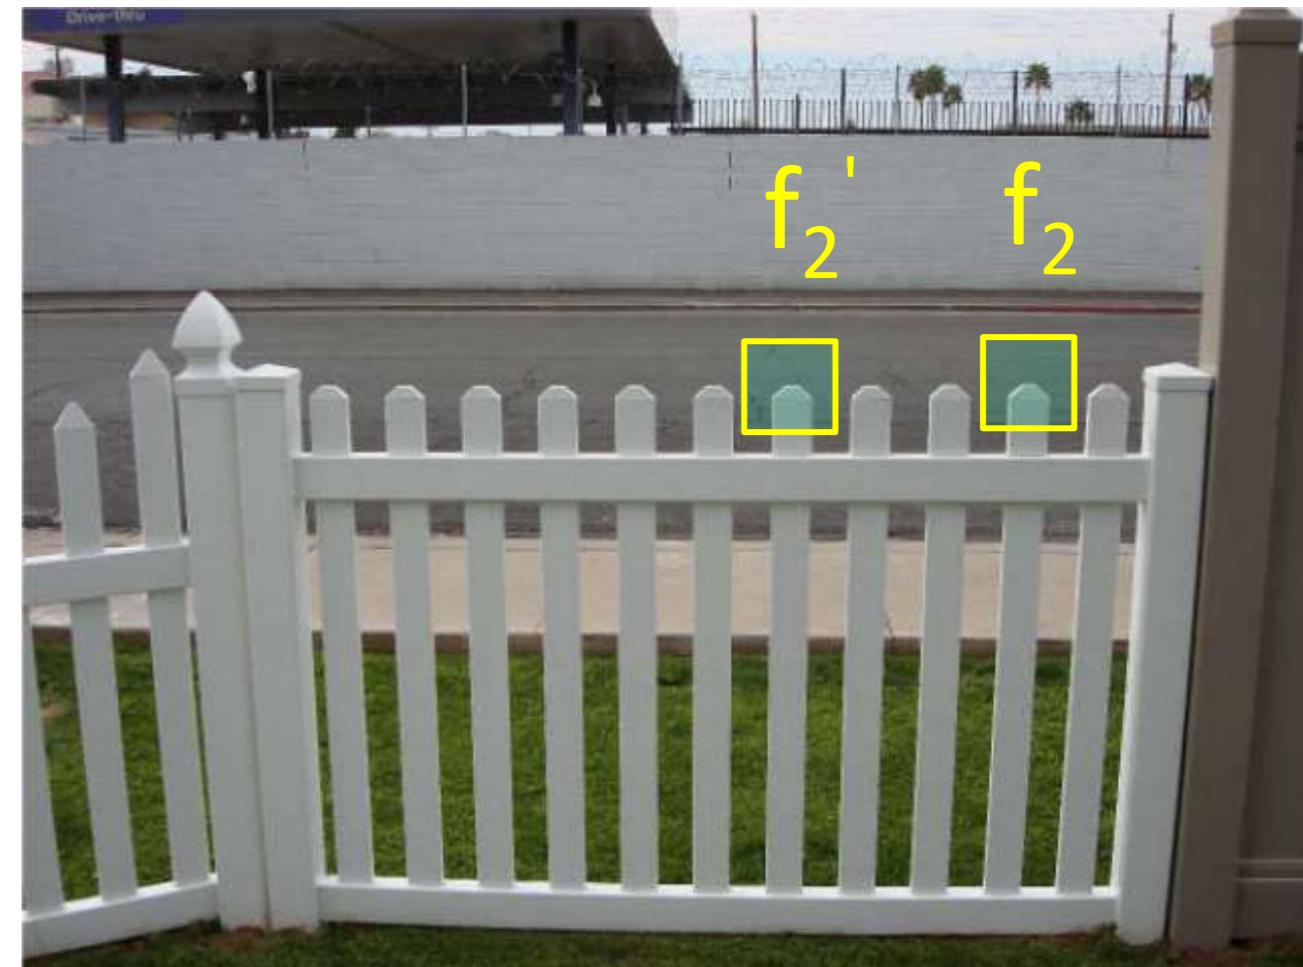
$I_2$

Credit: S. Szeliski

# Feature distance

How to define the difference between two features  $f_1, f_2$ ?

- Simple approach is  $\text{SSD}(f_1, f_2)$ 
  - **Sum of square differences** between entries of the two descriptors
  - Can give good scores to very ambiguous (bad) matches

 $|_1$  $|_2$ 

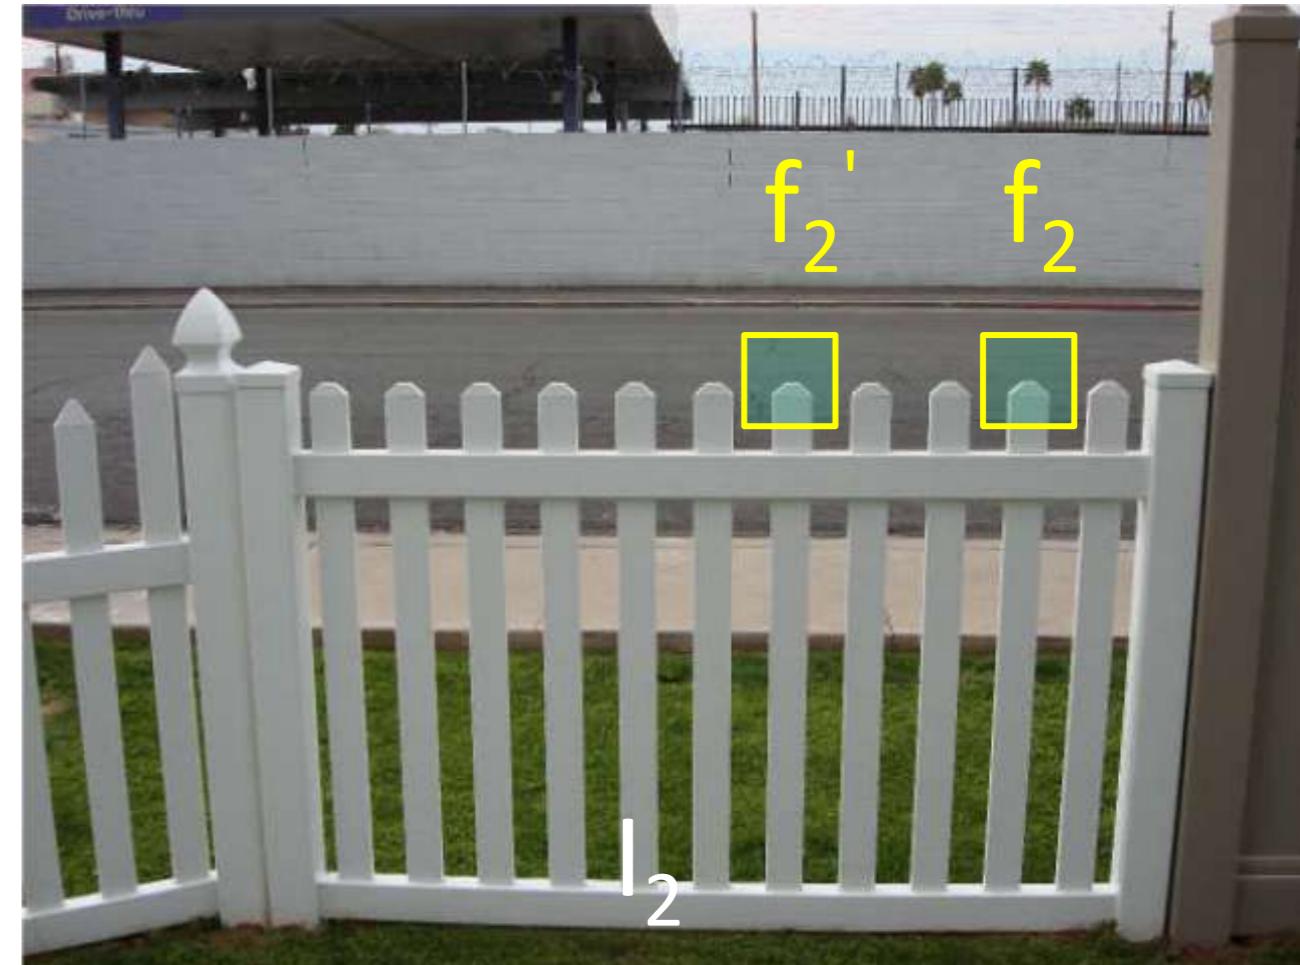
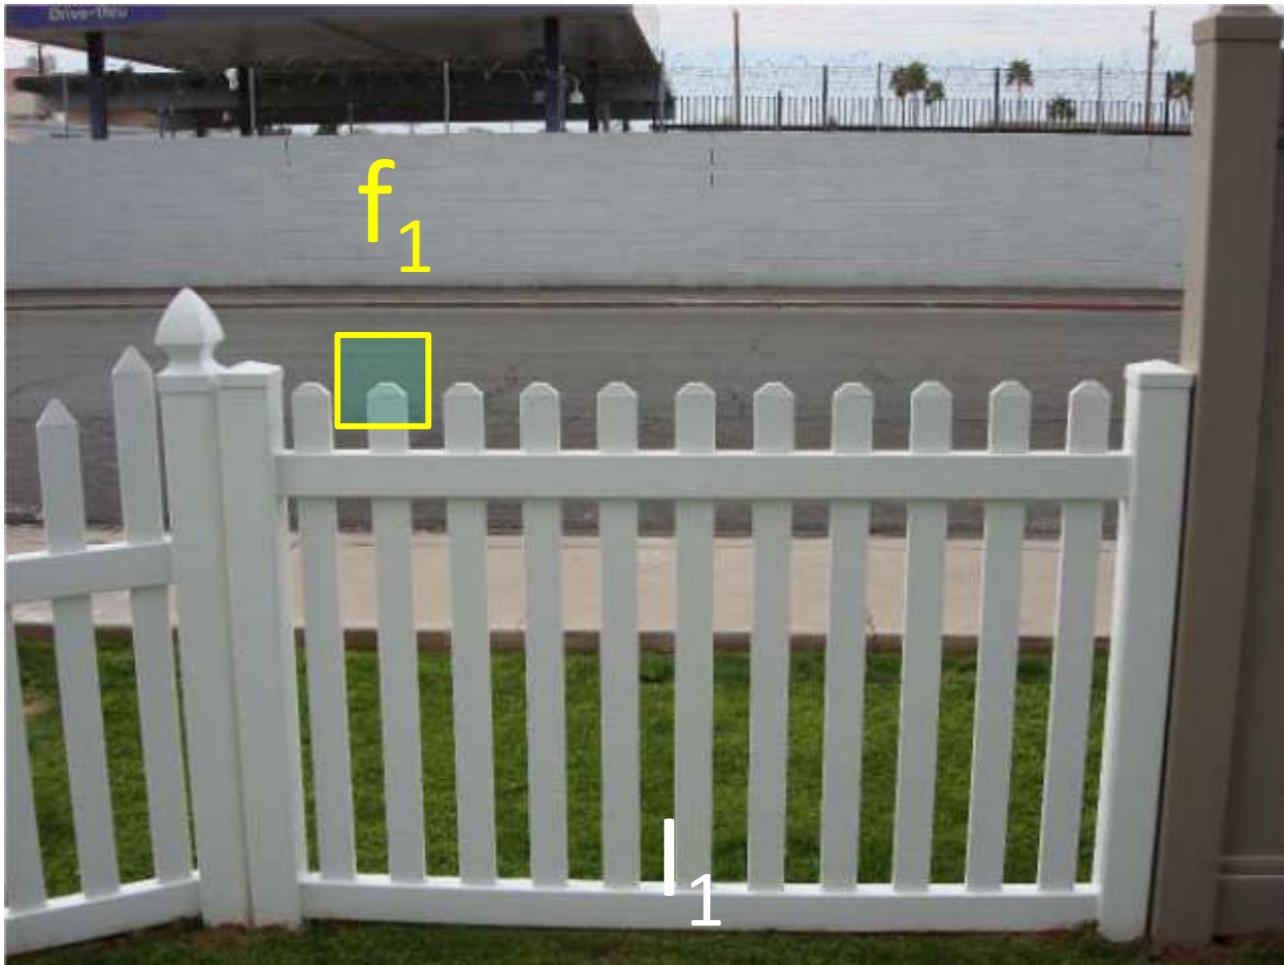
Credit: S. Szeliski

## Dual threshold method:

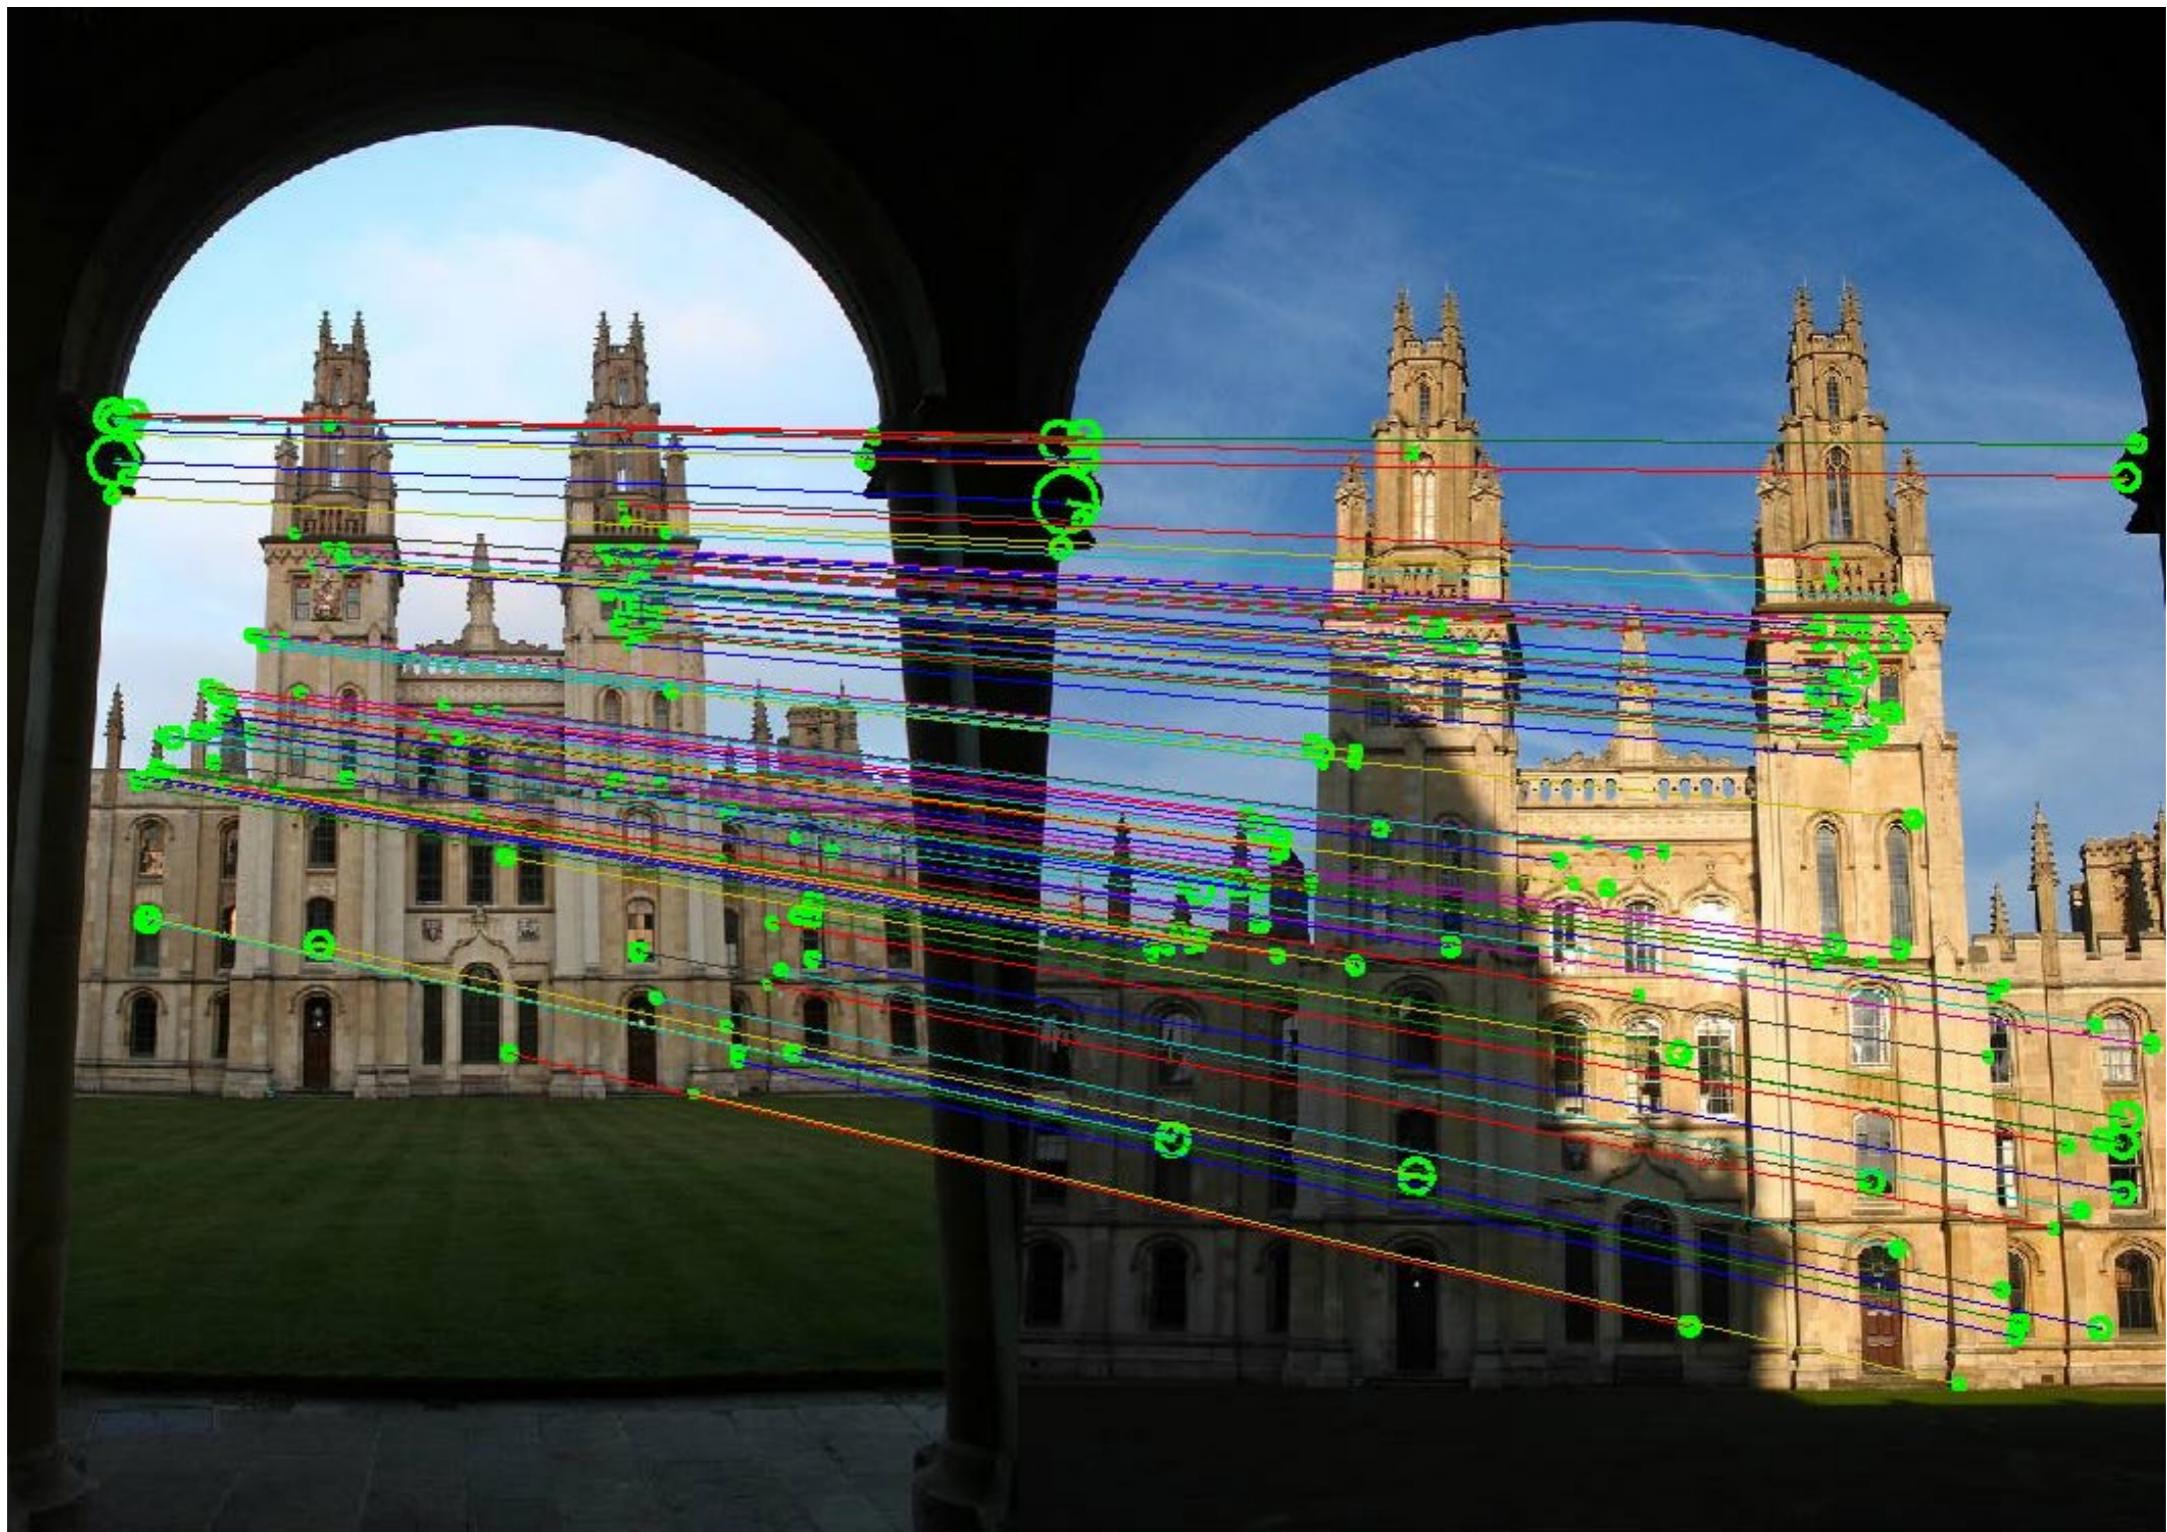
- Find the *two* nearest neighbors (Euclidean distance)
- Get their distances ( $d_1$  and  $d_2$ )
  - What is the possible value range of the distances?
- Threshold the “**ratio** of nearest to second nearest descriptor”:  $d_1/d_2$

# Dual Threshold method

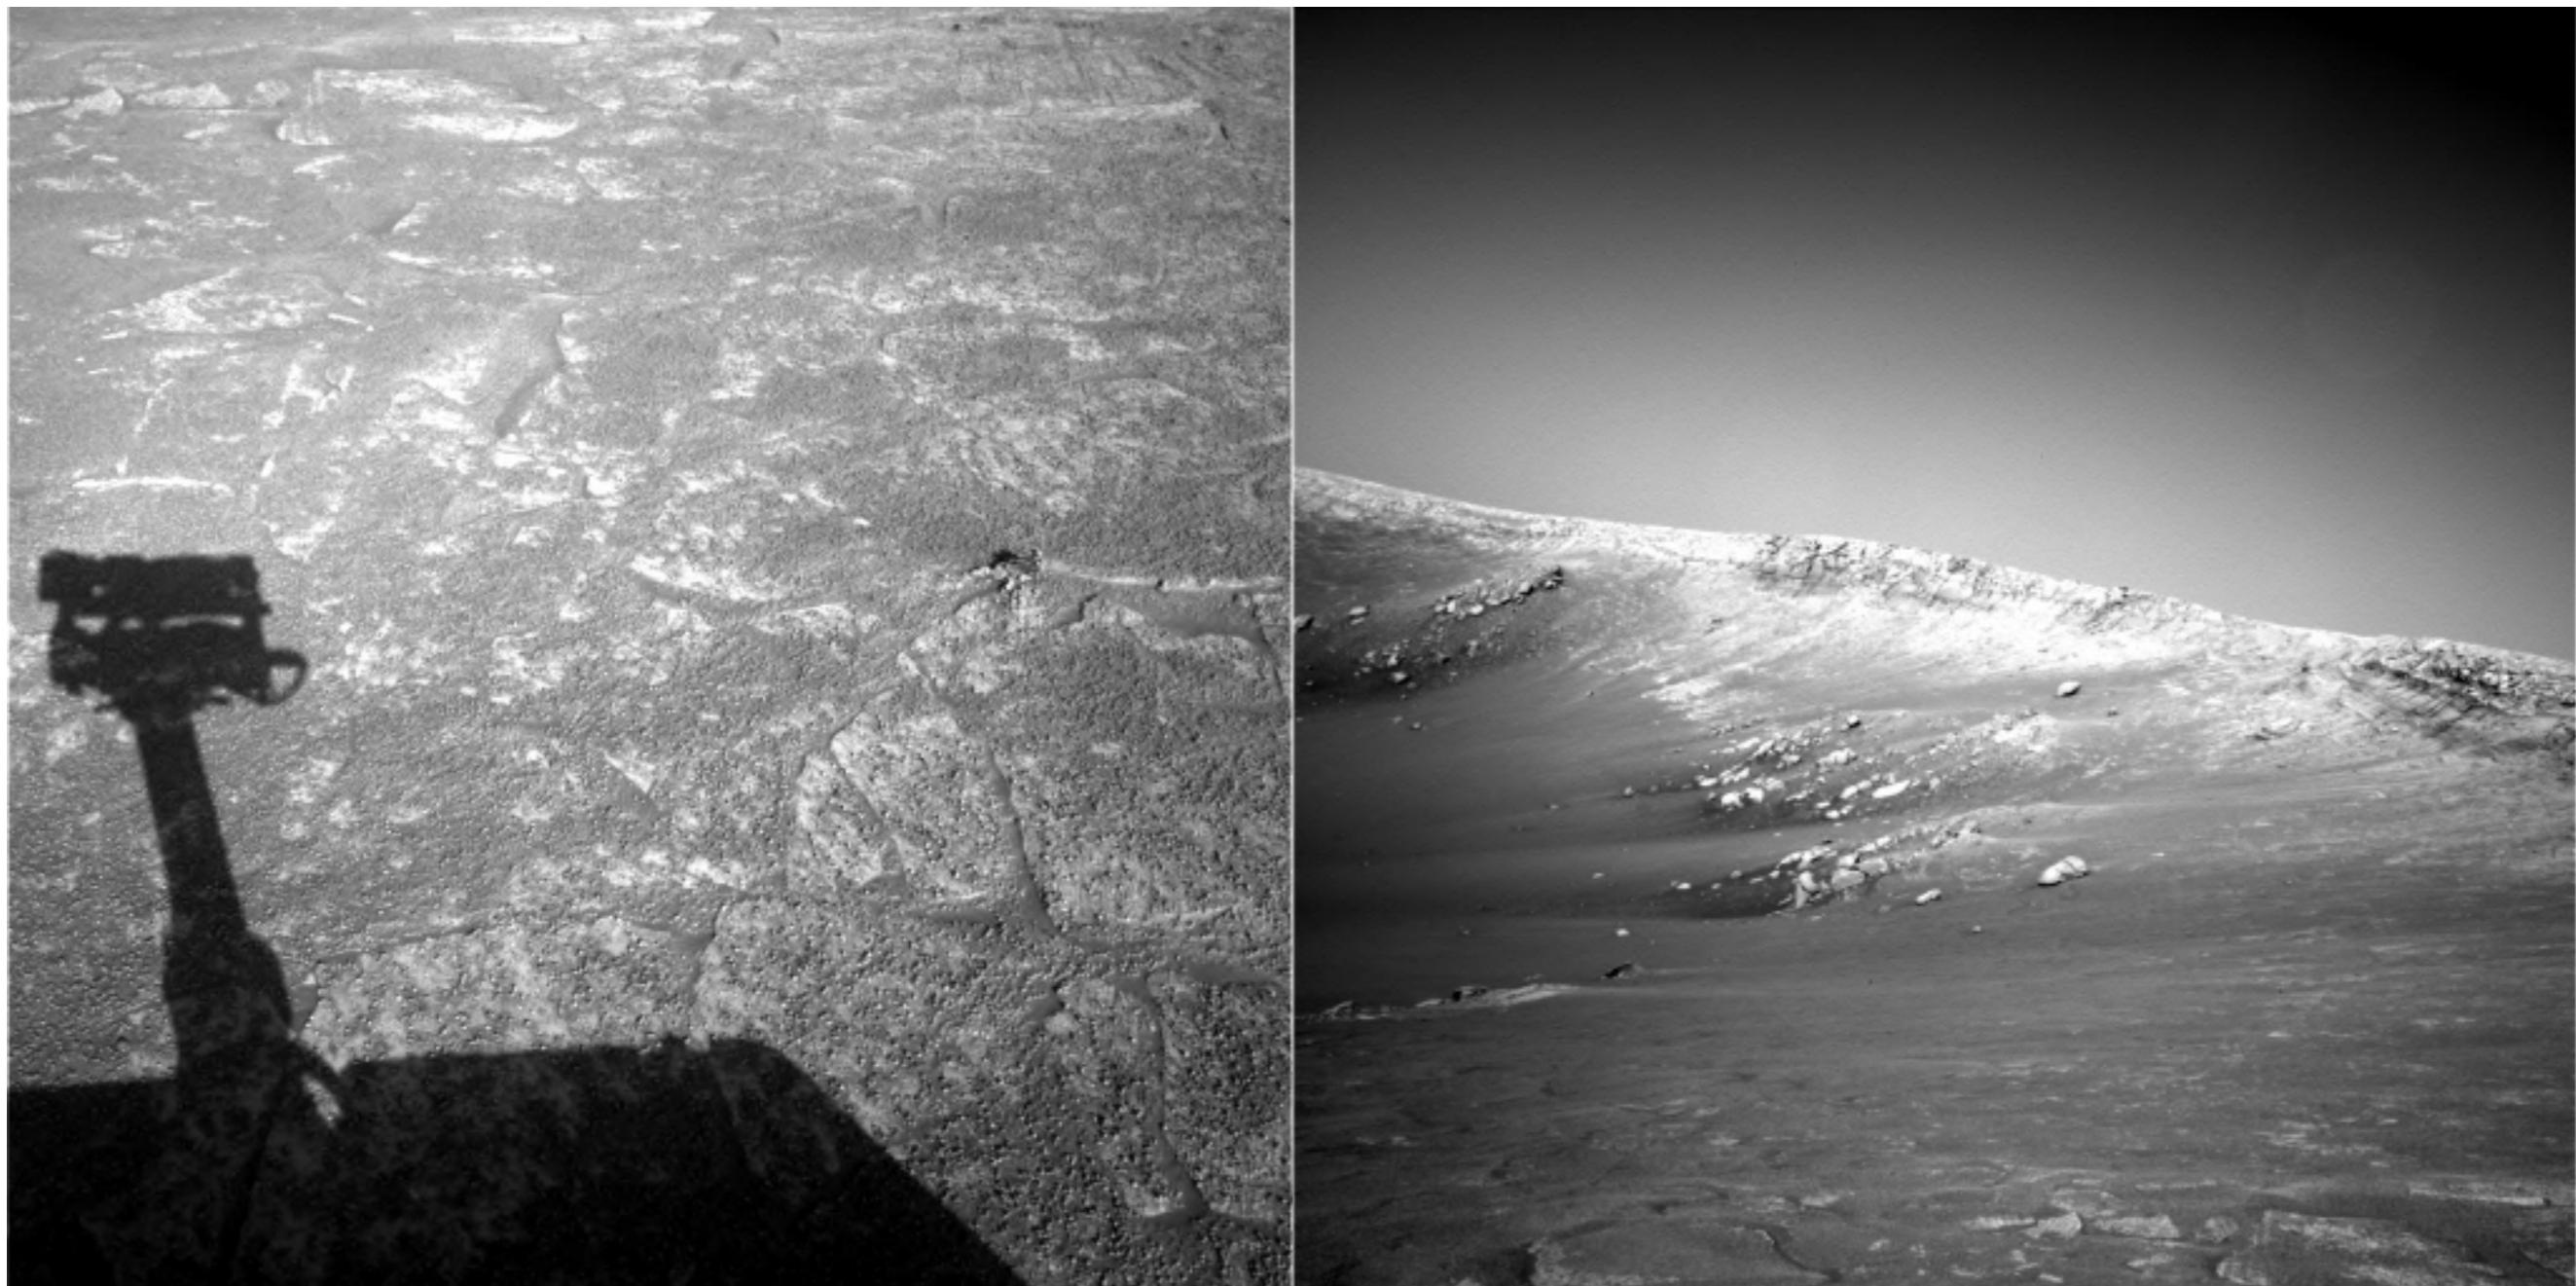
- Ratio distance =  $L2(f_1, f_2) / L2(f_1, f_2')$ 
  - $f_2$  is best L2 match to  $f_1$  in  $I_2$
  - $f_2'$  is 2nd best L2 match to  $f_1$  in  $I_2$
  - gives values close to 1 for ambiguous matches (e.g. 0.11/0.12) and small values for unique matches (e.g. 0.1/0.8)



# Matching Result

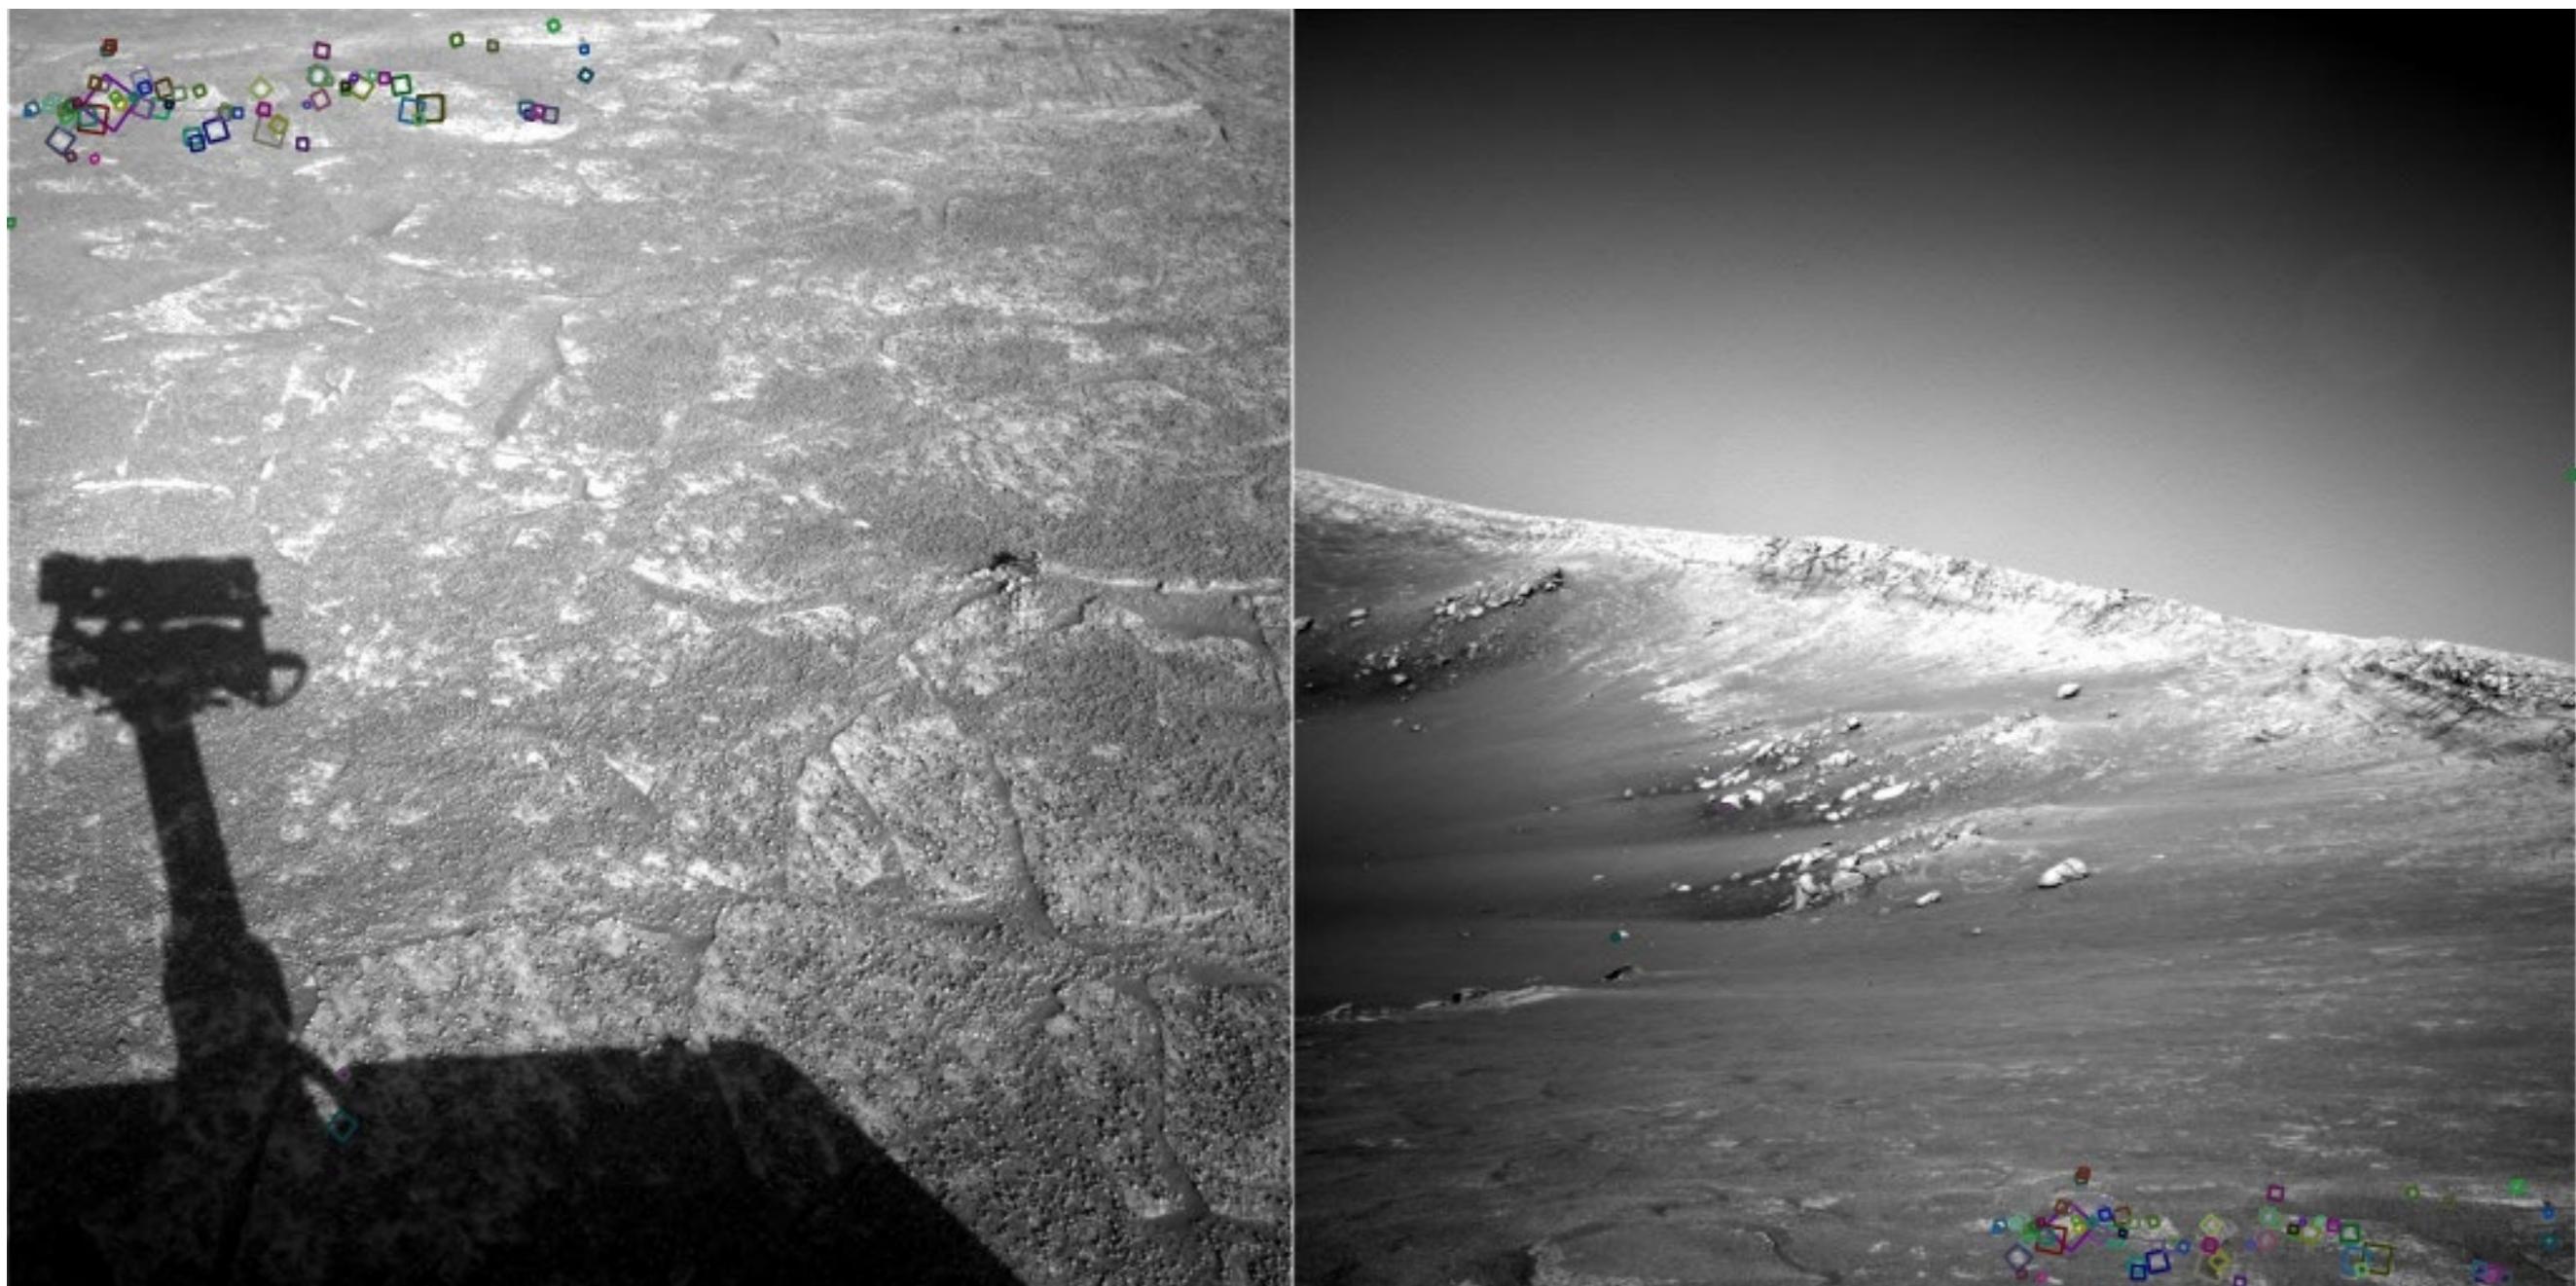


## Example – a hard case..



NASA Mars Rover images

## Example – a hard case..



NASA Mars Rover images  
with SIFT feature matches  
Figure by Noah Snavely

# Summary on Local Features

- Local features are **distinct and salient points** in an image, e.g.
  - Corners: Harris Corner Detector
  - Blobs: Laplacian of Gaussian Detector
- Feature descriptors represent the **neighborhood** of a local feature
  - Most: invariant to scale, rotation, translation, illumination
  - Some: invariant also to affine transforms
- Matching feature points = matching their descriptors
  - Distance measure: sum of squared differences (SSD) or L2 (Euclidean distance)
  - **Dual threshold approach works best in practice**
- **Note: every feature detector can be combined with every feature descriptor!**