

Introduction to Neural Networks

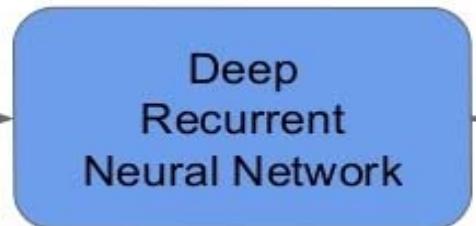
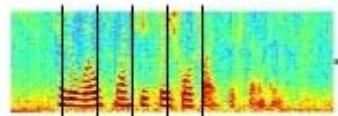
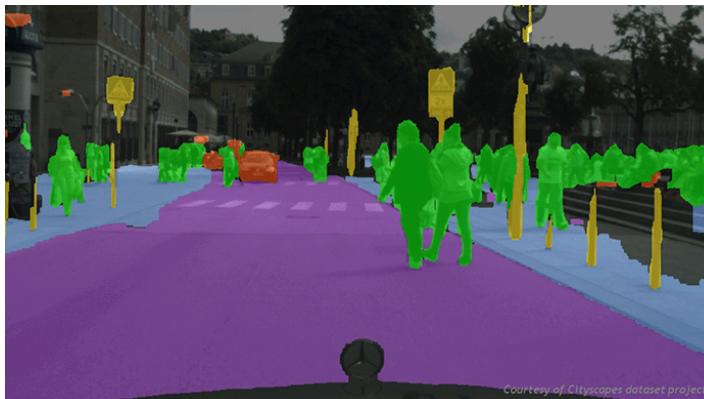
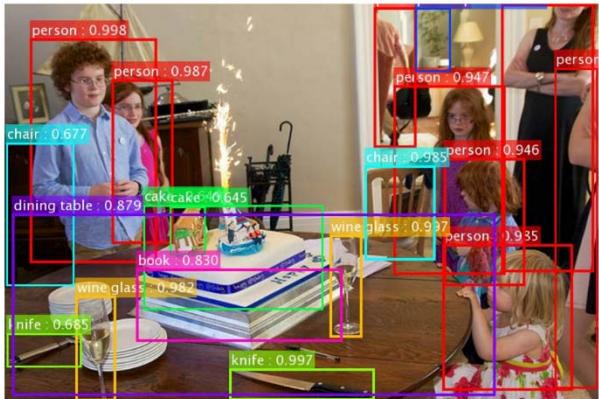
Matthias Zeppelzauer

Applied Computer Vision

MOTIVATION



Hey Siri

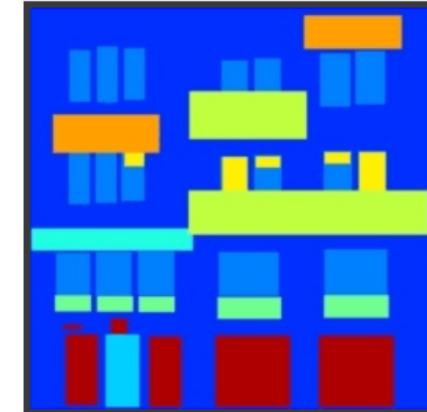


"How cold is it outside?"

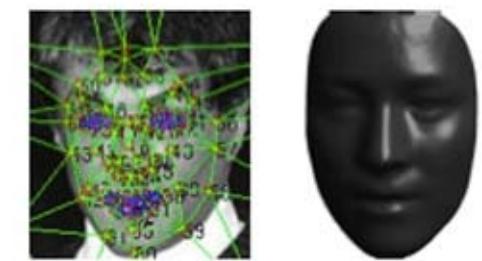
Text Output



INPUT



undo clear random



OUTPUT

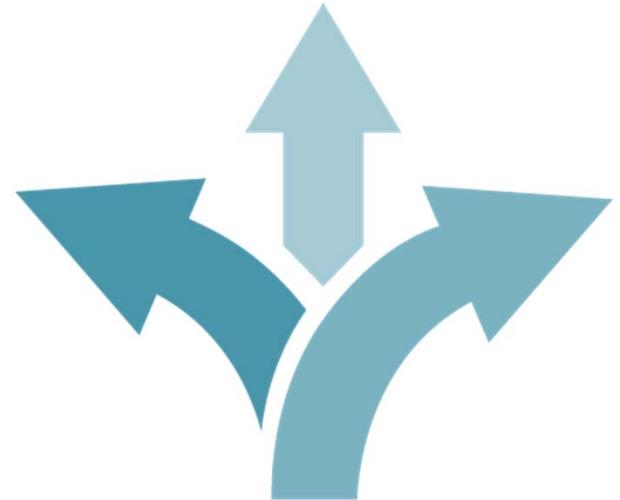
pix2pix
process



save

Strengths of Neural Networks

- › Neural networks can be used for many different purposes
 - › Different data types (image, time series, text etc.)
 - › Different problems and tasks (segmentation, classification, content generation)
 - › Networks can be combined, can have multiple inputs and multiple outputs
- › Neural networks are:
 - › Extremely flexible
 - › Powerful learning methodology
 - › State-of-the-art in many domains today



AGENDA

Introduction to neural networks

Neurons: the perceptron, the sigmoid neuron

Layers, multi-layer perceptron (MLP)

Basic intuition of networks training

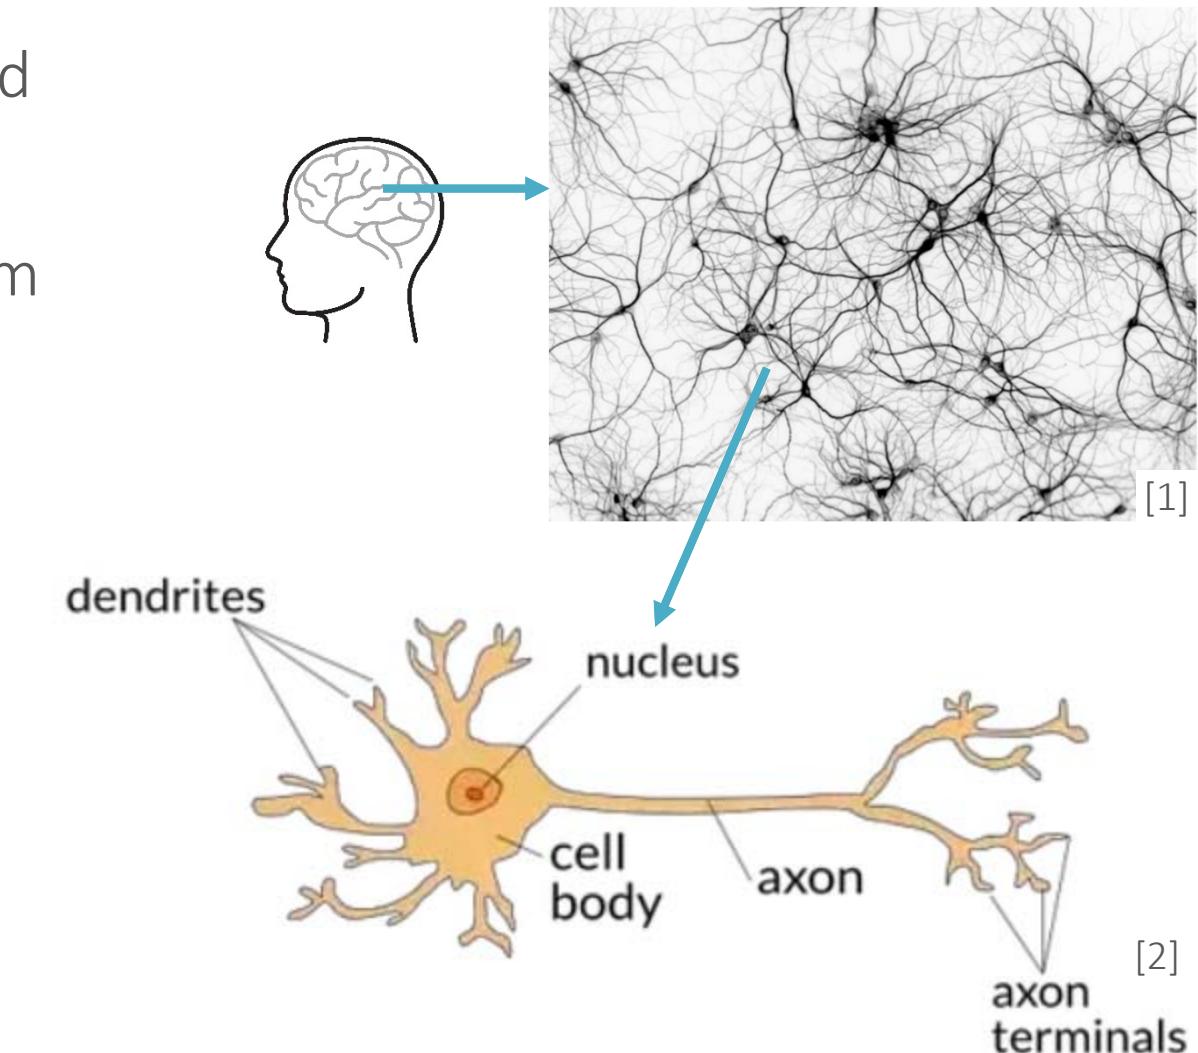
Example: Designing a network

Outlook

INTRODUCTION TO NEURAL NETWORKS

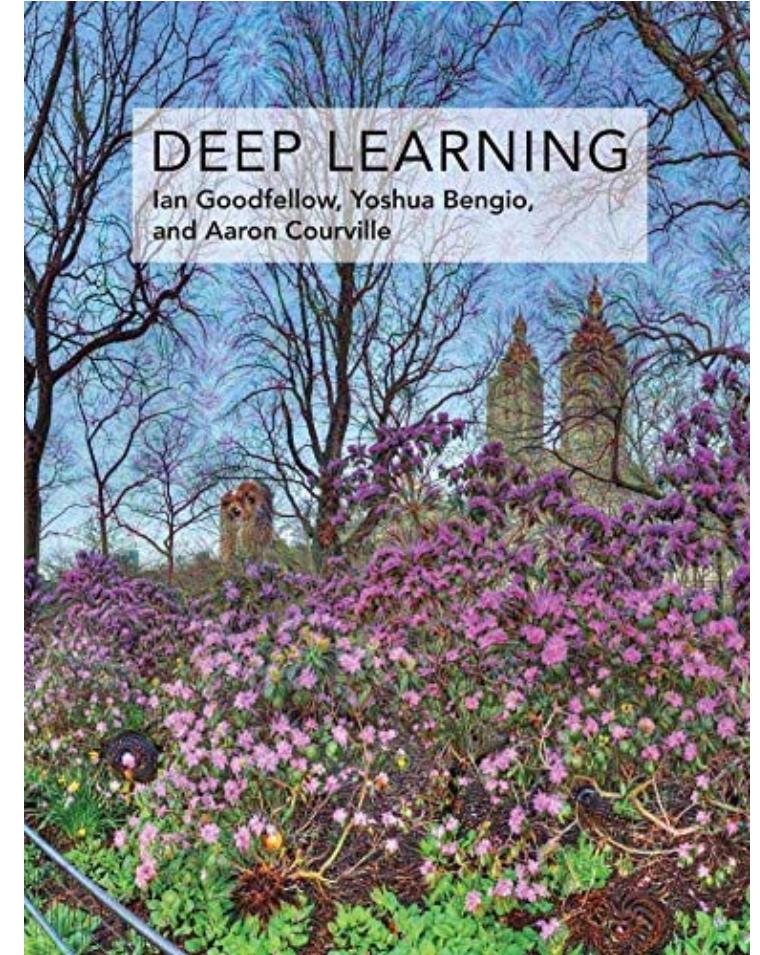
Neural Networks and the Human Brain

- › Neural network = machine learning method (among many others)
- › *Biologically-inspired* programming paradigm for learning from *data*
- › Model = human brain
 - › Extremely powerful and flexible
 - › Approx. 10^{11} neurons
 - › Processing power: 20Mbit/s (optical nerve)
 - › Neural network = *coarse approximation*



“Deep Learning”

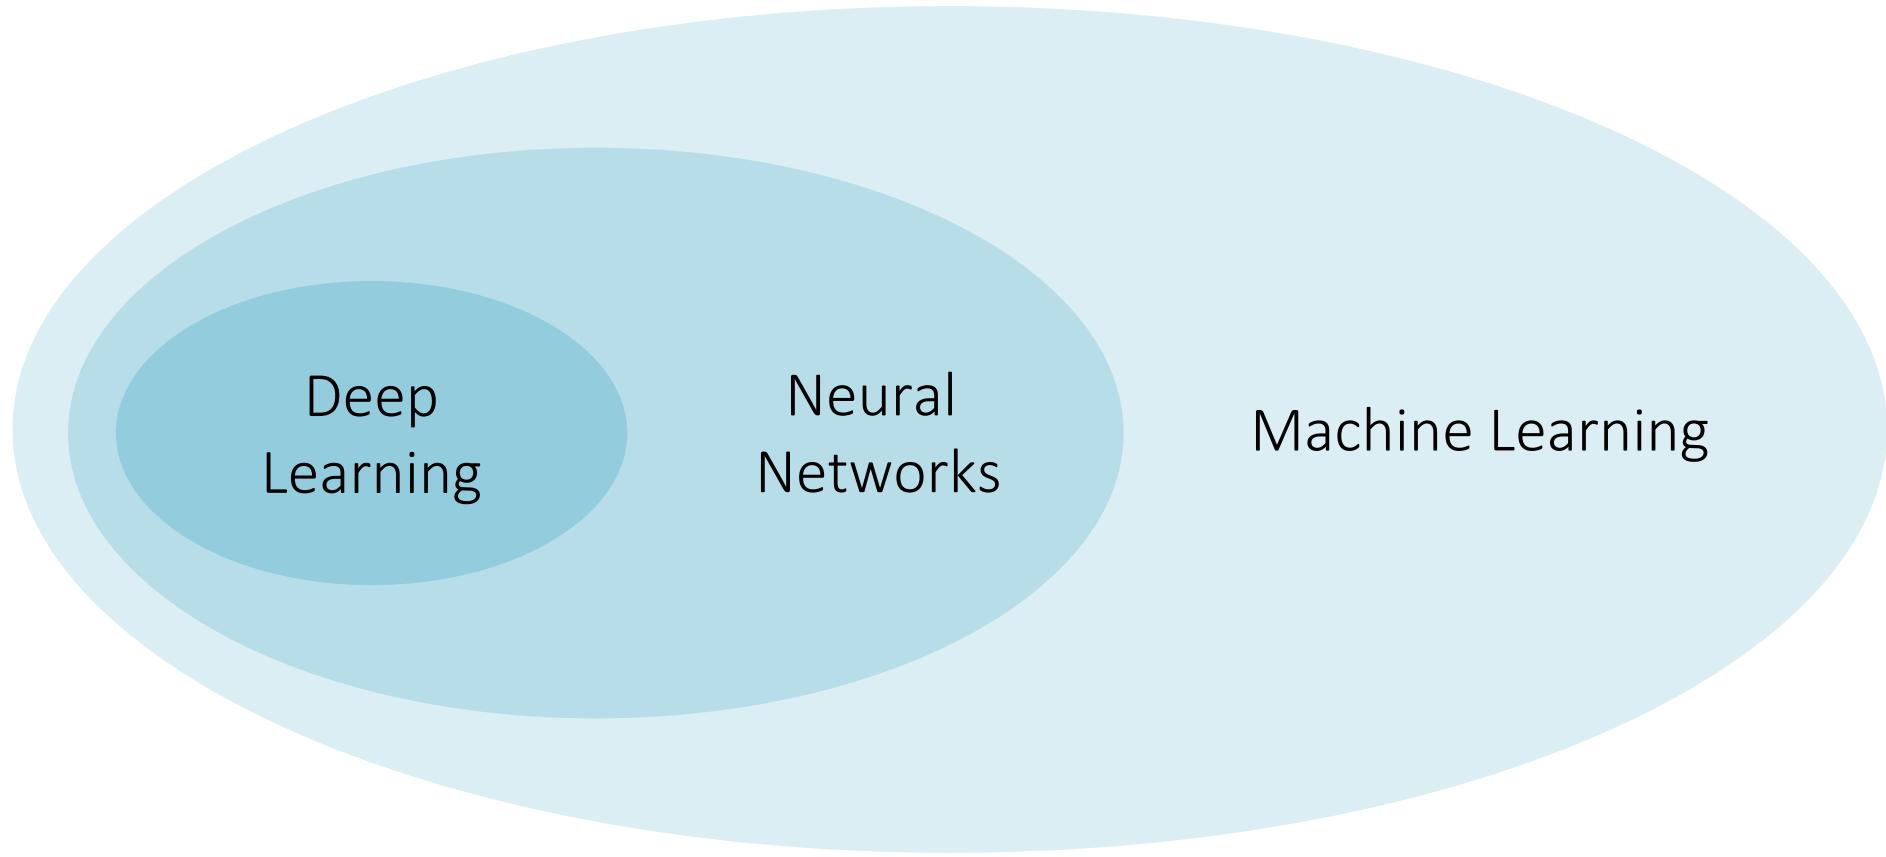
- › “Deep learning”: a powerful set of techniques for learning complex neural networks, i.e. with multiple layers
- › Deep learning = current state-of-the-art in
 - › Image recognition,
 - › Speech recognition,
 - › Natural language processing,
 - › ...
- › A network is considered deep when it has more than one *hidden layer*



Goodfellow et al., “Deep Learning”, MIT Press 2016
<http://www.deeplearningbook.org>

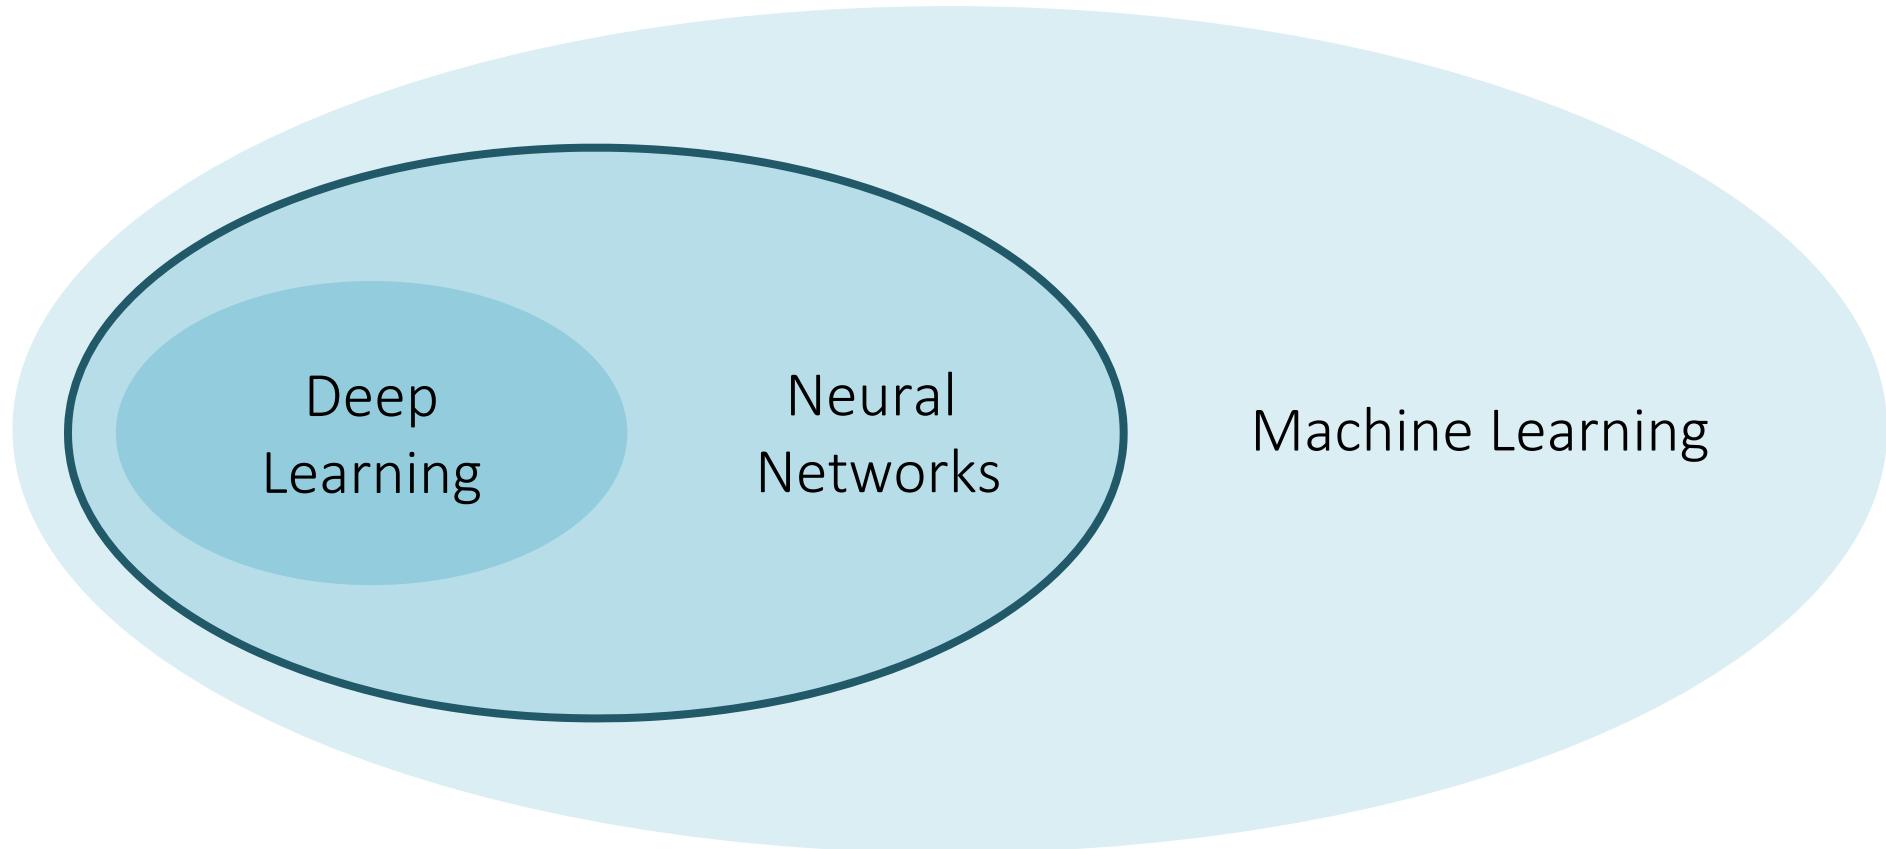
Methodological Categorization

Machine learning vs. neural networks vs. deep learning



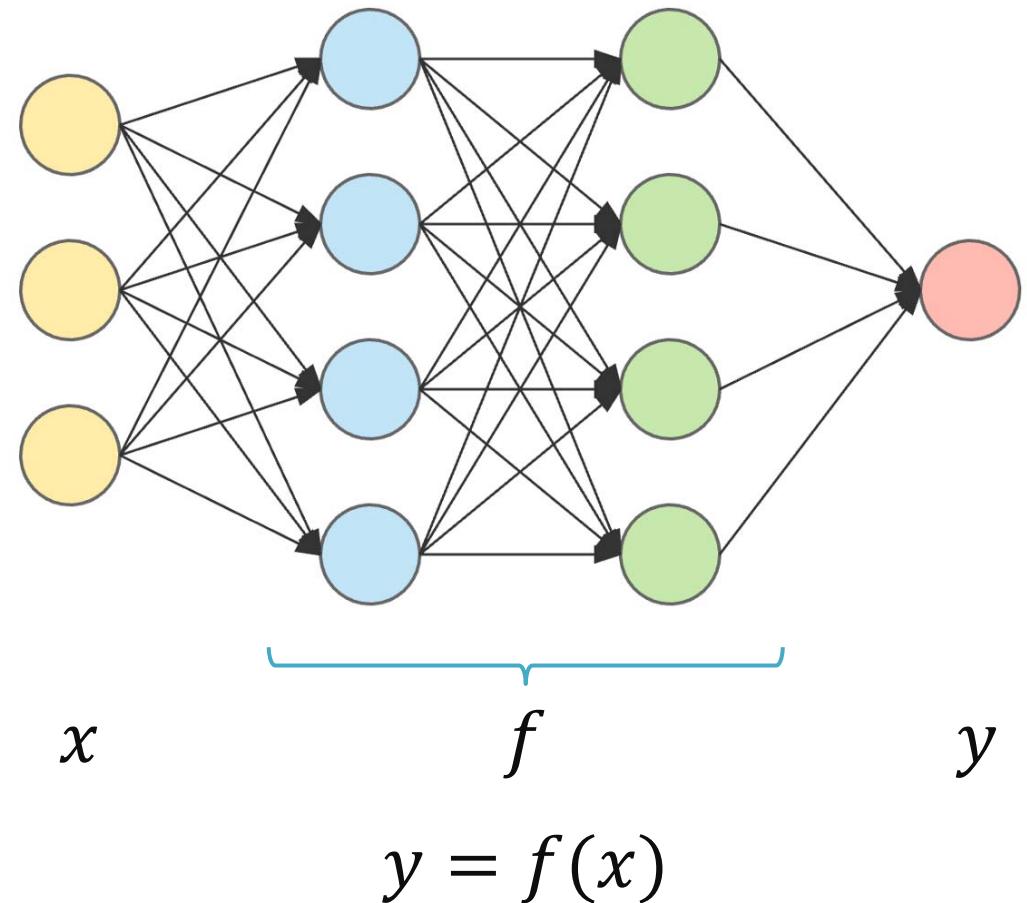
Methodological Categorization

Machine learning vs. neural networks vs. deep learning



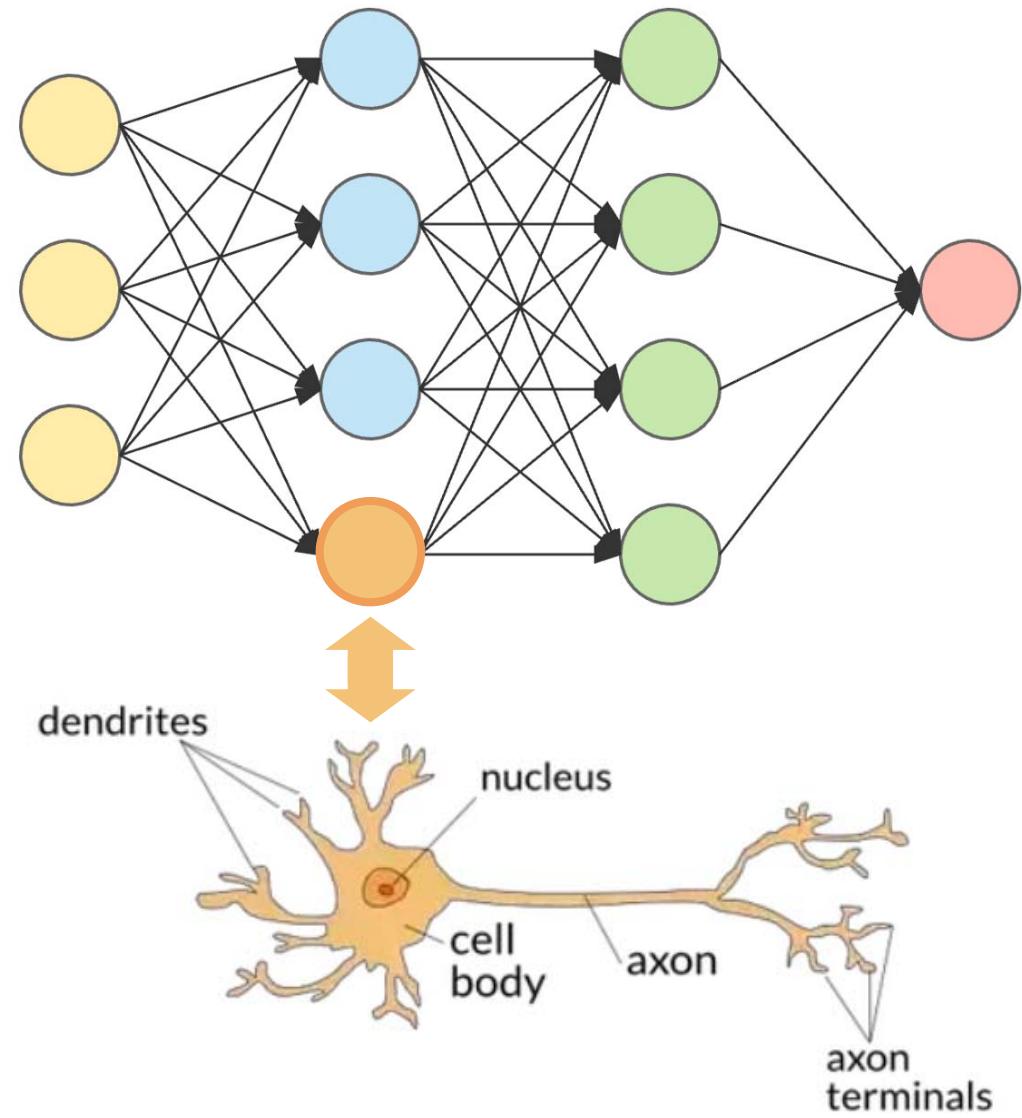
Neural Networks

- › Basic building blocks:
 - › Individual neurons
 - › Connections between neurons
 - › Layers (groups of neurons)
- › Formally: map input x to output y by a function f that represents our neural network



Building Blocks of Networks: the Neuron

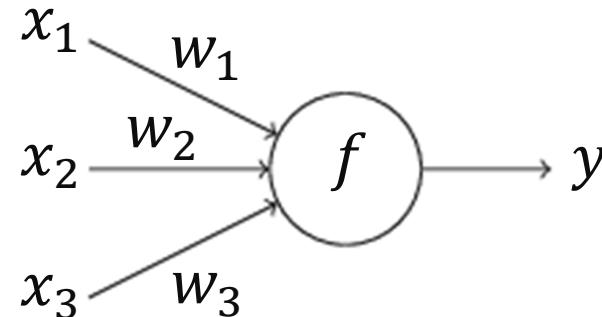
- › A network is built of *artificial neurons*
- › Different models of neurons exist (e.g. the *perceptron*)
- › Much simpler than the biological equivalent in our brain
- › Simplest networks: all neurons structurally the same



The Perceptron

- › Developed by Frank Rosenblatt in 1950s*

- › Early artificial neuron model



[https://en.wikipedia.org/
wiki/Frank_Rosenblatt](https://en.wikipedia.org/wiki/Frank_Rosenblatt)

- › A perceptron takes several inputs, x_1, x_2, \dots , and produces a single binary output y ($0|1$)
- › Weights, w_1, w_2, \dots , (real numbers) express the importance of the respective inputs

› Neuron function f :

$$y = \begin{cases} 0 & \text{if } \sum_j w_j x_j \leq \text{threshold} \\ 1 & \text{if } \sum_j w_j x_j > \text{threshold} \end{cases}$$

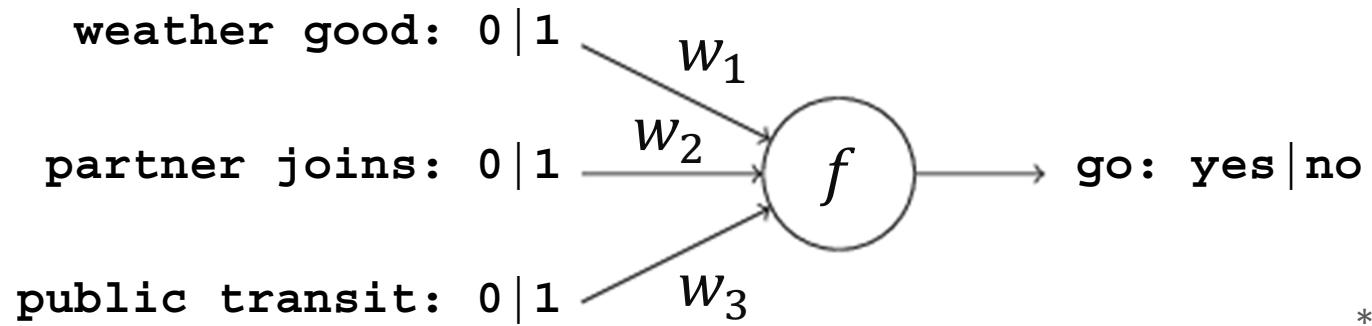
* The perceptron - a probabilistic model for information storage and organization in the brain.
F. Rosenblatt. *Psychological Review* 65, 1958.

The Perceptron...

... can be considered a decision maker

Shall I go to the Cheese Festival?*

- › The weekend is coming up, and you've heard that there's going to be a cheese festival in your city. You like cheese, and are trying to decide whether or not to go to the festival.
- › You might make your decision by weighing up three factors:
 1. Is the weather good? Yes | No
 2. Does your partner want to accompany you? Yes | No
 3. Is the festival near public transit (you don't own a car)? Yes | No
- › Build your own perceptron that helps you with the decision making:



How to set the weights?

Choose your preference..

* <http://neuralnetworksanddeeplearning.com/chap1.html>

Shall I go to the Cheese Festival?*

› Example

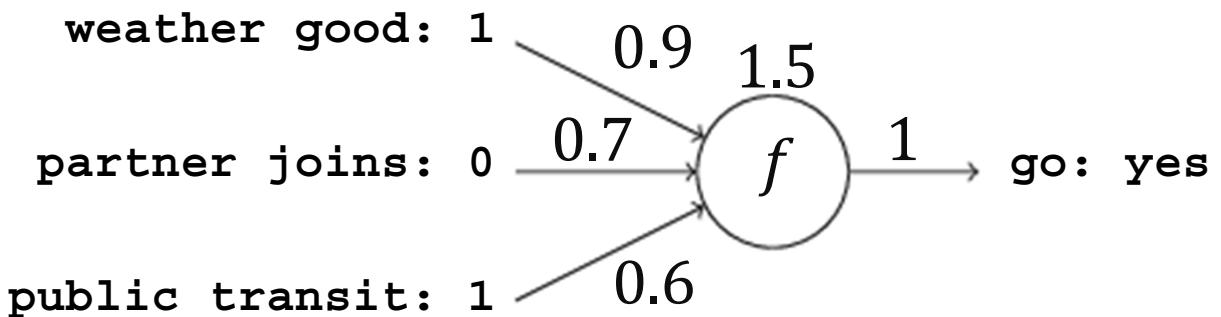
› I hate rain: $w_1 = 0.9$

› I prefer accompany: $w_2 = 0.7$ $w_j = 0$ means I don't care

› I am lazy: $w_3 = 0.6$

› threshold = 1

› Given:



$$y = 1 * 0.9 + 0 * 0.7 + 1 * 0.6 = 1.5 \Rightarrow f(x) = 1$$



$$y = \begin{cases} 0 & \text{if } \sum_j w_j x_j \leq \text{threshold} \\ 1 & \text{if } \sum_j w_j x_j > \text{threshold} \end{cases}$$

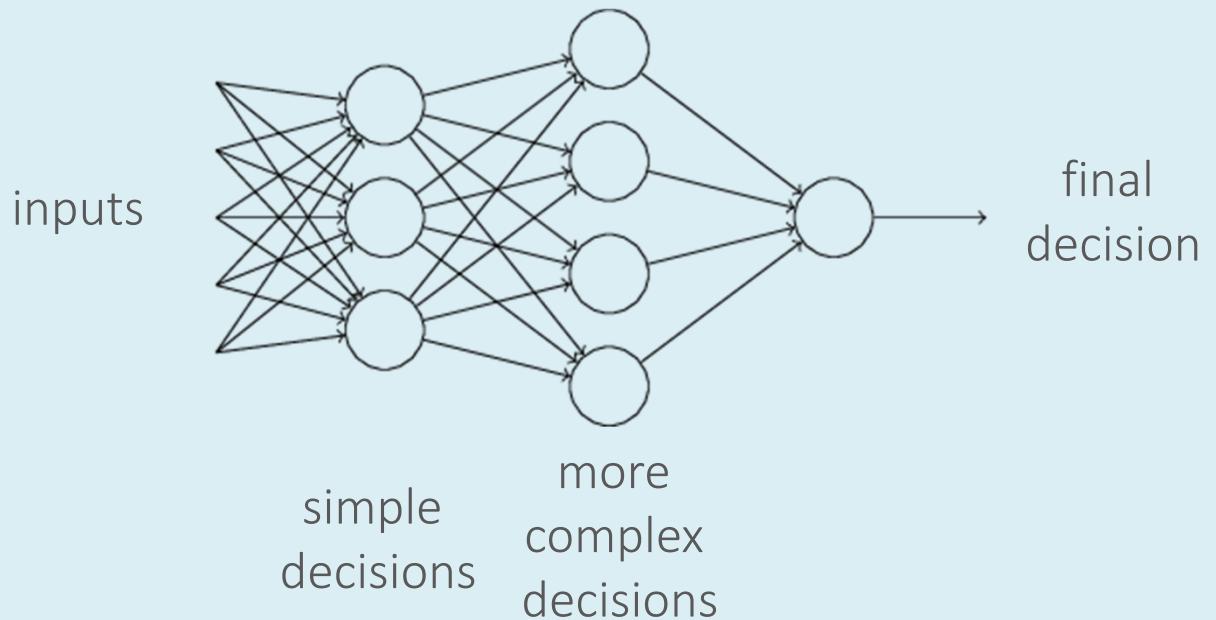
* <http://neuralnetworksanddeeplearning.com/chap1.html>

This was a Simple Decision..

... what about more complex decisions?

Example: whole group of people wants to jointly decide to go to the cheese festival

Connect multiple perceptrons to each other



Role of the Threshold

- › Question: how to set the threshold?

› Recap:

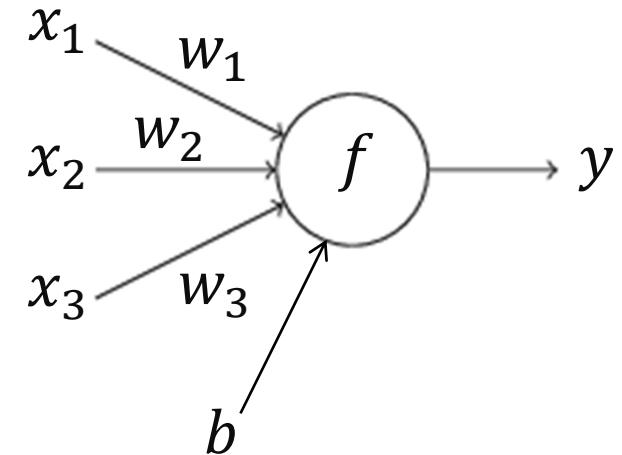
$$y = \begin{cases} 0 & \text{if } \sum_j w_j x_j \leq \text{threshold} \\ 1 & \text{if } \sum_j w_j x_j > \text{threshold} \end{cases}$$

- › Answer: learn it from data (same for weights)
- › Replace sum by dot product and threshold by variable “b”:

$$y = \begin{cases} 0 & \text{if } \mathbf{w} \cdot \mathbf{x} + b \leq 0 \\ 1 & \text{if } \mathbf{w} \cdot \mathbf{x} + b > 0 \end{cases}$$

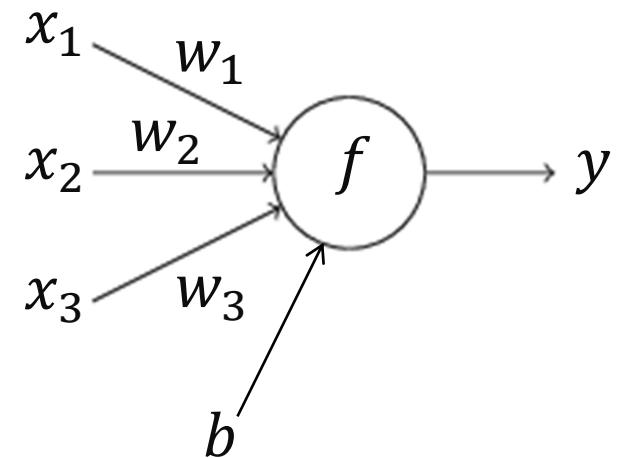
common notation in the literature

- › b is called “bias”, note: $b = -\text{threshold}$
- › A large bias means: it’s likely that the neuron outputs a 1 (i.e. it “fires”)



The Perceptron as a Classifier

- › Output: binary
- › Decision function: linear



Perceptron is a binary linear classifier

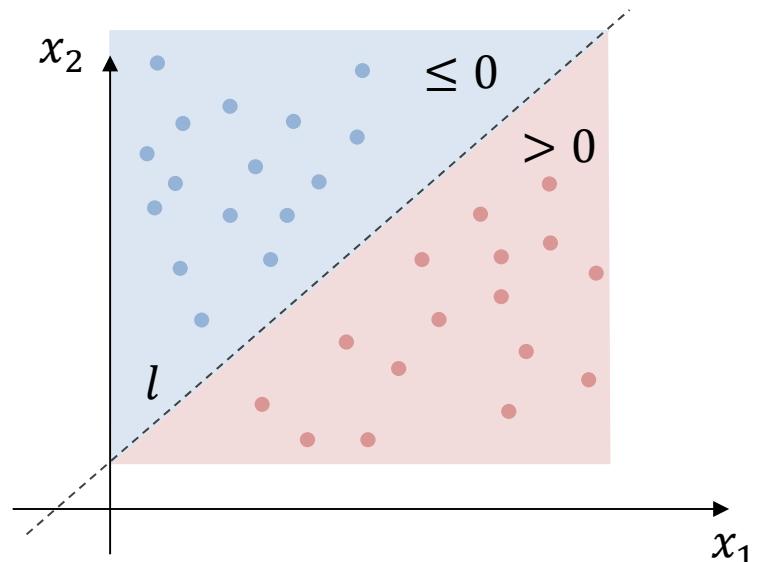
The Perceptron: Example in 2D

- › Perceptron is a binary linear classifier
- › Interpretation in 2D:
 - › Input: 2 variables x_1 and x_2
 - › 2 weights w_1 and w_2 , bias b
 - › Decision function f :

$$y = \begin{cases} 0 & \text{if } w_1 \cdot x_1 + w_2 \cdot x_2 + b \leq 0 \\ 1 & \text{if } w_1 \cdot x_1 + w_2 \cdot x_2 + b > 0 \end{cases}$$

- › Equations describe 2 regions separated by line:

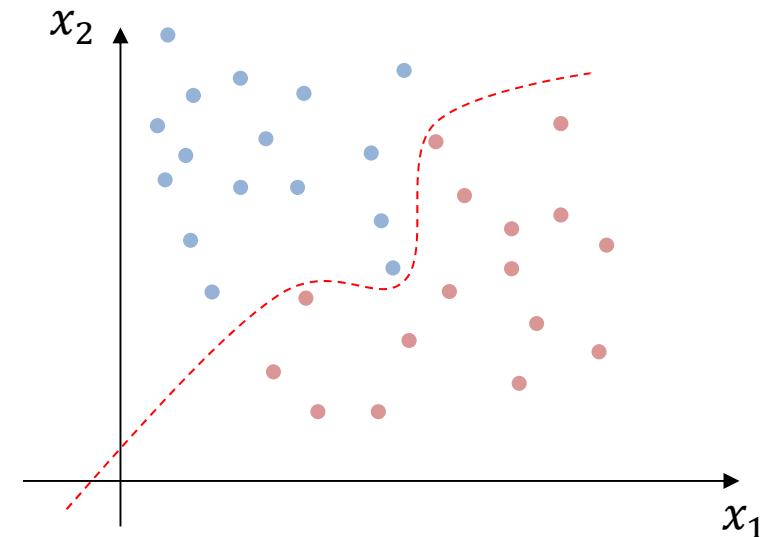
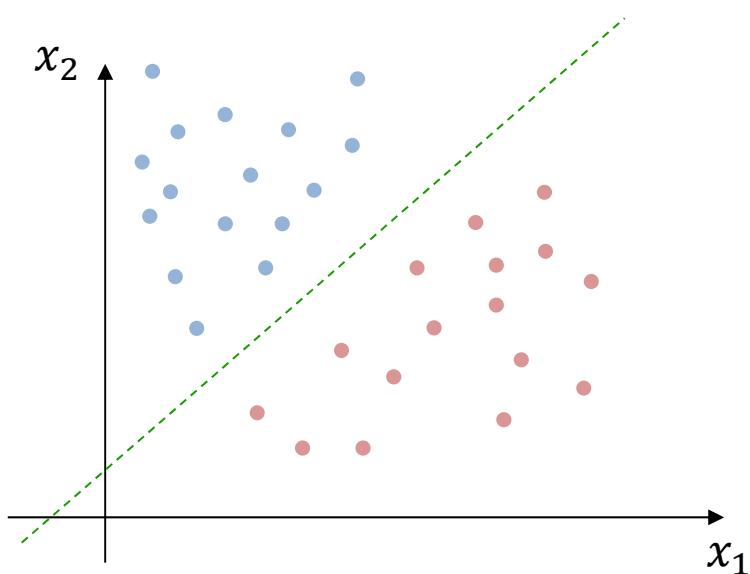
$$l: w_1 \cdot x_1 + w_2 \cdot x_2 + b = 0$$



$w = \begin{pmatrix} w_1 \\ w_2 \end{pmatrix}$ is the normal vector of the line

Limitations of the Perceptron

- › Perceptron can solve only linearly separable problems



The XOR Problem

- › Perceptron can solve only linearly separable problems
- › Simple logical operations:

$\begin{array}{c cc} 1 & 0 & 1 \\ 0 & 0 & 0 \end{array}$	$\begin{array}{c cc} 1 & 1 & 1 \\ 0 & 0 & 1 \end{array}$	$\begin{array}{c cc} 1 & 1 & 0 \\ 0 & 1 & 0 \end{array}$
AND	OR	NOT



- › The XOR Problem is not linearly separable:

$\begin{array}{c cc} 1 & 1 & 0 \\ 0 & 0 & 1 \end{array}$
XOR



led to the first AI winter (1970s)



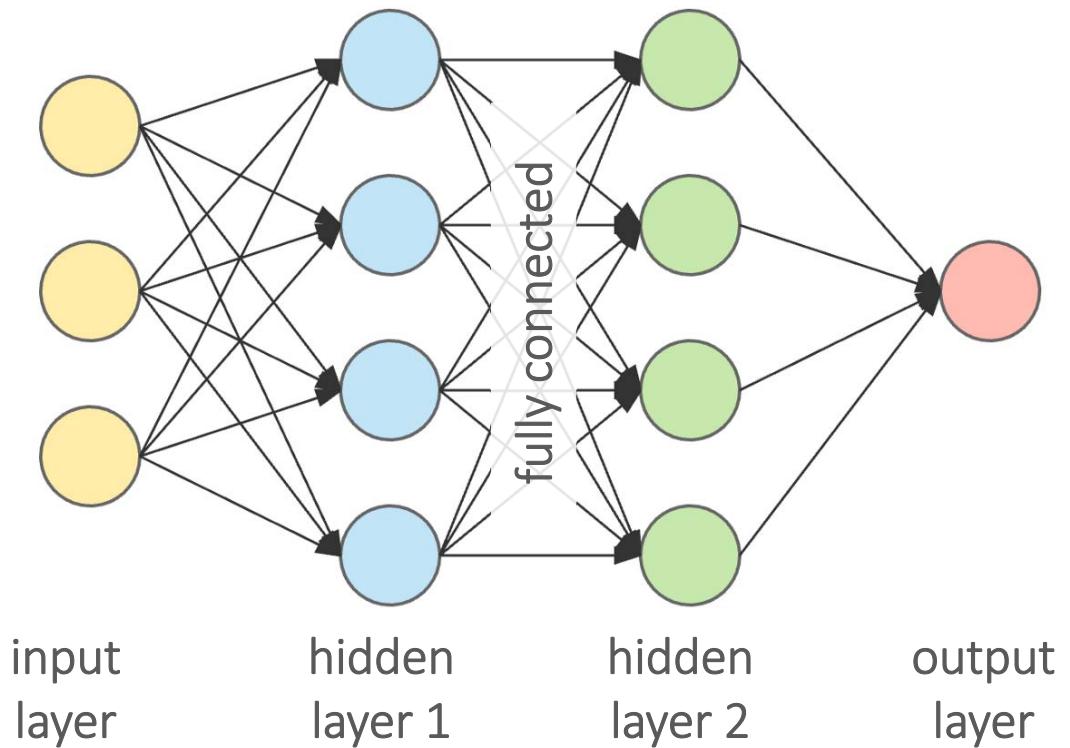
Minsky, M., & Papert, S. A. (1969). Perceptrons: An introduction to computational geometry. MIT press.

One Perceptron cannot do...

... but several can!

The Multi-Layer Perceptron

- › Combine multiple perceptrons to a network
- › Arrange perceptrons in layers
 - › Input layer
 - › Hidden layers
 - › Output layer
- › Connections between neurons
 - › full connectedness → fully connected layers
- › Feed-forward architecture



The Multi-Layer Perceptron

- › Good news:
 - › Can solve the XOR Problem (by learning multiple linear decision boundaries)
 - › MLP with more than one layer of perceptrons (with binary activation) can learn any Boolean function
- › Bad news:
 - › Perceptrons cannot learn non-linear relationships
 - › Inefficient for training



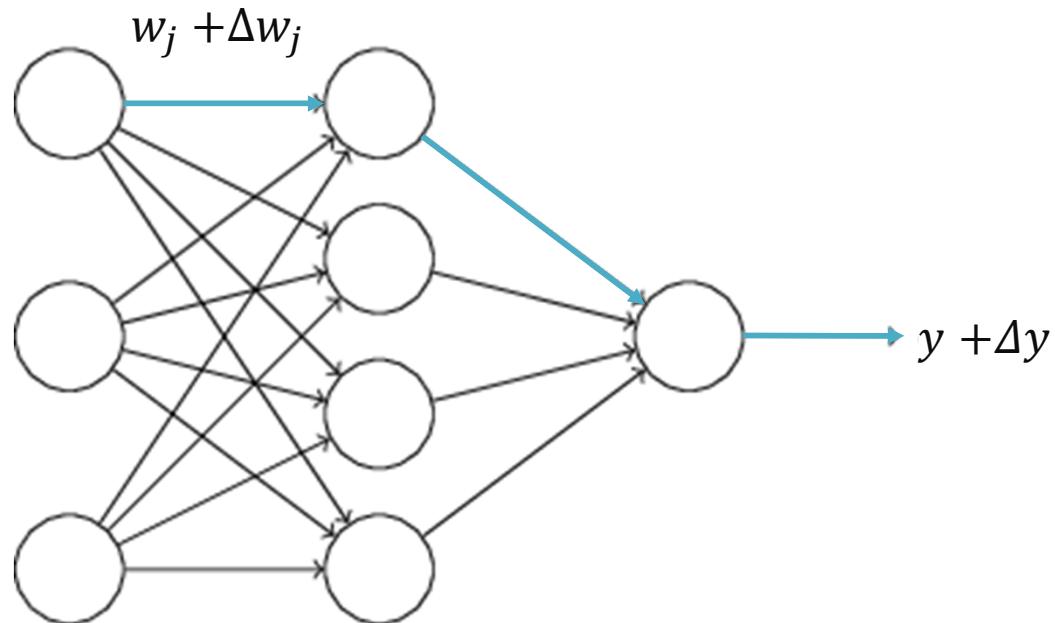
Strong limitation in practice!

What's wrong with the Perceptron?

- › Consider we want to train an MLP
- › Manually setting weights and biases = tedious!
 - › Even simple MLPs easily have several thousands of weights and biases
- › Better: estimate weights and biases from the data (and its labels)
- › Basic idea:
 - › Start with random values for weights and biases
 - › Change weights and biases so that desired output is achieved
- › Basic requirement: continuity

Pre-requisite for Learning

- › Learning requires continuity (no “jumps”)
- › Small change in learning parameter → small change in output



Not the case for the perceptron (output may flip from 0 to 1 abruptly), remember:

$$y = \begin{cases} 0 & \text{if } \sum_j w_j x_j \leq \text{threshold} \\ 1 & \text{if } \sum_j w_j x_j > \text{threshold} \end{cases}$$

- › Difficult to assess if change by Δw_j is beneficial
- › Danger to “forget” already learned patterns

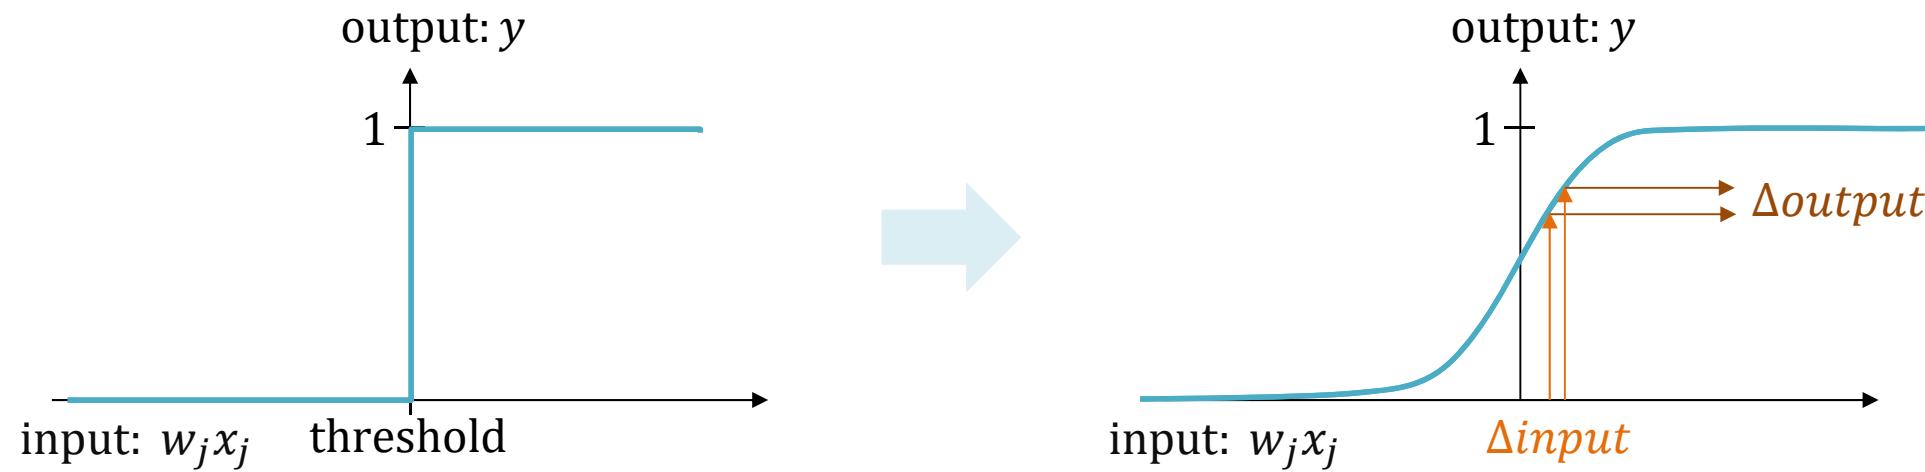
What's wrong with the Perceptron?

The binary output!

We need a neuron with continuous output!

Smooth the Activation Function

- Replace step-function of perceptron with a continuous “step-like” function



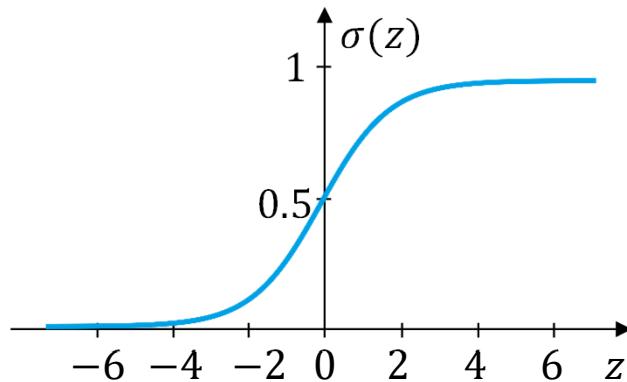
$$y = \begin{cases} 0 & \text{if } \sum_j w_j x_j \leq \text{threshold} \\ 1 & \text{if } \sum_j w_j x_j > \text{threshold} \end{cases}$$

Small change in input leads to
small change in output!

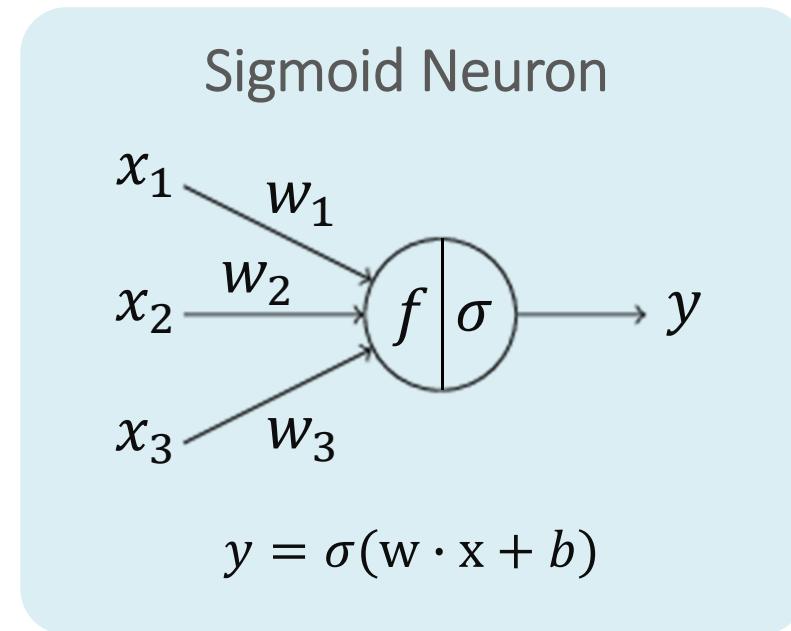
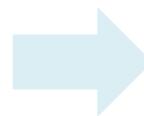
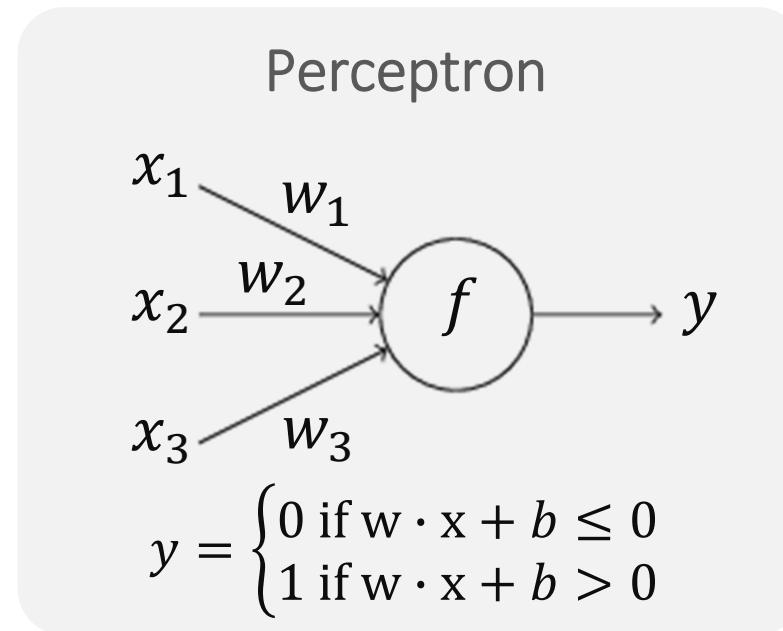
The Sigmoid Neuron

- › The sigmoid function:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

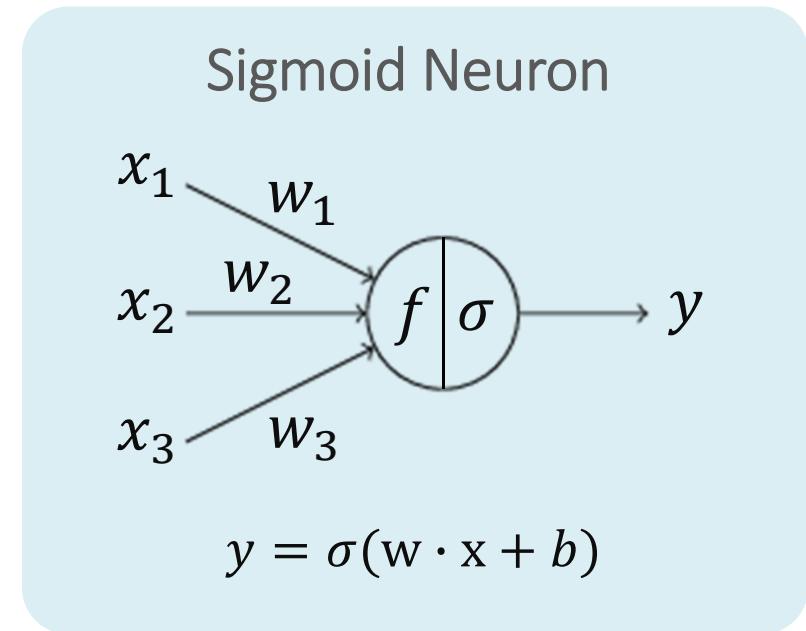
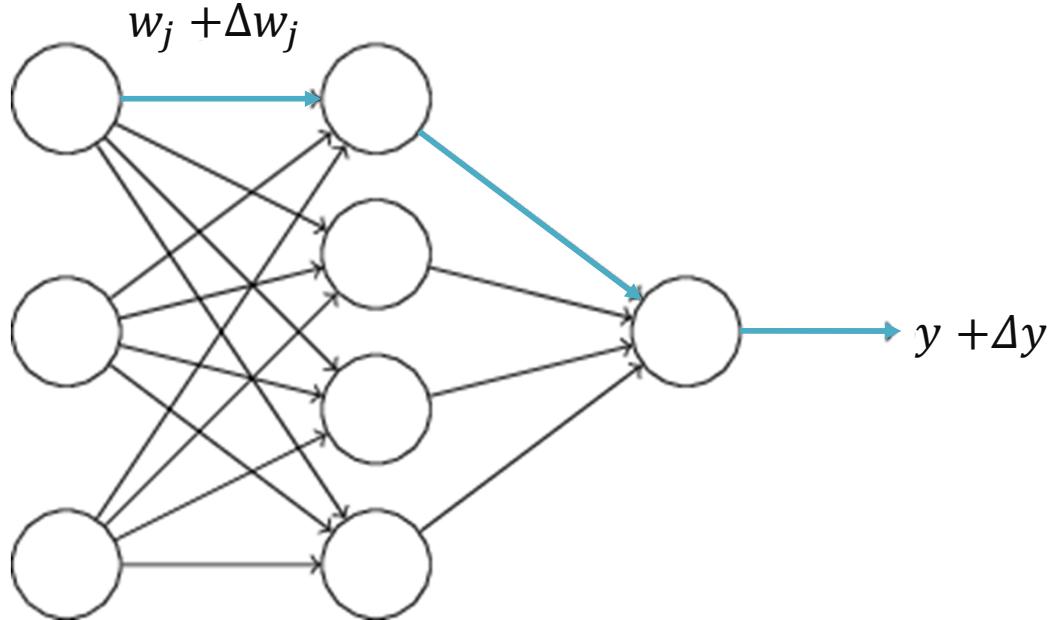


- › Scales inputs to 0..1
- › For large inputs: same as step function of perceptron
- › For moderate inputs: non-linear scaling



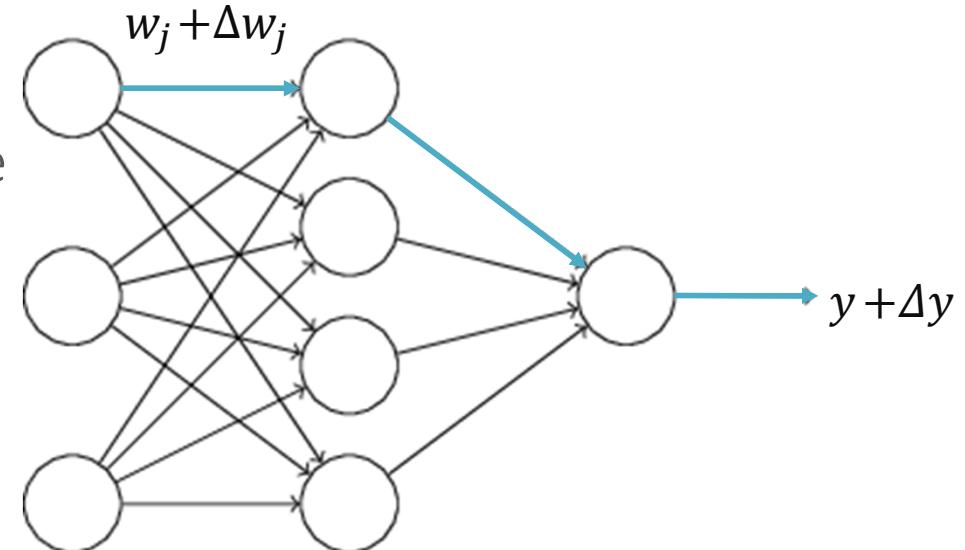
Learning with Sigmoid Neurons

- › Sigmoid is smooth: small change in parameter \rightarrow small change in output



Learning with Sigmoid Neurons

- › Learn from data $X = \{x^{(1)}, x^{(2)}, \dots, x^{(N)}\}$ to predict target Y given by labels $\{\hat{y}^{(1)}, \hat{y}^{(2)}, \dots, \hat{y}^{(N)}\}$
- › Intuition behind learning
 - › Start with random weights & biases
 - › Adjust weights by small Δw so that outputs $y^{(i)}$ of the network gets closer to desired outputs $\hat{y}^{(i)}$
- › How to adjust the weights to get closer to the desired output?

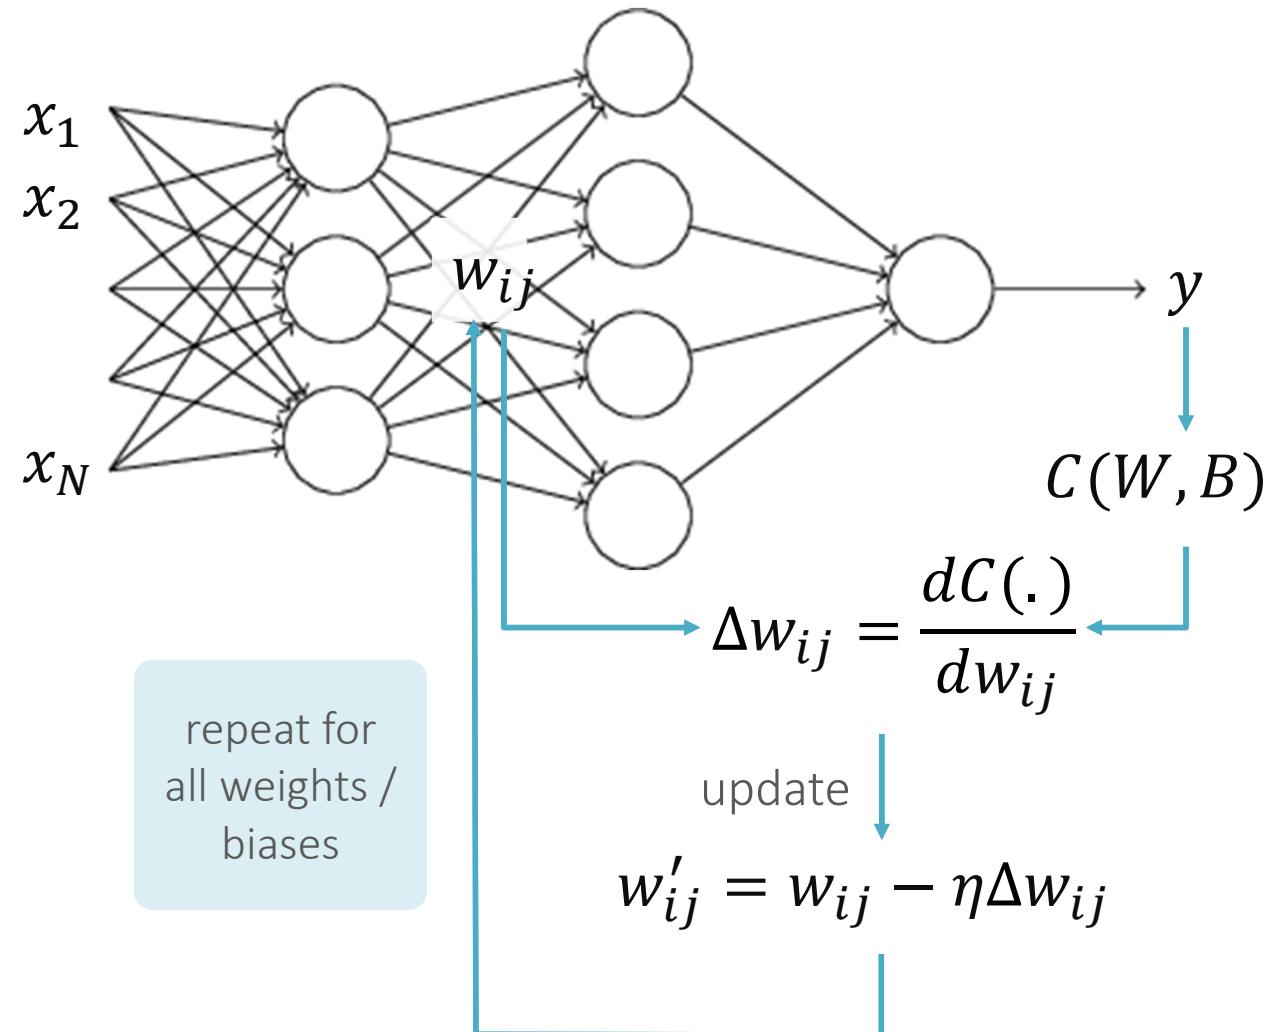


The Backpropagation Algorithm – An Intuition

- Estimate output quality by a cost function:

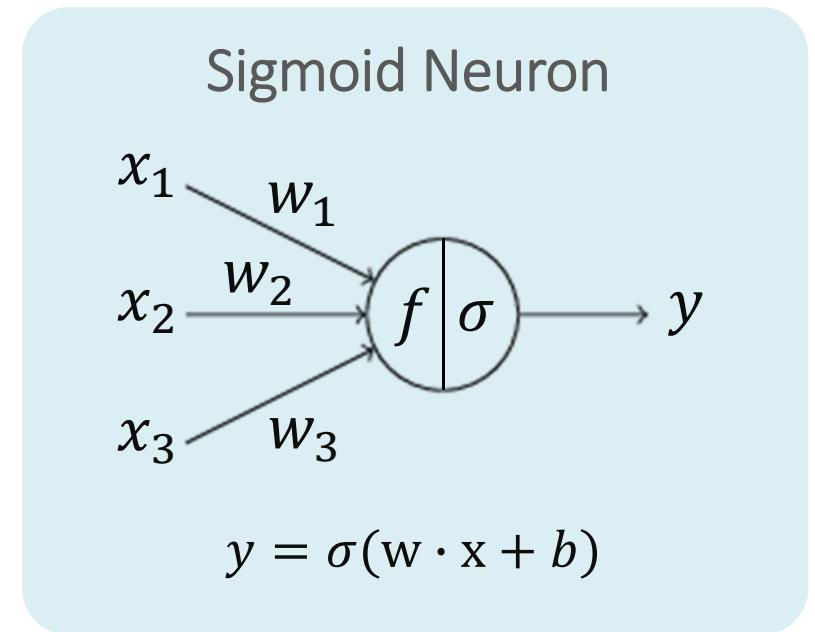
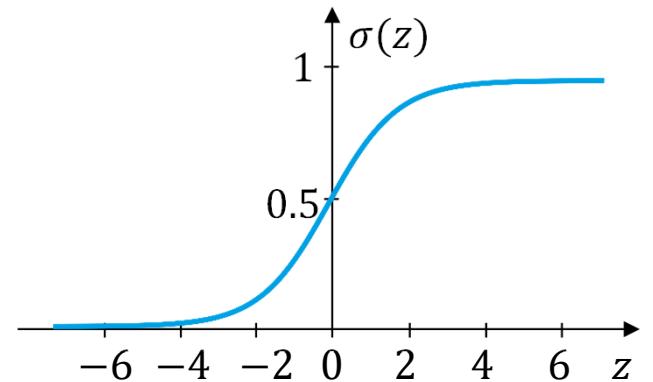
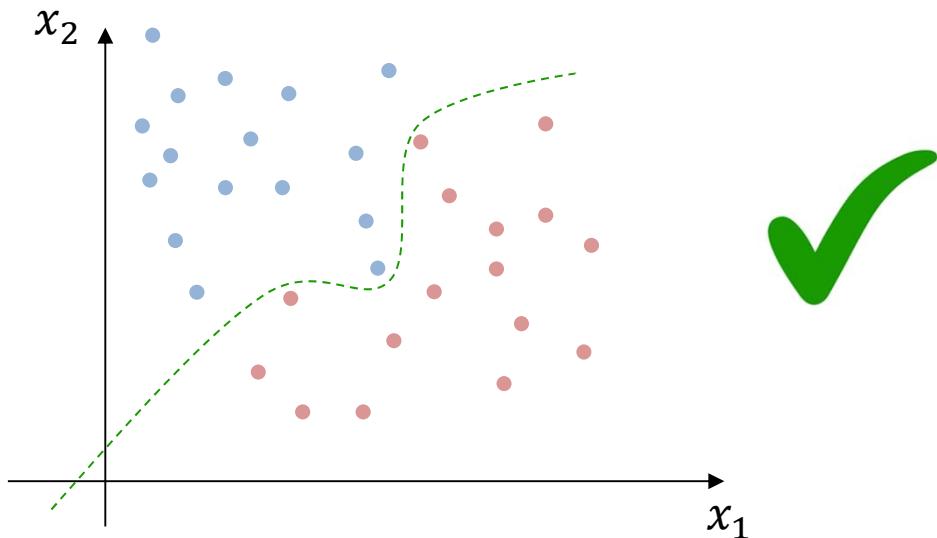
$$C(W, B) = \frac{1}{2N} \sum_i \|y - \hat{y}\|$$

- Goal: find Δw that decreases cost function
- Backpropagation algorithm provides an efficient solution for this.
 - Idea: estimate influence of parameter w_{ij} on output (=partial derivative w.r.t. cost function)
 - This tells us, how to change parameter w_{ij} to get closer to y' (and thereby to minimize C)
 - η is the learning rate



Making Non-Linear Decision

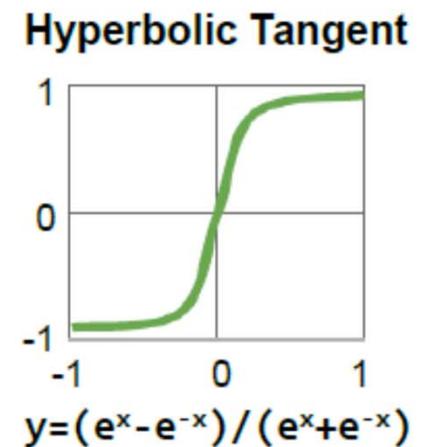
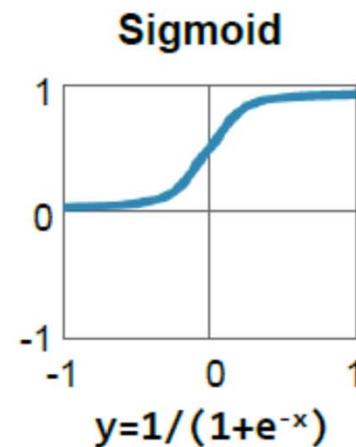
- › The sigmoid neuron is not only smooth, it is also non-linear → allows to model non-linear relationships
- › MLP based on sigmoid neuron is a non-linear classifier which can approximate nonlinear decision boundaries!



Activation Functions

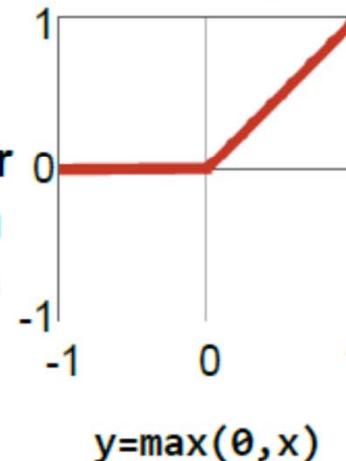
- › The sigmoid function in case of the sigmoid neuron is the “activation function” of the neuron
- › Many other activation functions have been designed
- › Important: they should be continuous and differentiable to facilitate learning!

Traditional Non-Linear Activation Functions

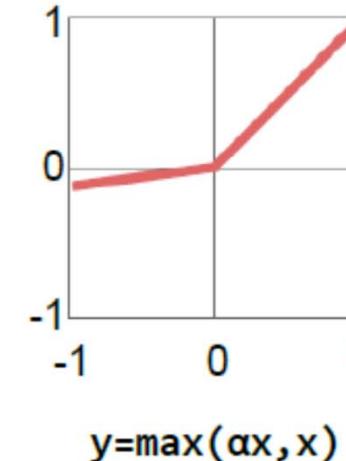


Modern Non-Linear Activation Functions

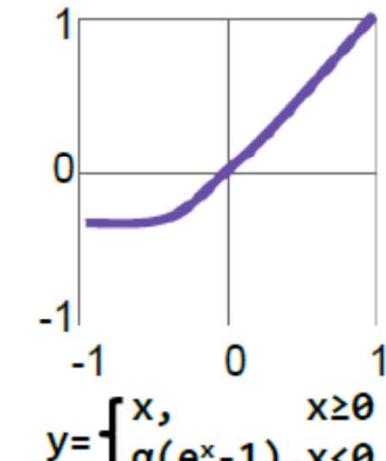
Rectified Linear Unit (ReLU)



Leaky ReLU



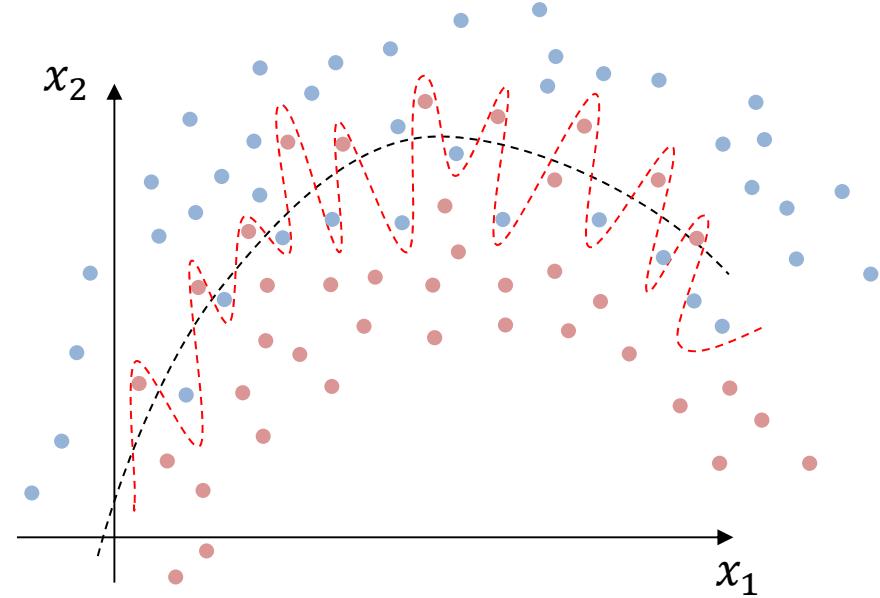
Exponential LU



$\alpha = \text{small const. (e.g. 0.1)}$

Practical Considerations in Network Training

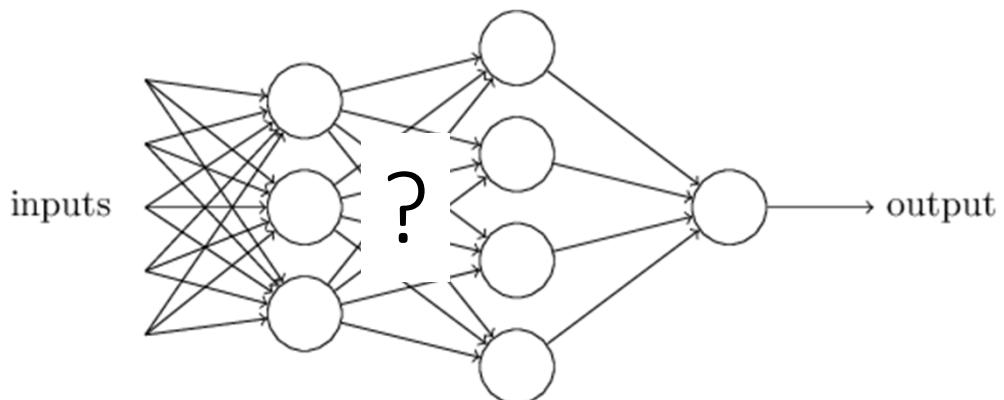
- › MLP based on sigmoid neurons can model highly-complex data and decision boundaries
- › Tendency to *overfit* on the data
- › No guarantee that an optimal solution is found because the search space for finding a *good* function is huge!
- › *Good* means a smooth function that fits the statistical *distribution* of the input data
- › Data always limited: a neural network is only as good as its training data allows



Model individual points
→ poorly generalizes to new data

Designing a Neural Network

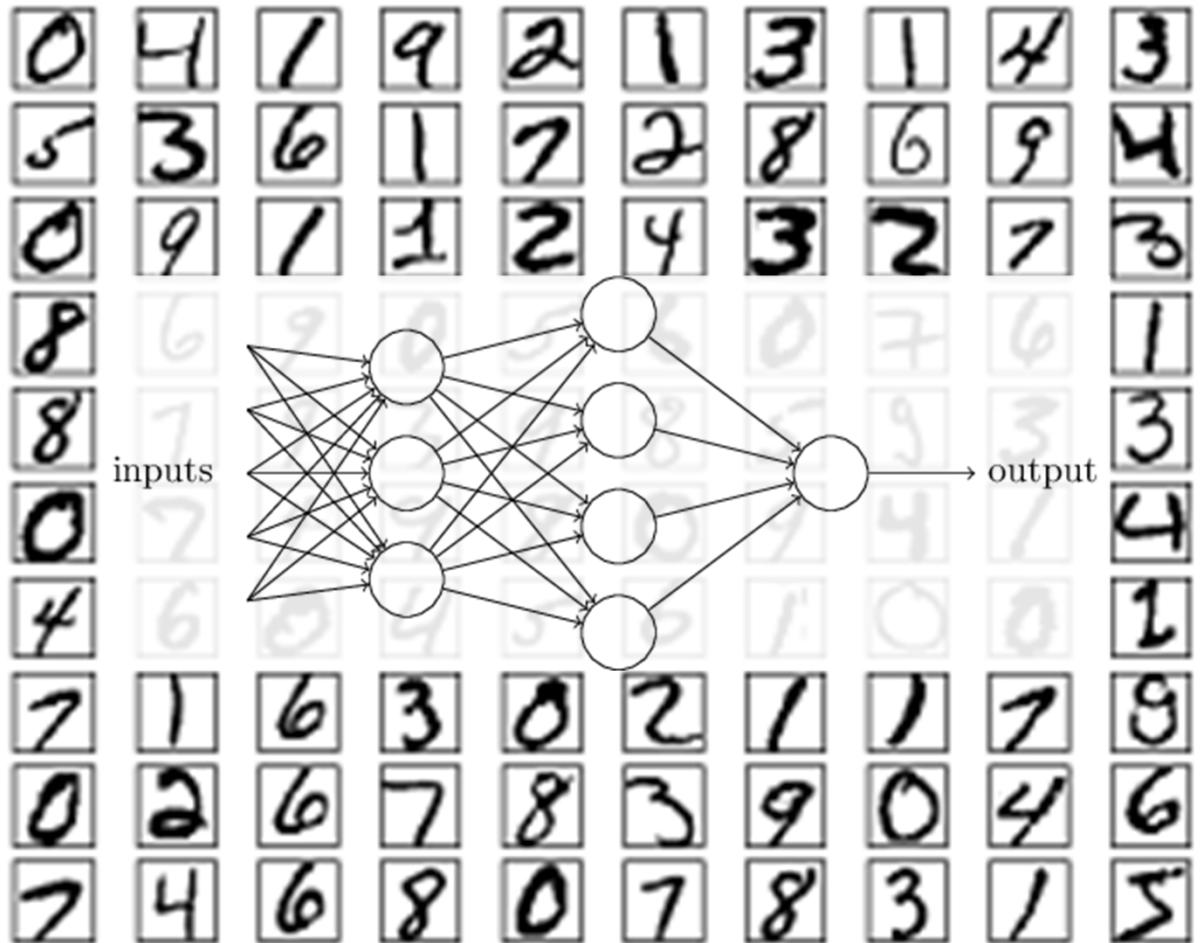
- › Task: predict hand written digits
- › Input: image (size $w \times h$), here $w=28$, $h=28$
- › Output: class $(0, 1, \dots, 9)$
- › How could an MLP solution look like?



MNIST Dataset: <http://yann.lecun.com/exdb/mnist/>

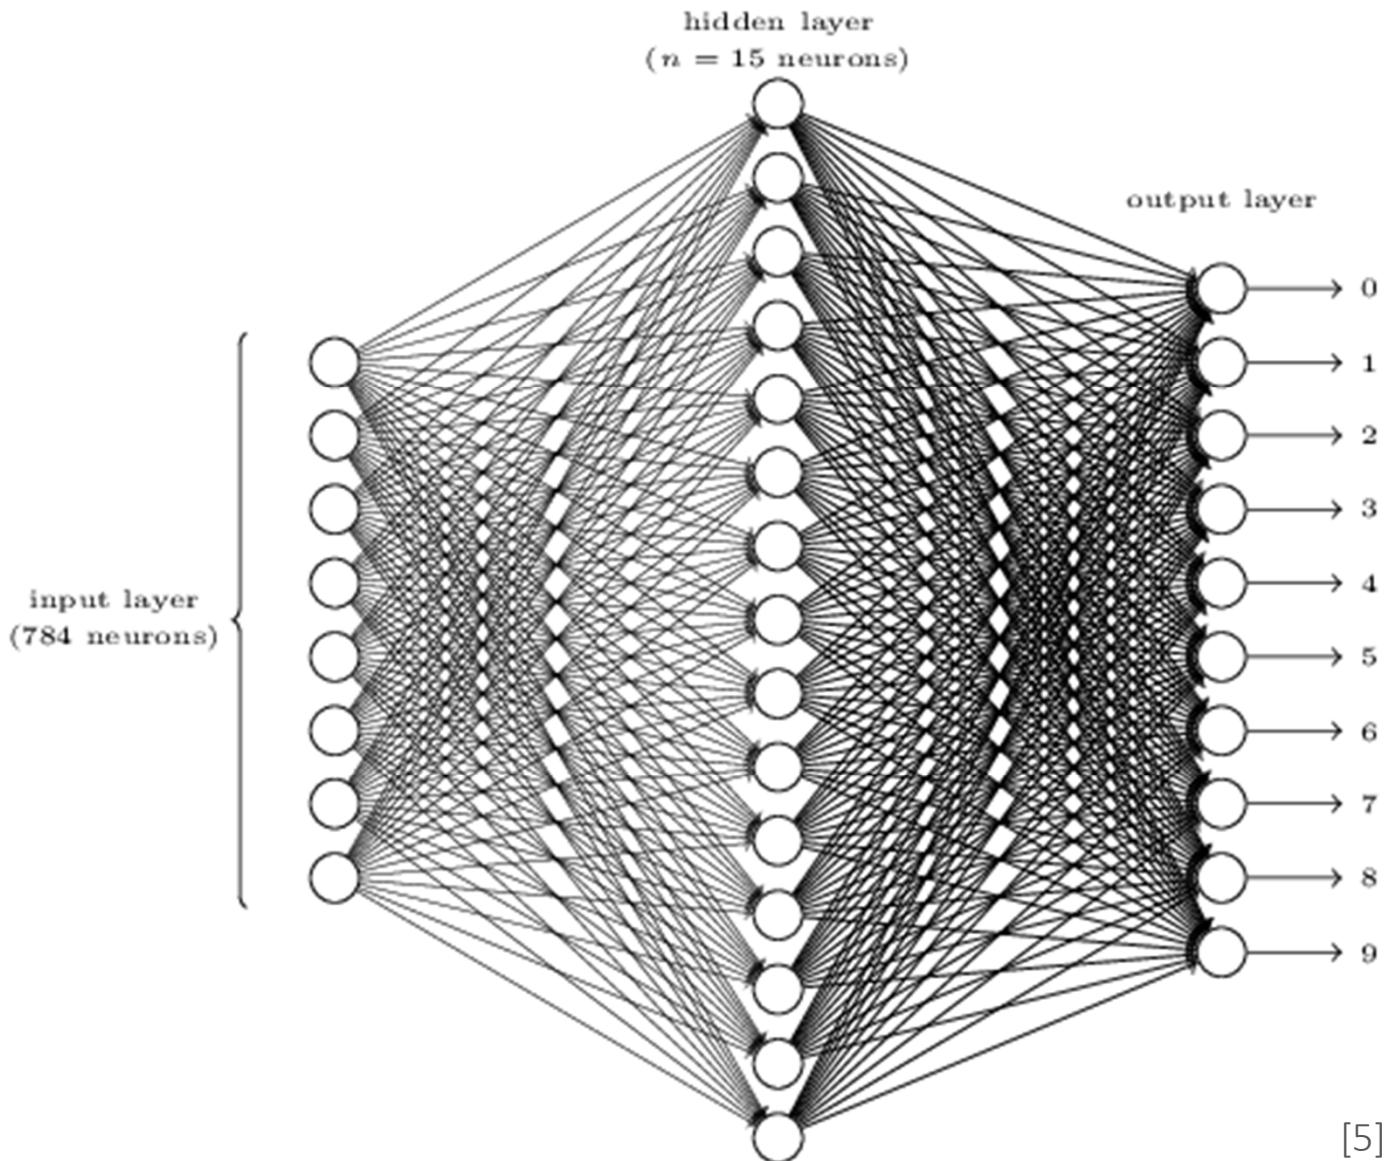
Designing a Neural Network

- › Input layer:
 - › $w \times h$ inputs ($28 \times 28 = 784$)
 - › Grayscale images: 1.0=white, 0.0=black
 - › Input layer simply provides the raw data – no neurons
 - › Here: inputs = raw image pixel values
- › Output layer:
 - › 10 Outputs: 0, 1, ..., 9 (one per class)
- › Hidden layers
 - › How many layers? / neurons per layer?
 - › Which connection pattern?



Designing a Neural Network

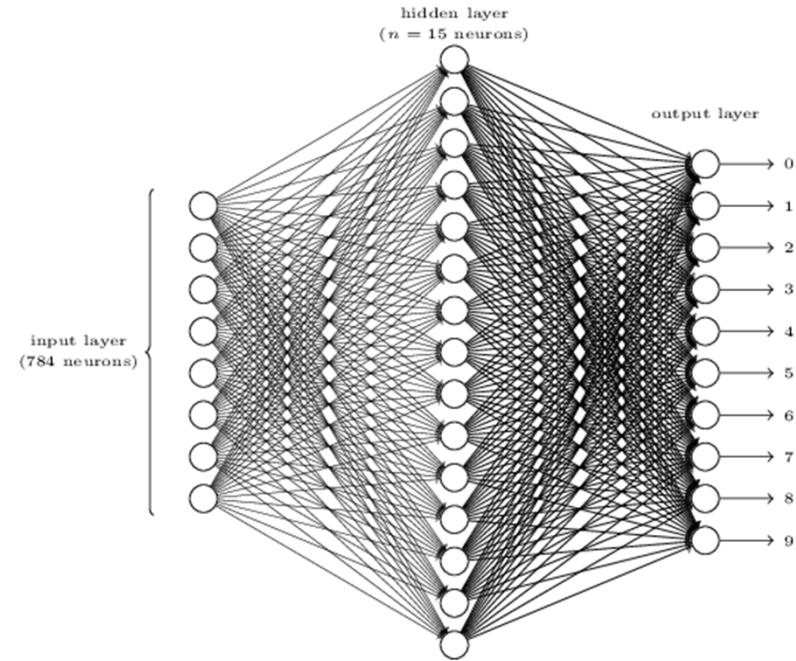
- › Example:
 - › 1 hidden layer
 - › 15 neurons
 - › Fully connected with input and output layer
- › Number of parameters:
 - › Weights: $784 * 15 + 15 * 10 = 11910$
 - › Biases: $15 + 10 = 25$
- › Result:
 - › Accuracy > 90% on MNIST dataset



[5]

Design Considerations

- › Open questions:
 - › Why 10 output neurons?
 - › Why not 15 hidden neurons?
 - › Why not 2 hidden layers
- › Non-trivial problem, one of the most difficult problems with neural networks
- › No rule of thumb for selecting the number of layers and the sizes of the layers (*hyper parameters*)
- › Network design is one of the core challenges today!



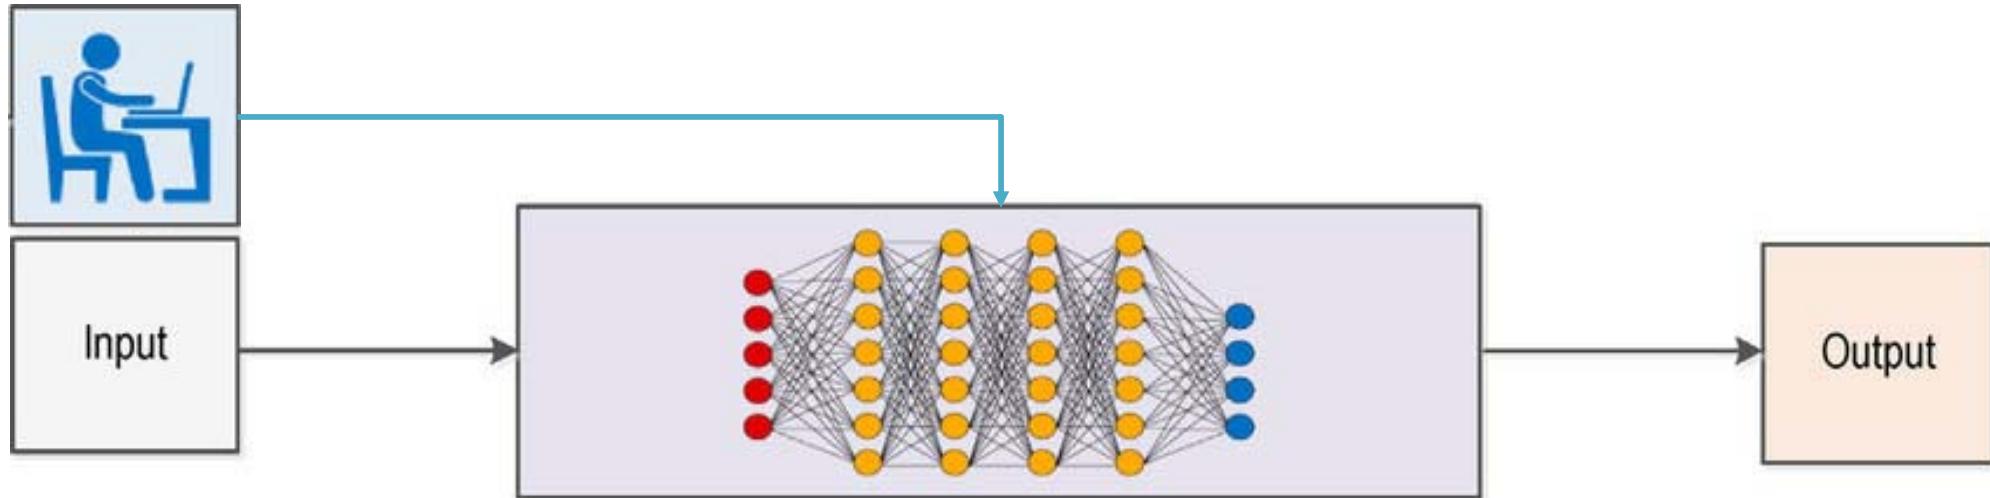
Design Tradeoff



- › Traditional design approach: *hand-crafted features*



- › Neural network design approach: *hand crafted-networks*



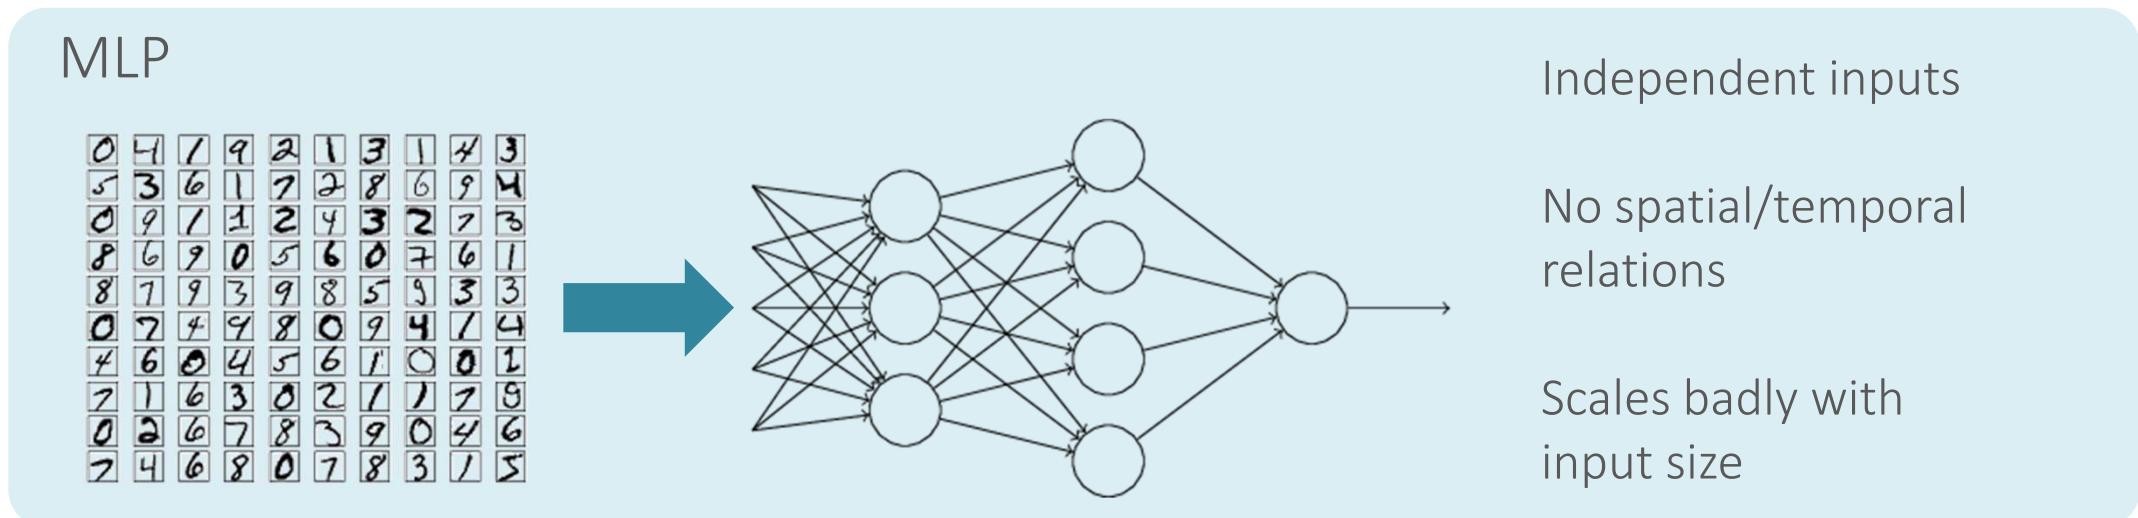
OUTLOOK

Limitations of MLP

Further network architectures

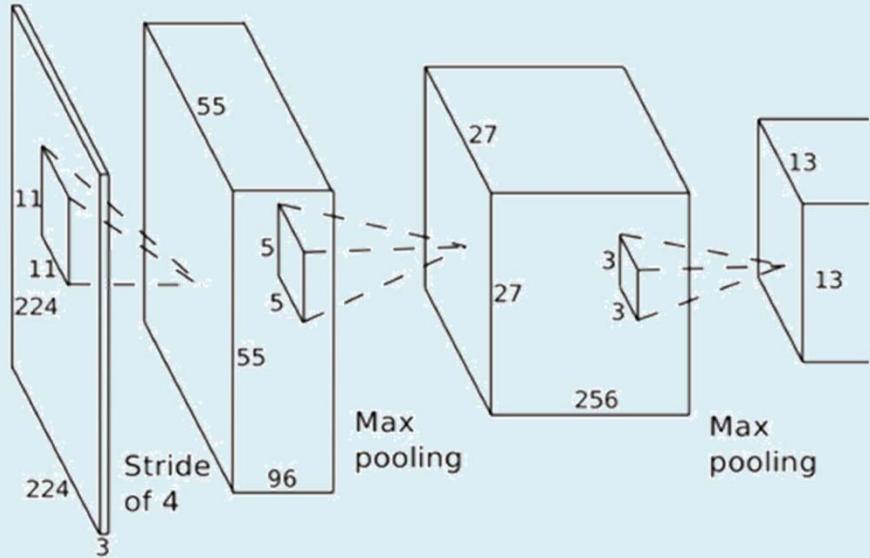
Limitations of MLP

- › MLP architecture
- › Good for structured data (e.g. demographic data – predict income)
- › Suboptimal for unstructured data: images, videos, audio/time series



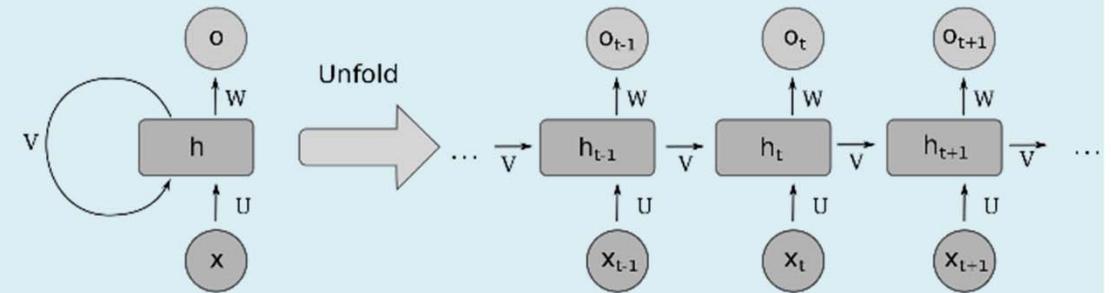
Outlook: More Powerful Network Architectures

Convolutional Neural Networks (CNNs)



Spatial data (e.g. image, video)

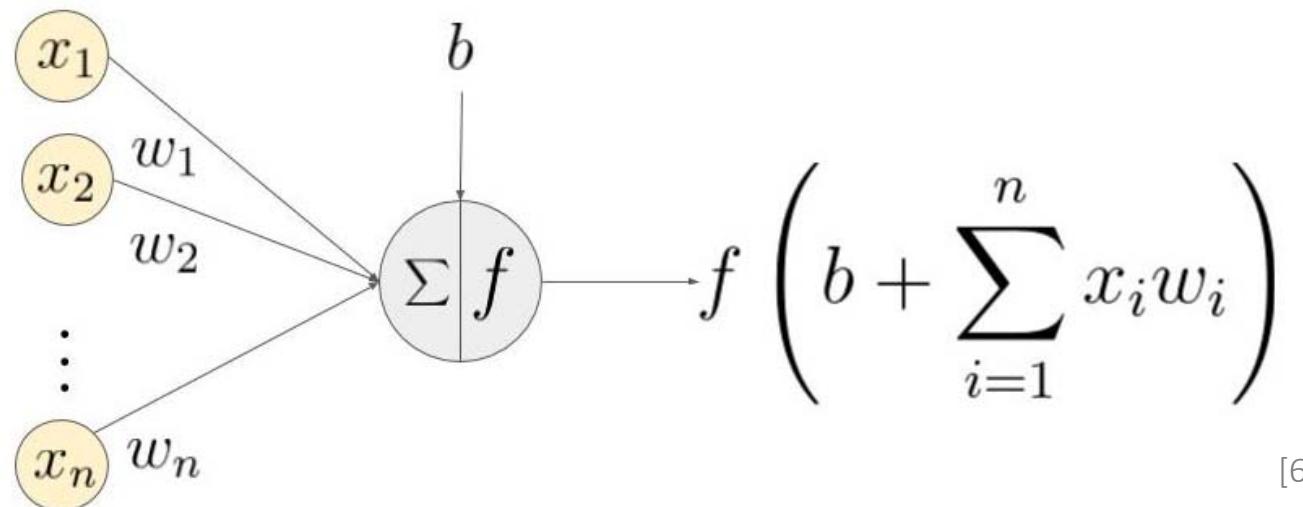
Recurrent Neural Networks (RNNs)



Temporal data (e.g. audio, time series)

Conclusion

- › Introduction to neural networks
- › Neurons: the perceptron, the sigmoid neuron
- › Basic intuition of networks training
- › Layers, multi-layer perceptron (MLP)
- › Example: Designing a network
- › Outlook



[6]

Contact:

matthias.zeppelzauer@fhstp.ac.at

Credits & References

- [1] <https://fractalfoundation.org/OFC/OFC-1-6.html>
- [2] <https://towardsdatascience.com/the-differences-between-artificial-and-biological-neural-networks-a8b46db828b7>
- [3] <https://www.sciencedirect.com/science/article/abs/pii/S0278612518300037>
- [4] <https://pixabay.com/es/users/cherebus0-1373501/>
- [5] <http://neuralnetworksanddeeplearning.com/chap1.html>
- [6] image credit: [Aditya Sharma](#)